

Lesson 7

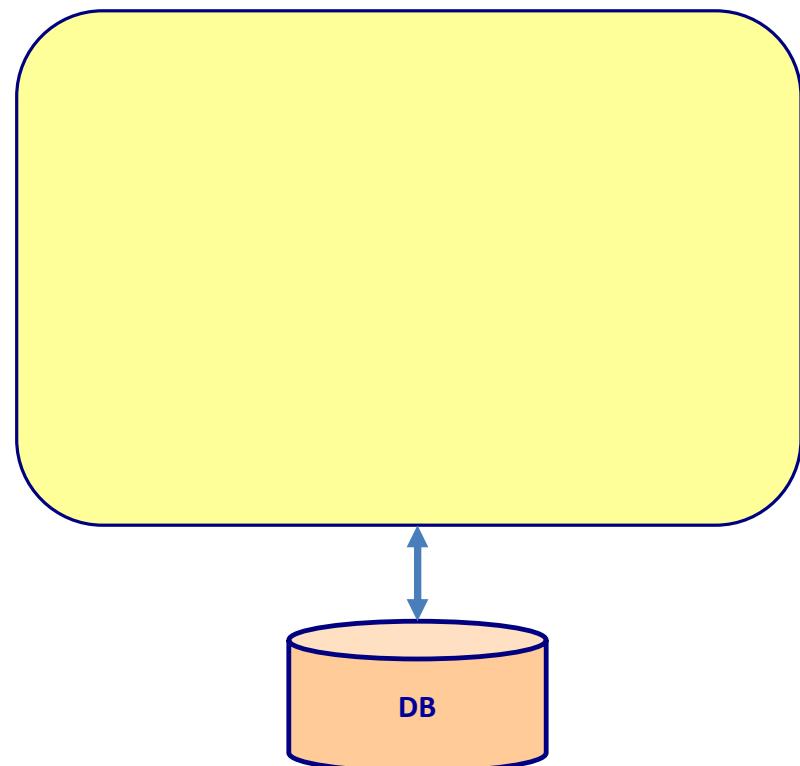
# **MICROSERVICES**

# **MONOLITH ARCHITECTURE**

# Monolith architecture

---

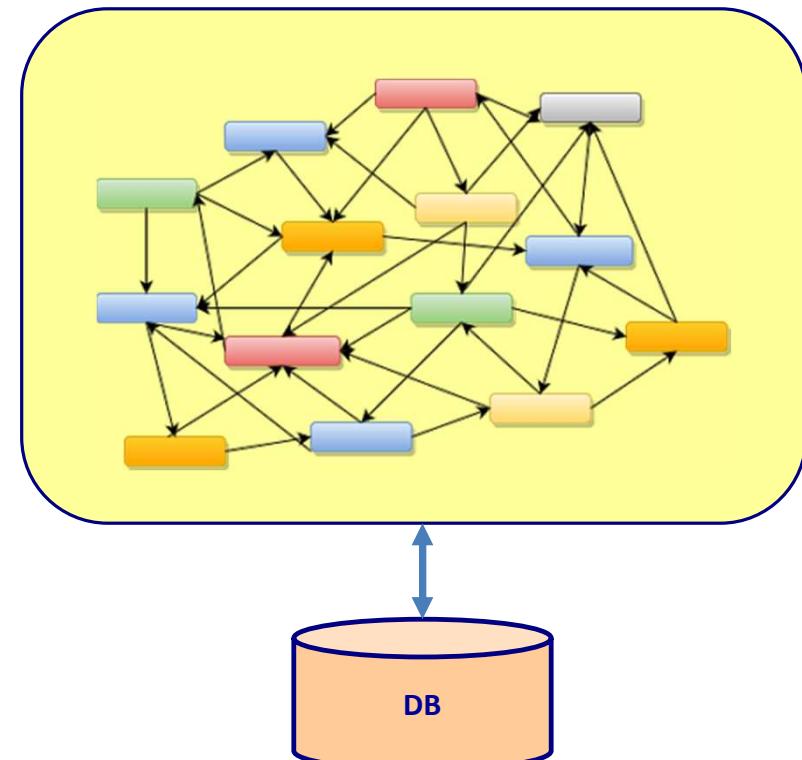
- Everything is implemented in one large system



# Monolith architecture

---

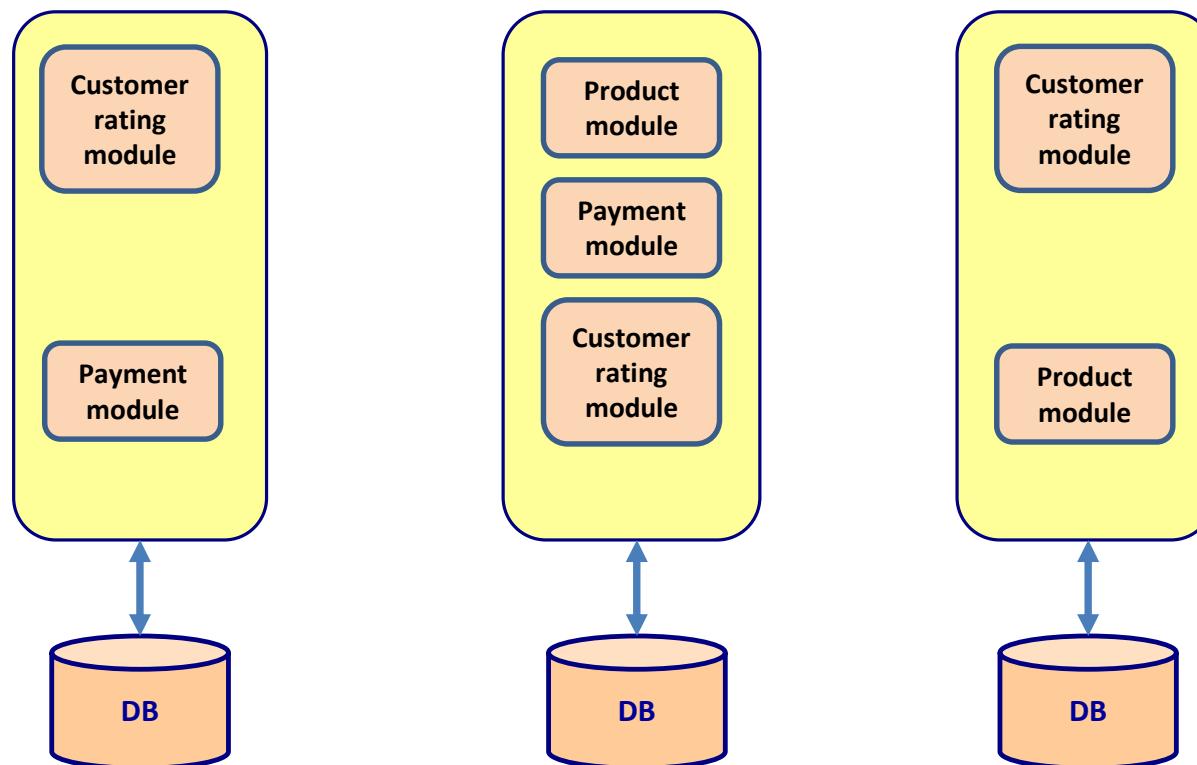
- Can evolve in a big ball of mud
  - Large complex system
    - Hard to understand
    - Hard to change



# Monolith architecture

---

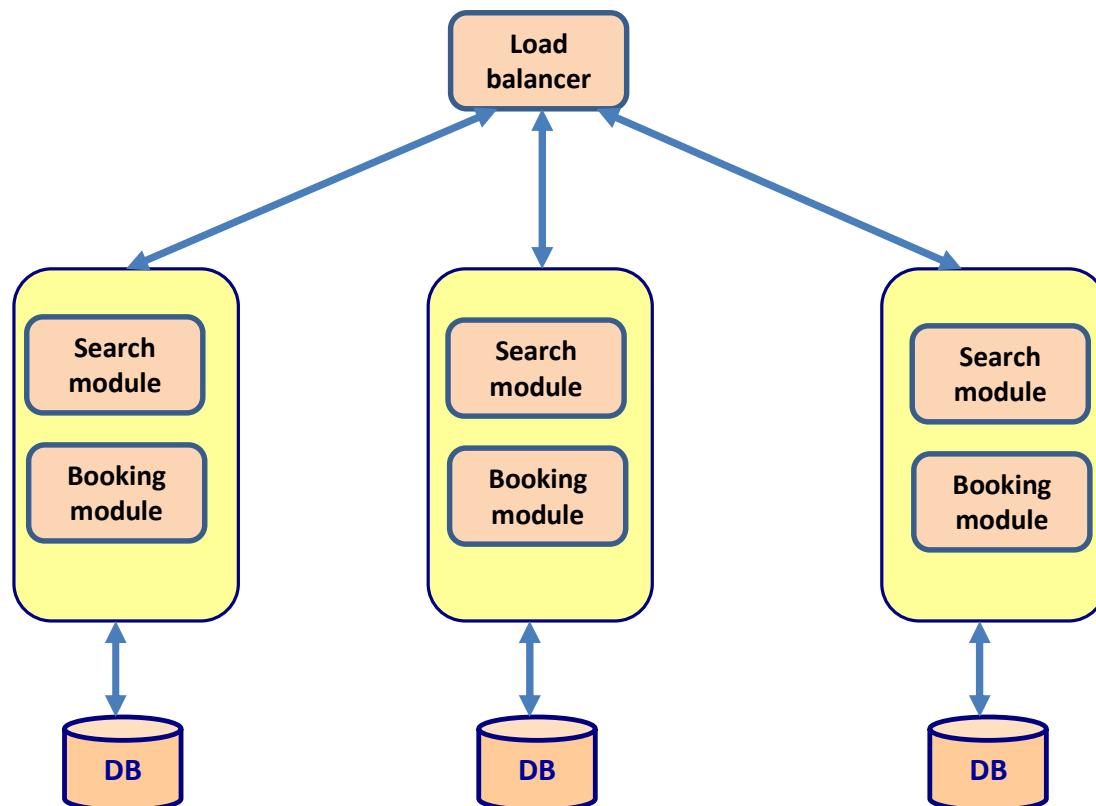
- Limited re-use is realized across monolithic applications



# Monolith architecture

---

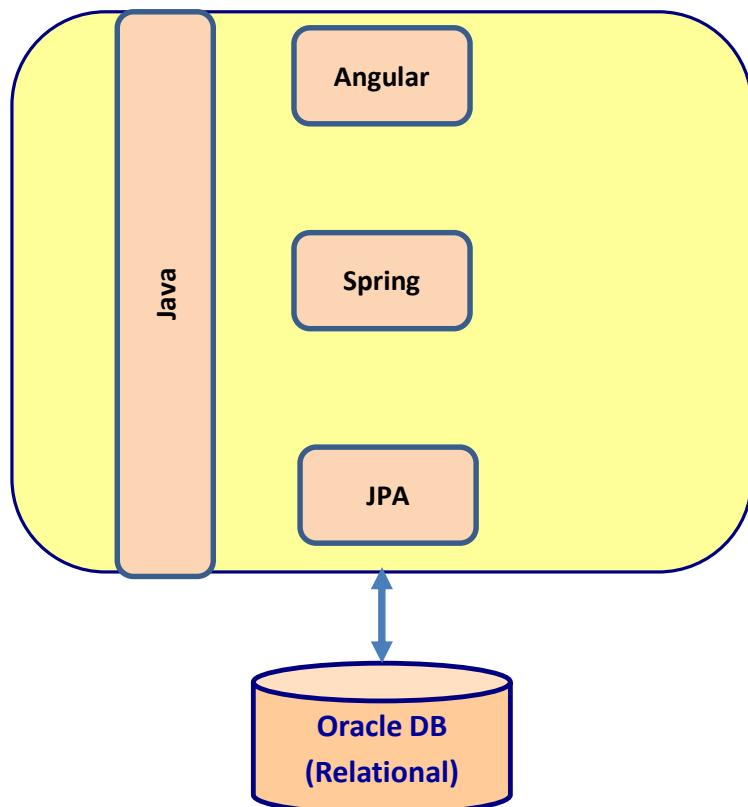
- All or nothing scaling
  - Difficult to scale separate parts



# Monolith architecture

---

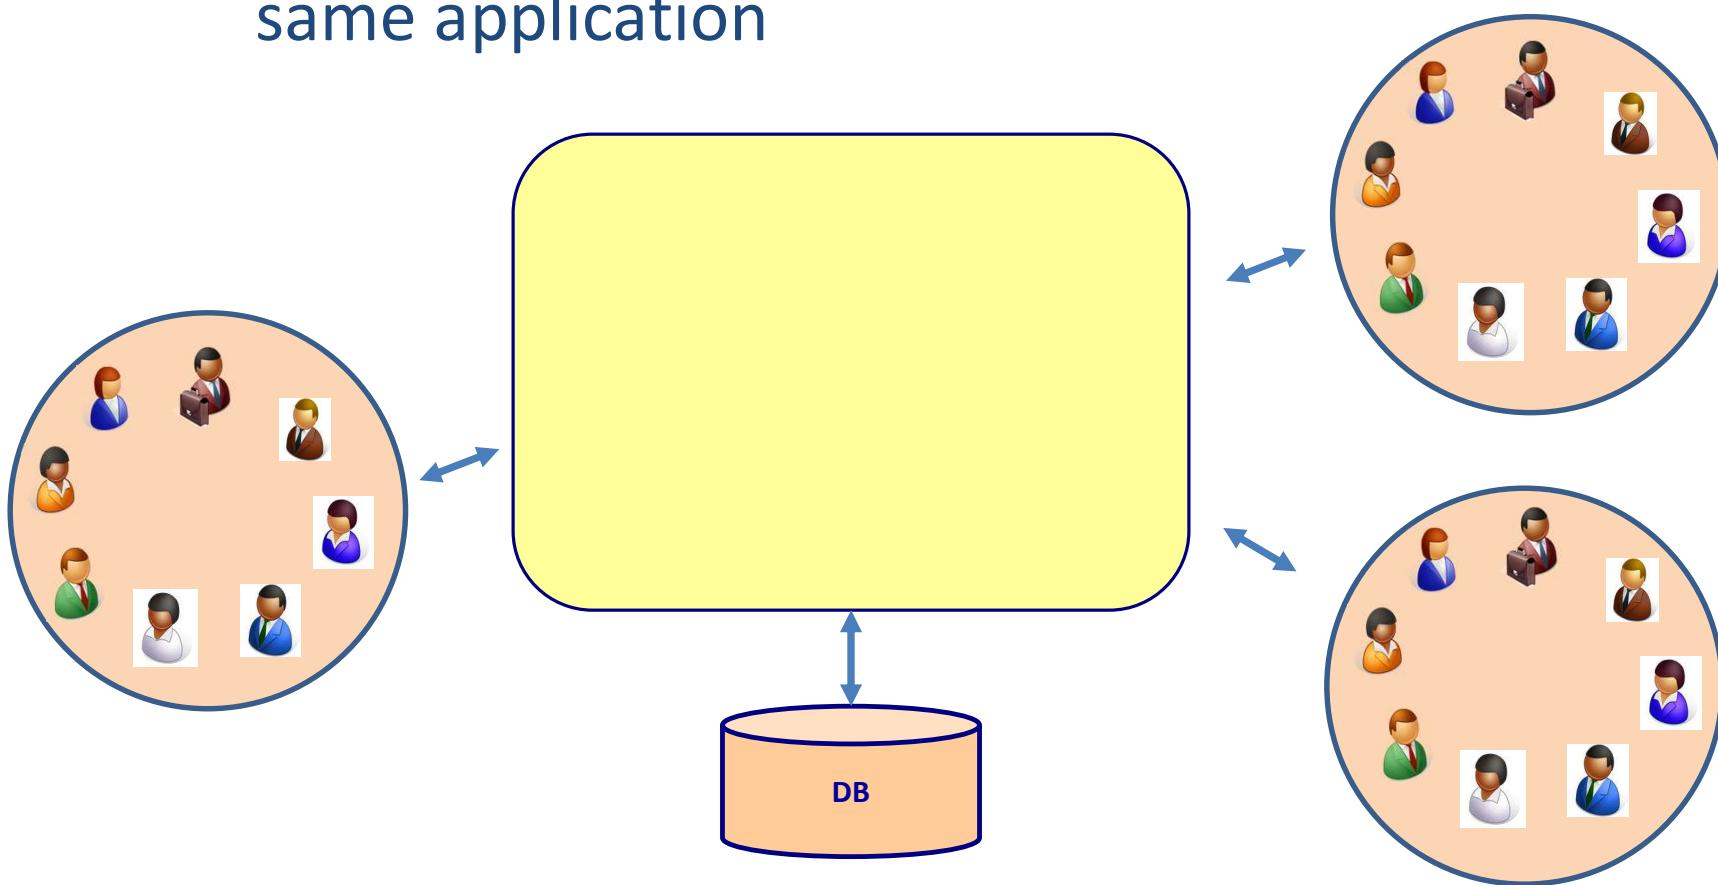
- Single development stack
  - Hard to use “the right tool for the job.”



# Monolith architecture

---

- Does not support small agile scrum teams
  - Hard to have different agile teams work on the same application



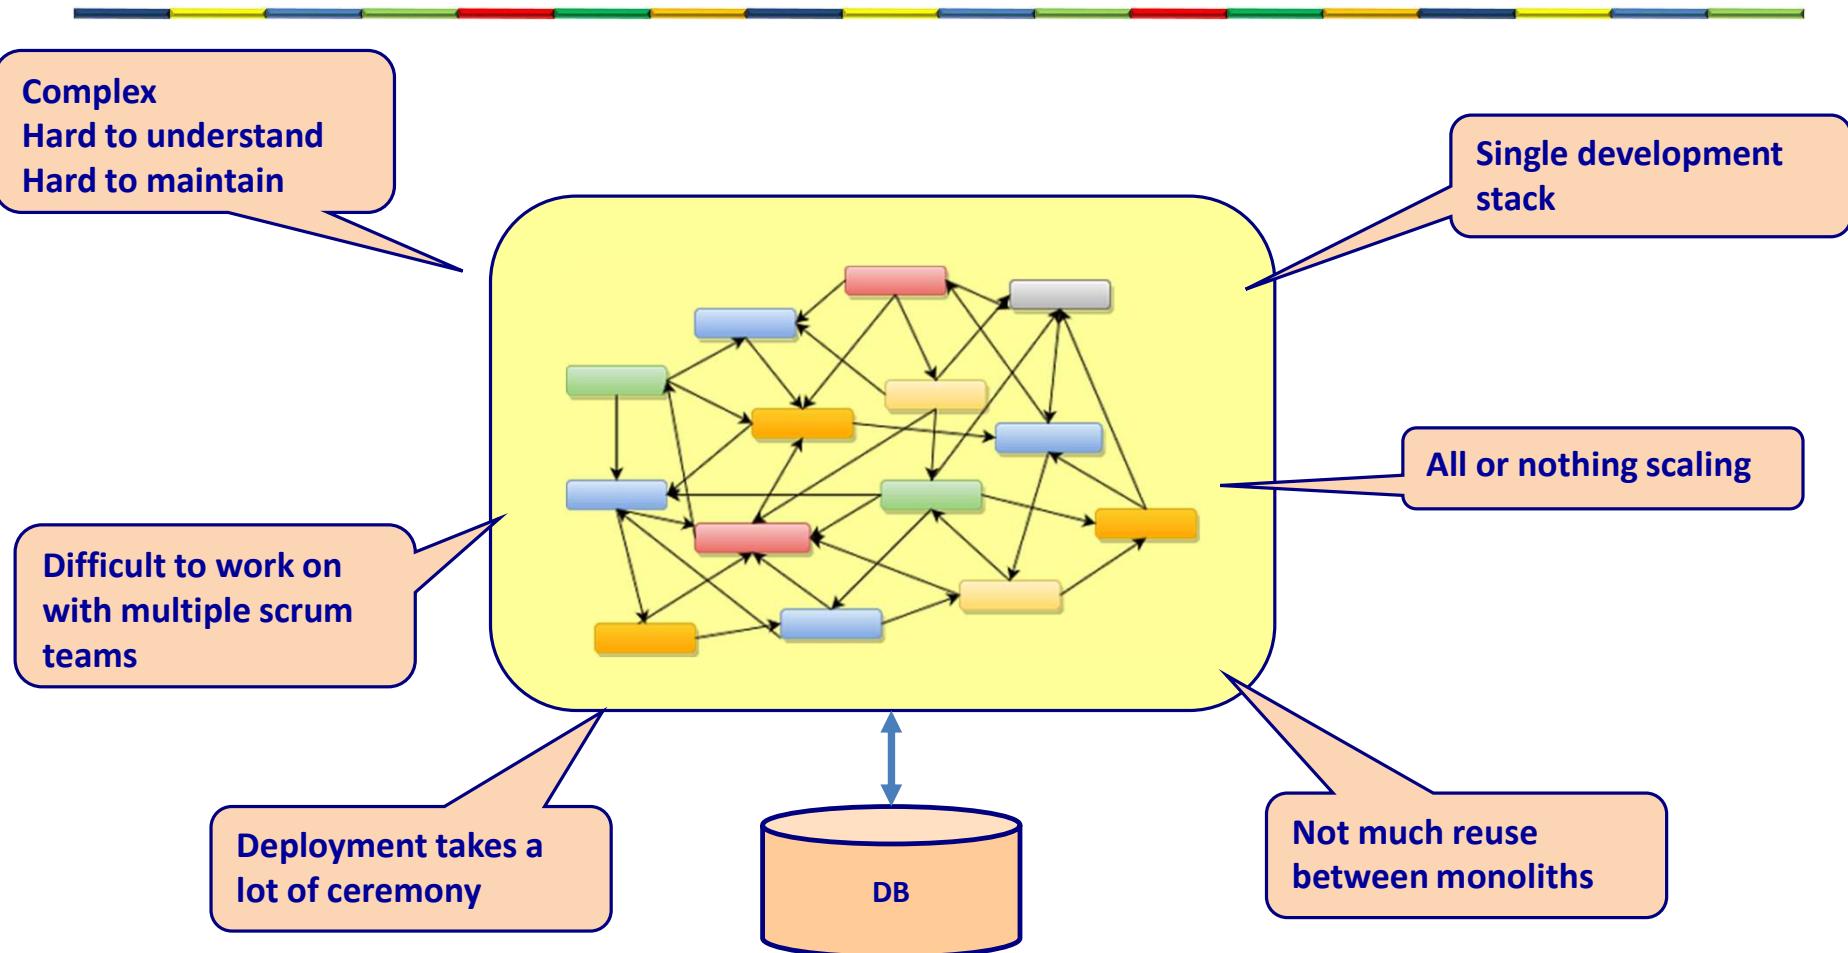
# Monolith architecture

---

- Deploying a monolith takes a lot of ceremony
  - Every deployment is of high risk
  - I cannot deploy very frequently
  - Long build-test-release cycles

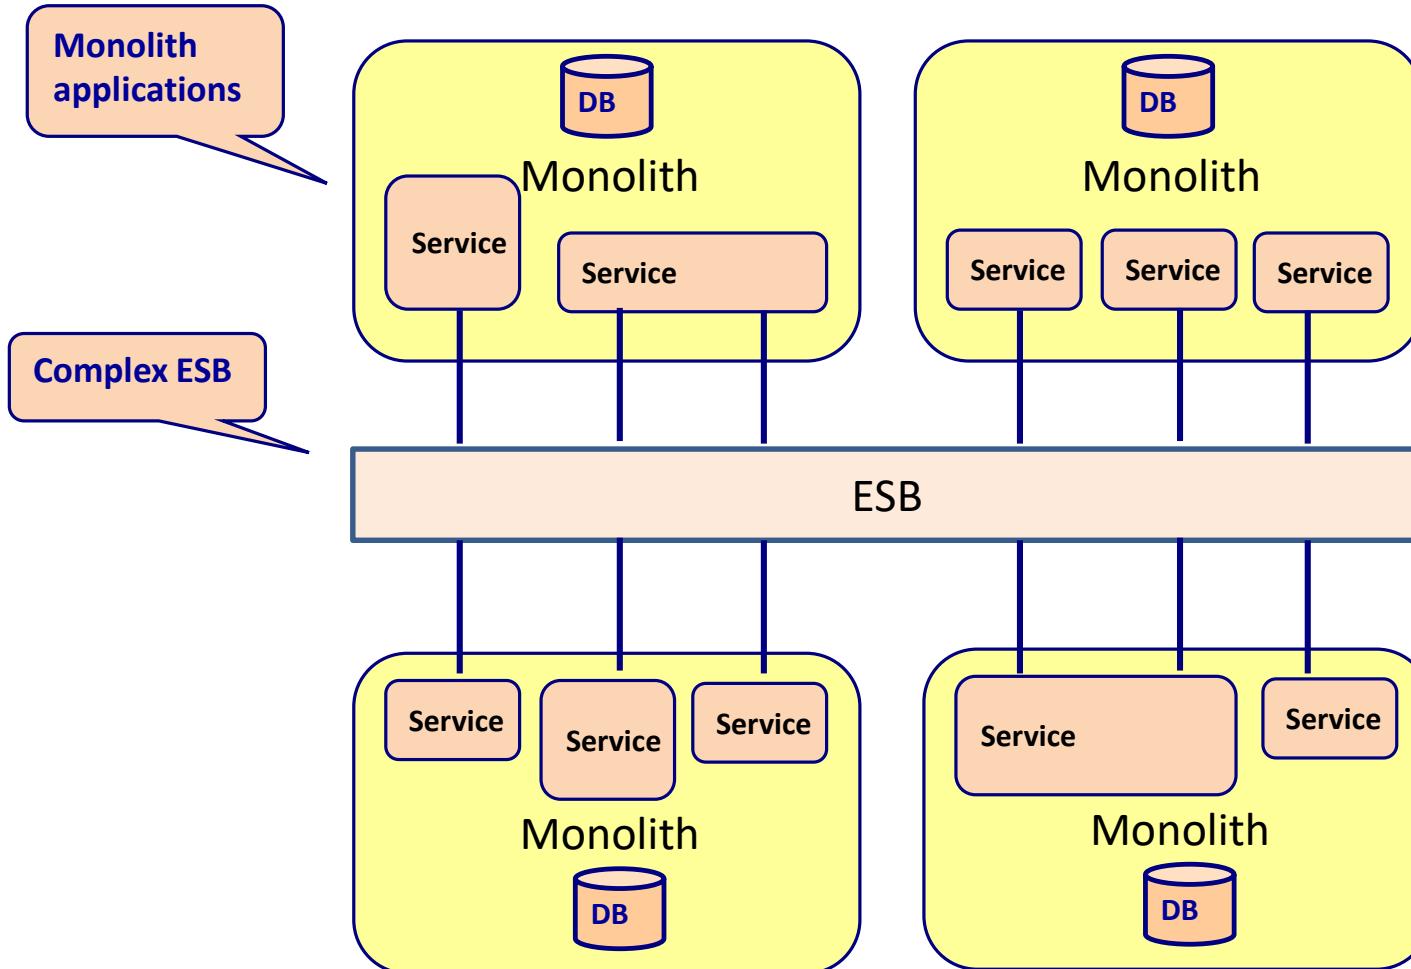


# Problems with a monolith architecture



# Problems with SOA

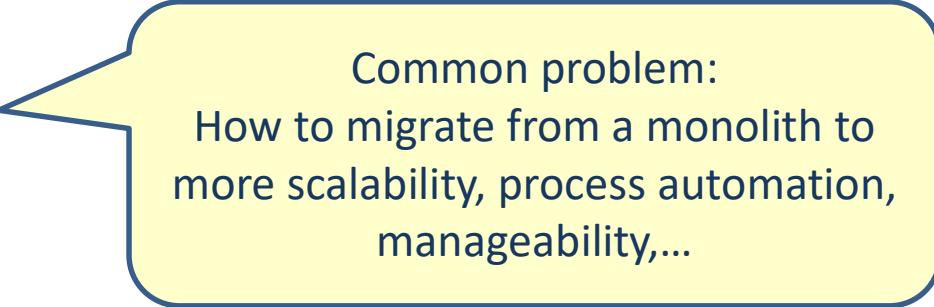
---



# Microservice early adopters

---

- Netflix
- Uber
- Airbnb
- Orbiz
- eBay
- Amazon
- Twitter
- Nike

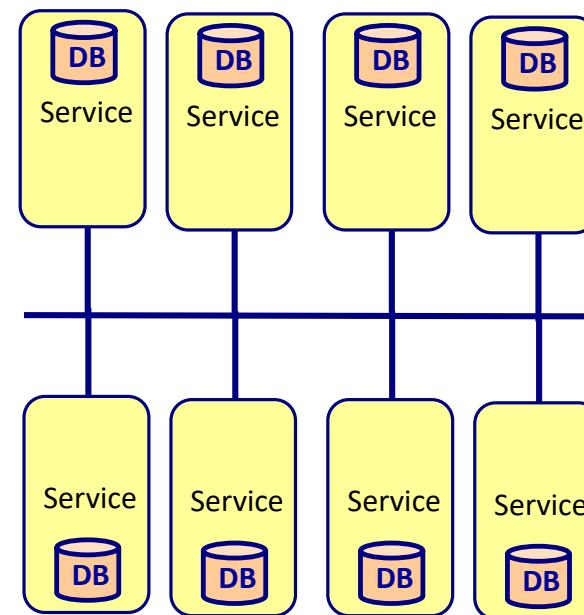


Common problem:  
How to migrate from a monolith to  
more scalability, process automation,  
manageability,...

# Microservices

---

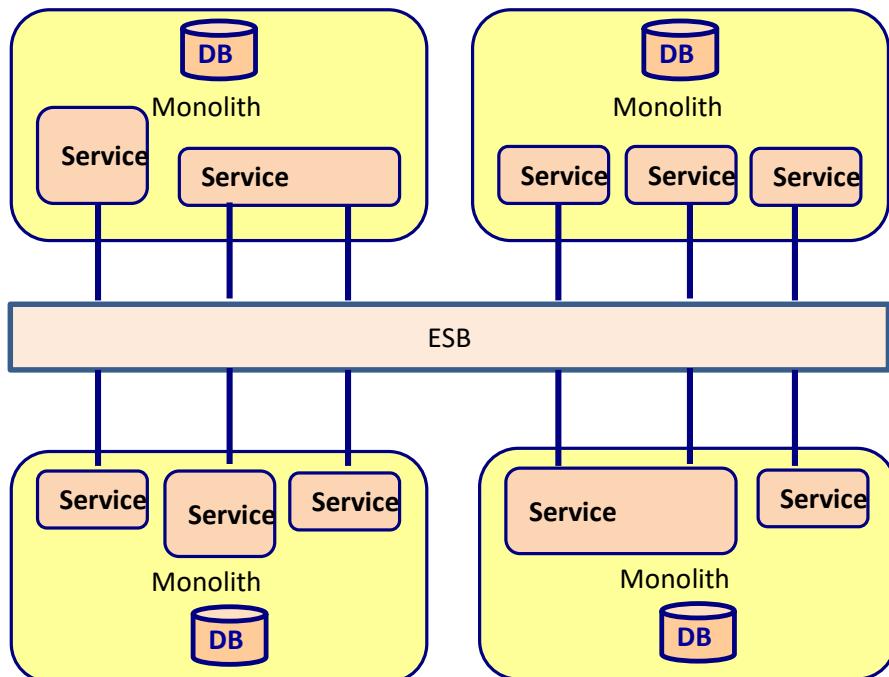
- Small independent services
  - Simple and lightweight
  - Runs in an independent process
  - Language agnostic
  - Decoupled



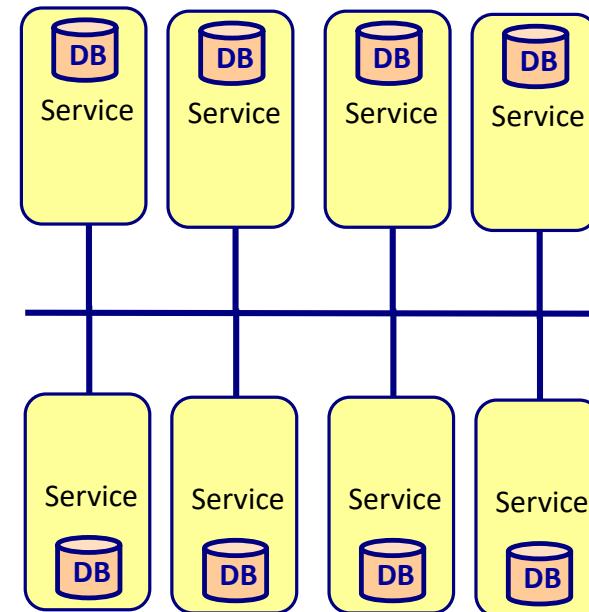
# SOA vs Microservice

---

SOA



Microservice

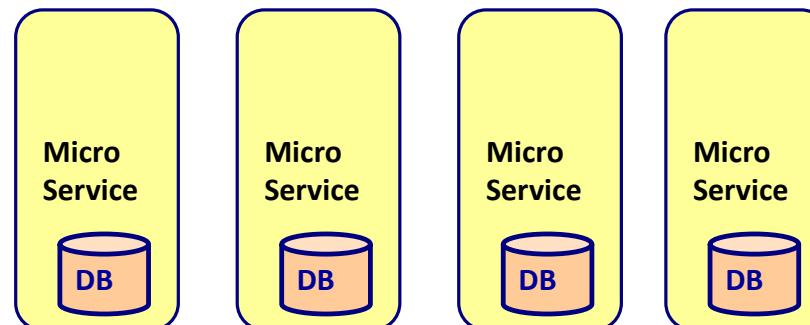


# **CHARACTERISTICS OF A MICROSERVICE**

# Microservices

---

- Small independent services
  - Simple and lightweight
  - Runs in an independent process
  - Technology agnostic
  - Decoupled

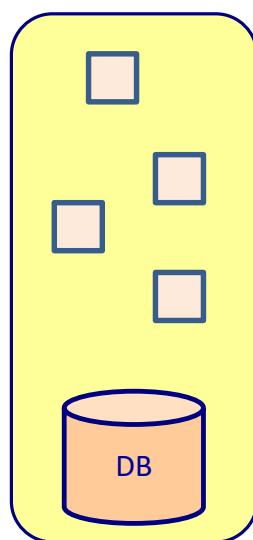


# Simple and lightweight

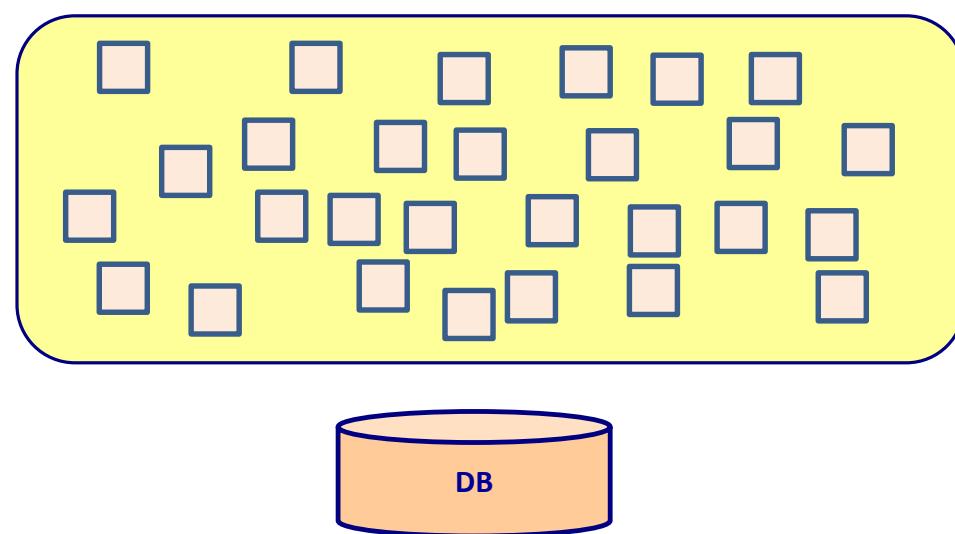
---

- Small and simple
- Can be build and maintained by 1 agile team

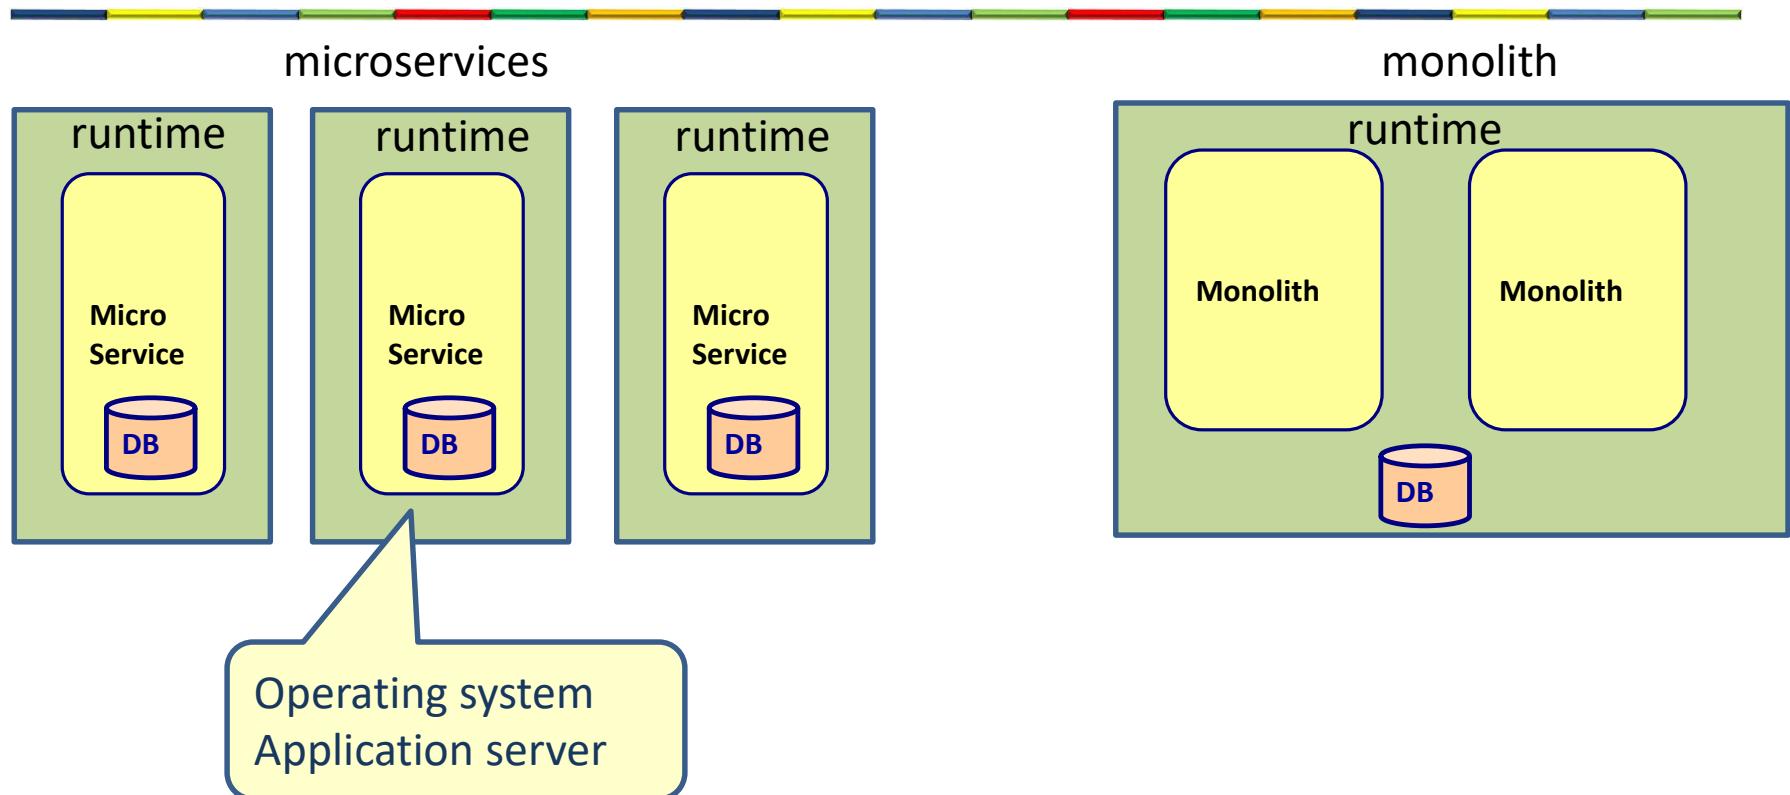
microservices



monolith



# Runs in an independent process



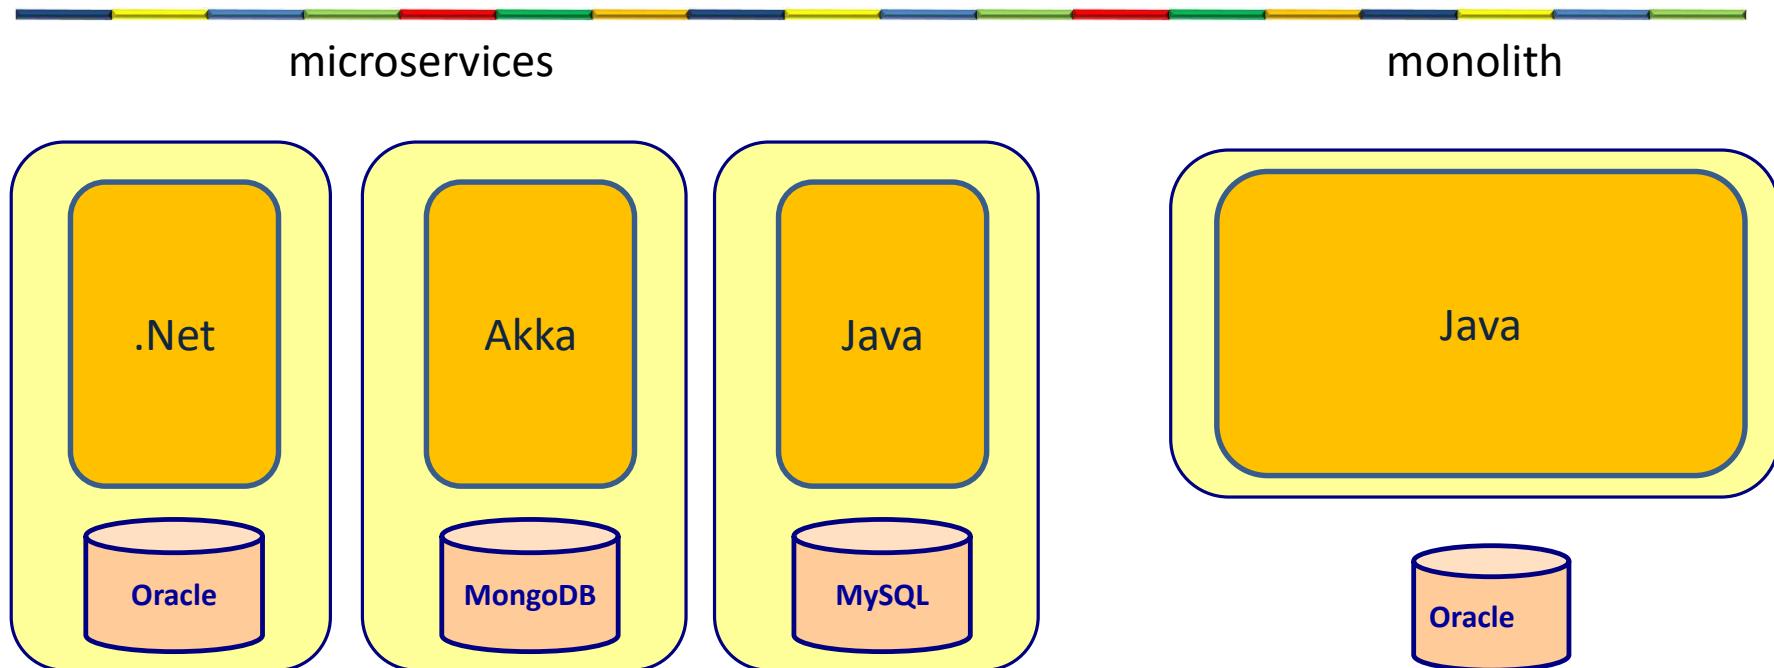
## Advantages

- Runtime can be small
  - Only add what you need
- Runtime can be optimized
- Runtime can start and stop fast
- If runtime goes down, other services will still run

## Disadvantages

- We need to manage many runtimes

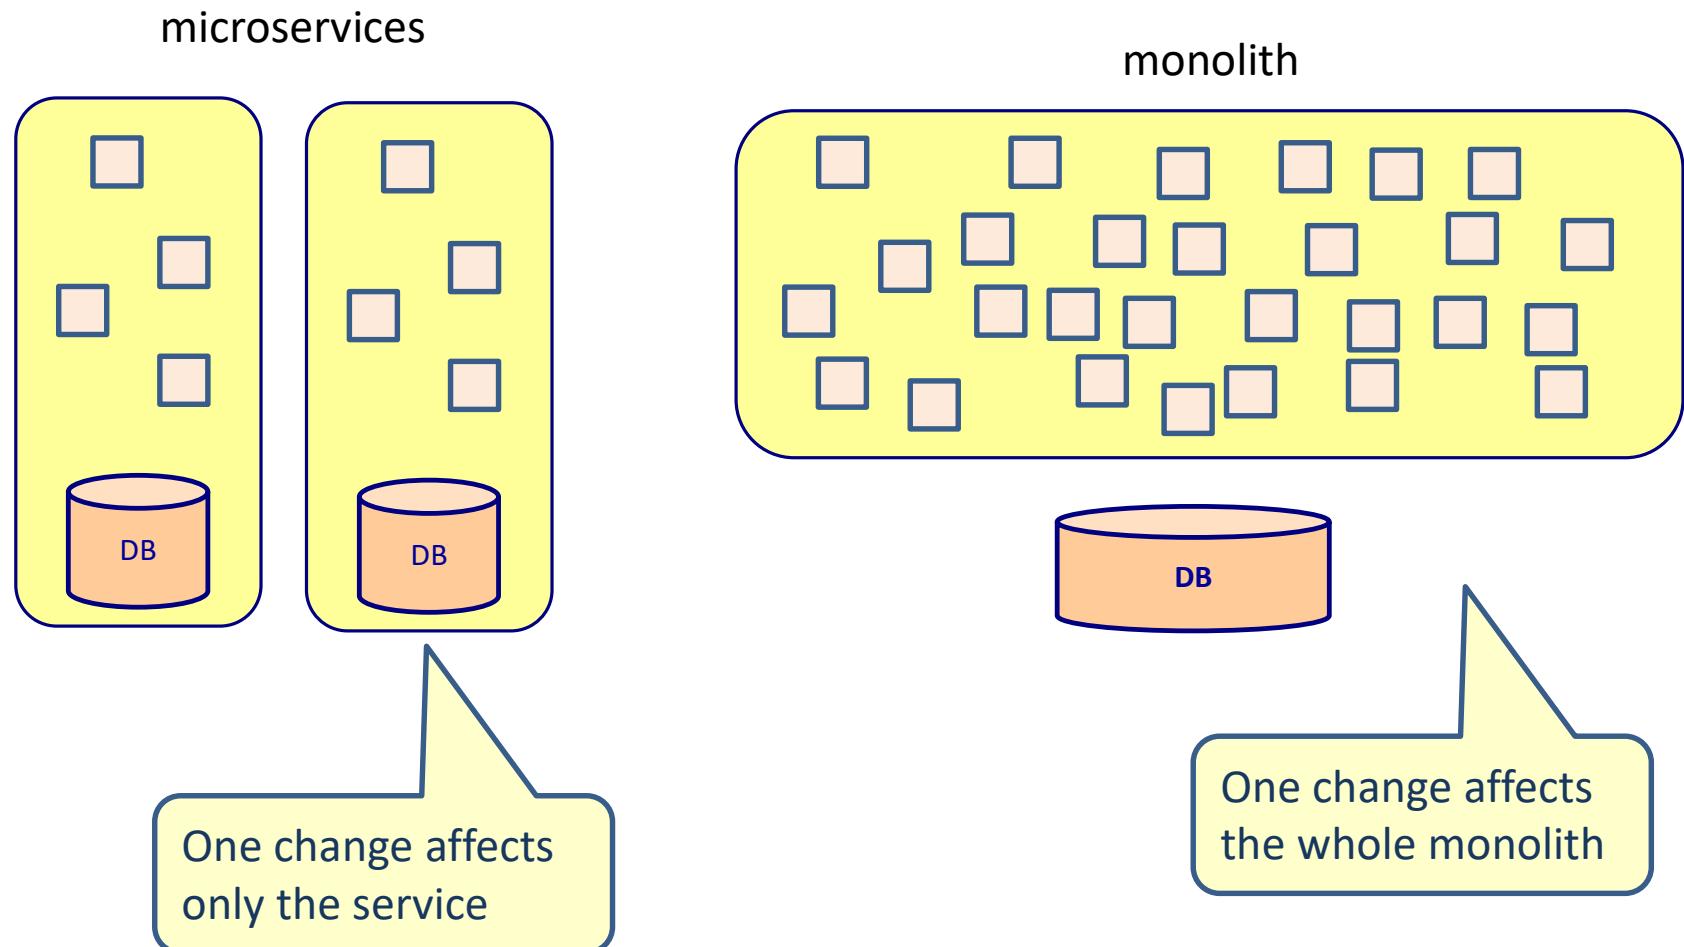
# Technology agnostic



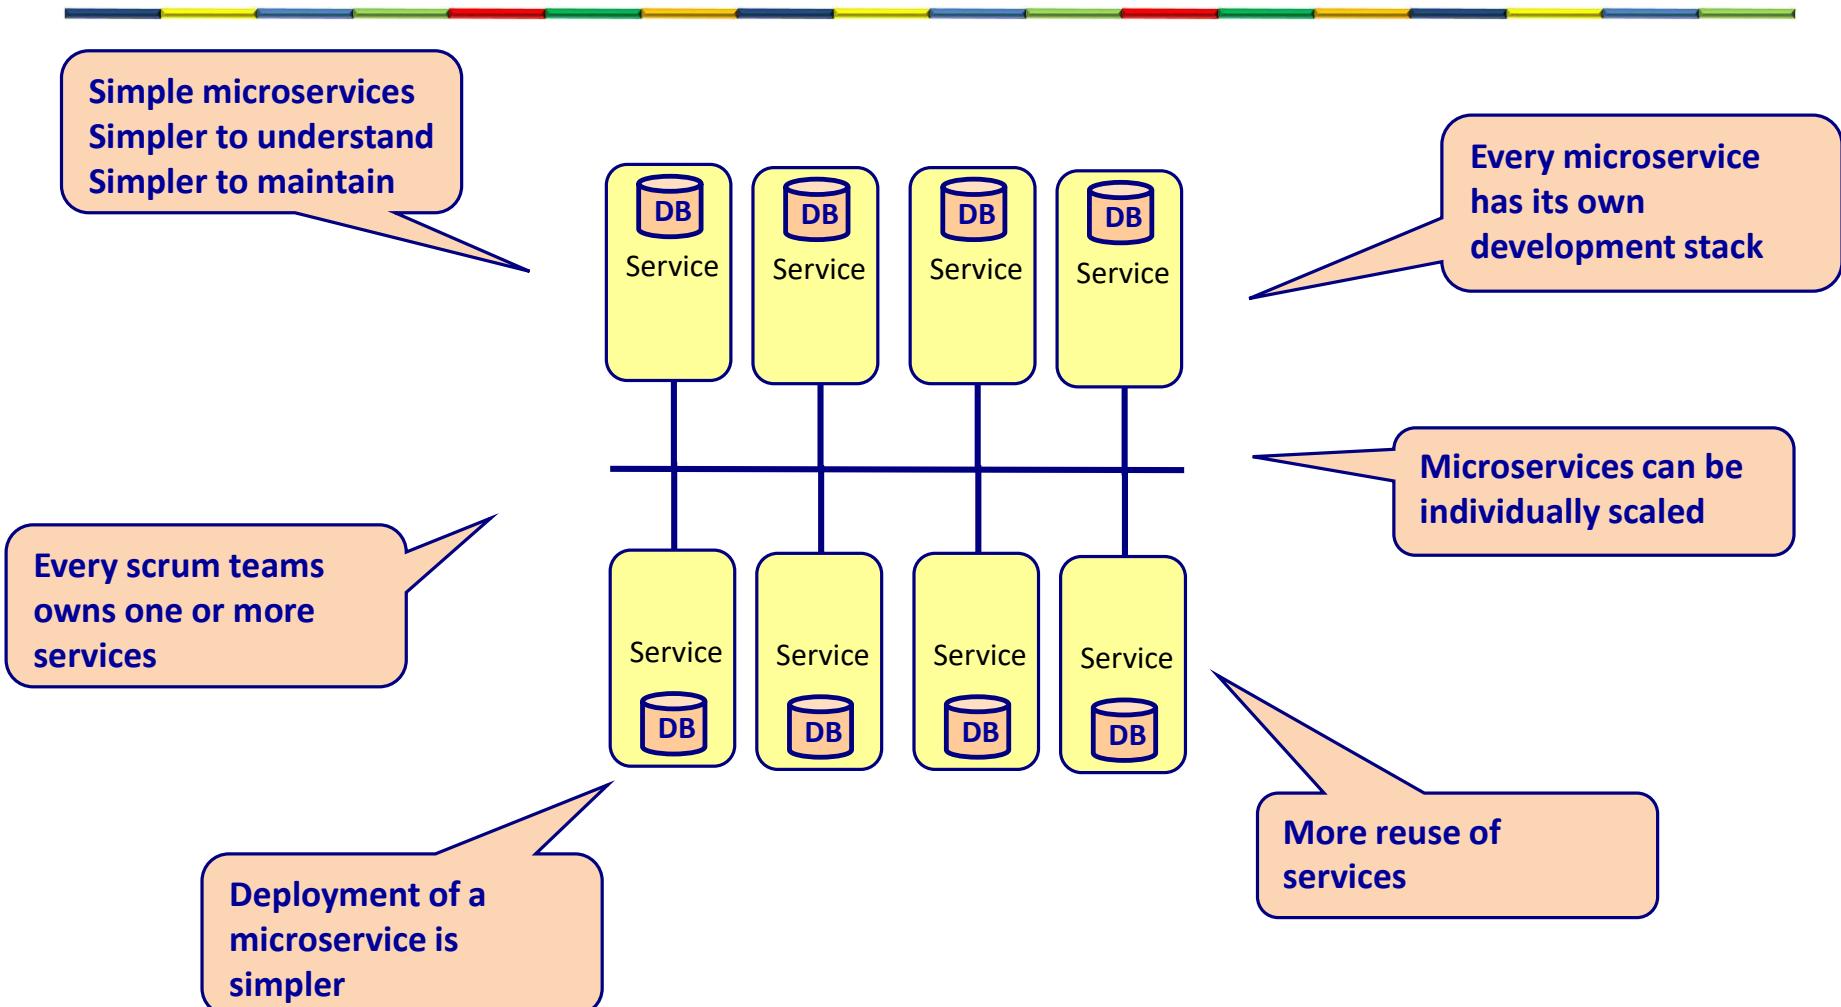
- Use the architecture and technologies that fits the best for this particular microservice

# Decoupled

---

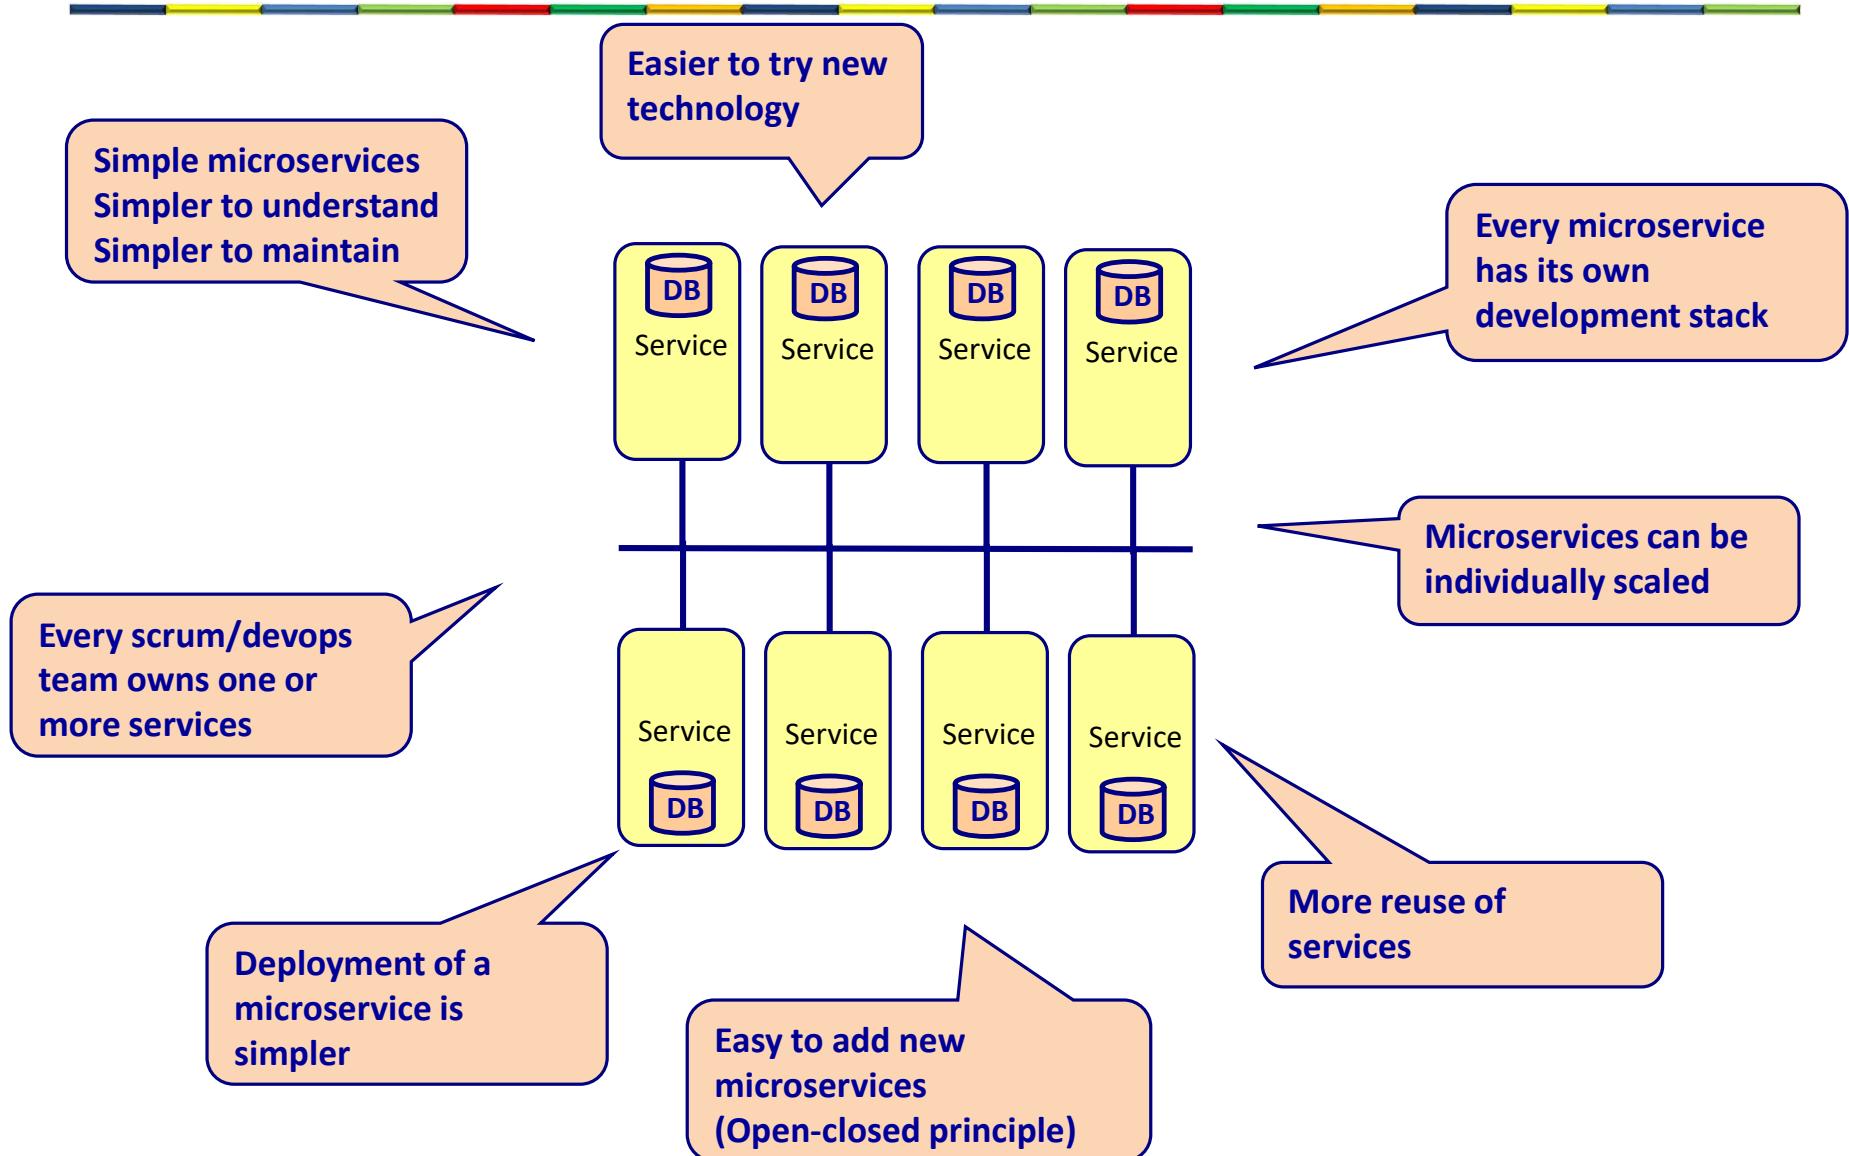


# Microservice architecture

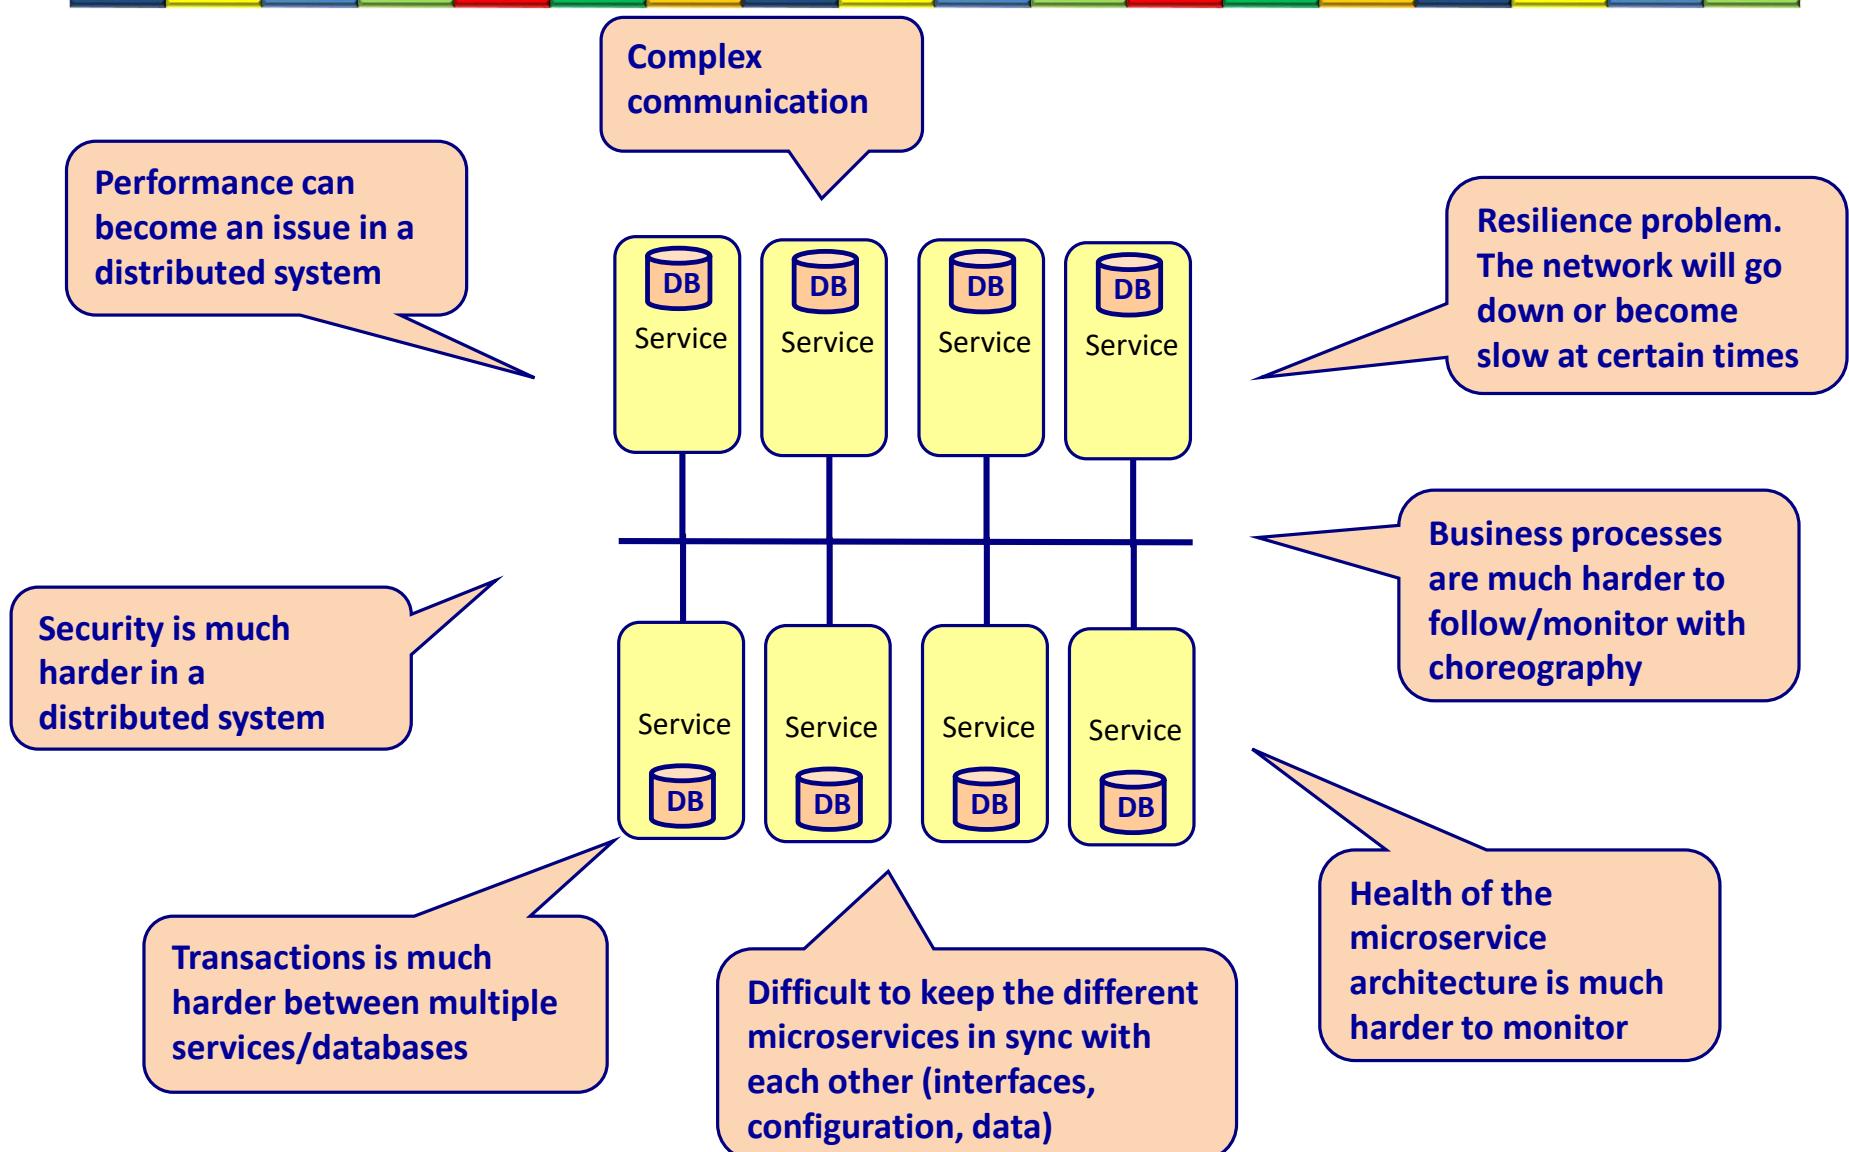


# **ADVANTAGES AND DISADVANTAGES OF A MICROSERVICE ARCHITECTURE**

# Advantages



# Disadvantages



# Challenges of a microservice architecture

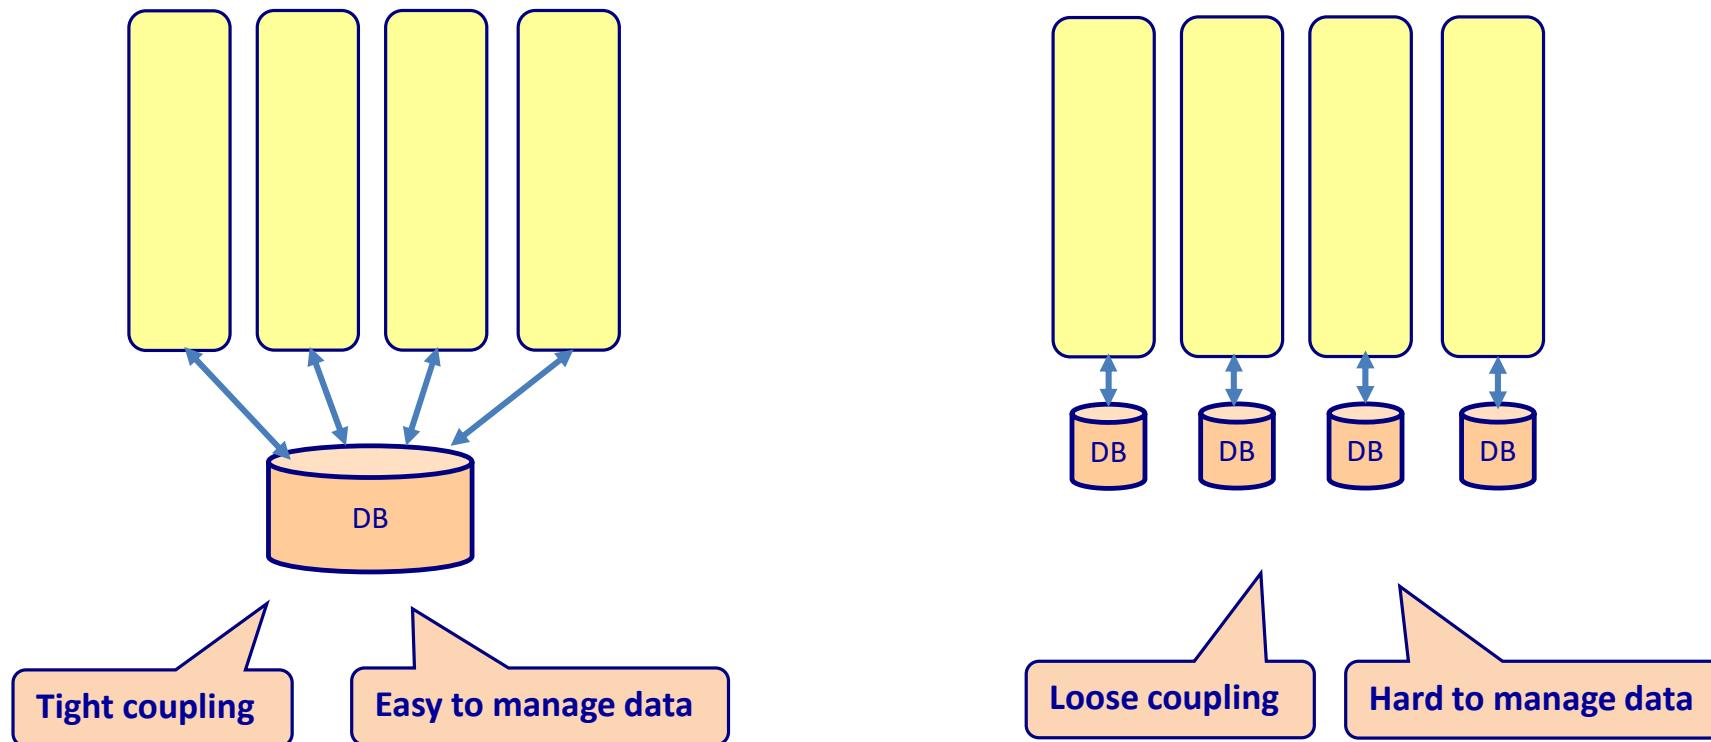
---

Challenge	Solution
Complex communication	
Performance	
Resilience	
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	

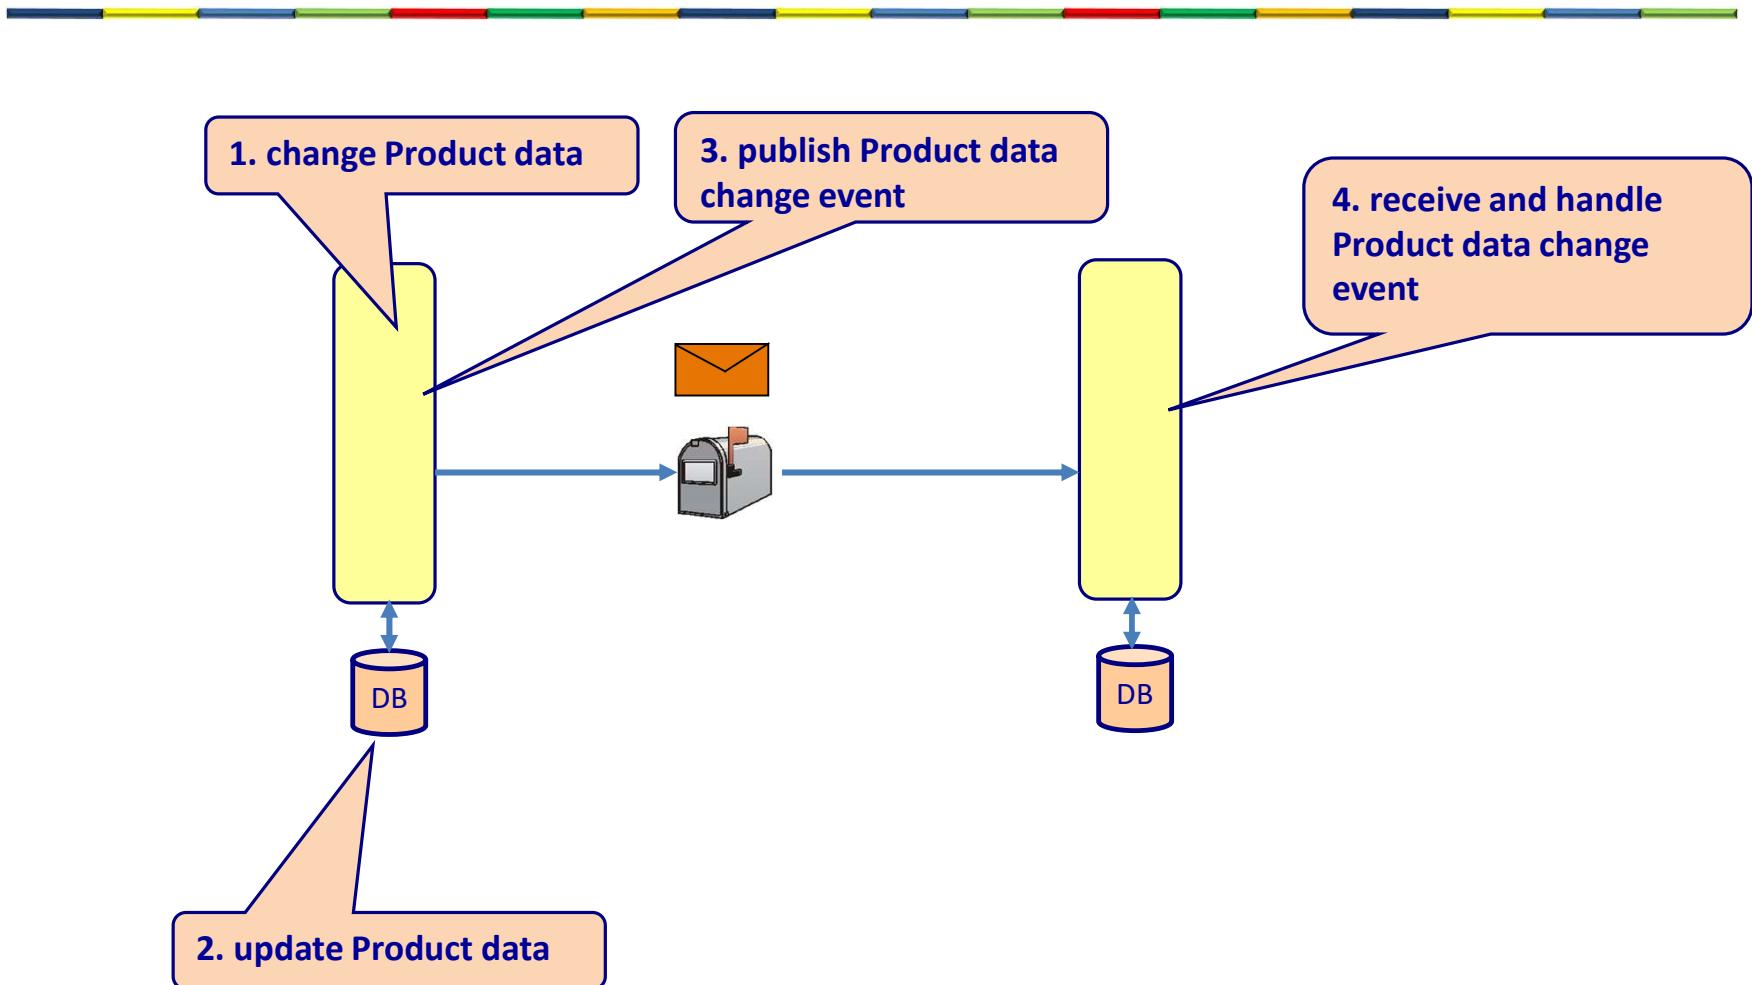
# **MICROSERVICE AND DATABASES**

# Every service manages its own data

---



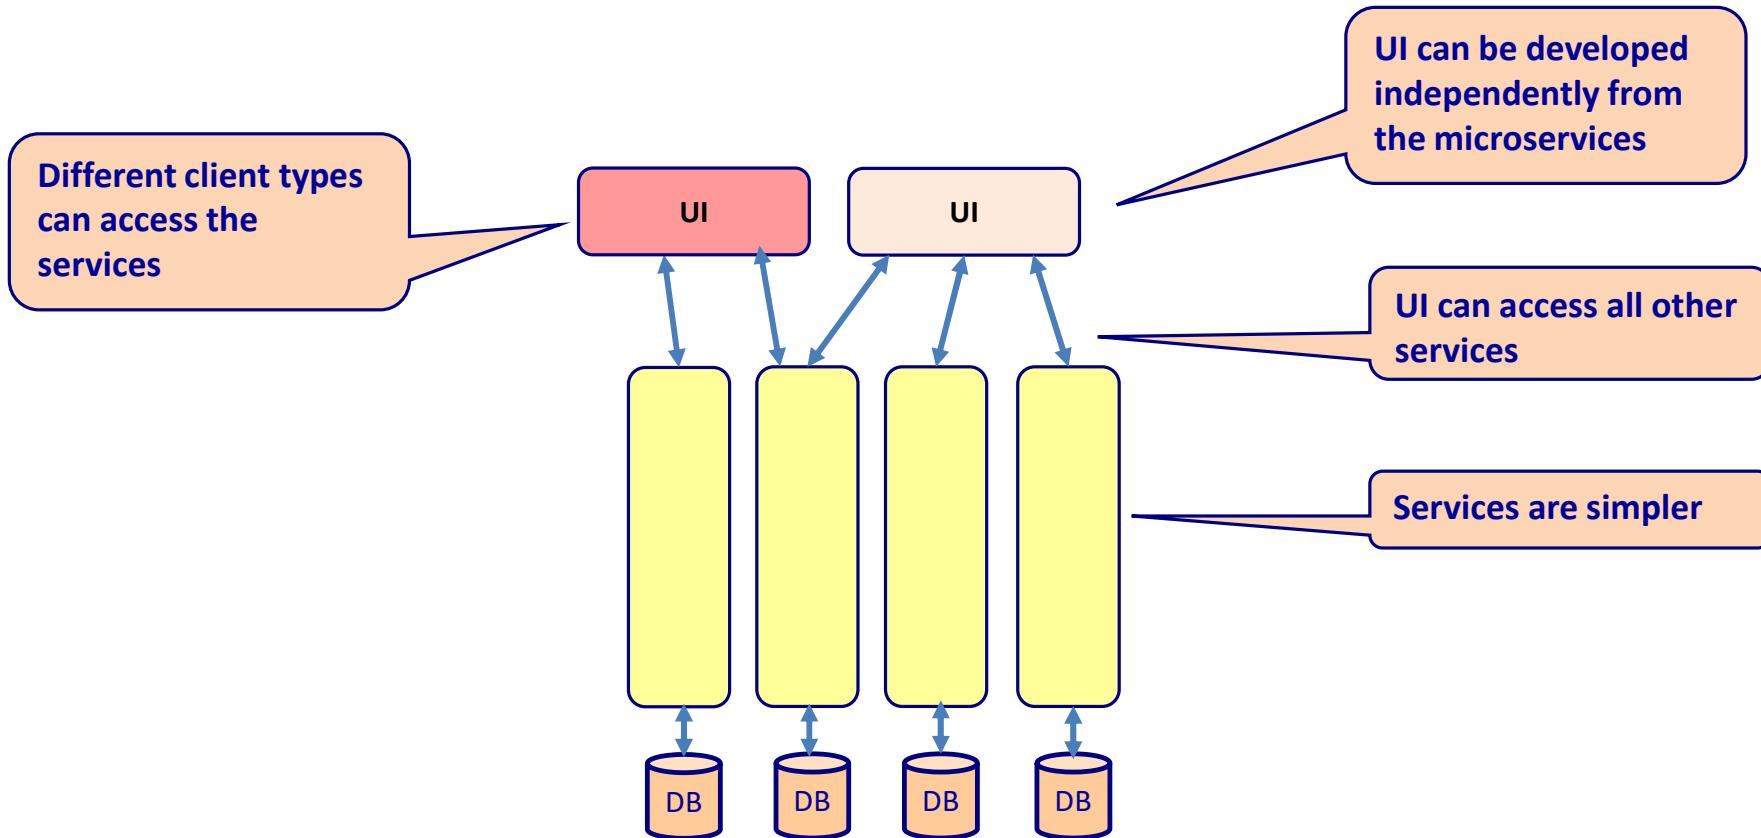
# Data consistency



# **UI AND MICROSERVICE**

# Split front-end and back-end

---



# **MICROSERVICE BOUNDARIES**

# Appropriate boundaries

---

- DDD bounded context
  - Isolated domains that are closely aligned with business capabilities
- Autonomous functions
  - Accept input, perform its logic and return a result
    - Encryption engine
    - Notification engine
    - Delivery service that accept an order and informs a trucking service

# Appropriate boundaries

---

- Size of deployable unit
  - Manageable size
- Most appropriate function or subdomain
  - What is the most useful component to detach from the monolith?
  - Hotel booking system: 60-70% are search request
    - Move out the search function
- Polyglot architecture
  - Functionality that needs different architecture
    - Booking service needs transactions
    - Search does not need transactions

# Appropriate boundaries

---

- Selective scaling
  - Functionality that needs different scaling
    - Booking service needs low scaling capabilities
    - Search needs high scaling capabilities
- Small agile teams
  - Specialist teams that work on their expertise
- Single responsibility

# Appropriate boundaries

---

- Replicability or changeability
  - The microservice is easy detachable from the overall system
  - What functionality might evolve in the future?
- Coupling and cohesion
  - Avoid chatty services
  - Too many synchronous request
  - Transaction boundaries within one service

# Appropriate boundaries

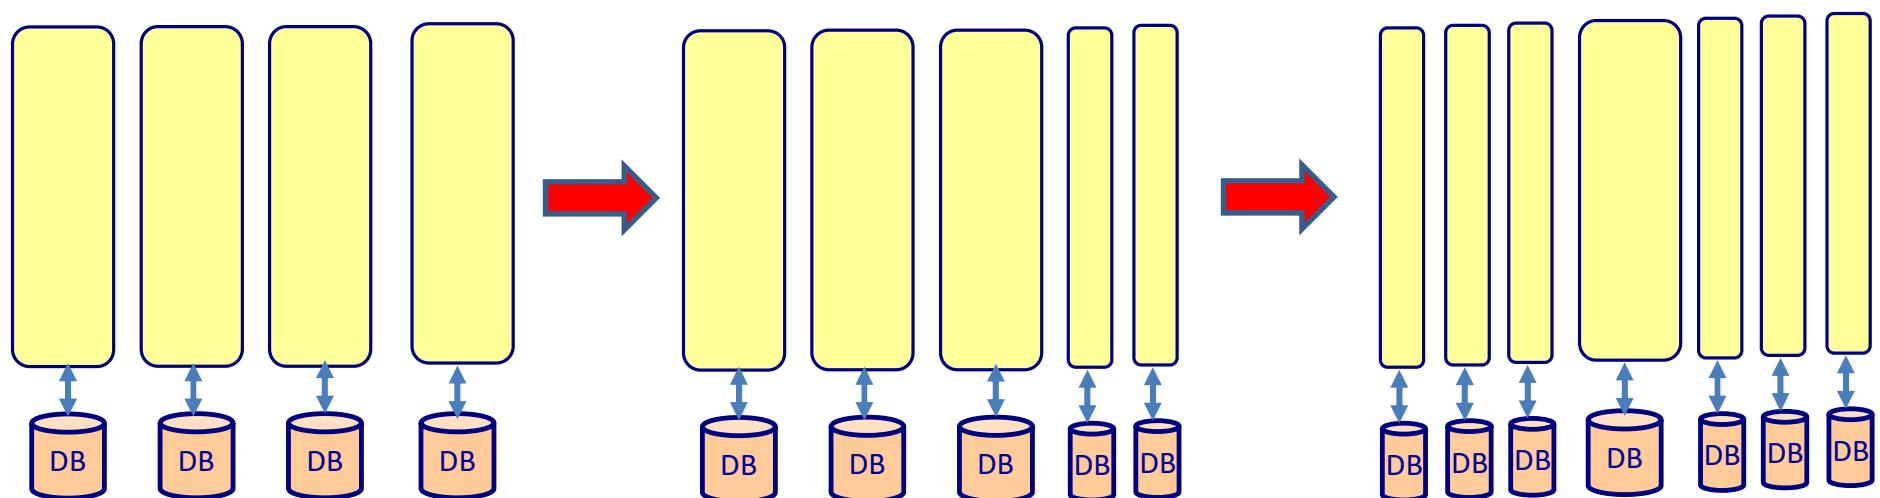
---

- DDD bounded context
- Autonomous functions
- Size of deployable unit
- Most appropriate function or subdomain
- Polyglot architecture
- Selective scaling
- Small agile teams
- Single responsibility
- Replicability or changeability
- Coupling and cohesion

# Microservice boundaries

---

- Start with a few services and then evolve to more services



# Domains

---

- Core subdomain
  - This is the reason you are writing the software.
- Supporting subdomain
  - Supports the core domain
- Generic subdomain
  - Very generic functionality
    - Email sending service
    - Creating reports service

# Distilling the domain

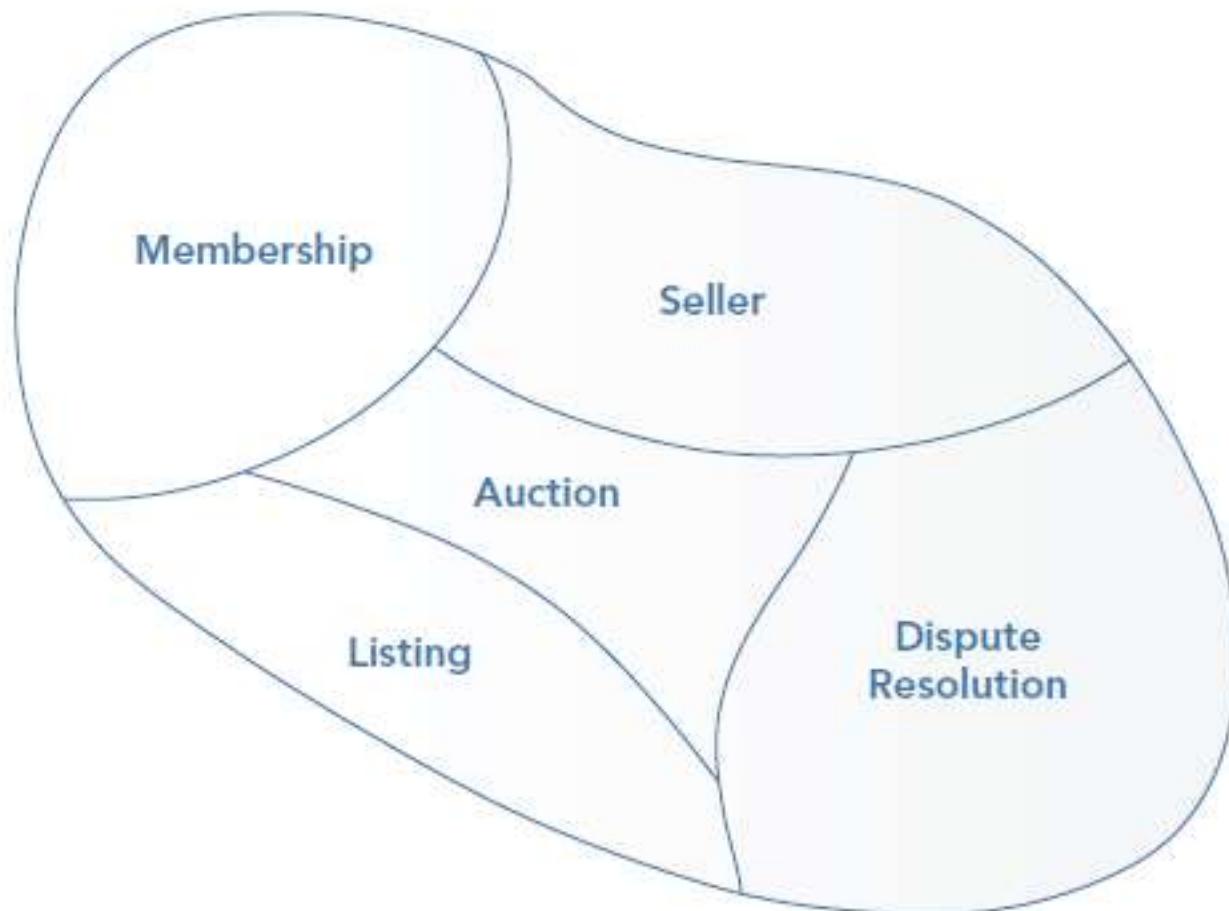
---

- The large domain of online auction



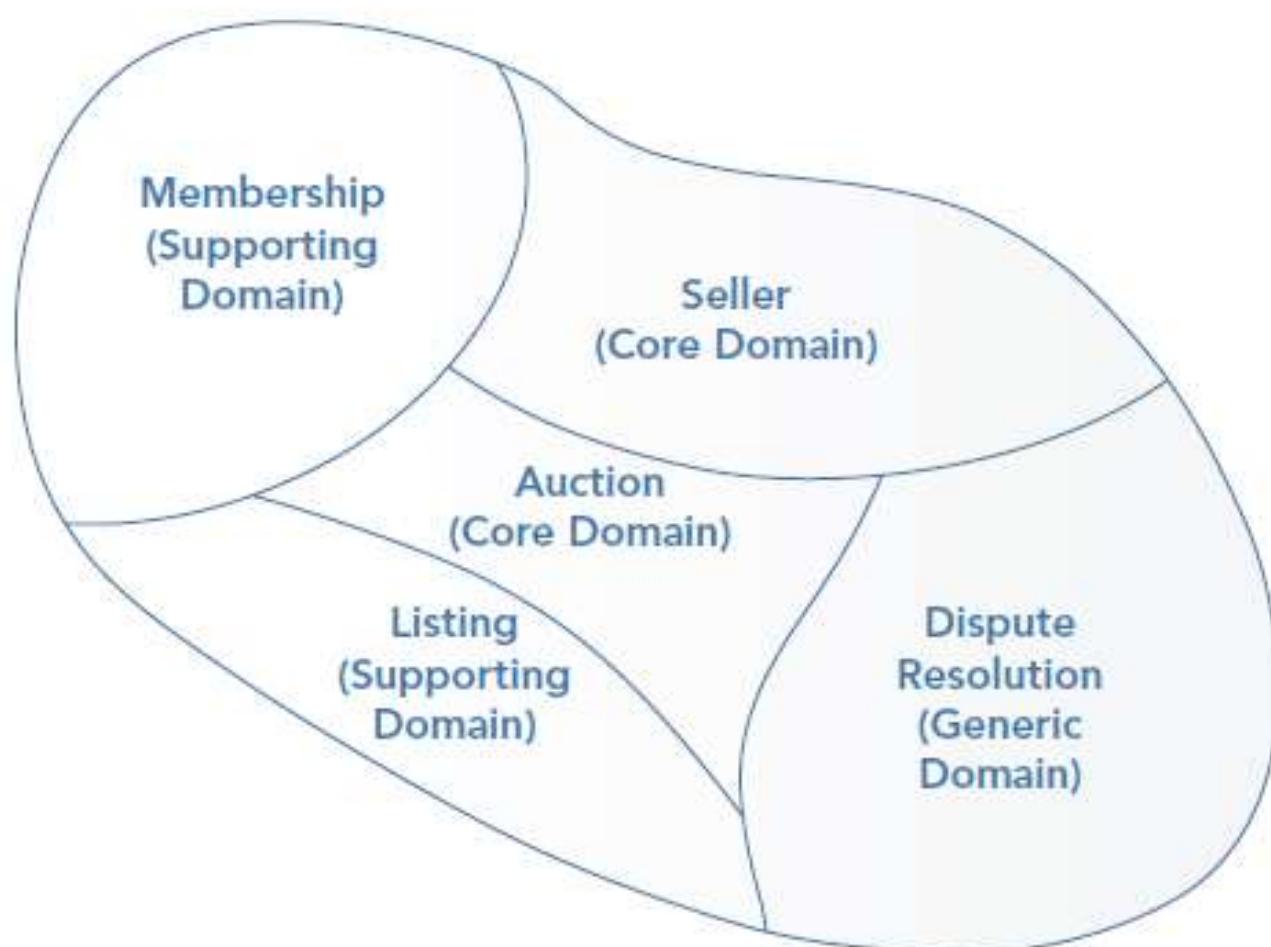
# Find the subdomains

---

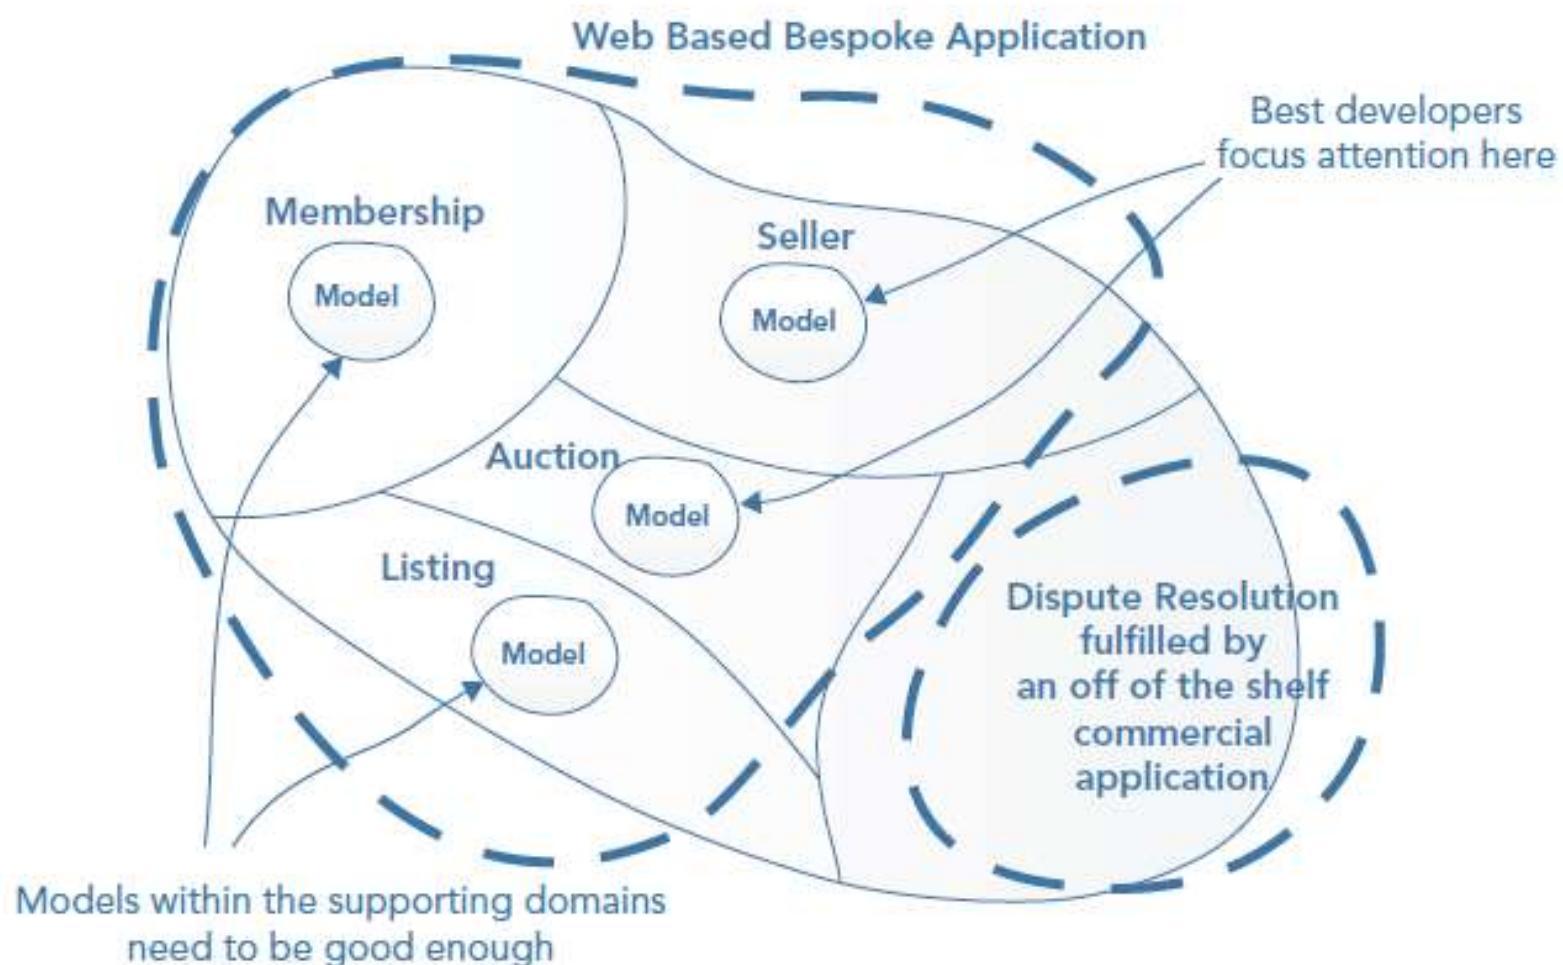


# Identify the core domain

---

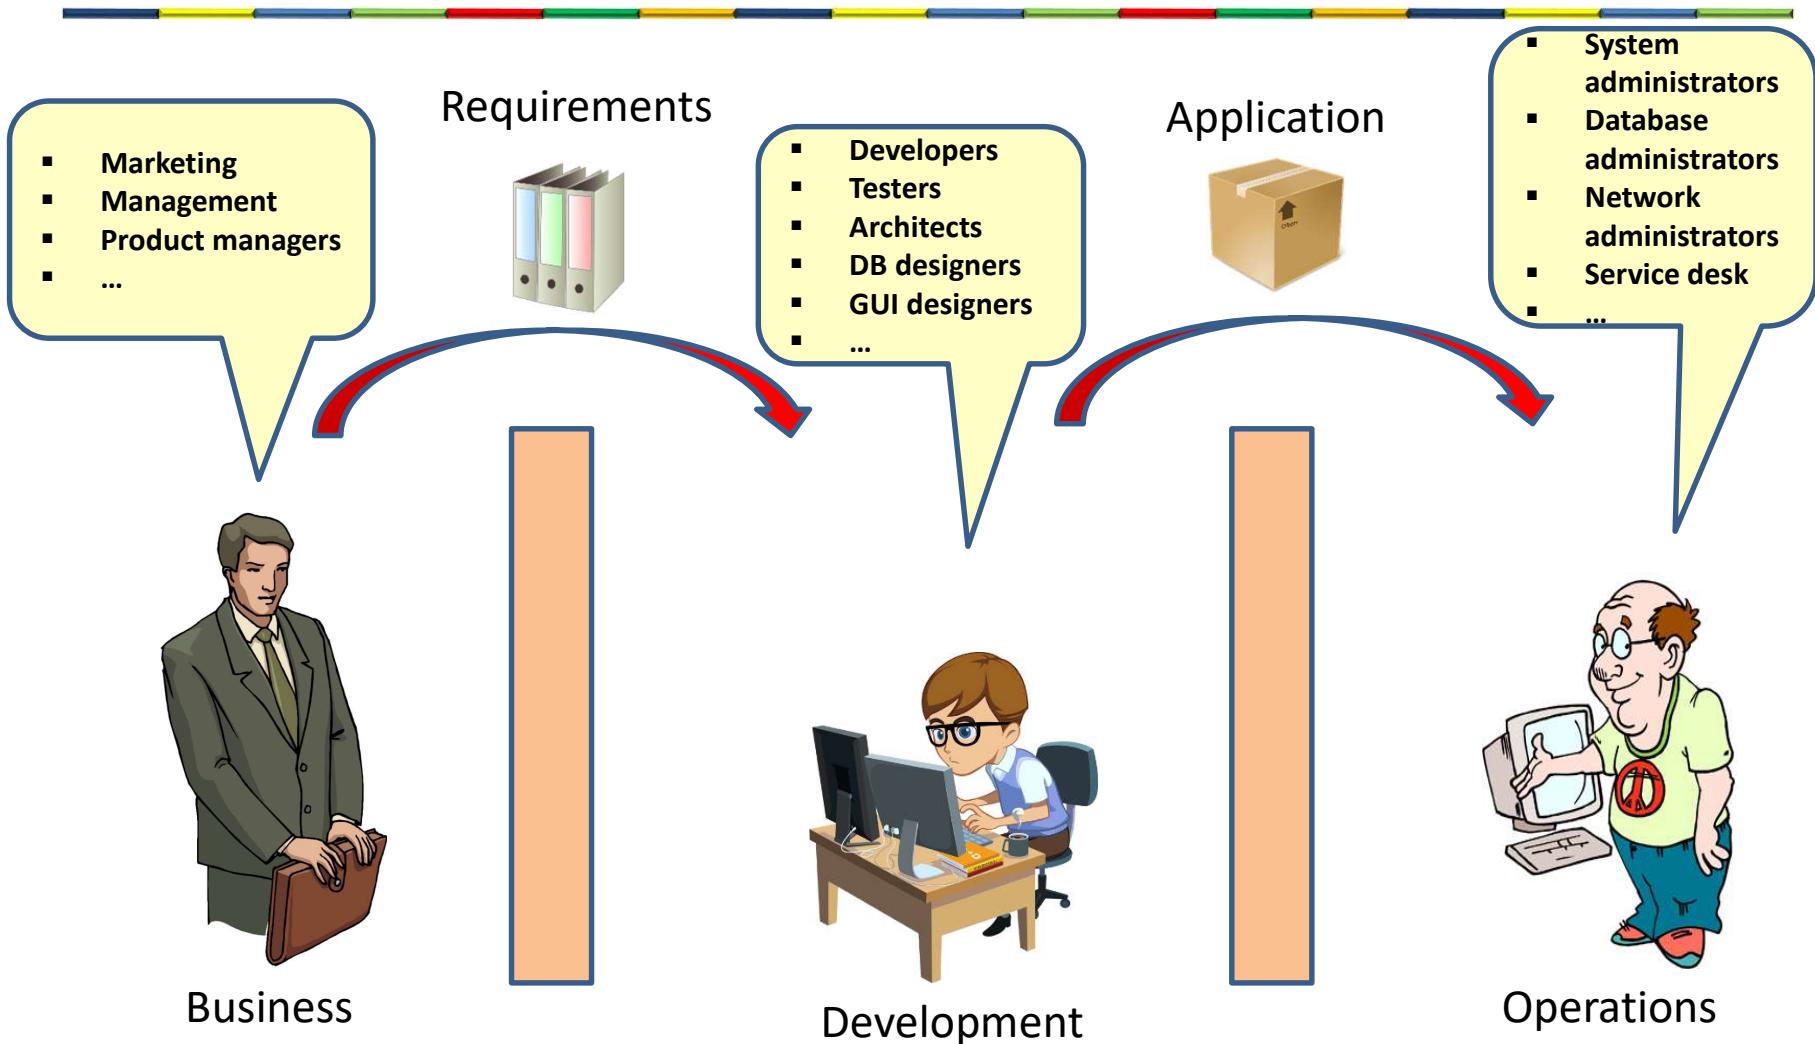


# Subdomains shape the solution

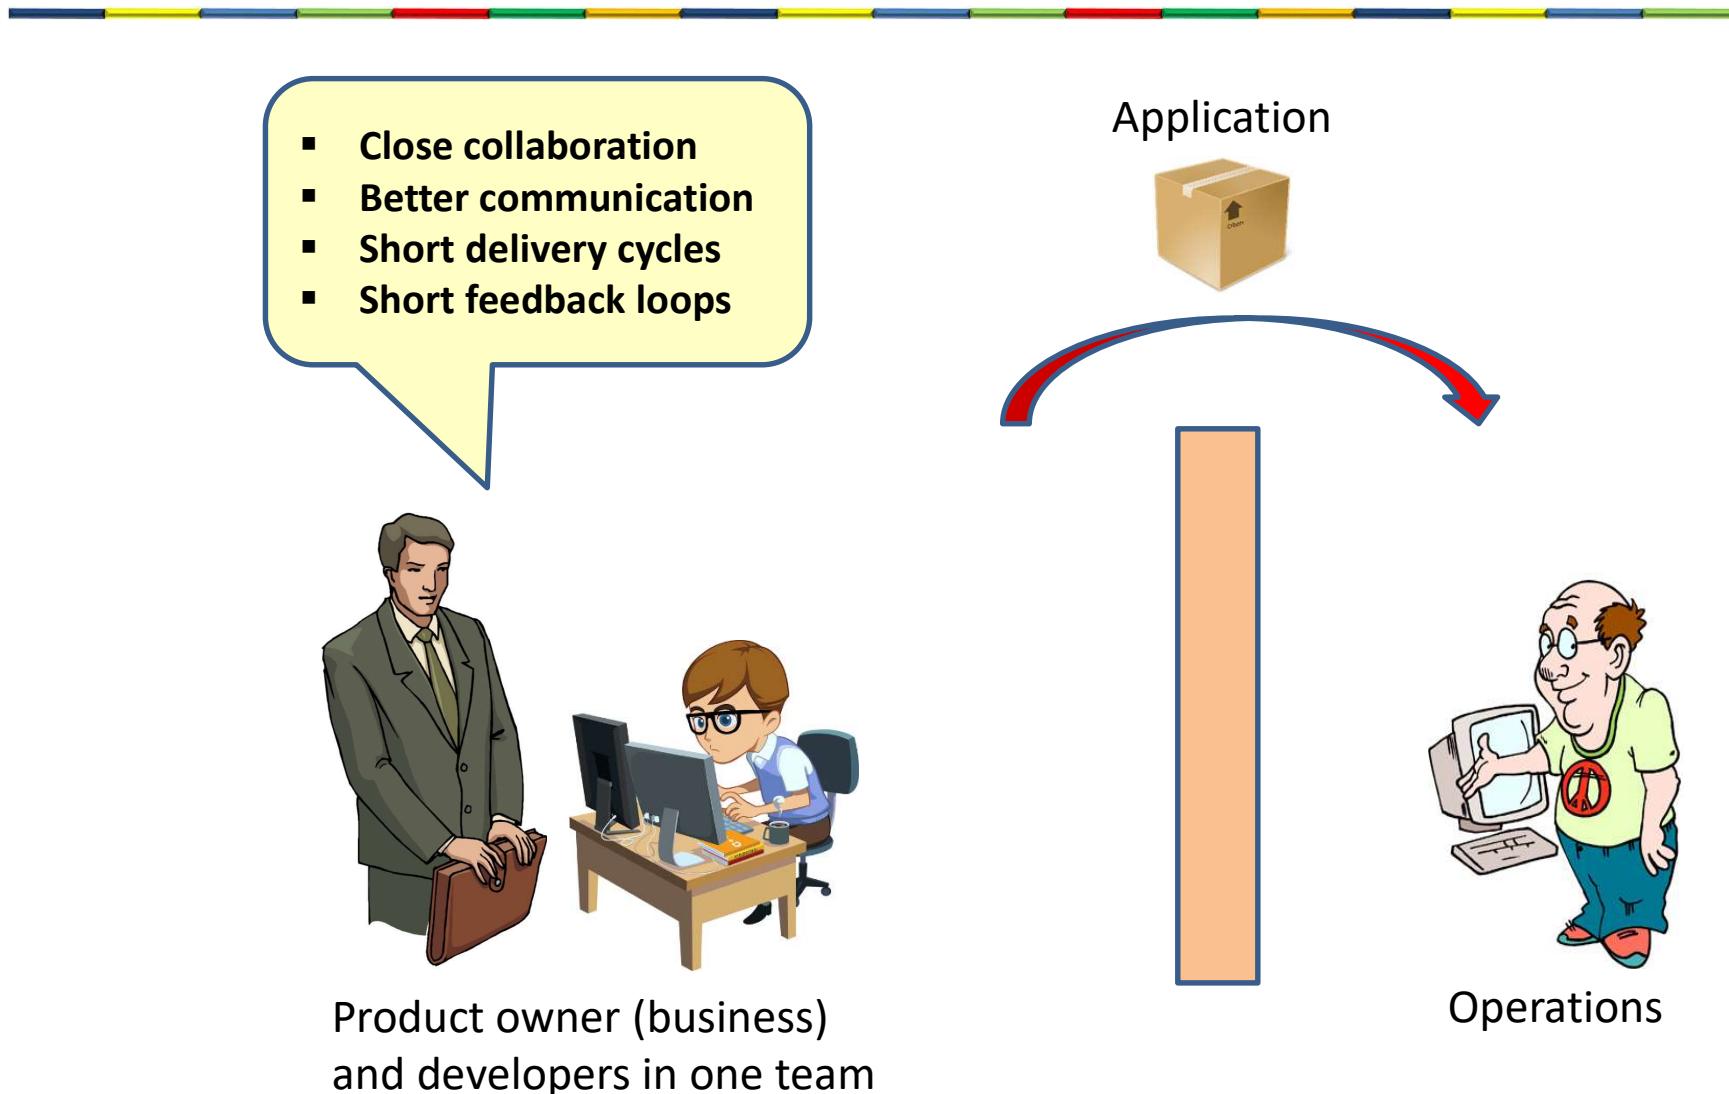


# **MICROSERVICES IN THE ORGANIZATION**

# Traditional software development



# Agile software development: Scrum



# DevOps

- Close collaboration between developers and operations
- Streamlines the delivery process of software from business requirements to production
- Better communication
- Identical development and production environment
- Shared tools
  - Automate everything
  - Monitor everything

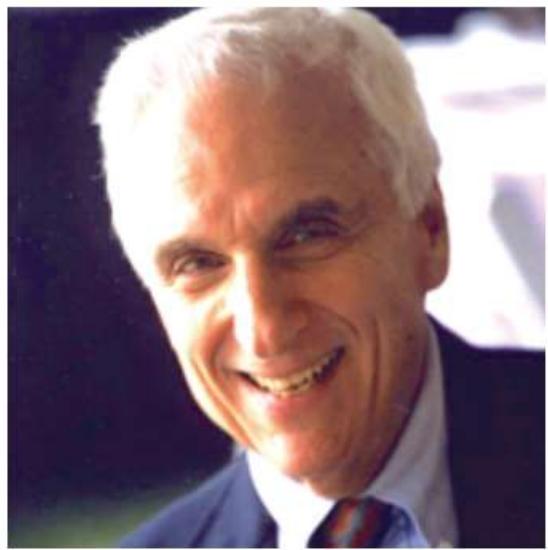


Product owner (business)  
and developers in one team

Operations

# Conways law

---



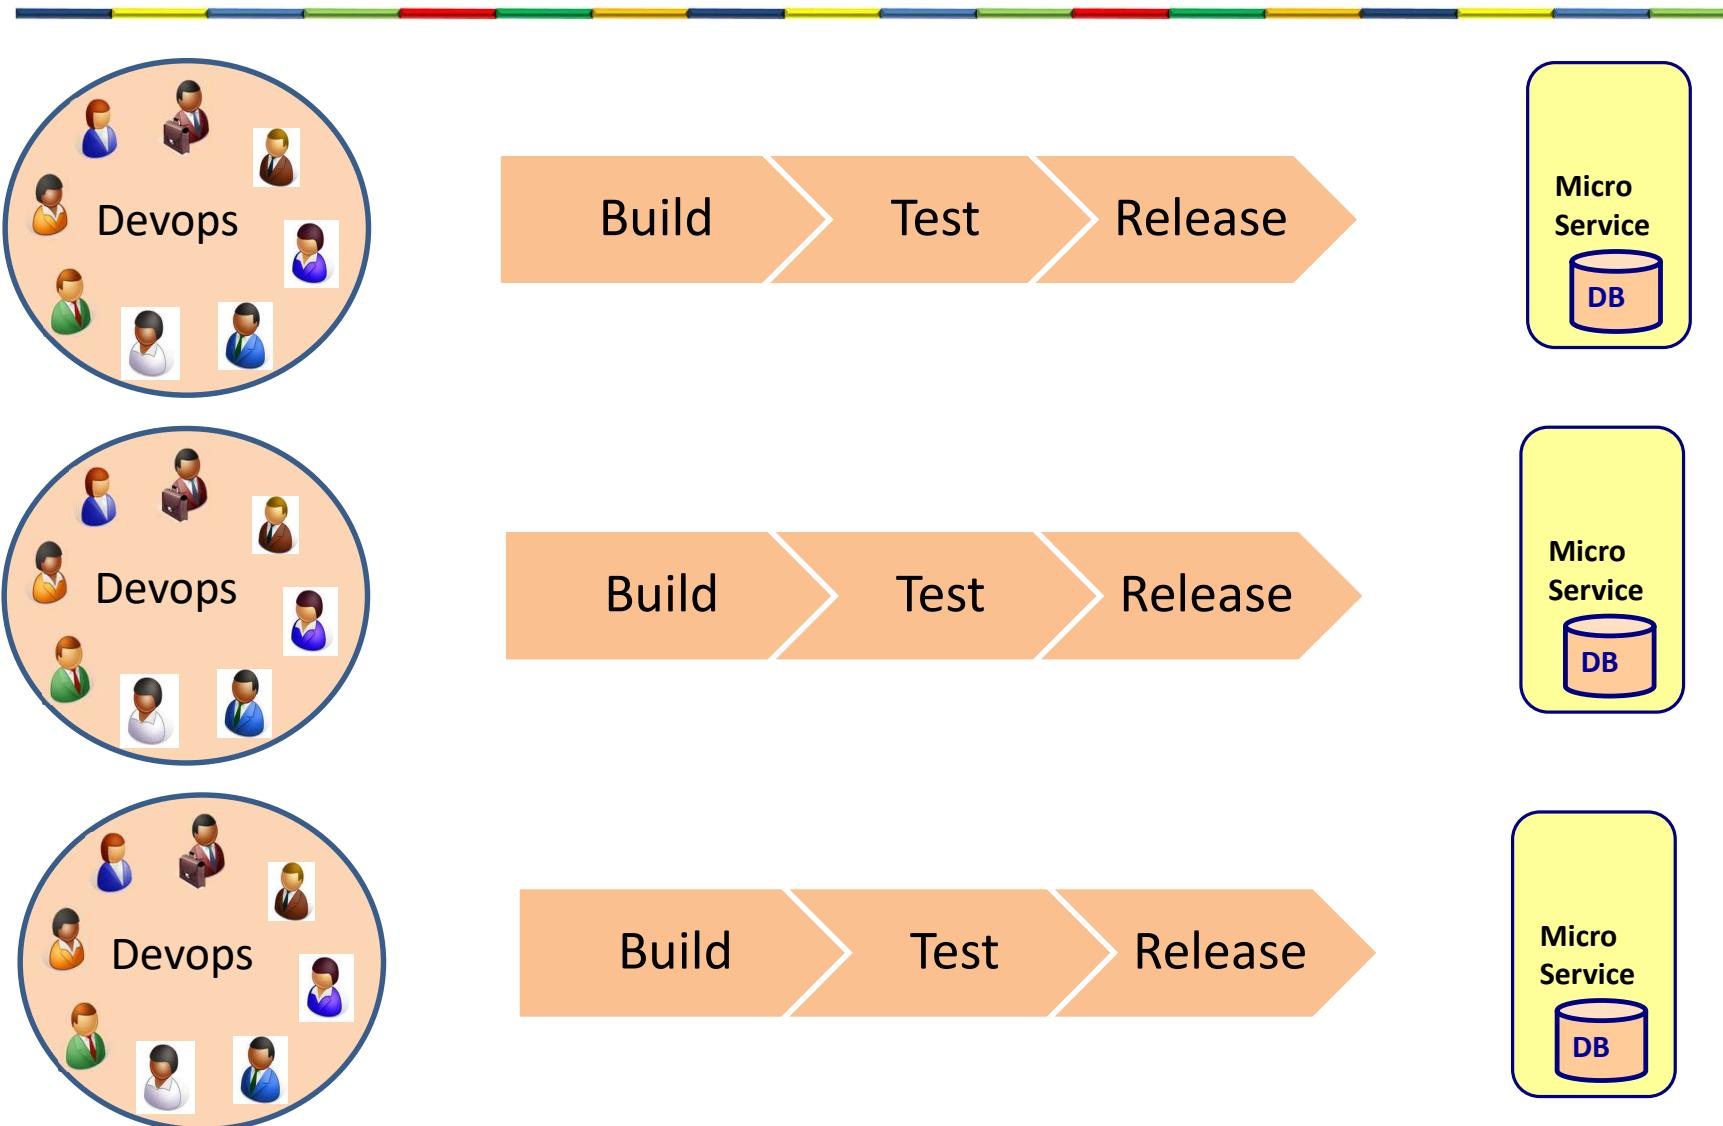
*"If you have four groups  
working on a compiler, you'll  
get a 4-pass compiler"*

*—Eric S Raymond*

*"organizations which design  
systems ... are constrained to  
produce designs which are copies  
of the communication structures  
of these organizations "*

*—Melvin Conway*

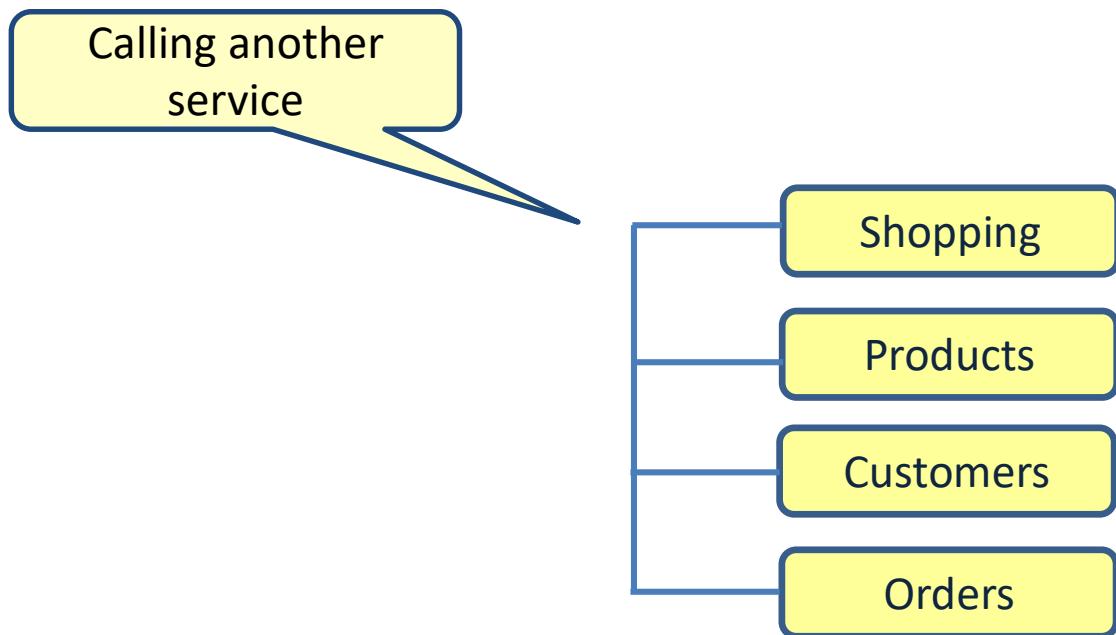
# Microservice organization



# **CALLING ANOTHER MICROSERVICE: FEIGN**

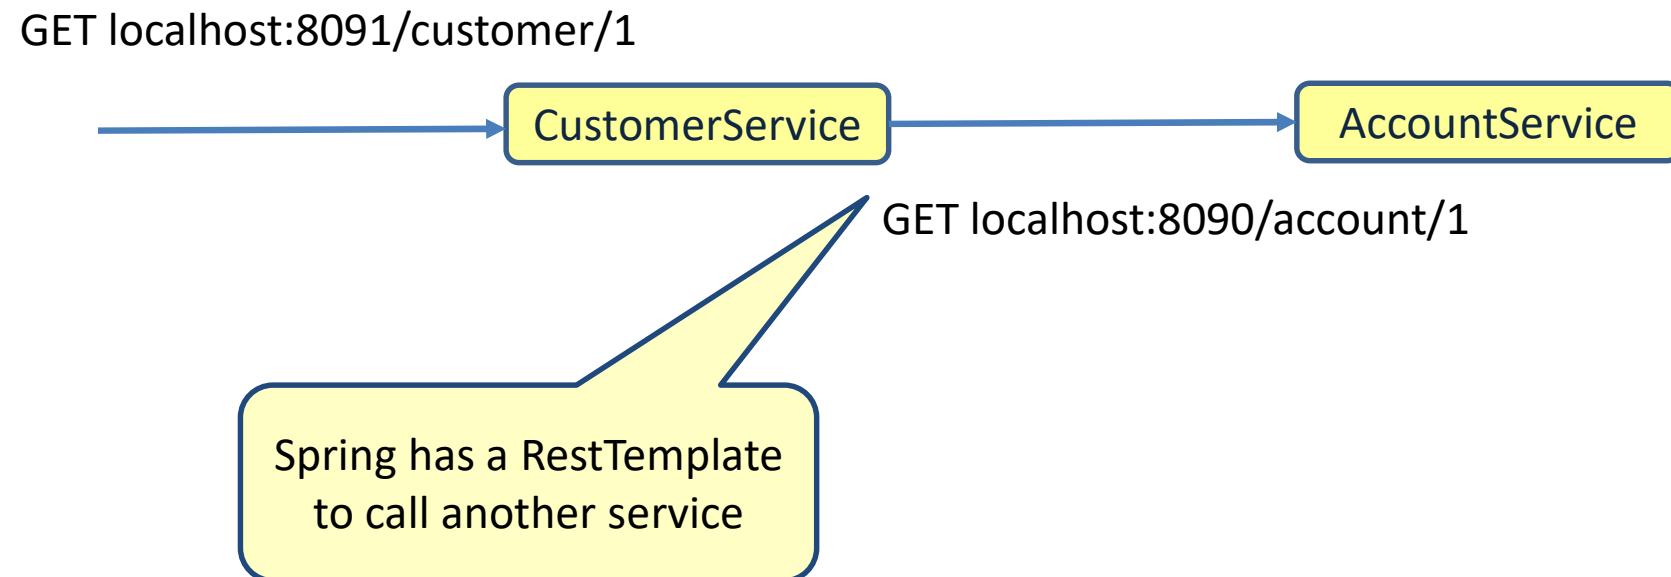
# Implementing microservices

---



# Calling another service

---



# RestTemplate

```
@Component
public class RestClient {
    @Autowired
    private RestOperations restTemplate;

    public void callRestServer(){
        Greeting greeting =
            restTemplate.getForObject("http://localhost:8080/greeting", Greeting.class);
        System.out.println("Receiving message:"+greeting.getContent());
    }
}
```

RestTemplate has to be configured.  
Developer has to know REST details

```
@Configuration
public class AppConfig {
    @Bean
    RestTemplate restTemplate(){
        return new RestTemplate();
    }
}
```

# Feign

---

- Declarative HTTP client
  - Simplify the HTTP client
- You only need to declare and annotate the interface

# AccountService

```
@RestController
public class AccountController {
    @RequestMapping("/account/{customerid}")
    public Account getName(@PathVariable("customerid") String customerId) {
        return new Account("1234", "1000.00");
    }
}
```

```
public class Account {
    private String accountNumber;
    private String balance;
    ...
}
```

application.yml

```
server:
  port: 8090
```

```
@SpringBootApplication
public class AccountServiceApplication {

    public static void main(String[] args) {
        SpringApplication.run(AccountServiceApplication.class, args);
    }
}
```

# Properties files and yml files

---

application.properties

```
Mapping single properties
myapp.mail.to=frank@hotmail.com
myapp.mail.host=mail.example.com
myapp.mail.port=250

#Mapping list or array
myapp.mail.cc=mike@gmail.com,david@gmail.com
myapp.mail.bcc=john@hotmail.com,admin@acme.com

#Mapping nested POJO class
myapp.mail.credential.user-name=john1234
myapp.mail.credential.password=xyz@1234
```

application.yml

```
myapp:
  mail:
    to: frank@hotmail.com
    host: mail.example.com
    port: 250
    cc:
      - mike@gmail.com
      - david@gmail.com
    bcc:
      - john@hotmail.com
      - admin@acme.com
  credential:
    user-name: john1234
    password: xyz@1234
```

# CustomerService

```
@SpringBootApplication  
@EnableFeignClients  
public class CustomerServiceApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(AccountServiceApplication.class, args);  
    }  
}
```

Use Feign

## application.yml

```
server:  
  port: 8091
```

```
<dependency>  
  <groupId>org.springframework.cloud</groupId>  
  <artifactId>spring-cloud-starter-openfeign</artifactId>  
</dependency>
```

# CustomerService: the controller

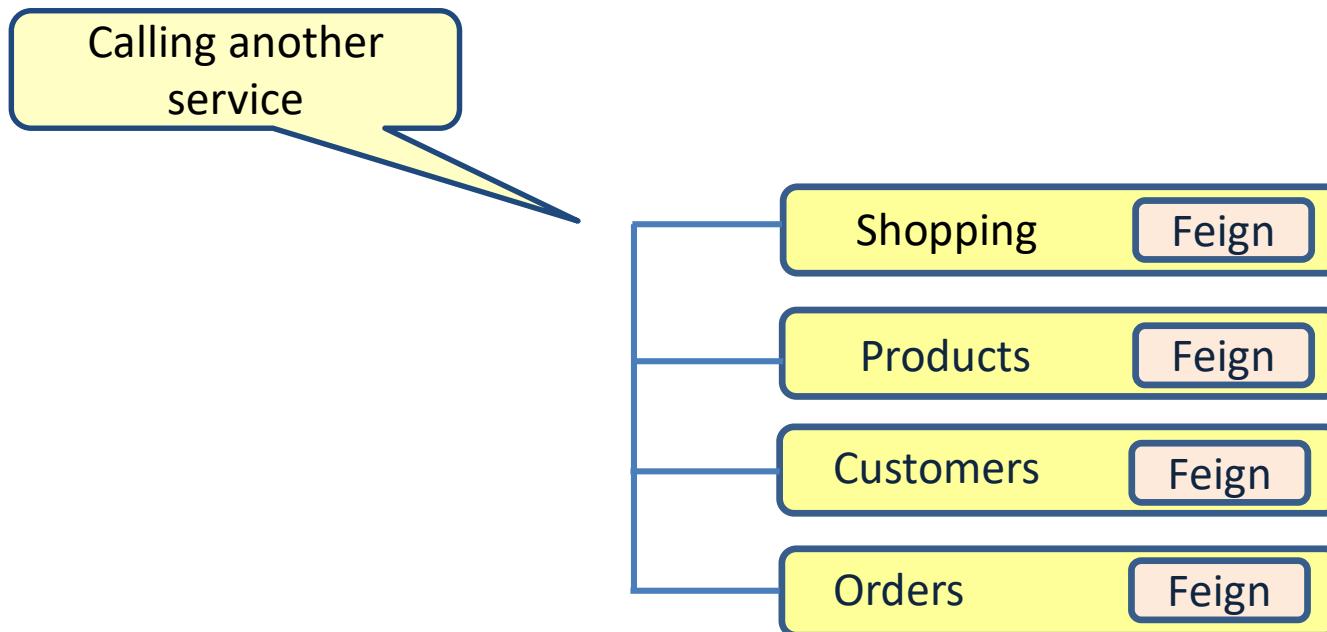
```
@RestController  
public class CustomerController {  
    @Autowired  
    AccountFeignClient accountClient;  
  
    @RequestMapping("/customer/{customerId}")  
    public Account getName(@PathVariable("customerId") String customerId) {  
        Account account = accountClient.getName(customerId);  
        return account;  
    }  
  
    @FeignClient(name = "account-service", url = "http://localhost:8090")  
    interface AccountFeignClient {  
        @RequestMapping("/account/{customerId}")  
        public Account getName(@PathVariable("customerId") String customerId);  
    }  
}
```

Autowire the client

Declare the interface, Spring creates the implementation

# Implementing microservices

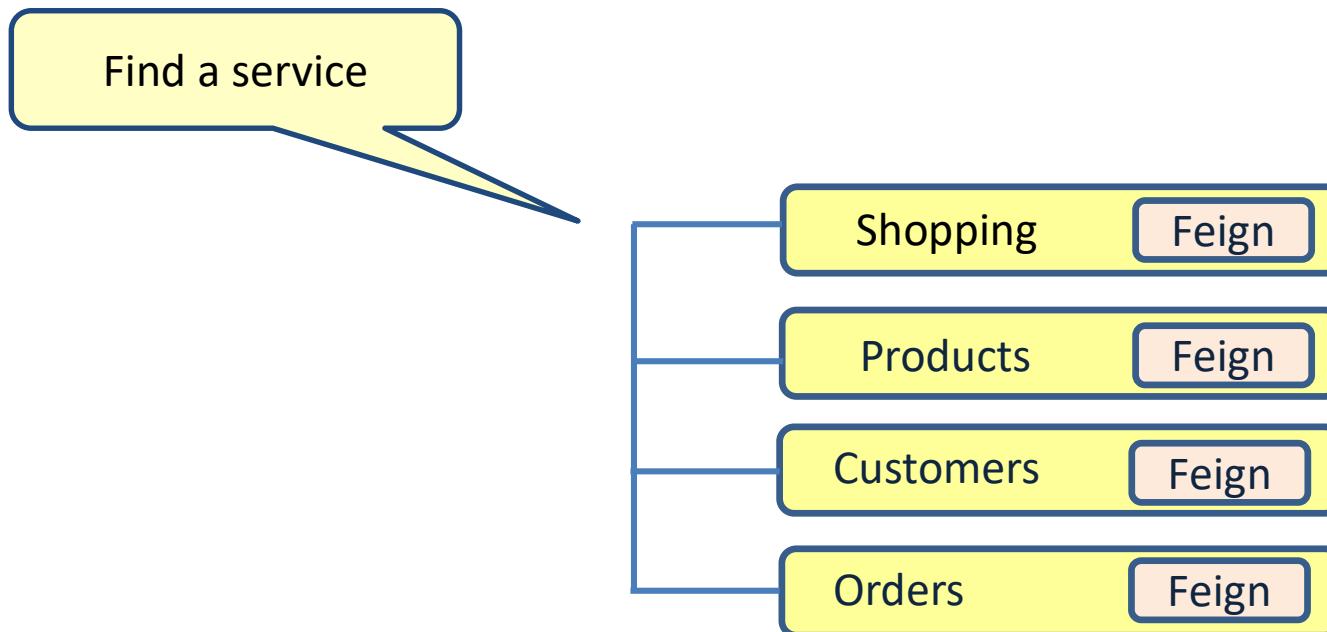
---



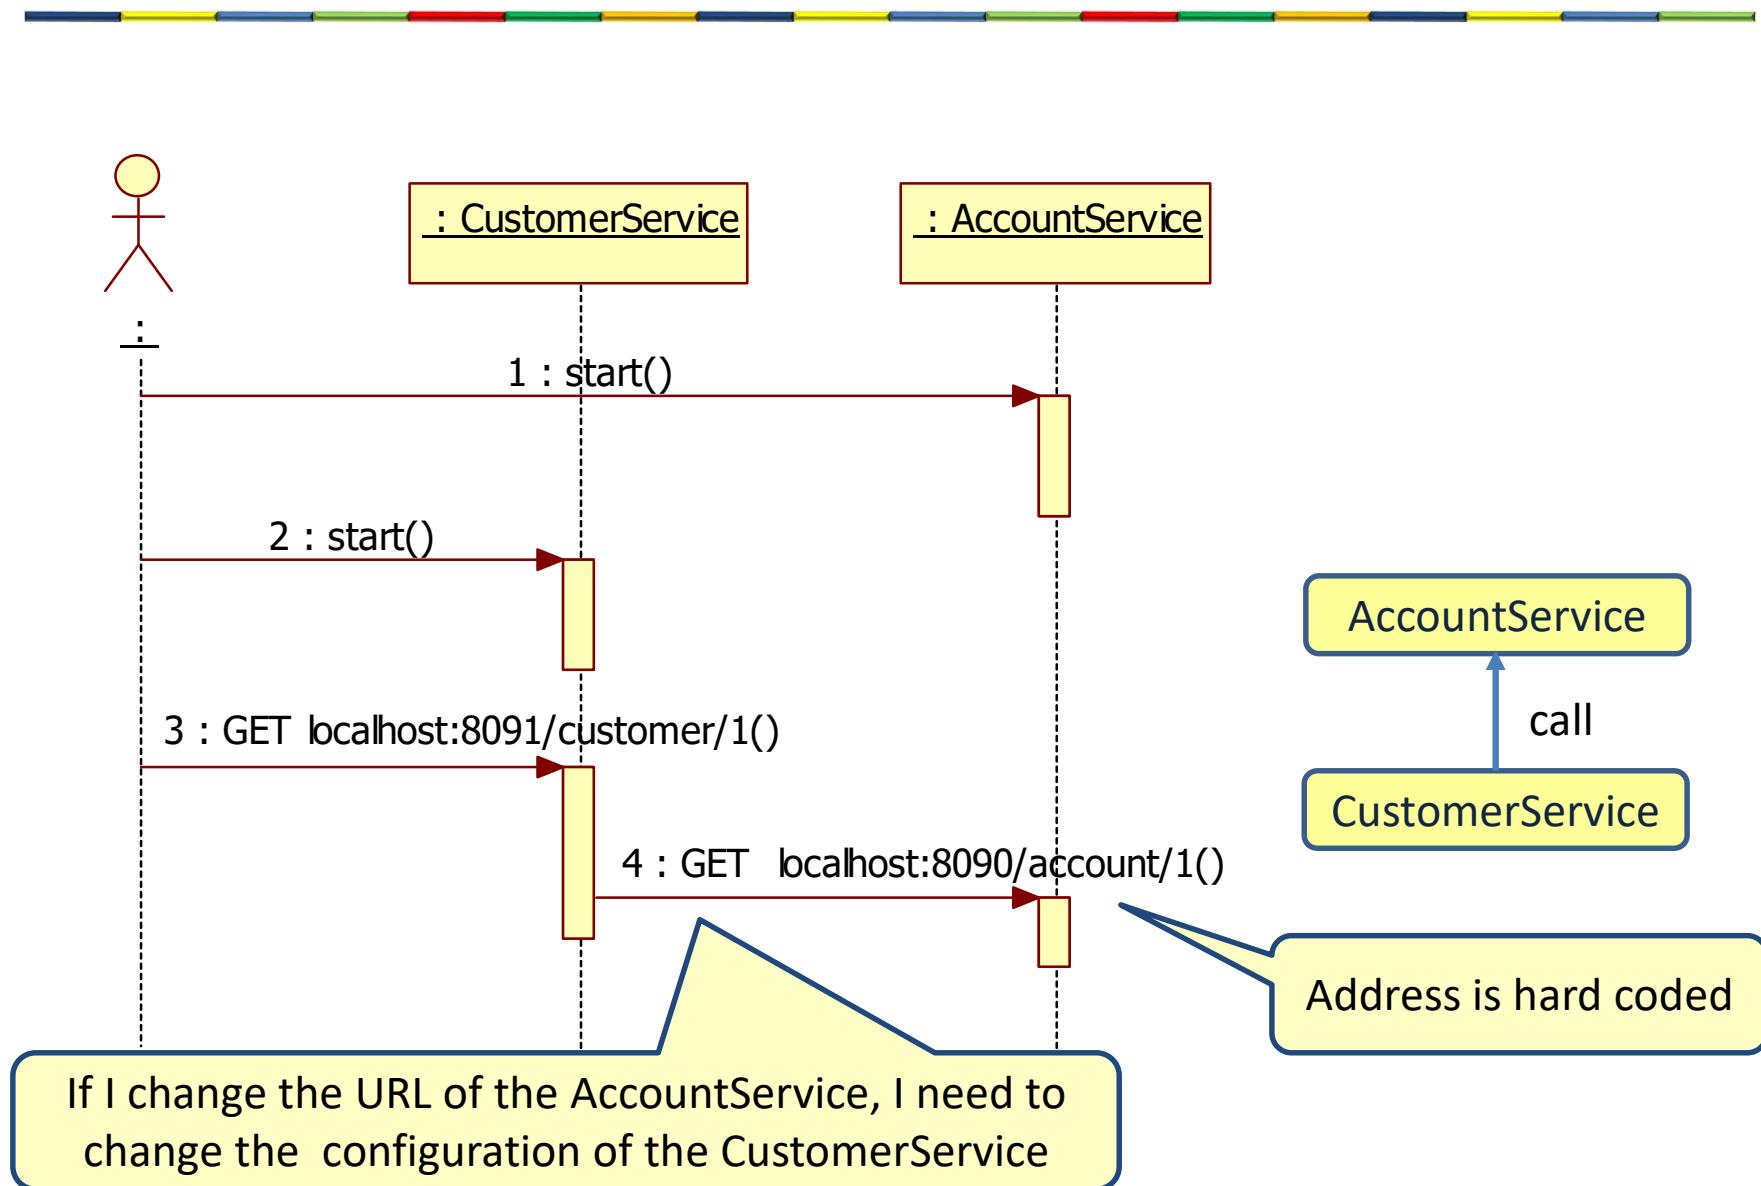
# **SERVICE REGISTRY: EUREKA**

# Implementing microservices

---



# One service calling another service

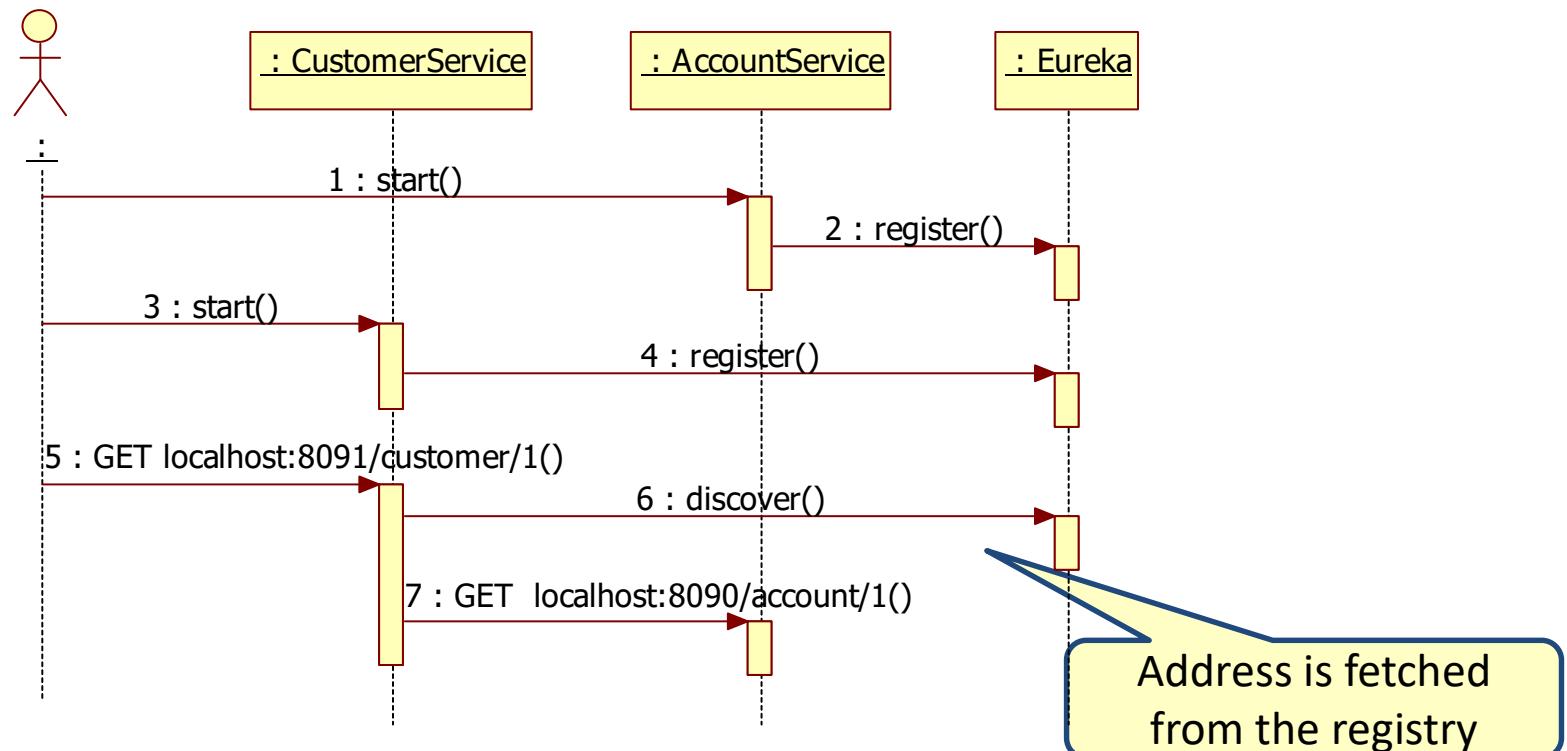
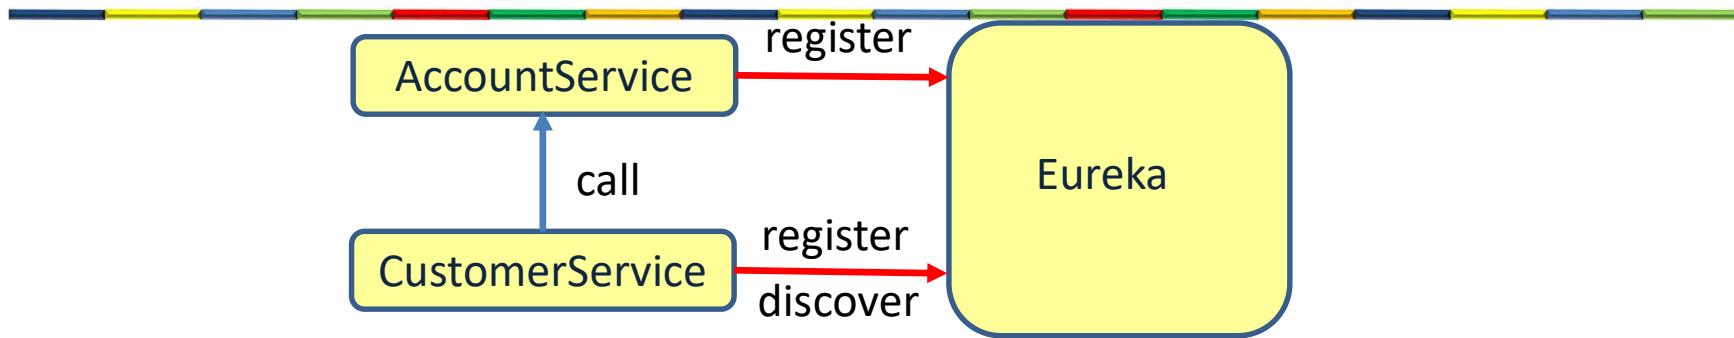


# Service Registry

---

- Like the phone book for microservices
  - Services register themselves with their location and other meta-data
  - Clients can lookup other services
- Netflix Eureka

# Using Eureka



# Why service registry/discovery?

---

## 1. Loosely coupled services

- Service consumers should not know the physical location of service instances.
  - We can easily scale up or scale down service instances

## 2. Increase application resilience

- If a service instance becomes unhealthy or unavailable, the service discovery engine will remove that instance from the list of available services.

# Eureka Server

```
@SpringBootApplication  
@EnableEurekaServer  
public class EurekaServerApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(EurekaServerApplication.class, args);  
    }  
}
```

application.yml

```
server:  
  port: 8761  
  
eureka:  
  client:  
    registerWithEureka: false      #telling the server not to register himself  
    fetchRegistry: false
```

```
<dependency>  
  <groupId>org.springframework.cloud</groupId>  
  <artifactId>spring-cloud-starter-netflix-eureka-server</artifactId>  
</dependency>
```

# Running Eureka



# AccountService

```
@SpringBootApplication
@EnableDiscoveryClient
public class AccountServiceApplication {

    public static void main(String[] args) {
        SpringApplication.run(AccountServiceApplication.class, args);
    }
}
```

## application.yml

```
server:
  port: 8090

eureka:
  client:
    serviceUrl:
      defaultZone: http://localhost:8761/eureka/

spring:
  application:
    name: AccountService
```

# AccountService



```
@RestController
public class AccountController {
    @RequestMapping("/account/{customerid}")
    public Account getName(@PathVariable("customerid") String customerId) {
        return new Account("1234", "1000.00");
    }
}
```

```
public class Account {
    private String accountNumber;
    private String balance;
    ...
}
```

```
<dependency>
<groupId>org.springframework.cloud</groupId>
<artifactId>spring-cloud-starter-netflix-eureka-client</artifactId>
</dependency>
```

# Running the AccountService

Eureka

localhost:8761

localhost

### Instances currently registered with Eureka

Application	AMIs	Availability Zones	Status
ACCOUNTSERVICE	n/a (1)	(1)	UP (1) - 192.168.1.104:AccountService:8090

### General Info

Name	Value
total-avail-memory	67mb
environment	test
num-of-cpus	8

# CustomerService

```
@SpringBootApplication  
@EnableDiscoveryClient  
@EnableFeignClients  
public class CustomerServiceApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(AccountServiceApplication.class, args);  
    }  
}
```

Use Feign and Eureka

## application.yml

```
server:  
  port: 8091  
  
eureka:  
  client:  
    serviceUrl:  
      defaultZone: http://localhost:8761/eureka/  
  
spring:  
  application:  
    name: CustomerService
```

# CustomerService: the controller

---

```
@RestController
public class CustomerController {
    @Autowired
    AccountFeignClient accountClient;

    @RequestMapping("/customer/{customerId}")
    public Account getName(@PathVariable("customerId") String customerId) {
        Account account = accountClient.getName(customerId);
        return account;
    }

    @FeignClient("AccountService")
    interface AccountFeignClient {
        @RequestMapping("/account/{customerId}")
        public Account getName(@PathVariable("customerId") String customerId);
    }
}
```

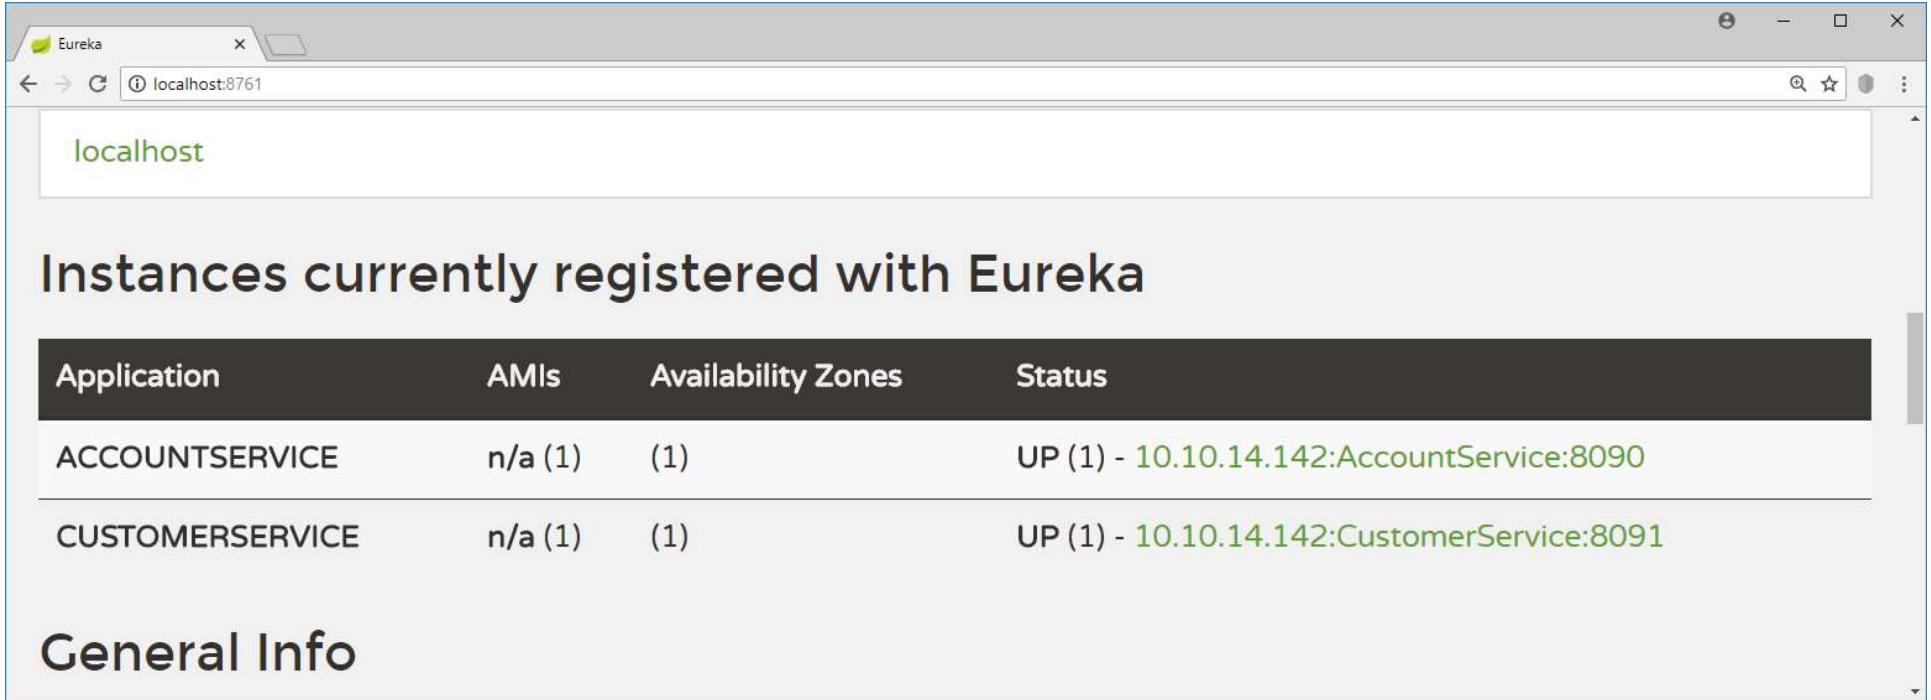
Name of the service instead of the URL

Feign works together  
with Eureka

application.yml

```
server:
  port: 8091
```

# Running the CustomerService



Instances currently registered with Eureka

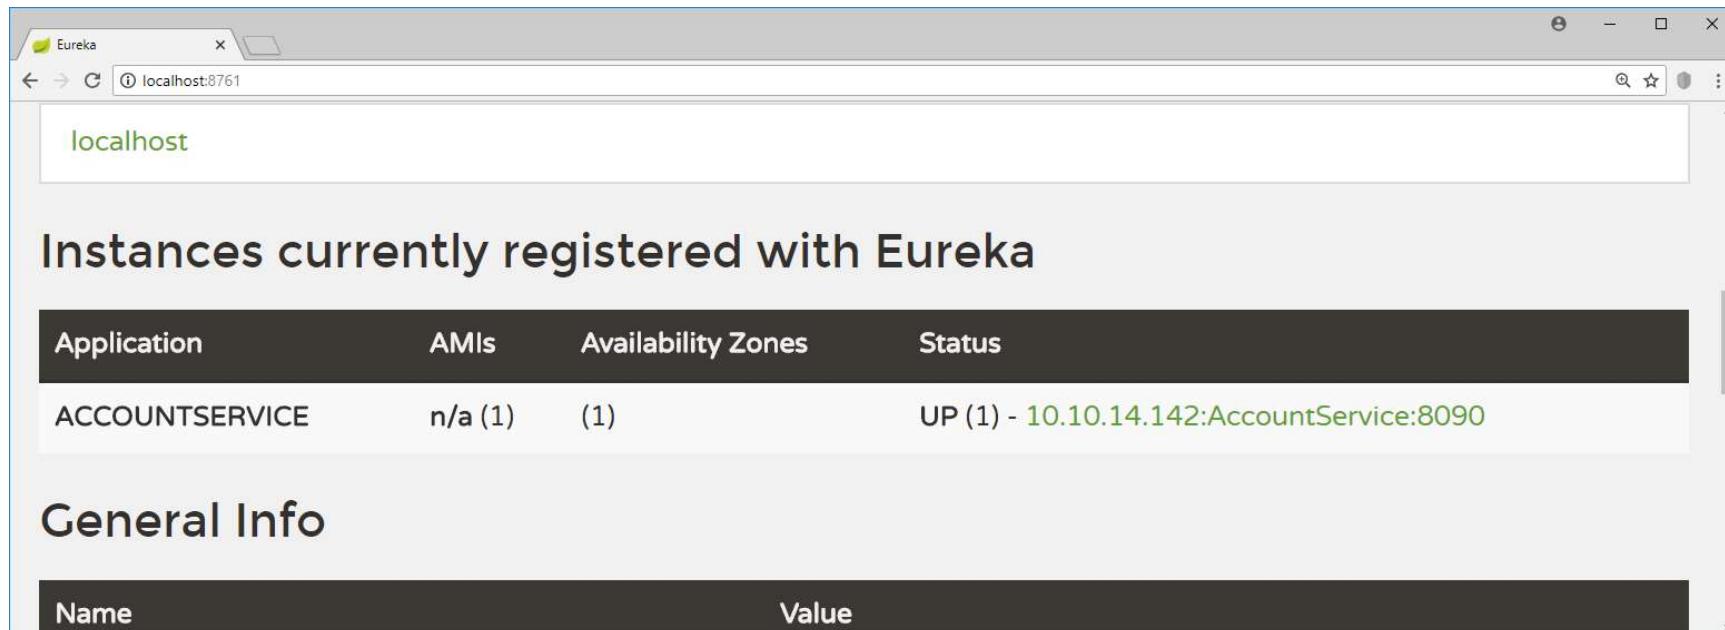
Application	AMIs	Availability Zones	Status
ACCOUNTSERVICE	n/a (1)	(1)	UP (1) - <a href="#">10.10.14.142:AccountService:8090</a>
CUSTOMERSERVICE	n/a (1)	(1)	UP (1) - <a href="#">10.10.14.142:CustomerService:8091</a>

General Info

# Stopping the CustomerService

---

- Eureka monitors the health of registered services.
- If we stop the CustomerService, Eureka will notice that automatically



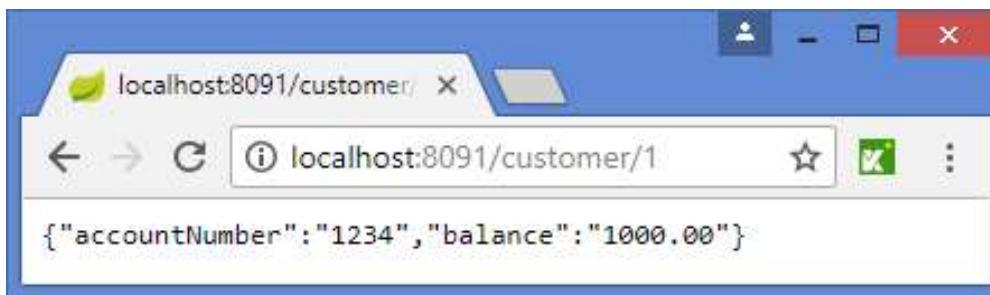
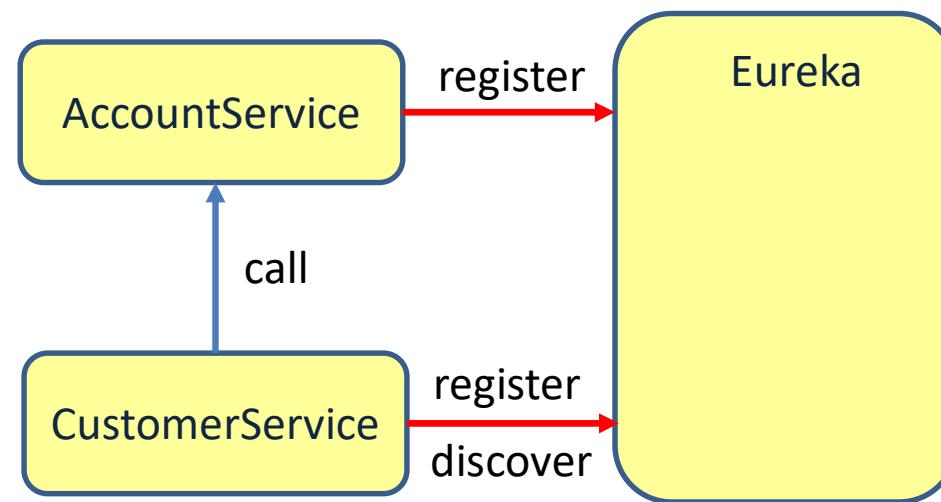
The screenshot shows the Eureka user interface running in a browser window. The title bar says "Eureka". The address bar shows "localhost:8761". The main content area has a heading "Instances currently registered with Eureka". Below it is a table:

Application	AMIs	Availability Zones	Status
ACCOUNTSERVICE	n/a (1)	(1)	UP (1) - 10.10.14.142:AccountService:8090

Below the table is a section titled "General Info" with a table:

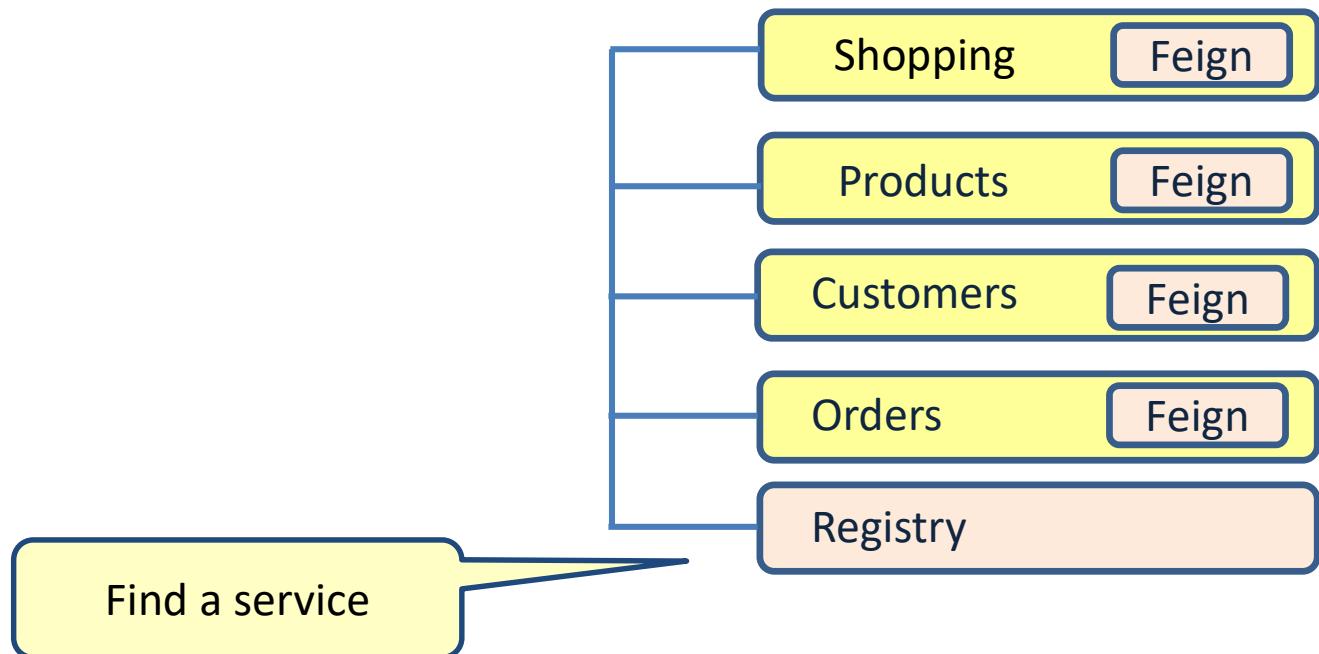
Name	Value

# Using Eureka



# Implementing microservices

---



Lesson 9

# **MICROSERVICES**

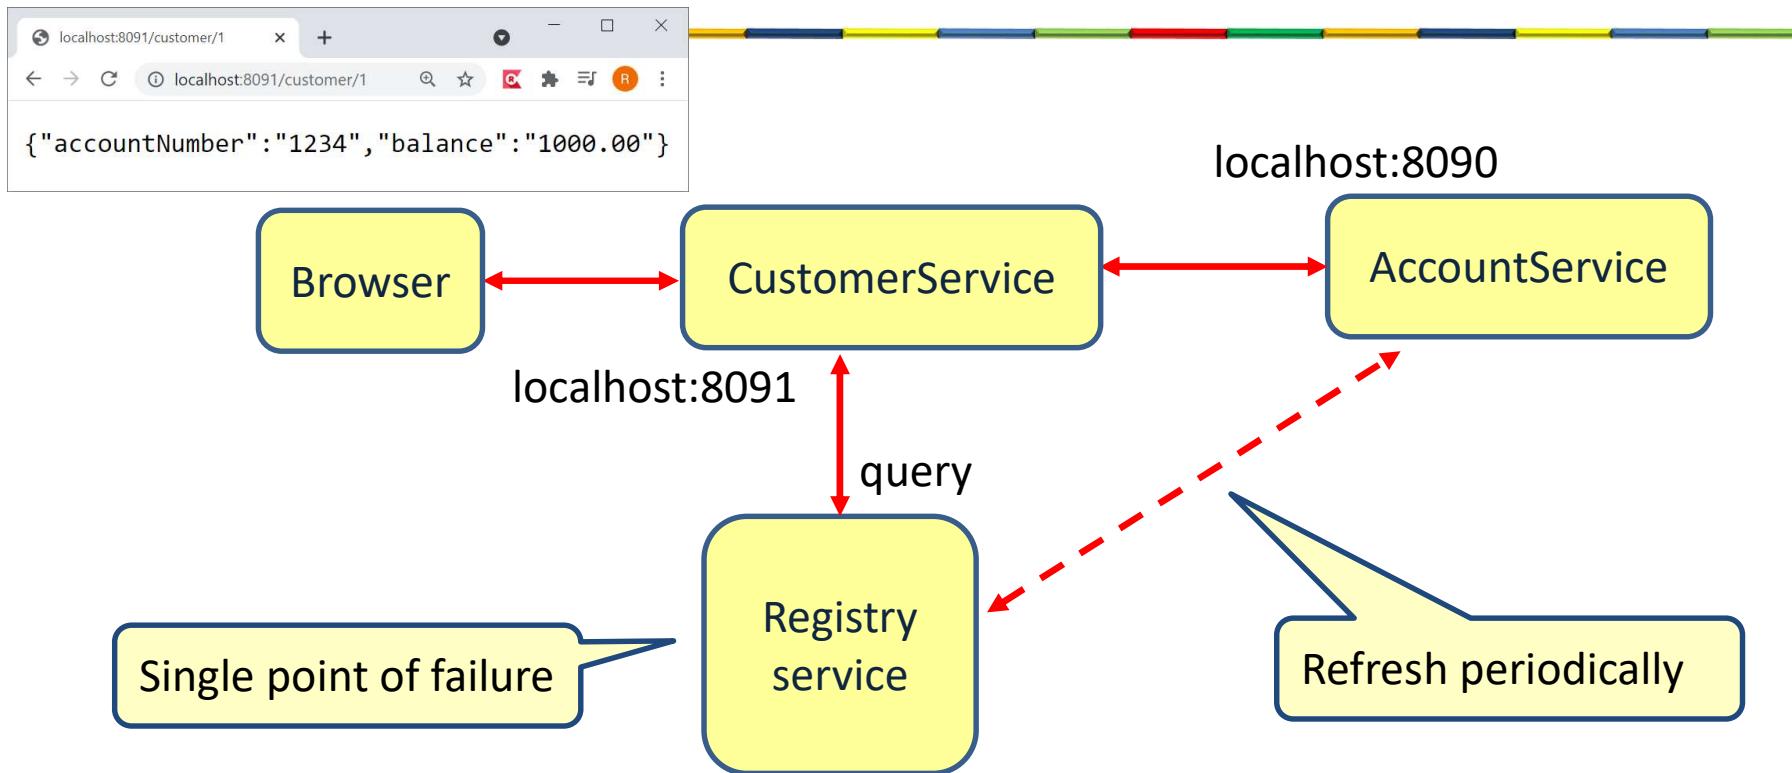
# Challenges of a microservice architecture

---

Challenge	Solution
Complex communication	Feign Registry
Performance	
Resilience	
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	

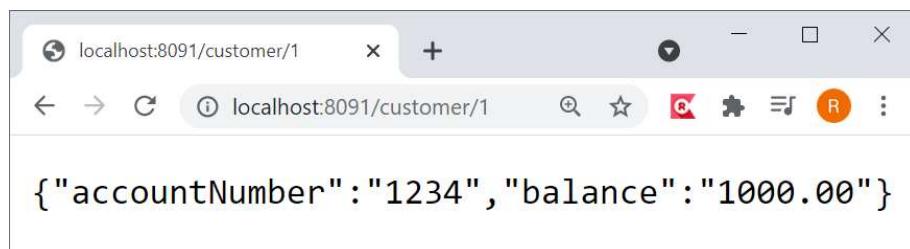
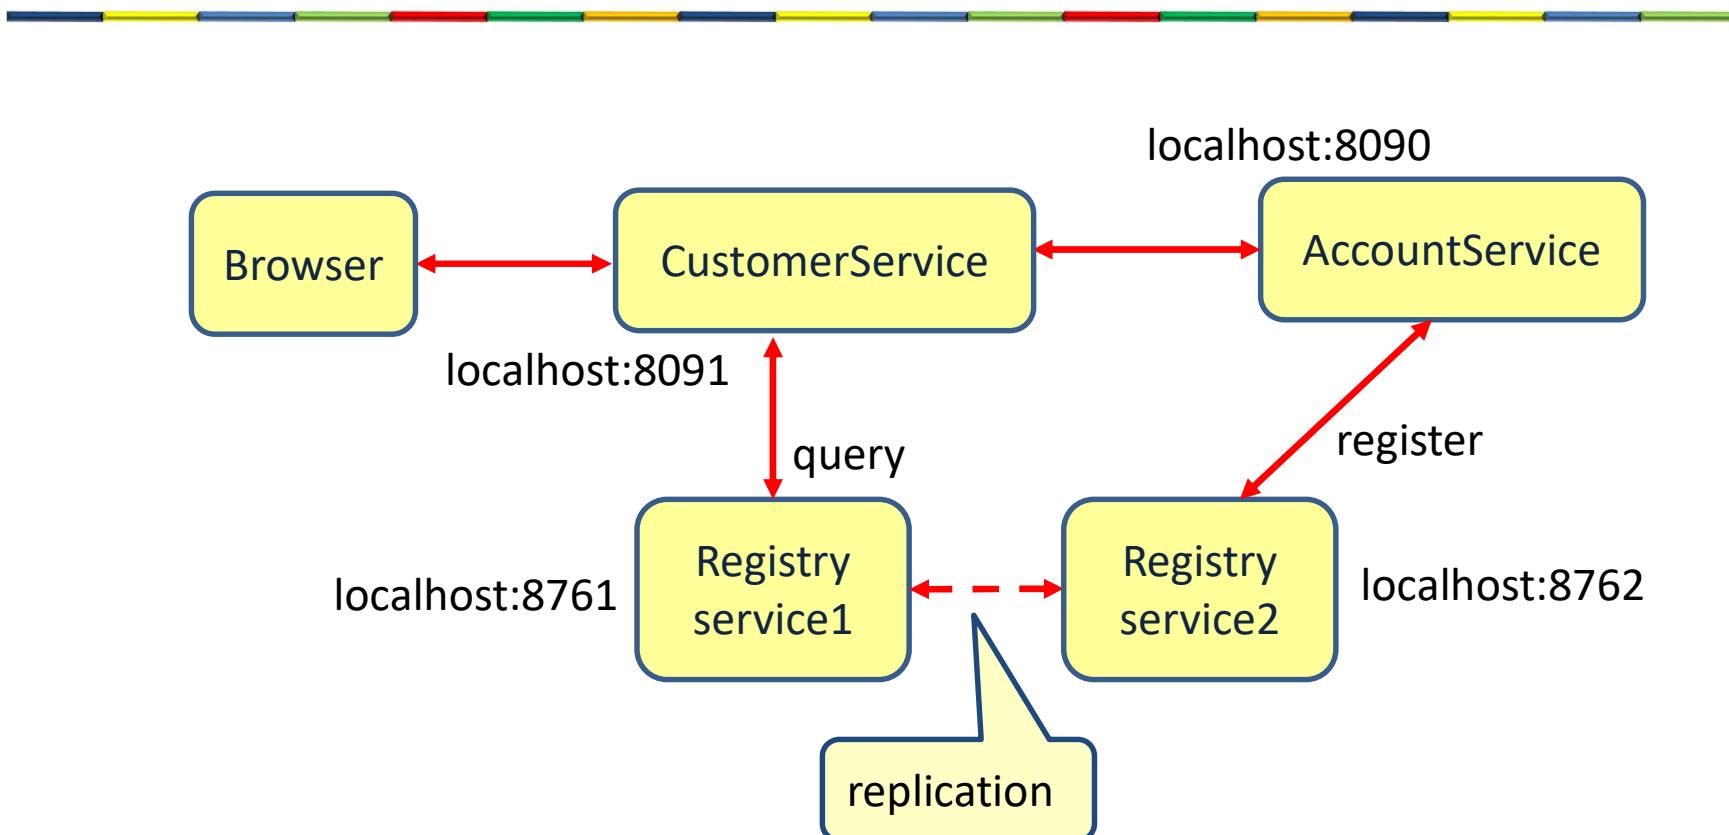
# **EUREKA FAILOVER**

# Eureka registry

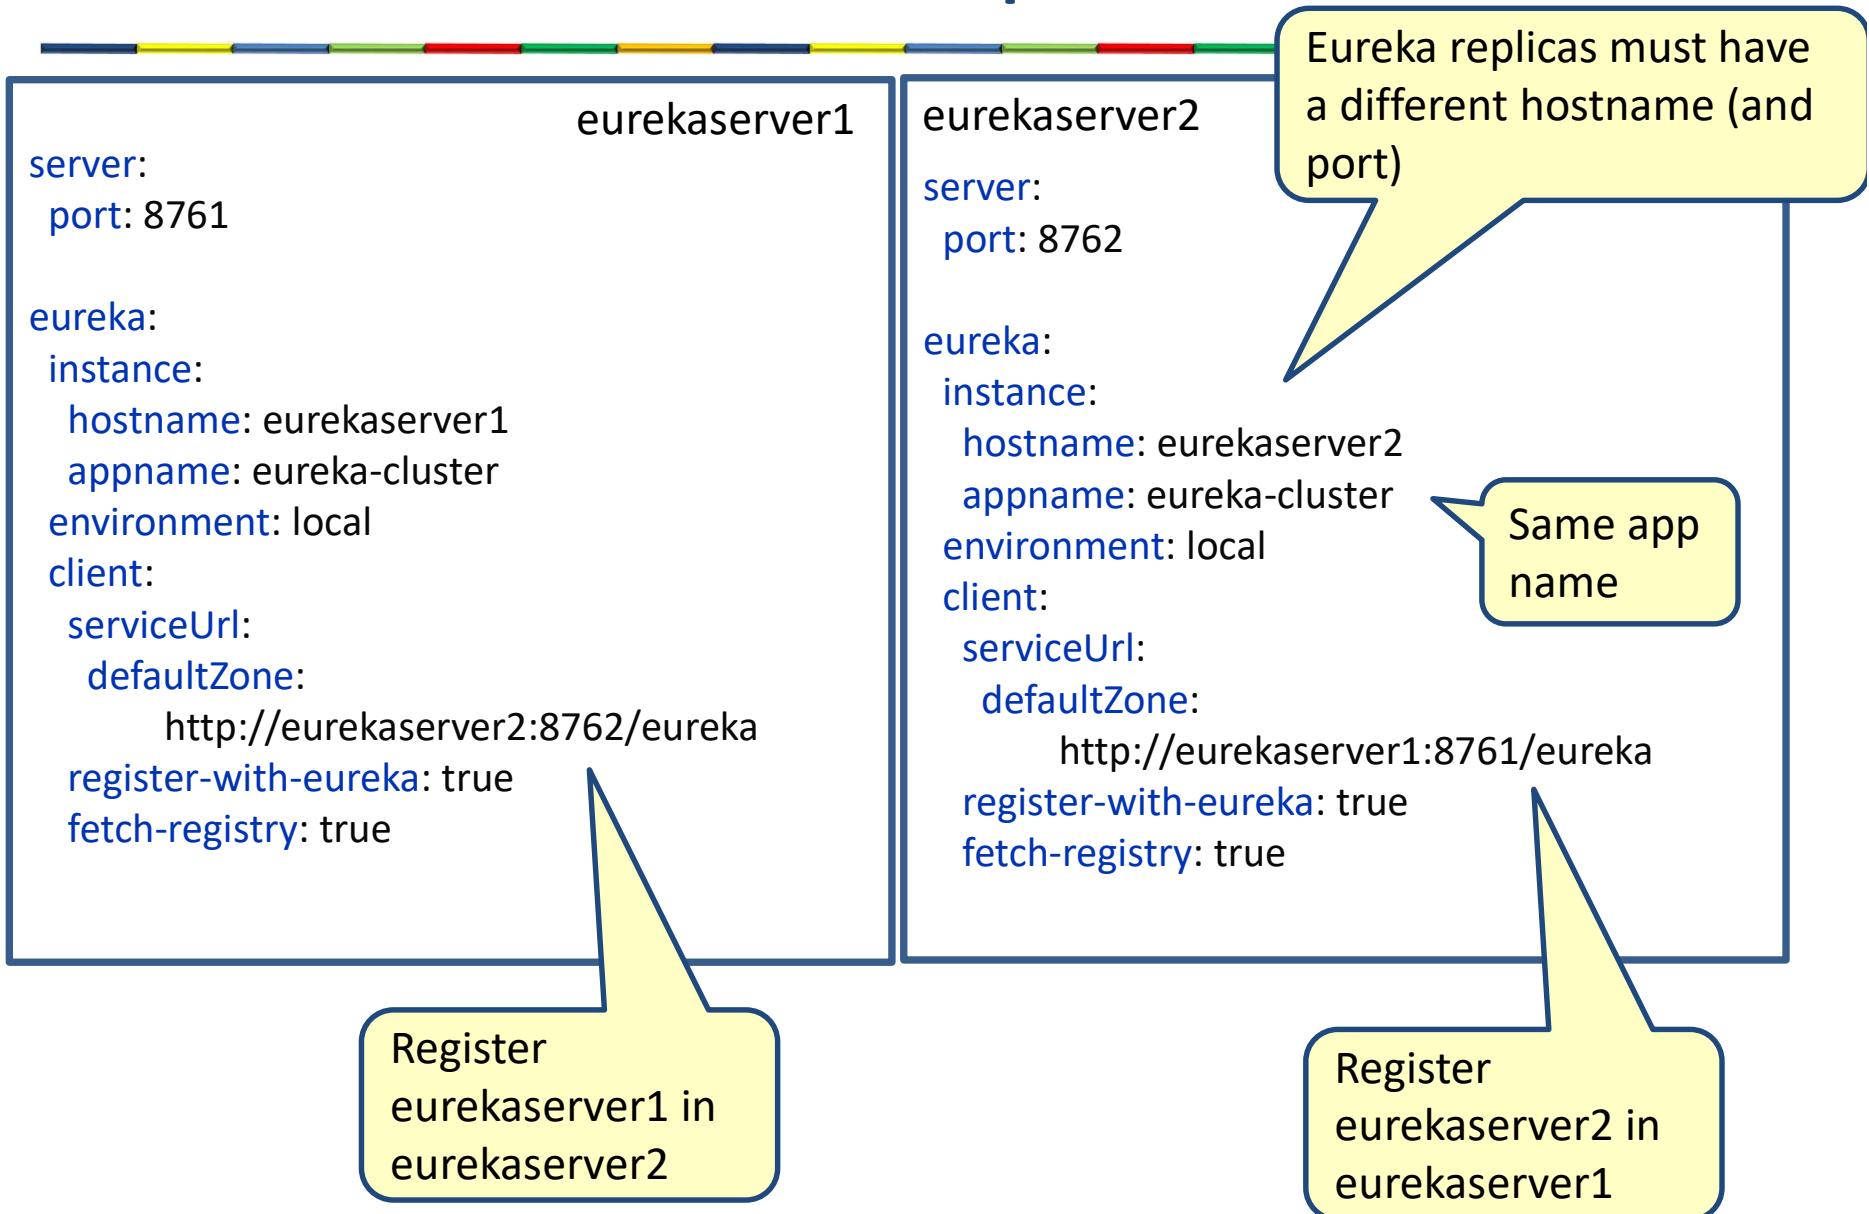


ACCOUNTSERVICE	n/a (1)	(1)	UP (1) - 10.10.14.142:AccountService:8090
CUSTOMERSERVICE	n/a (1)	(1)	UP (1) - 10.10.14.142:CustomerService:8091

# Registry replication



# Eureka replicas



# Hosts file

---

- Window:  
c:\Windows\System32\Drivers\etc\hosts
- Linux : /etc/hosts

```
# localhost name resolution is handled within DNS itself.  
# 127.0.0.1      localhost  
# ::1            localhost  
127.0.0.1      eurekaserver1  
127.0.0.1      eurekaserver2
```

Map host names  
to machine  
addresses

# EurekaServer1

The screenshot shows the Spring Eureka UI running at `localhost:8761`. The interface includes sections for System Status, DS Replicas, Instances currently registered with Eureka, General Info, and Instance Info.

**System Status**

Environment	N/A	Current time	2021-07-29T03:38:31 -0500
Data center	N/A	Uptime	00:00
		Lease expiration enabled	false
		Renews threshold	3
		Renews (last min)	0

**DS Replicas**

eurekaserver2
---------------

**Instances currently registered with Eureka**

Application	AMIs	Availability Zones	Status
EUREKA-CLUSTER	n/a (1)	(1)	UP (1) - DESKTOP-BVHRK6K.home:8762

**General Info**

Name	Value
total-avail-memory	256mb
num-of-cpus	8
current-memory-usage	115mb (44%)
server-uptime	00:00
registered-replicas	<a href="http://eurekaserver2:8762/eureka/">http://eurekaserver2:8762/eureka/</a>
unavailable-replicas	
available-replicas	<a href="http://eurekaserver2:8762/eureka/">http://eurekaserver2:8762/eureka/</a> ,

**Instance Info**

# EurekaServer2

The screenshot shows the Spring Eureka UI running at `localhost:8762`. The interface includes sections for System Status, DS Replicas, Instances currently registered with Eureka, General Info, and Instance Info.

**System Status**

Environment	N/A	Current time	2021-07-29T03:40:16 -0500
Data center	N/A	Uptime	00:02
		Lease expiration enabled	true
		Renews threshold	1
		Renews (last min)	4

**DS Replicas**

eurekaserver1
---------------

**Instances currently registered with Eureka**

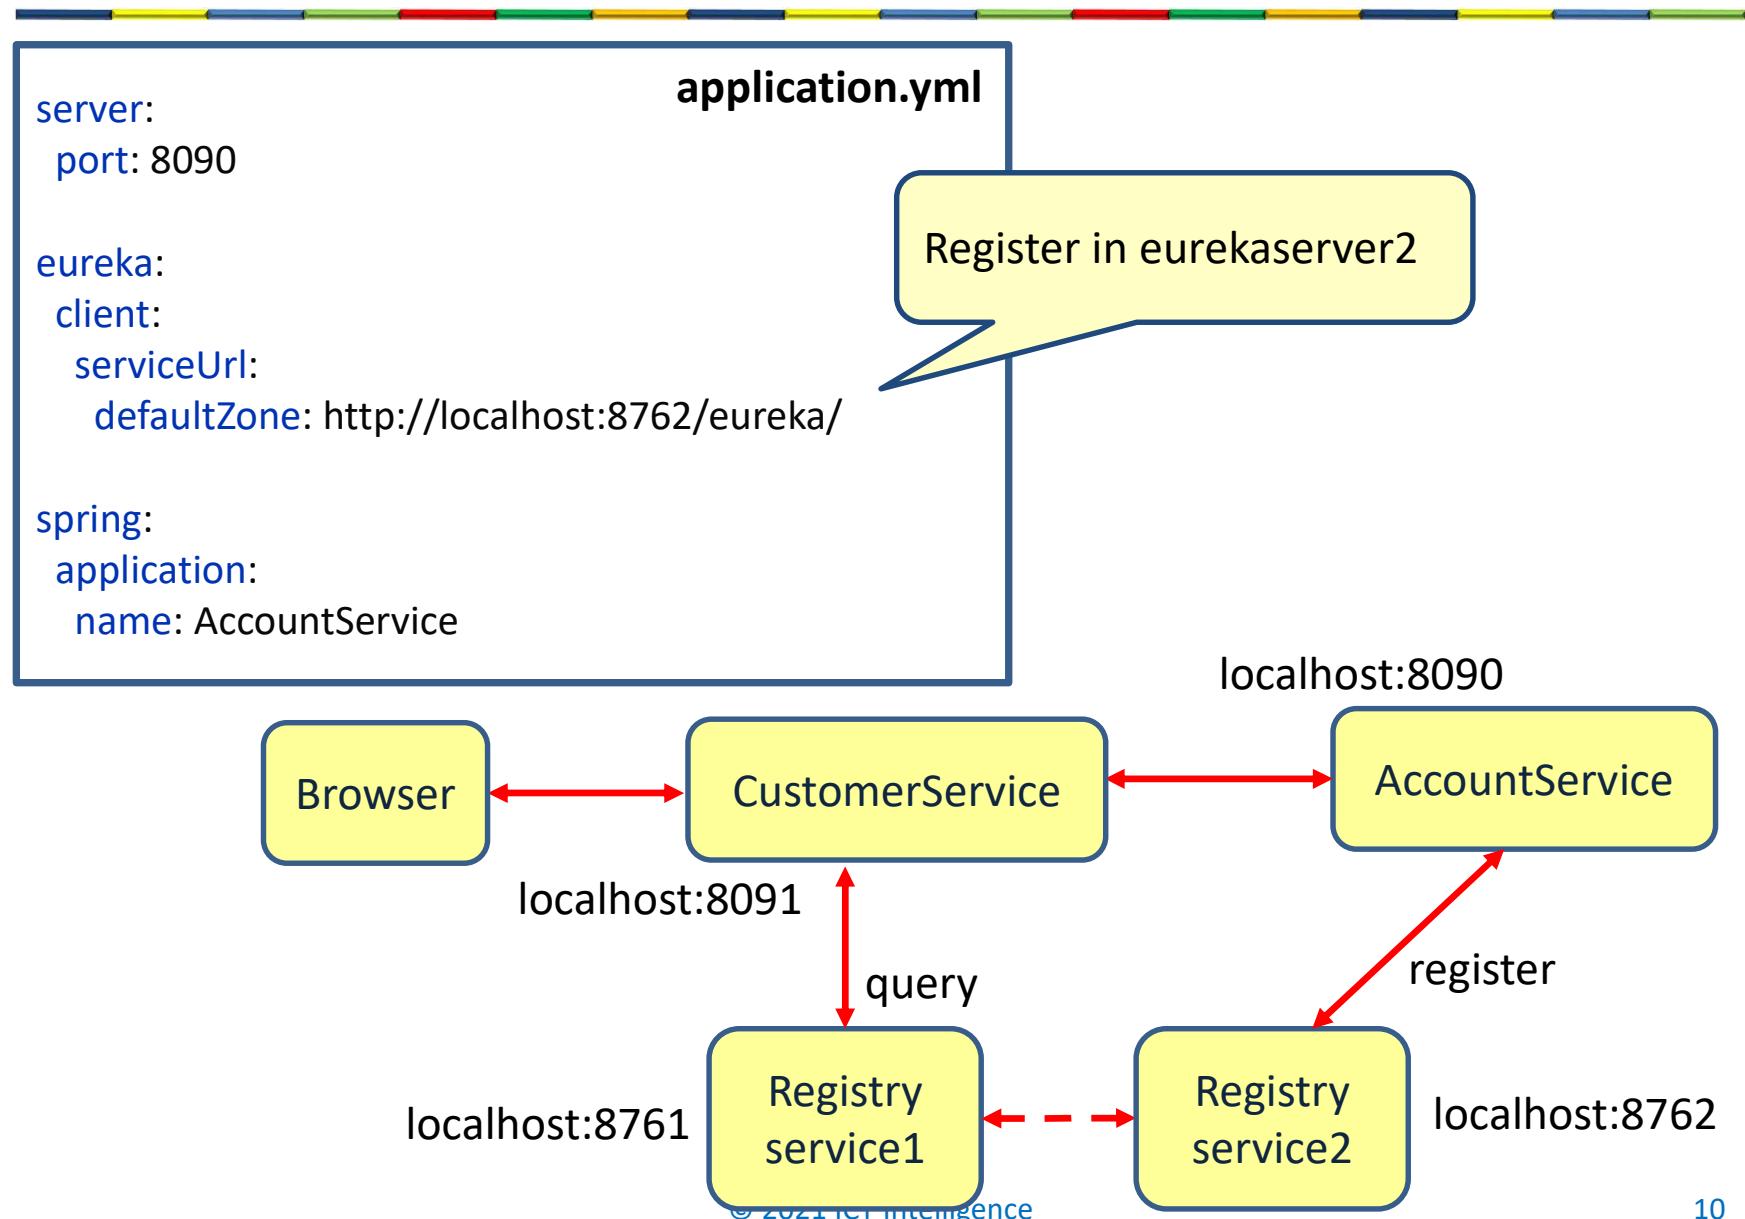
Application	AMIs	Availability Zones	Status
EUREKA-CLUSTER	n/a (2)	(2)	UP (2) - DESKTOP-BVHRK6K.home:8762 , DESKTOP-BVHRK6K.home:8761

**General Info**

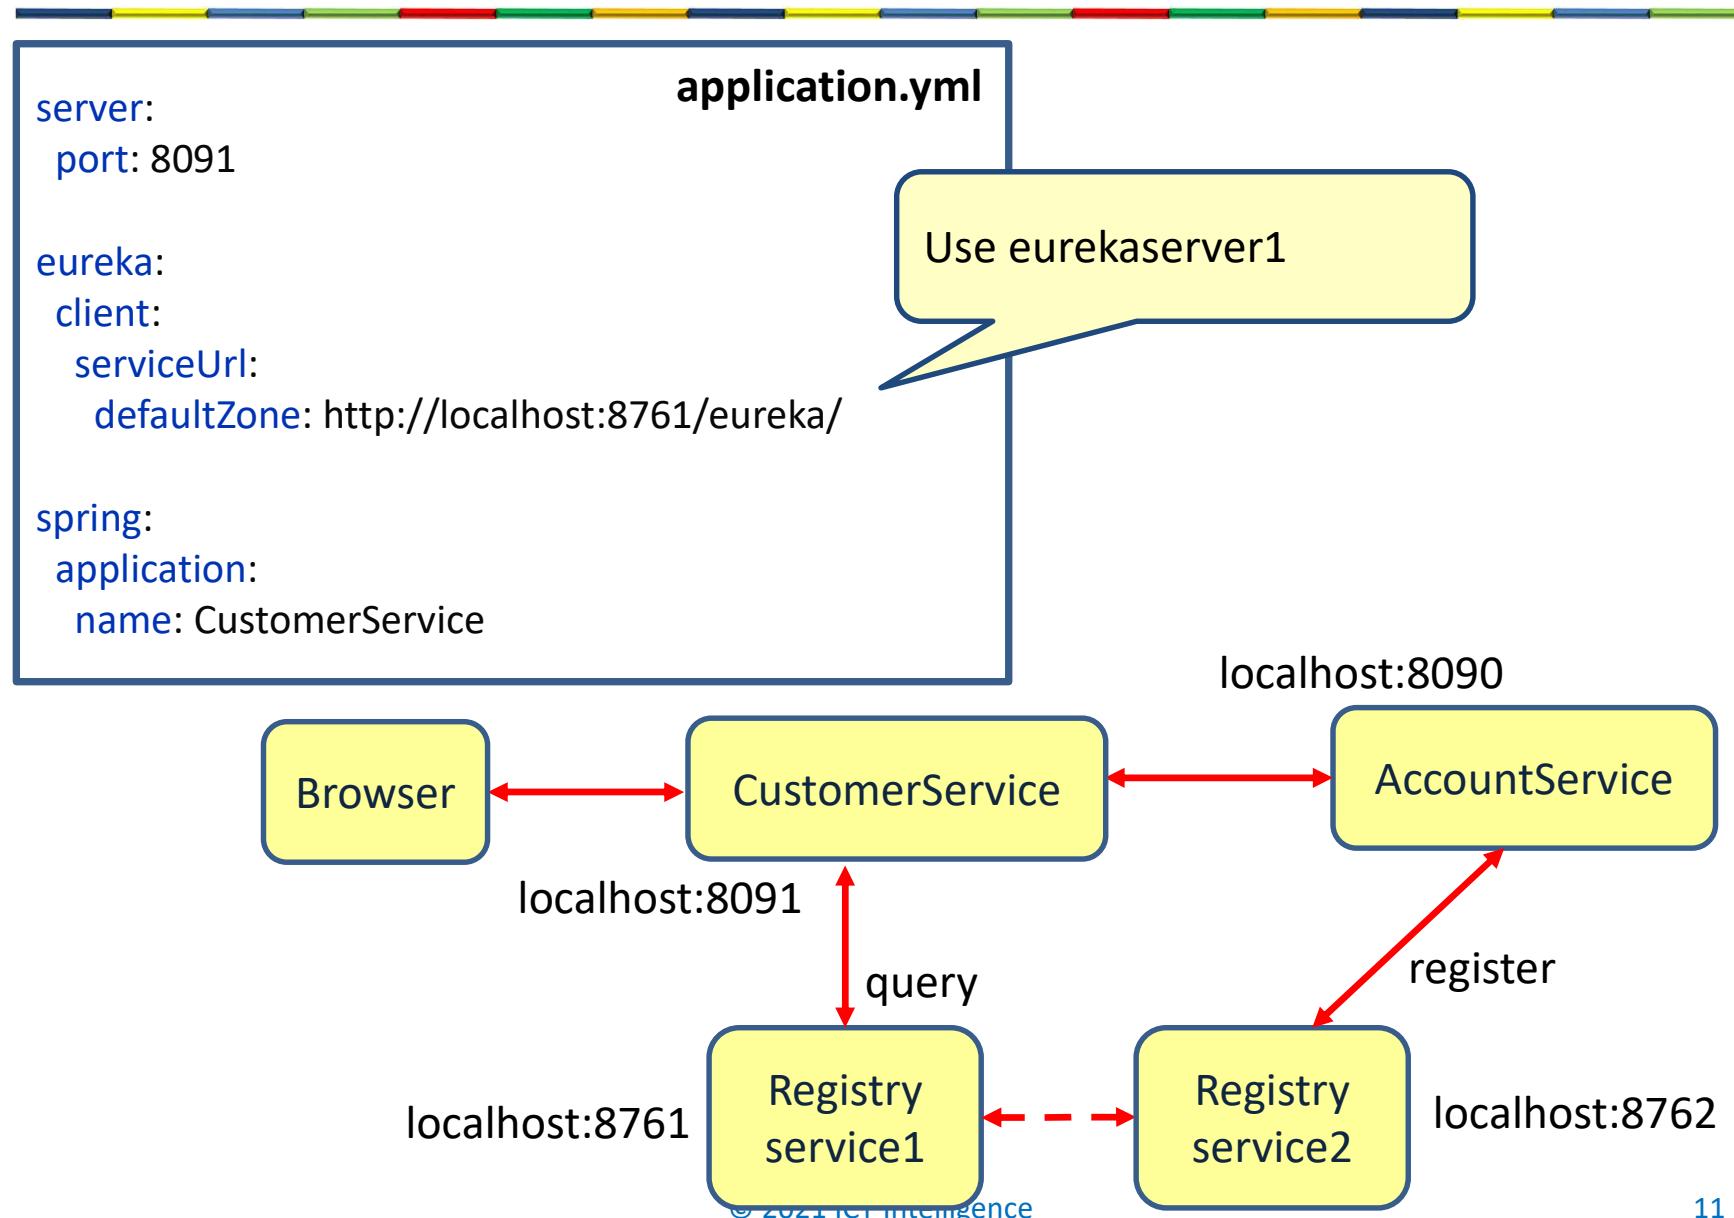
Name	Value
total-avail-memory	256mb
num-of-cpus	8
current-memory-usage	109mb (42%)
server-upptime	00:02
registered-replicas	<a href="http://eurekaserver1:8761/eureka/">http://eurekaserver1:8761/eureka/</a>
unavailable-replicas	
available-replicas	<a href="http://eurekaserver1:8761/eureka/">http://eurekaserver1:8761/eureka/</a> ,

**Instance Info**

# Accountservice



# CustomerService



# Eureka high availability

---

- In the client, multiple Eureka servers can be configured.

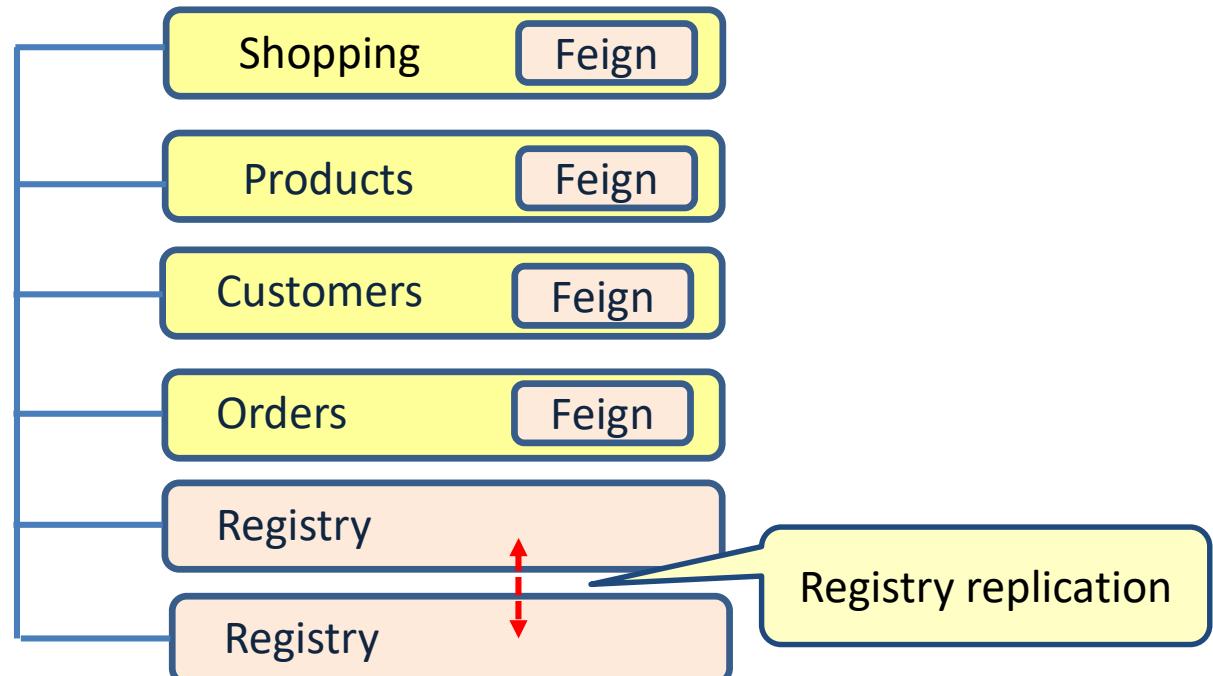
**application.yml**

```
server:  
  port: 8091  
  
eureka:  
  client:  
    serviceUrl:  
      defaultZone: http://eurekaserver1:8761/eureka/,  
                  http://eurekaserver2:8762/eureka/
```

This can be a comma separated list of Eureka instances.  
If the first instance does not respond, we try the next instance

# Implementing microservices

---



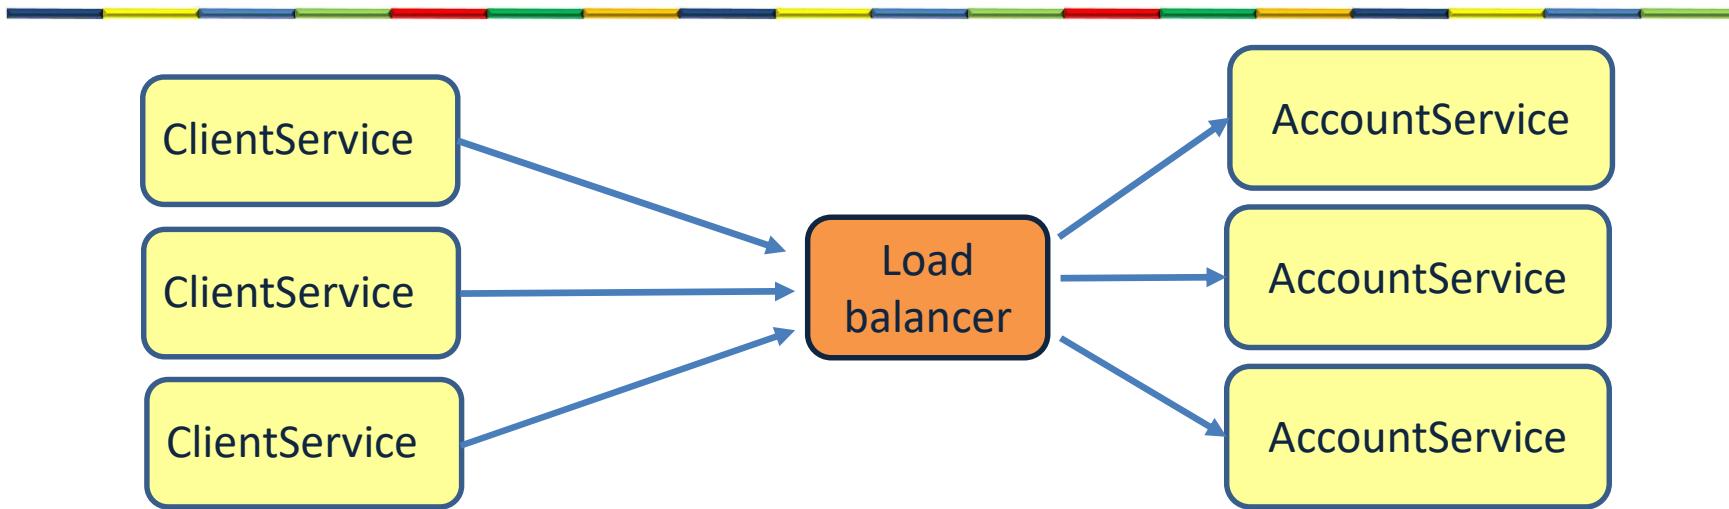
# Challenges of a microservice architecture

---

Challenge	Solution
Complex communication	Feign Registry
Performance	
Resilience	Registry replicas
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	

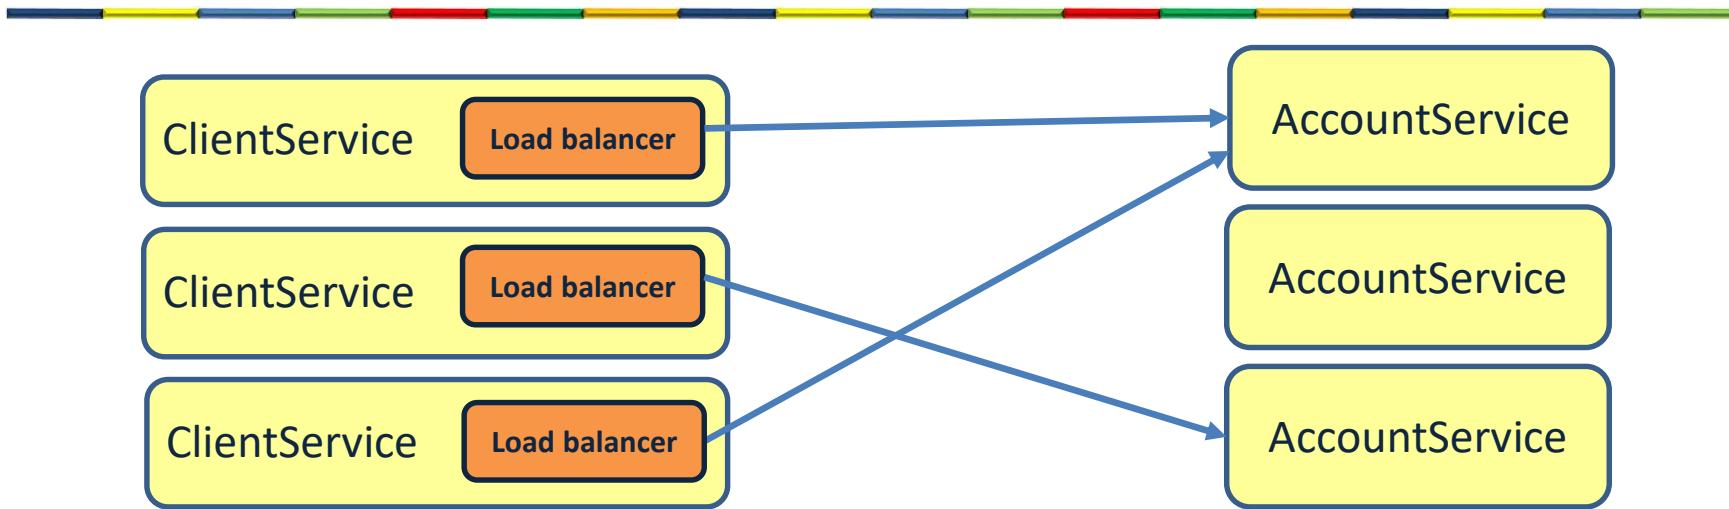
# **LOAD BALANCING: RIBBON**

# Server side load balancing



- Single point of failure
- If we add a new instance of AccountService, we need to reconfigure the load balancer
- Extra hop (performance)
- Every microservice needs its own load balancer
- Same load balance algorithm for every client
- Scaling limitation, load balance can handle only a certain number of requests

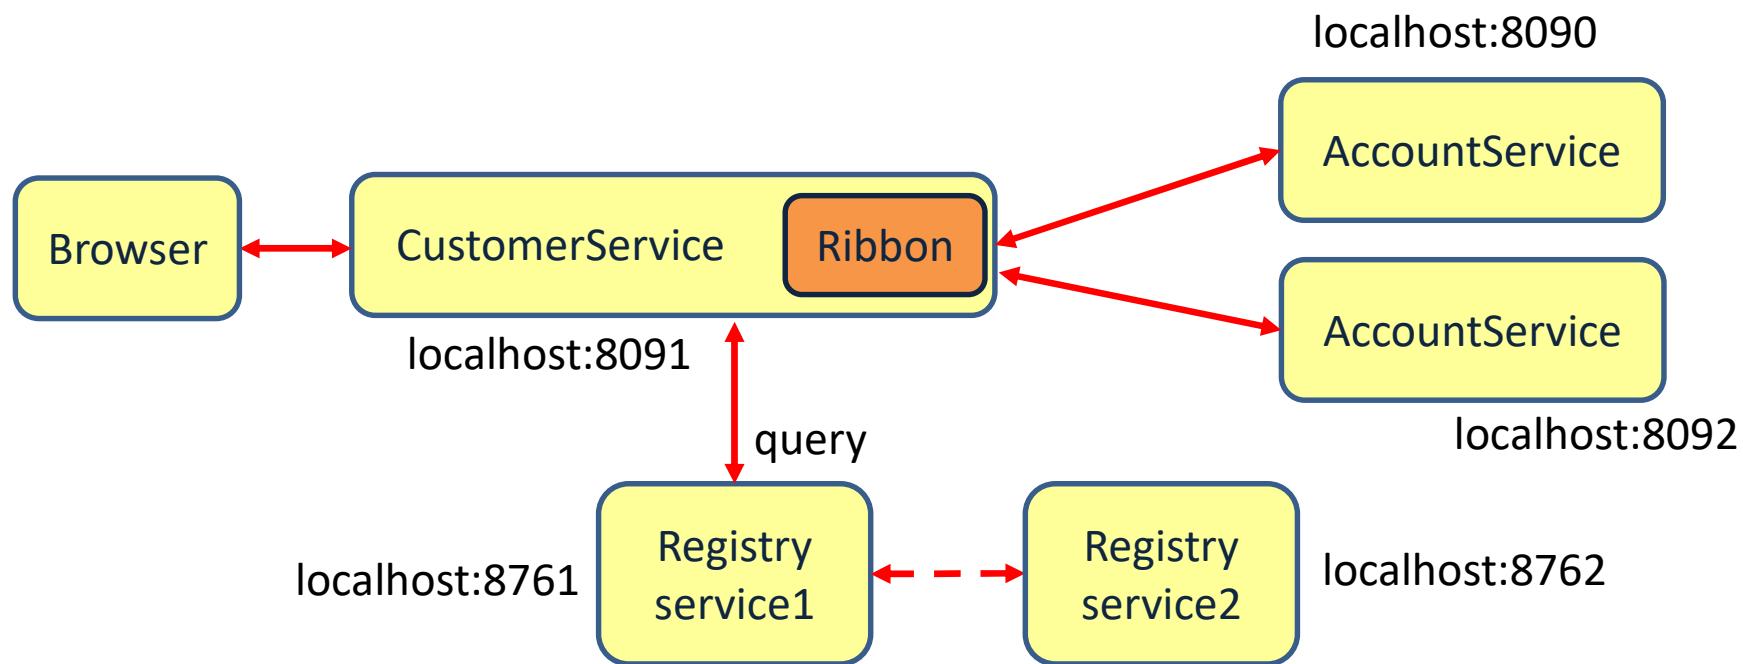
# Client side load balancing



- No single point of failure
- Simplifies service management
- Only one hop (performance)
- Auto discovery with registry based lookup (flexibility)
- Every client can use its own load balancing algorithm
- Unlimited scalable

# Load balancing using Ribbon

---



# AccountService

---

```
@RestController
public class AccountController {
    @GetMapping("/account/{customerid}")
    public Account getName(@PathVariable("customerid") String customerId) {
        return new Account("1234", "1000.00");
    }
}
```

server:  
port: 8090

eureka:  
client:  
serviceUrl:  
defaultZone: http://localhost:8761/eureka/,  
http://localhost:8762/eureka/

spring:  
application:  
name: AccountService

**application.yml**

# AccountService2

---

```
@RestController
public class AccountController {
    @GetMapping("/account/{customerid}")
    public Account getName(@PathVariable("customerid") String customerId) {
        return new Account("1234", "2000.00");
    }
}
```

server:  
port: 8092

eureka:  
client:  
serviceUrl:  
defaultZone: http://localhost:8761/eureka/,  
http://localhost:8762/eureka/

spring:  
application:  
name: AccountService

**application.yml**

# CustomerService: the controller

```
@RestController
public class CustomerController {
    @Autowired
    AccountFeignClient accountClient;

    @RequestMapping("/customer/{customerId}")
    public Account getName(@PathVariable("customerId") String customerId) {
        Account account = accountClient.getName(customerId);
        return account;
    }

    @FeignClient(name = "account-service")
    interface AccountFeignClient {
        @RequestMapping("/account/{customerId}")
        public Account getName(@PathVariable("customerId") String customerId);
    }
}
```

Feign automatically uses the  
Ribbon load balancer

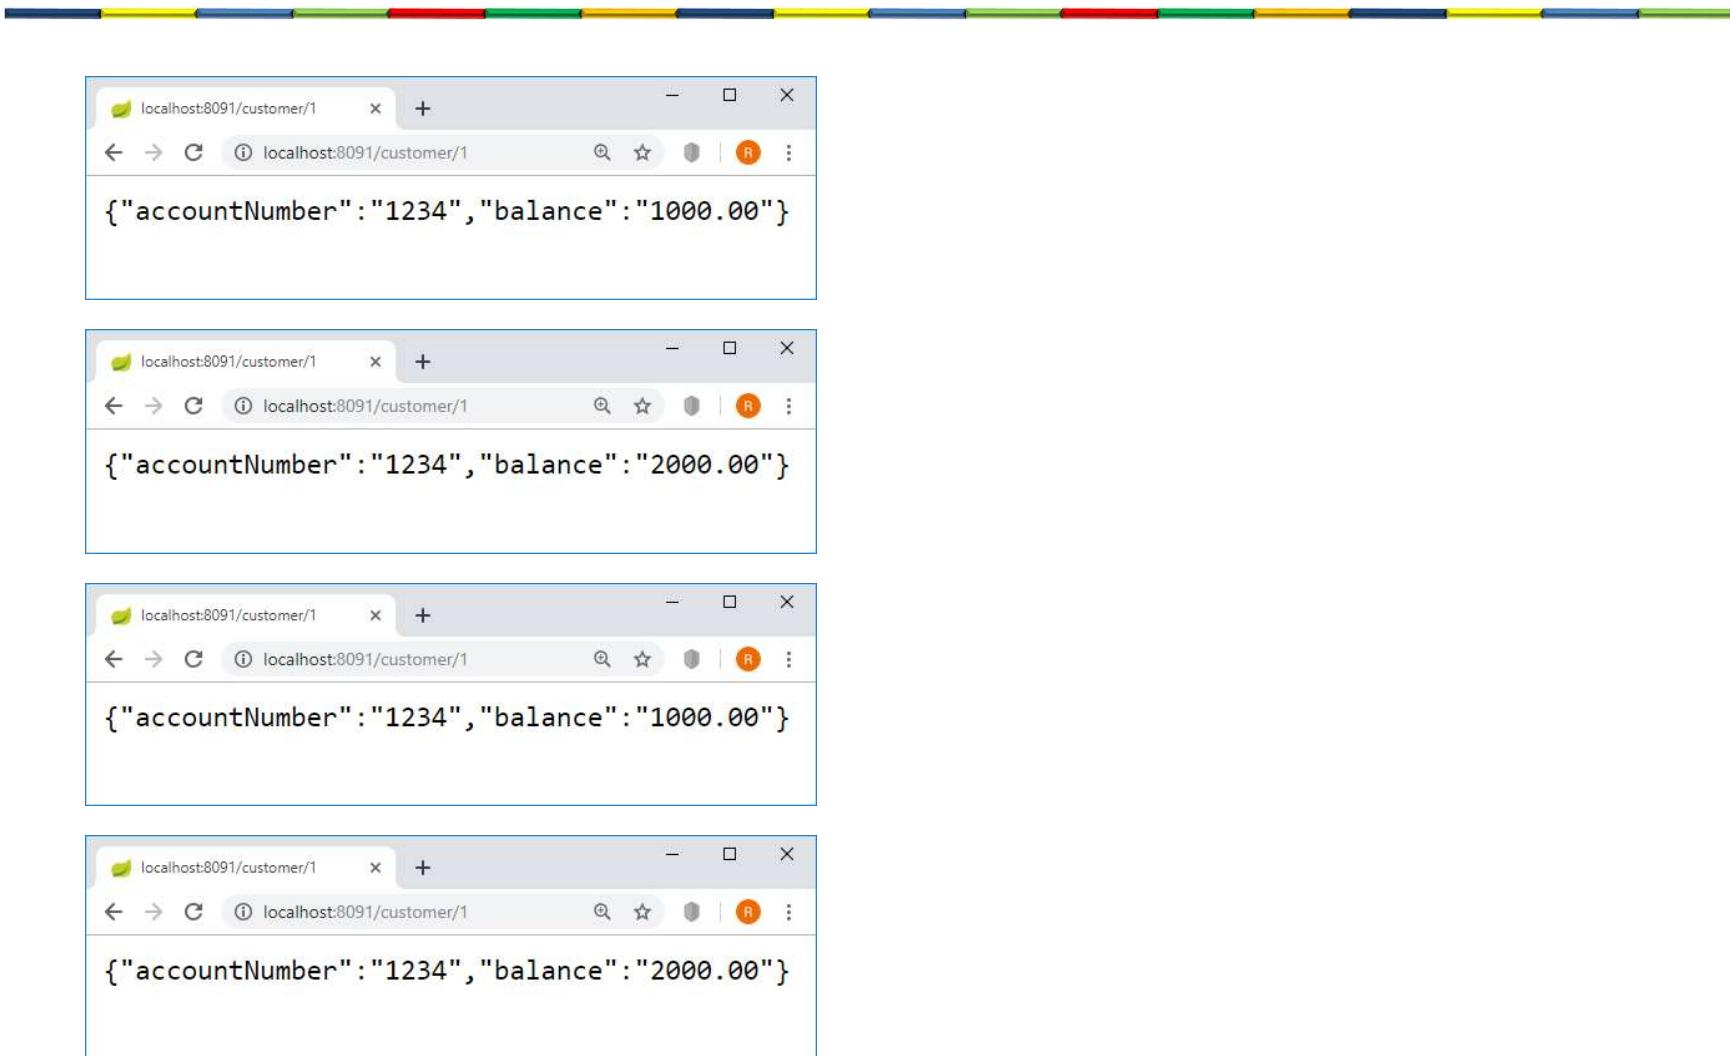
# Eureka



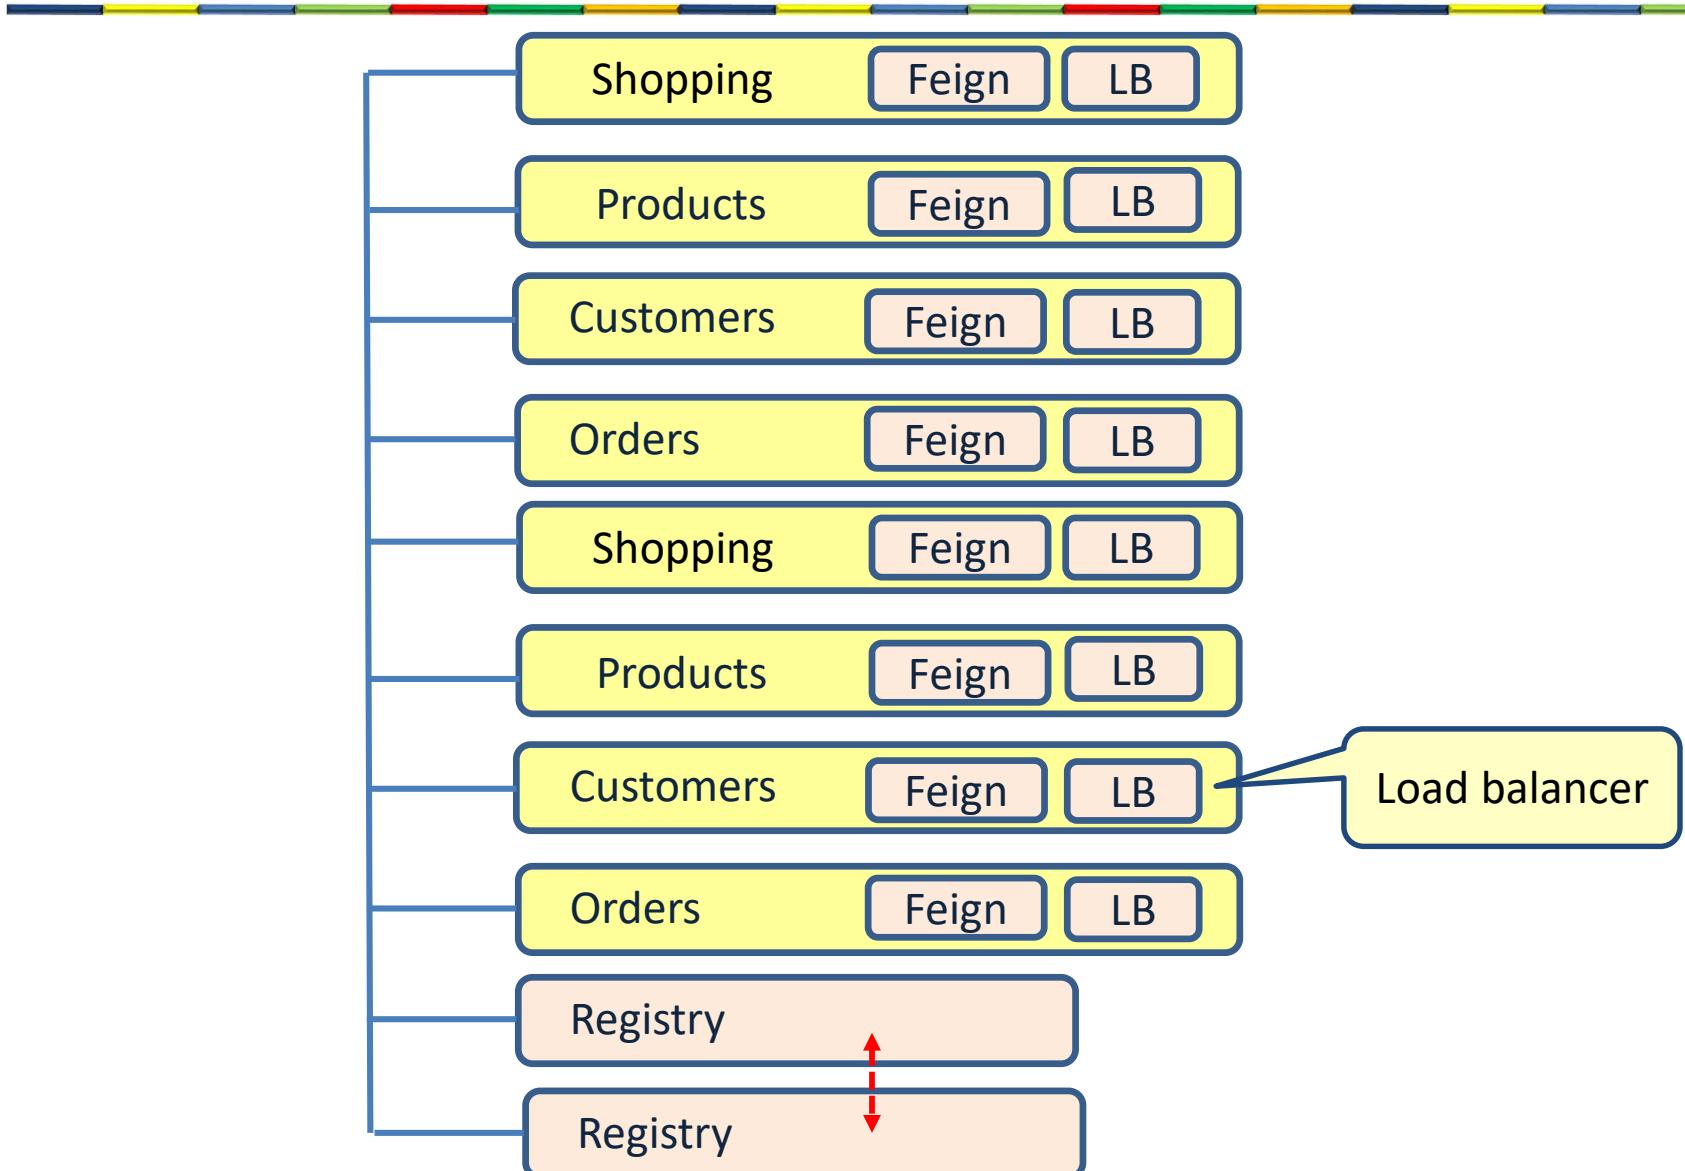
Instances currently registered with Eureka			
Application	AMIs	Availability	Status
		Zones	
ACCOUNTSERVICE	n/a (2)	(2)	UP (2) - DESKTOP-BVHRK6K.home:AccountService:8092 , DESKTOP-BVHRK6K.home:AccountService:8090
CUSTOMERSERVICE	n/a (1)	(1)	UP (1) - DESKTOP-BVHRK6K.home:CustomerService:8091
EUREKA-CLUSTER	n/a (2)	(2)	UP (2) - DESKTOP-BVHRK6K.home:8762 , DESKTOP-BVHRK6K.home:8761

2 instances of accountservice

# Round robin



# Implementing microservices

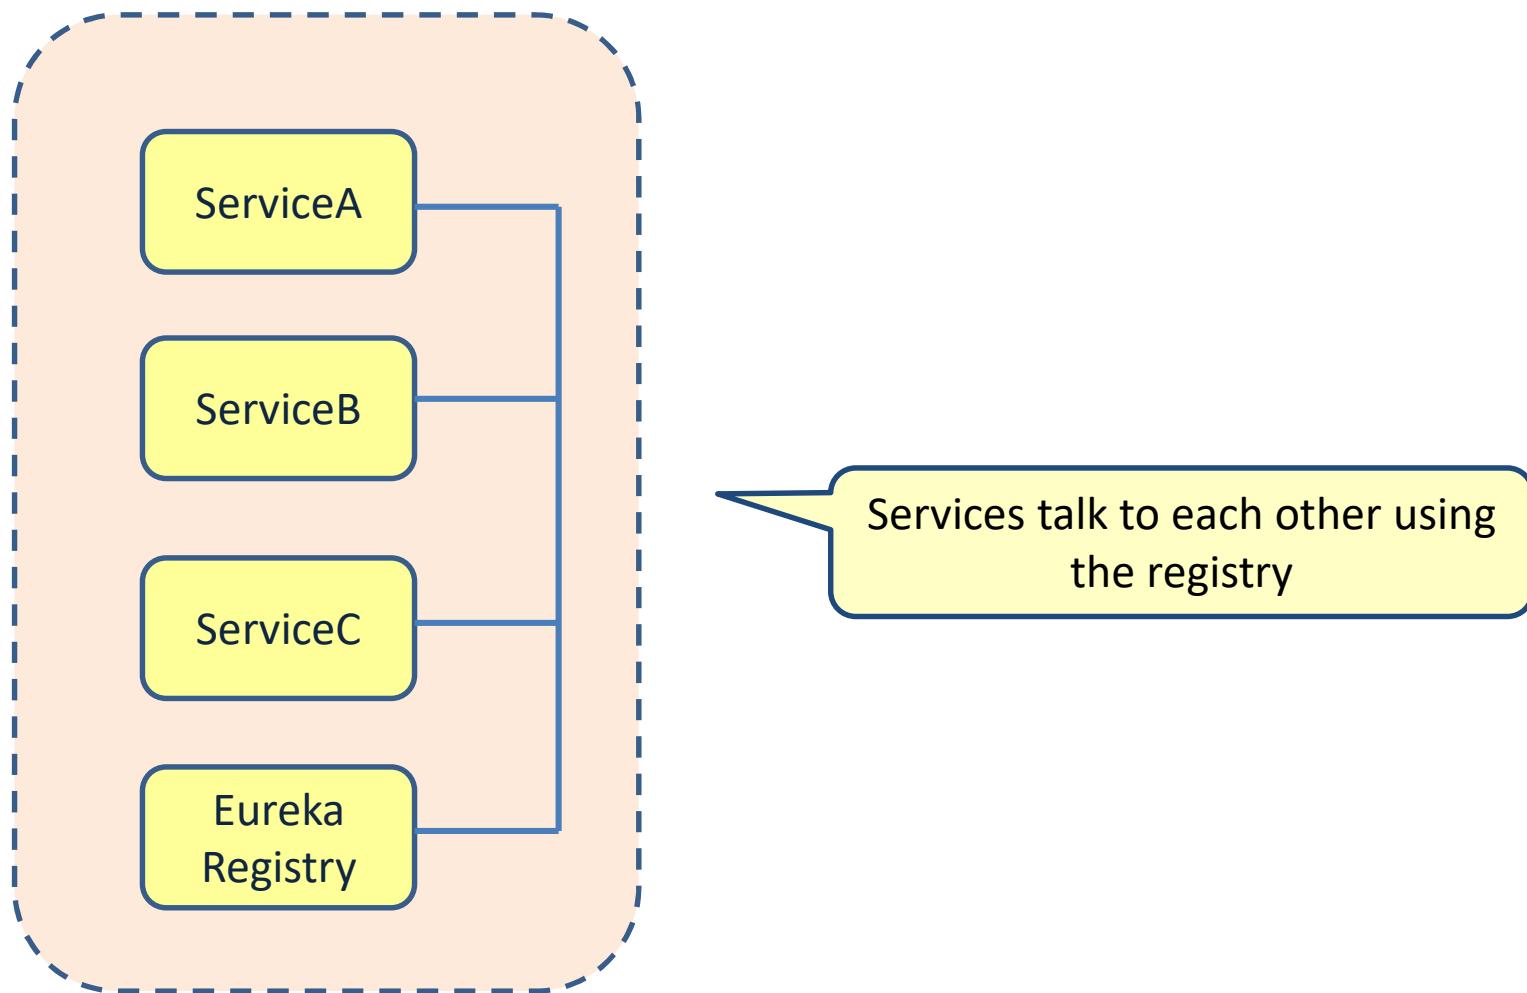


# Challenges of a microservice architecture

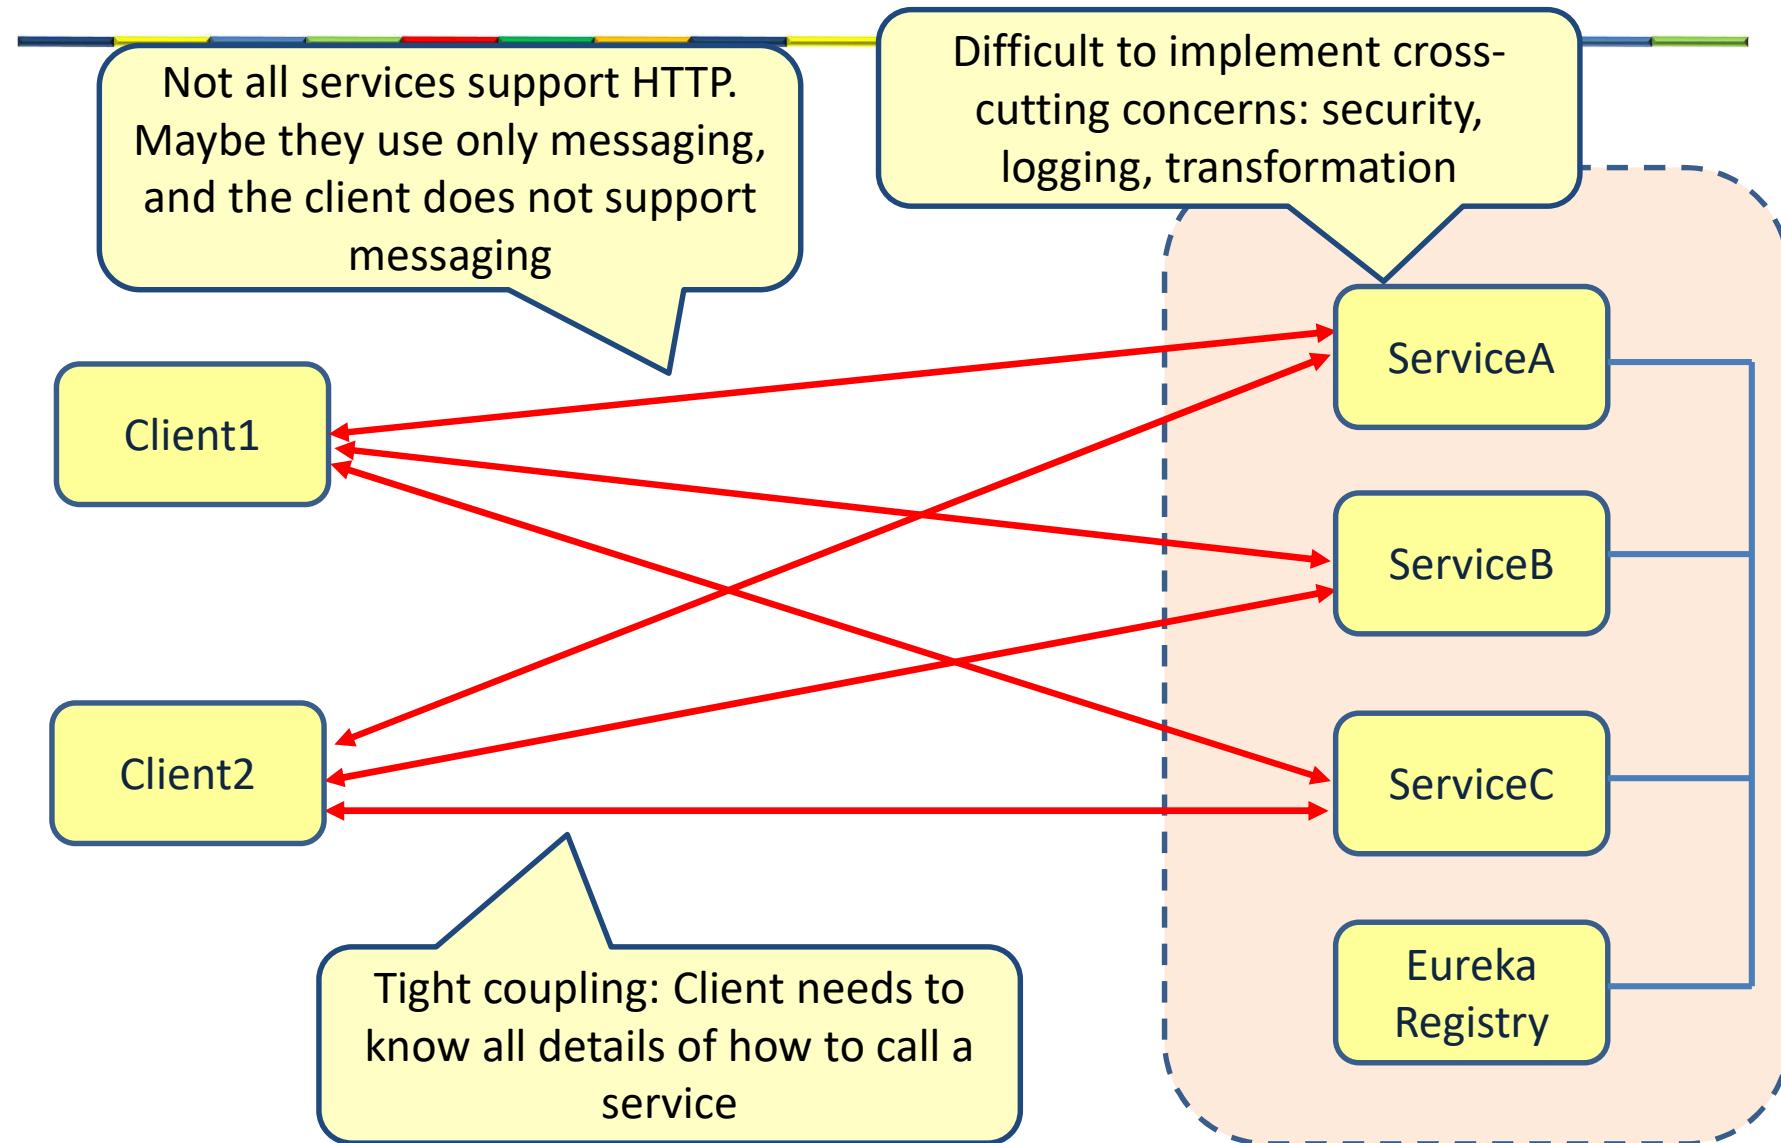
Challenge	Solution
Complex communication	Feign Registry
Performance	
Resilience	Registry replicas Load balancing between multiple service instances
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	

# **API GATEWAY**

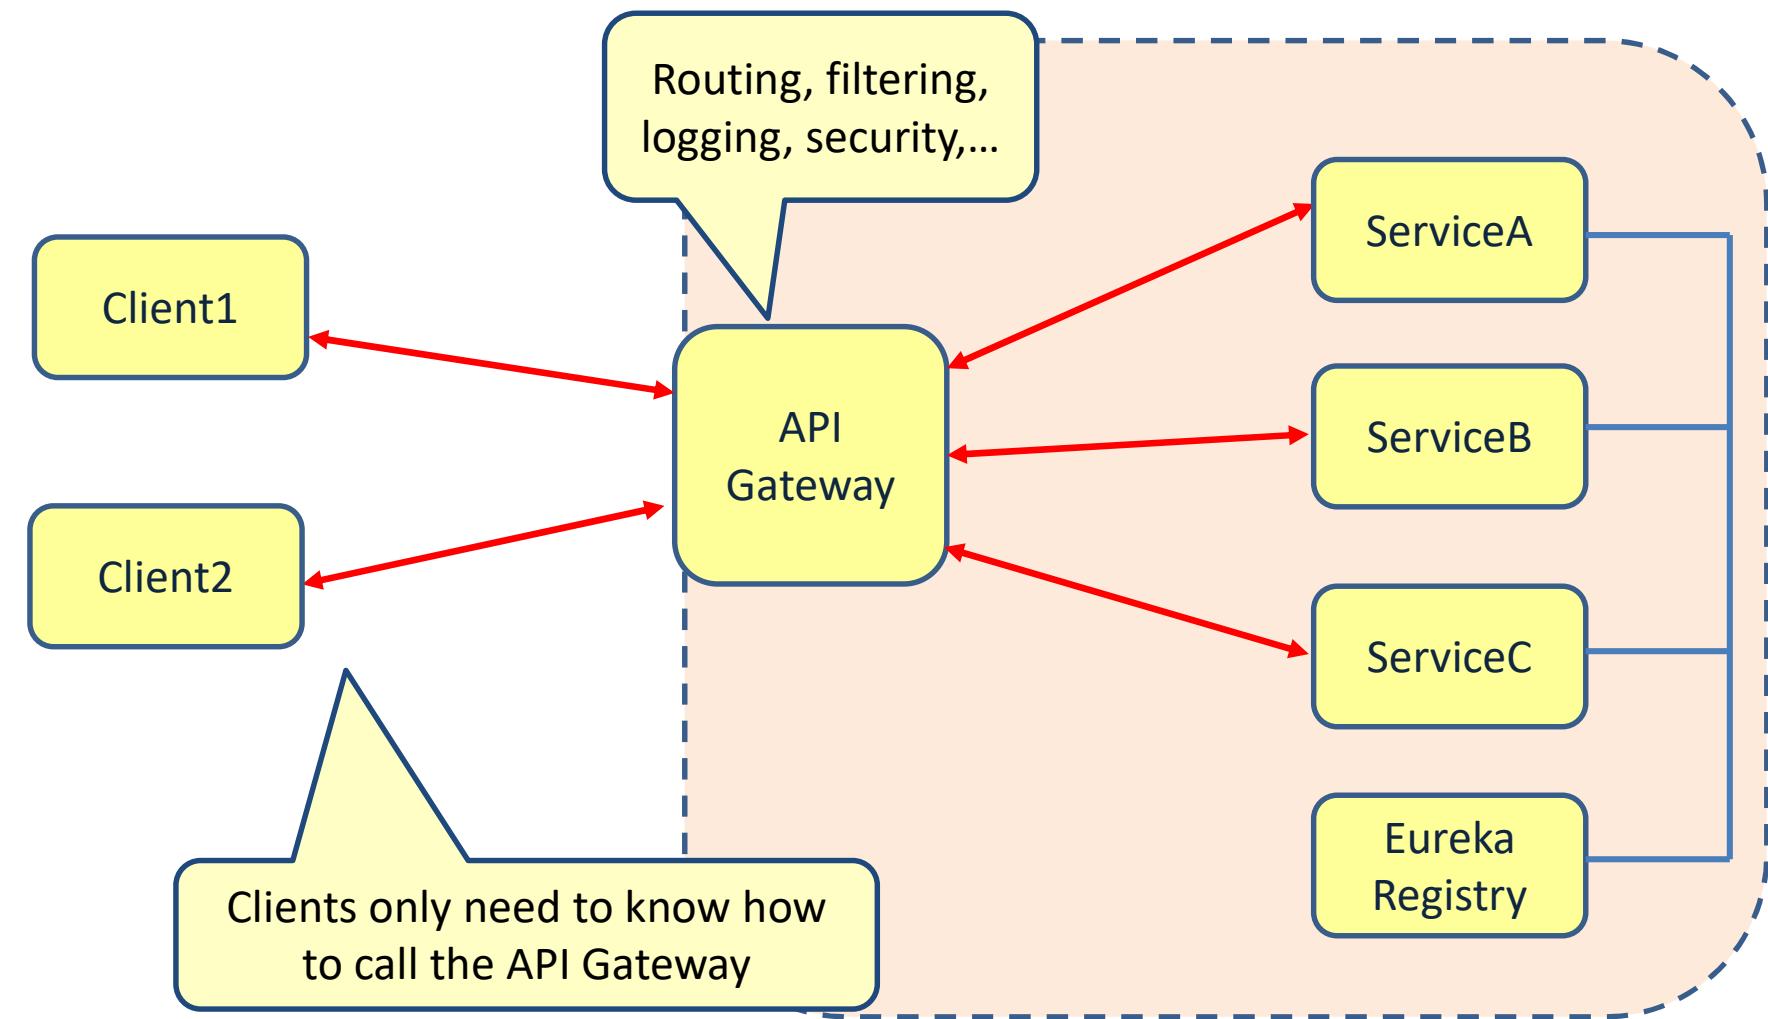
# Microservice architecture



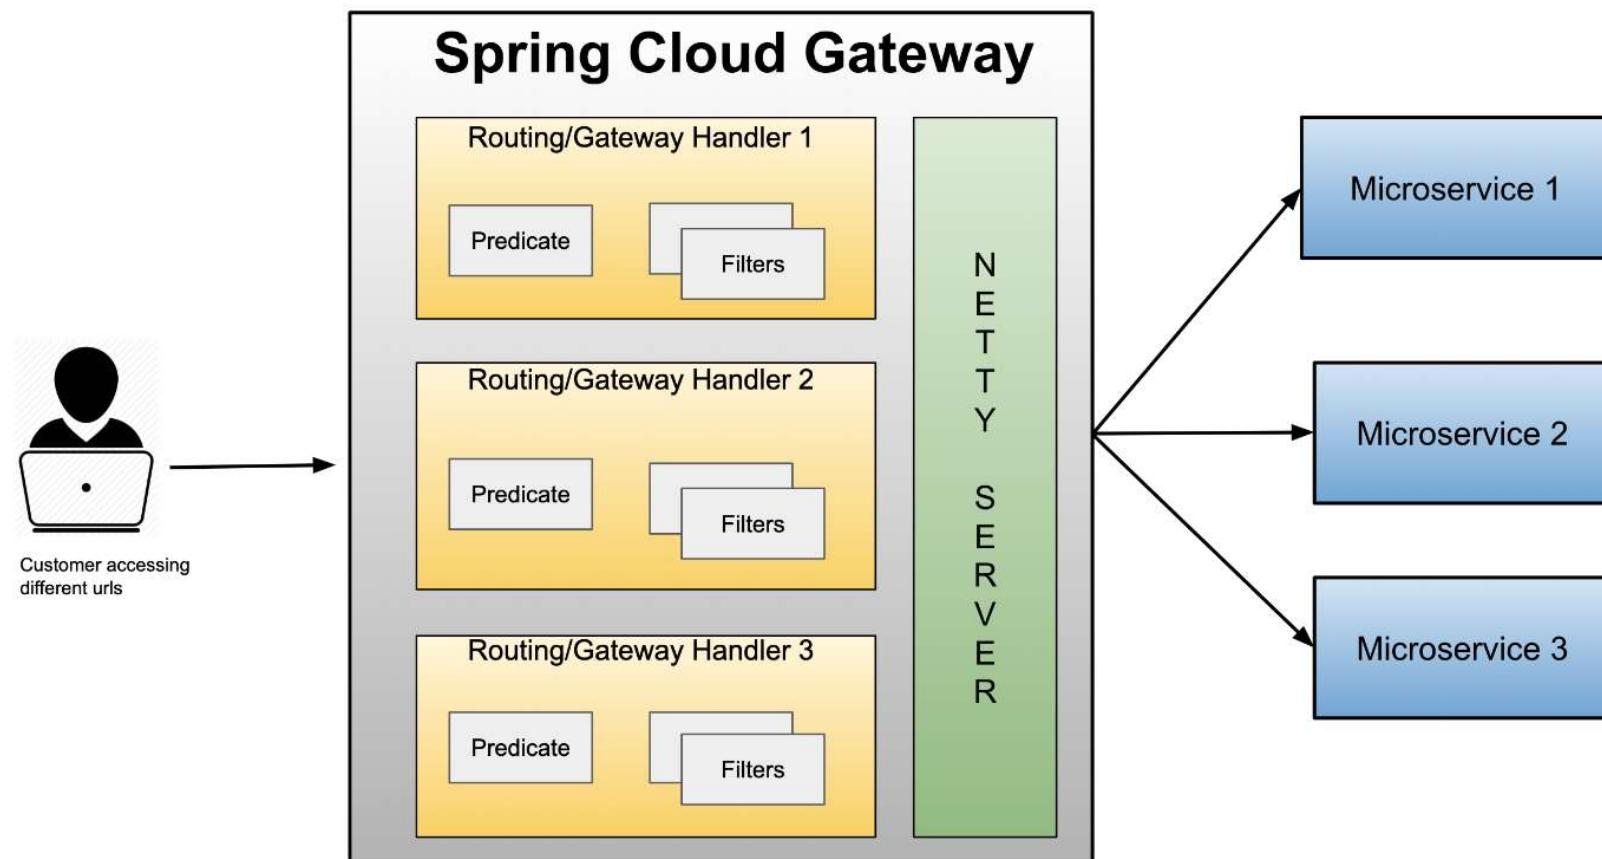
# Adding clients



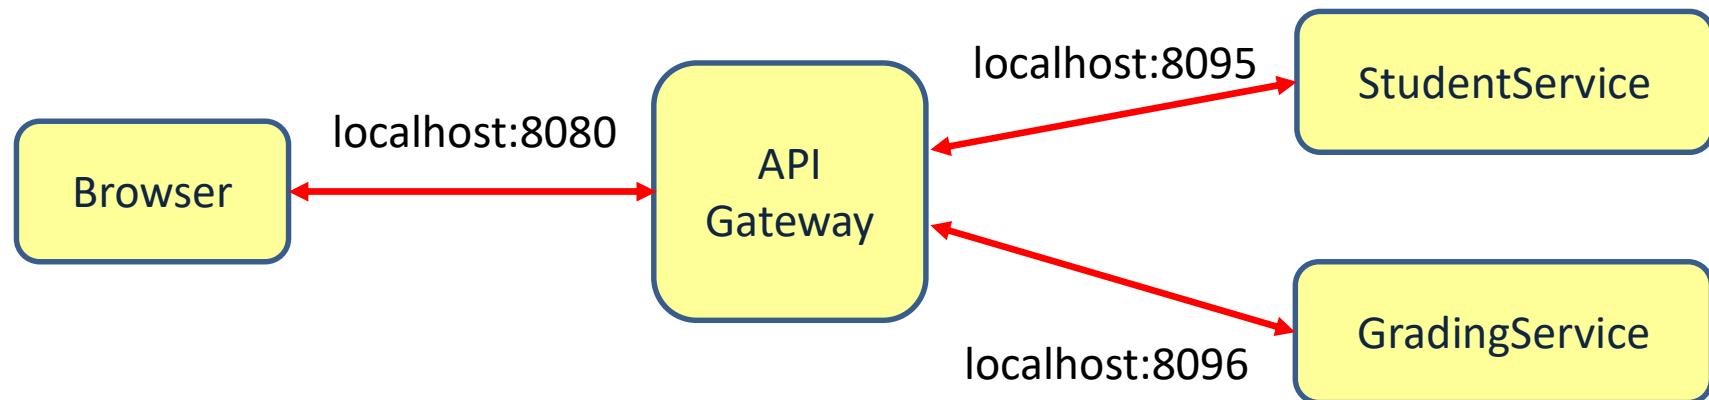
# Api Gateway



# Spring cloud gateway



# Api Gateway example



# StudentService

```
@SpringBootApplication  
@EnableDiscoveryClient  
public class StudentServiceApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(StudentServiceApplication.class, args);  
    }  
}
```

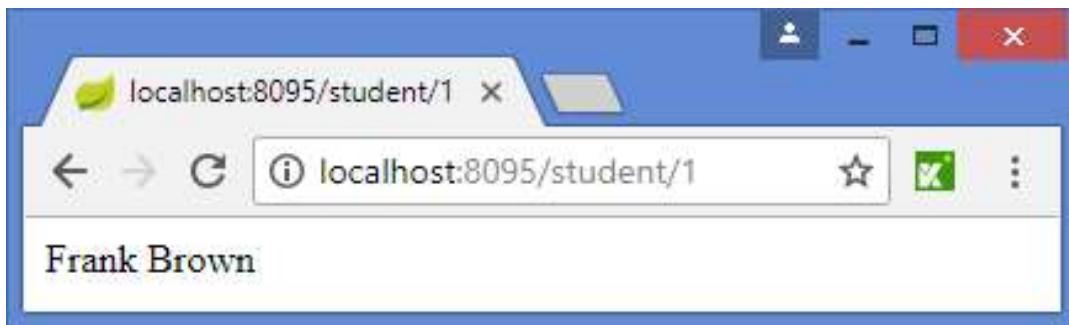
application.yml

```
server:  
  port: 8095  
  
spring:  
  application:  
    name: StudentService  
  
eureka:  
  client:  
    serviceUrl:  
      defaultZone: http://localhost:8761/eureka/
```

# StudentService: the controller

---

```
@RestController
public class StudentController {
    @GetMapping("/student/{studentid}")
    public String getName(@PathVariable("studentid") String studentid) {
        return "Frank Brown";
    }
}
```



# GradingService

```
@SpringBootApplication
@EnableDiscoveryClient
public class GradingServiceApplication {

    public static void main(String[] args) {
        SpringApplication.run(StudentServiceApplication.class, args);
    }
}
```

application.yml

```
server:
  port: 8096

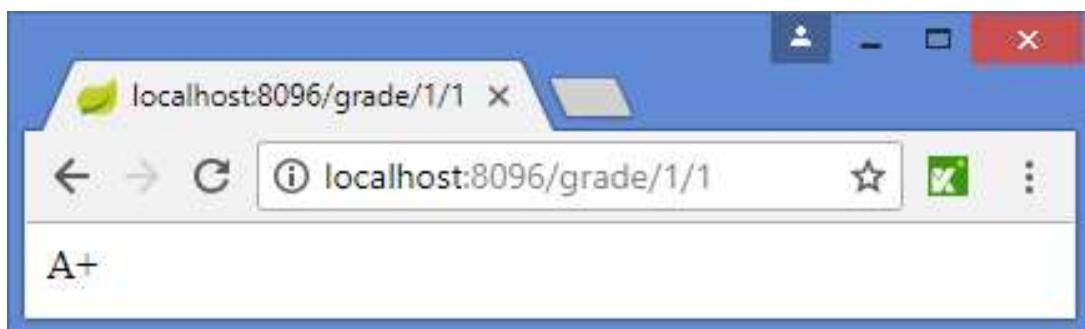
spring:
  application:
    name: GradingService

eureka:
  client:
    serviceUrl:
      defaultZone: http://localhost:8761/eureka/
```

# GradingService: the controller

---

```
@RestController
public class GradingController {
    @GetMapping("/grade/{studentid}/{courseid}")
    public String getGrade(@PathVariable("studentid") String studentid,
                          @PathVariable("courseid") String courseid) {
        return "A+";
    }
}
```



# Spring Cloud Gateway

```
@SpringBootApplication
public class CloudgatewayApplication {

    public static void main(String[] args) {
        SpringApplication.run(CloudgatewayApplication.class, args);
    }
}
```

## POM.xml

```
<dependency>
    <groupId>org.springframework.cloud</groupId>
    <artifactId>spring-cloud-starter-gateway</artifactId>
</dependency>
```

# Spring Cloud Gateway

---

## application.yml

```
spring:
  application:
    name: api-gateway
  cloud:
    gateway:
      routes:
        - id: studentModule
          uri: http://localhost:8095/
          predicates:
            - Path=/student/**
        - id: gradingModule
          uri: http://localhost:8096/
          predicates:
            - Path=/grade/**
server:
  port: 8080
```

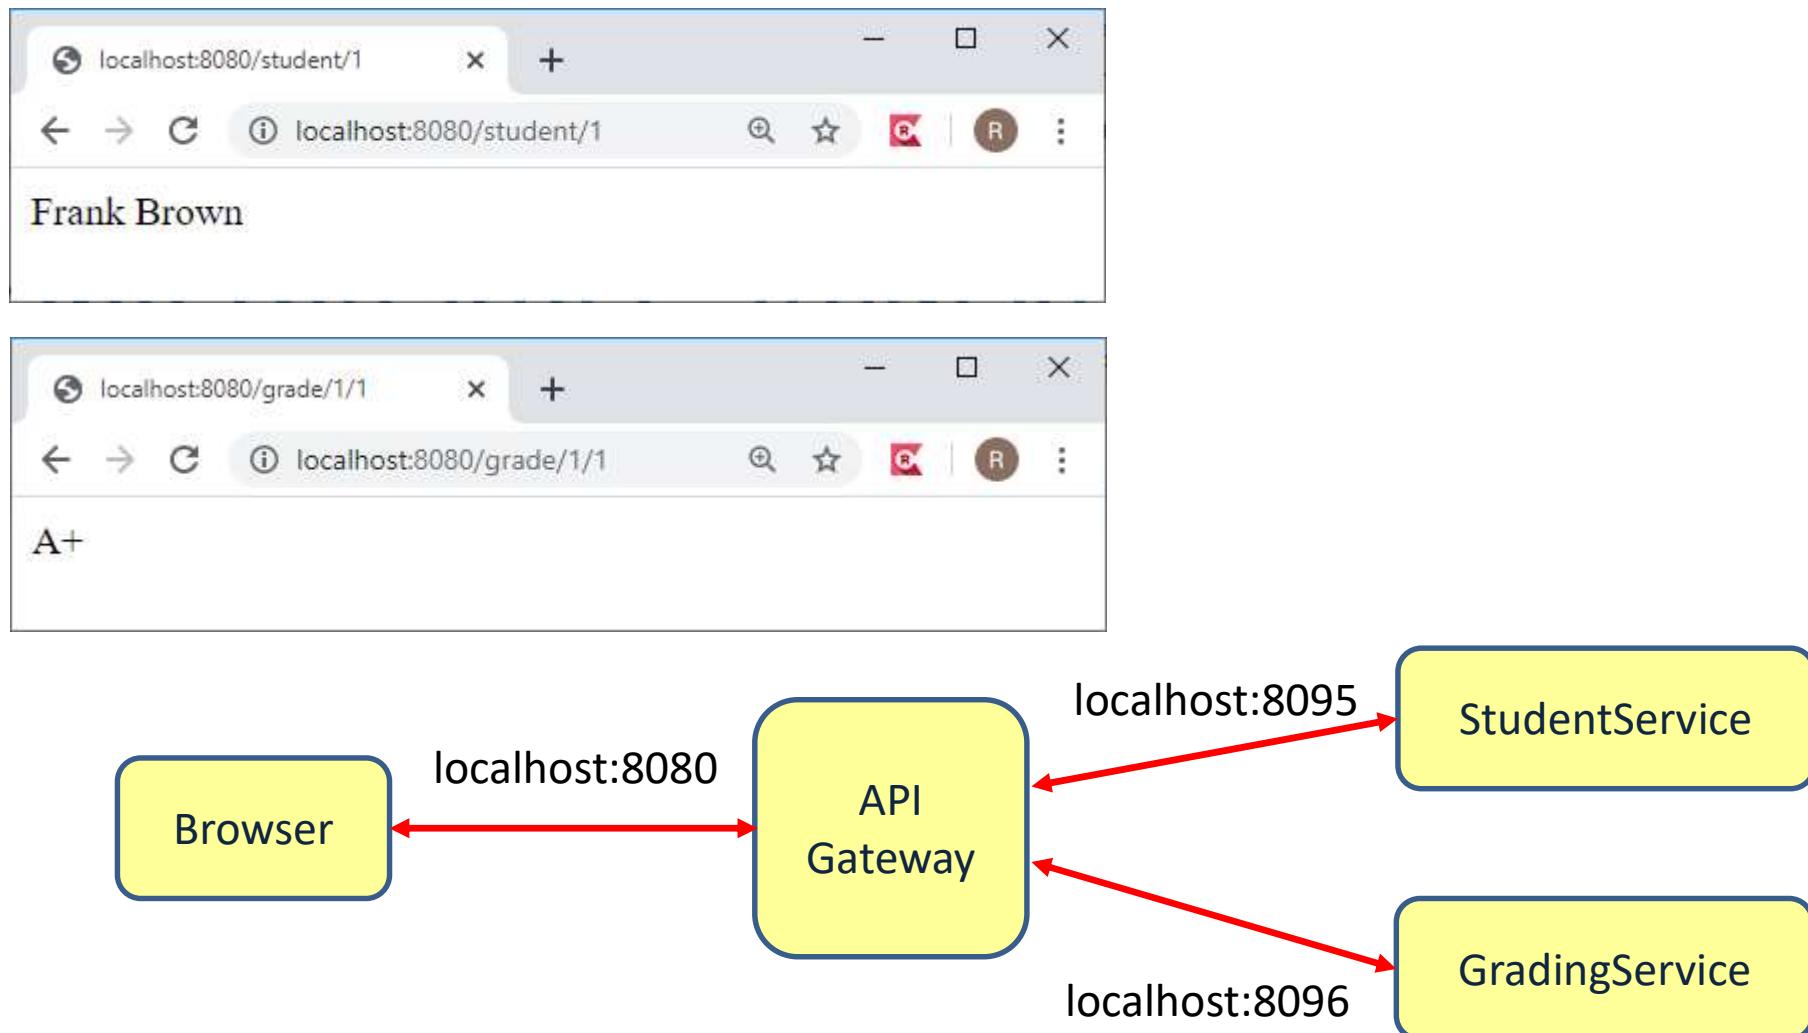
Id should be unique for every route

Route /student to localhost:8095

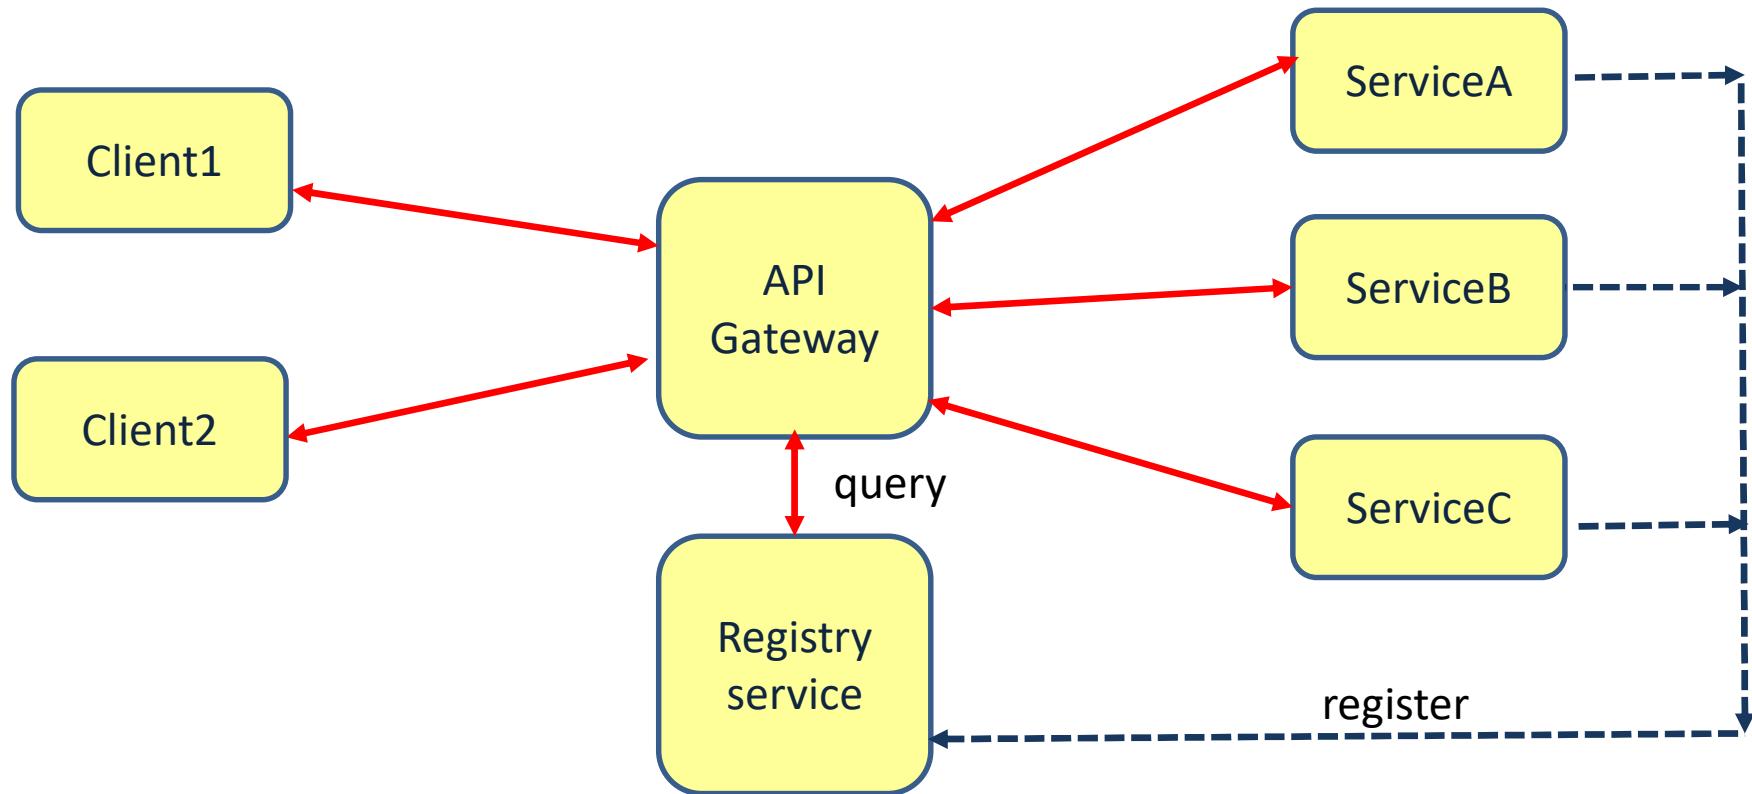
Route /grades to localhost:8096

A route is matched if predicate is true

# Using the API Gateway



# Api Gateway and registry service



# Spring cloud gateway with the registry

```
spring:  
  application:  
    name: api-gateway  
  cloud:  
    gateway:  
      routes:  
        - id: studentModule  
          uri: lb://StudentService  
          predicates:  
            - Path=/student/**  
        - id: gradingModule  
          uri: lb://GradingService  
          predicates:  
            - Path=/grade/**  
  
  server:  
    port: 8080  
  
eureka:  
  client:  
    serviceUrl:  
      defaultZone: http://localhost:8761/eureka/
```

application.yml

Route /student to the service with the name StudentService using the load balancer

Route /grade to the service with the name GradingService using the load balancer

URL to Eureka

# Java based config

```
@Configuration
public class BeanConfig {

    @Bean
    public RouteLocator gatewayRoutes(RouteLocatorBuilder builder) {
        return builder.routes()
            .route(r -> r.path("/student/**"))
                .uri("lb://StudentService")
                .id("studentModule"))

            .route(r -> r.path("/grade/**"))
                .uri("lb://GradingService")
                .id("gradesModule"))
        .build();
    }
}
```

# Build-in predicates

---

Name	Description	Example
After Route	It takes a date-time parameter and matches requests that happen after it	After=2017-11-20T...
Before Route	It takes a date-time parameter and matches requests that happen before it	Before=2017-11-20T...
Between Route	It takes two date-time parameters and matches requests that happen between those dates	Between=2017-11-20T..., 2017-11-21T...
Cookie Route	It takes a cookie name and regular expression parameters, finds the cookie in the HTTP request's header, and matches its value with the provided expression	Cookie=SessionID, abc.
Header Route	It takes the header name and regular expression parameters, finds a specific header in the HTTP request's header, and matches its value with the provided expression	Header=X-Request-Id, \d+
Host Route	It takes a hostname ANT style pattern with the . separator as a parameter and matches it with the Host header	Host=*.example.org
Method Route	It takes an HTTP method to match as a parameter	Method=GET
Path Route	It takes a pattern of request context path as a parameter	Path=/account/{id}
Query Route	It takes two parameters—a required param and an optional regexp and matches them with query parameters	Query=accountId, 1.
RemoteAddr Route	It takes a list of IP addresses in CIDR notation, like 192.168.0.1/16, and matches it with the remote address of a request	RemoteAddr=192.168.0.1/16

# Filters

---

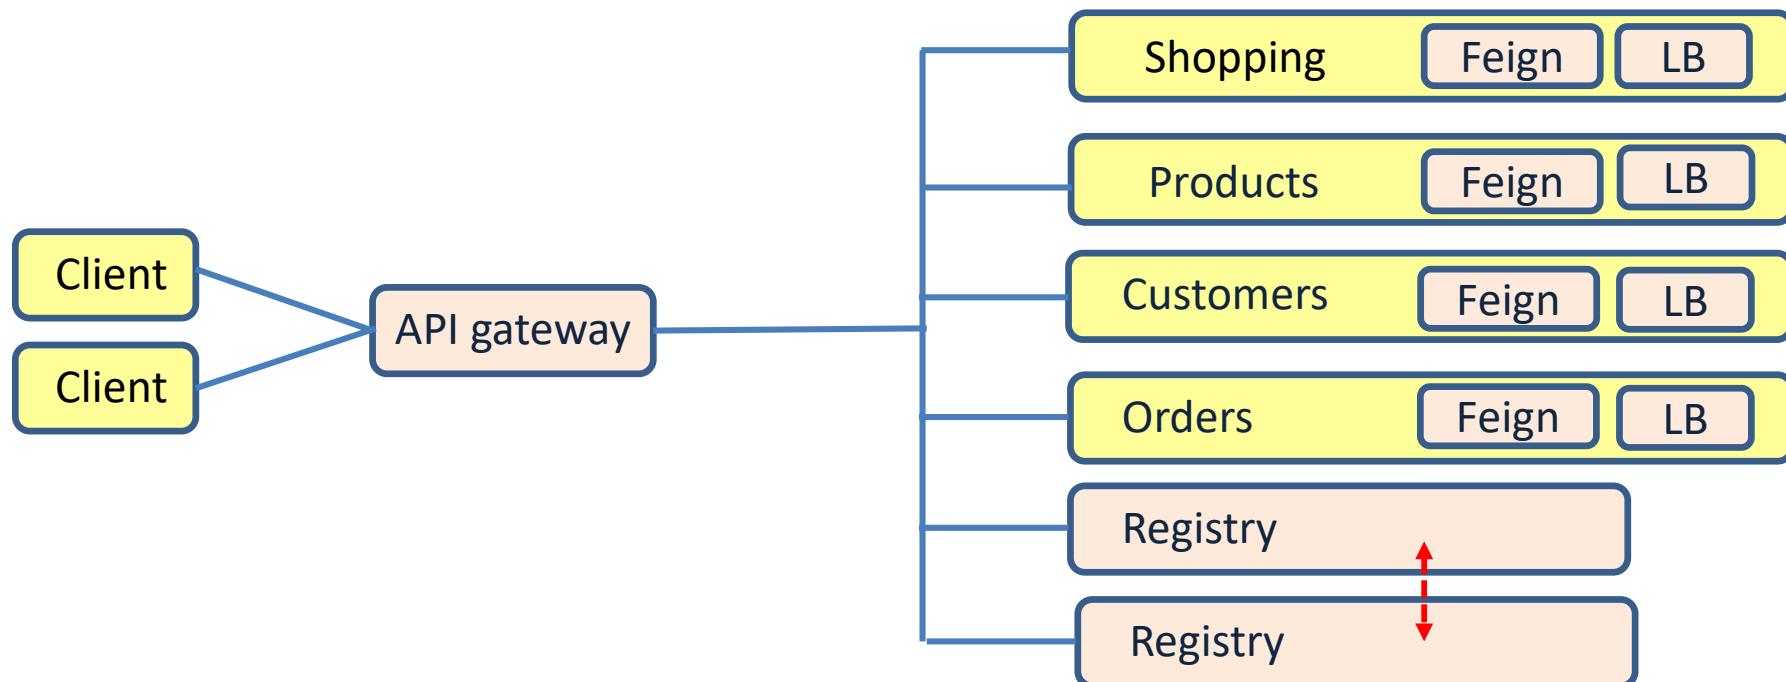
- Pre and post filters
  - Build-in filters
  - Custom filters
  - Global filters

```
cloud:  
  gateway:  
    routes:  
      - id: studentModule  
        uri: lb://StudentService  
        filters:  
          - AddRequestHeader=X-myHeader, Hello  
          - AddRequestParameter=name, John  
          - AddResponseHeader=X-someHeader, Hello World  
        predicates:  
          - Path=/student/**
```

Build-in filters

# Implementing microservices

---



# Challenges of a microservice architecture

---

Challenge	Solution
Complex communication	Feign Registry API gateway
Performance	
Resilience	Registry replicas Load balancing between multiple service instances
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	

Lesson 9

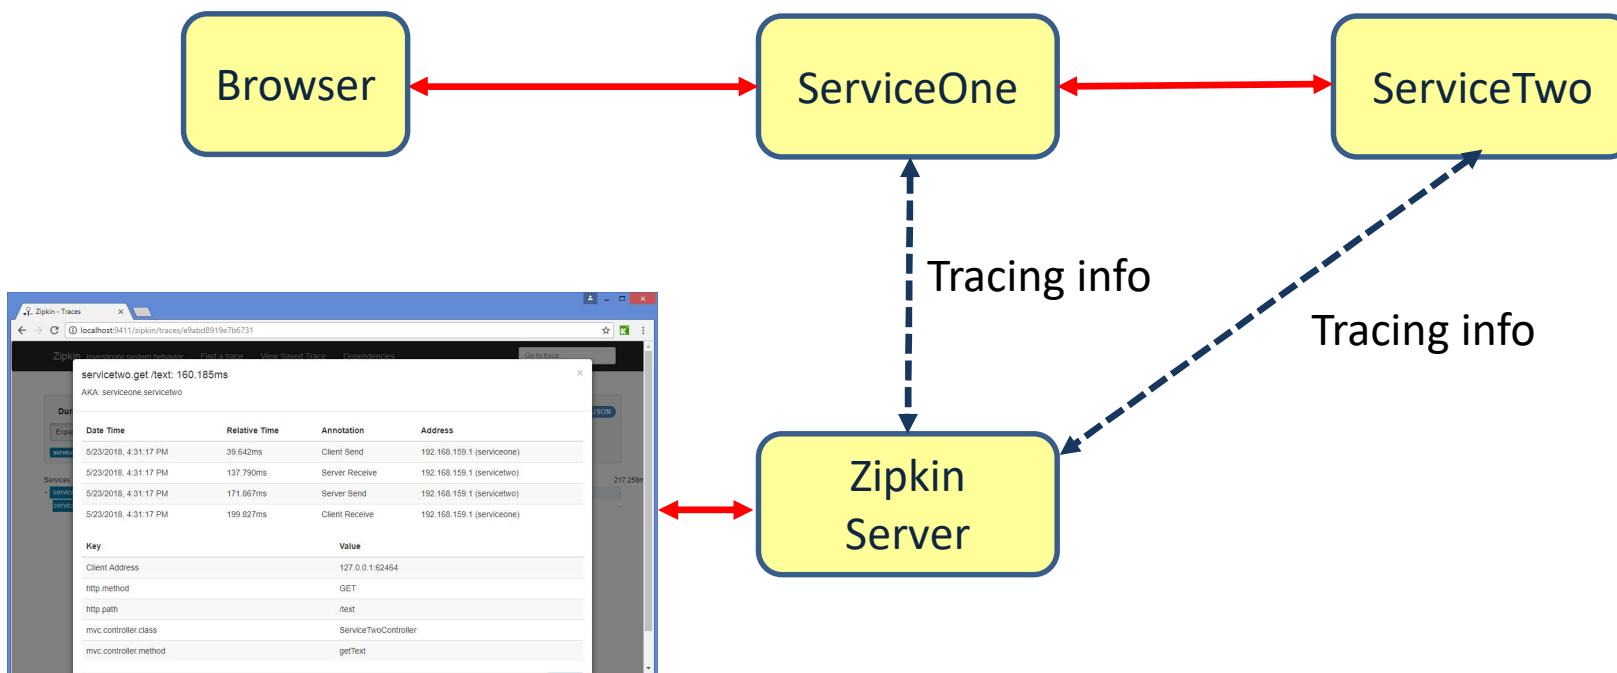
# **MICROSERVICES**



# DISTRIBUTED TRACING: ZIPKIN

# Distributed Tracing

- One central place where one can see the end-to-end tracing of all communication between services



# Spring cloud Sleuth

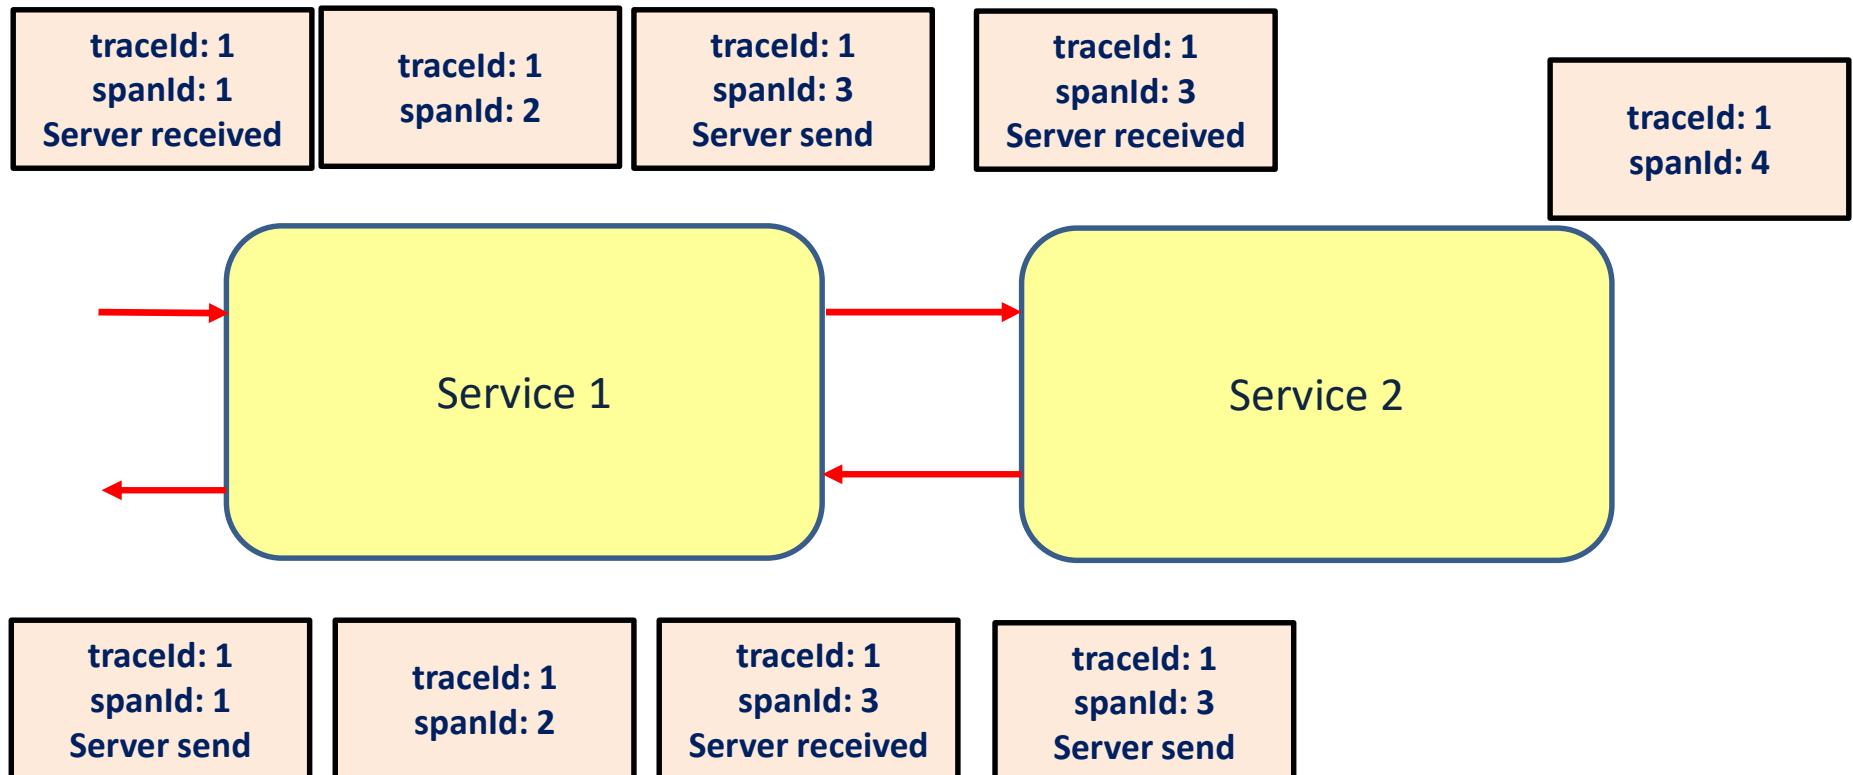
---

- Adds unique id's to a request so we can trace the request
  - Span id: id for an individual operation
  - Trace id: id for a set of spans
- Also embeds these unique id's to log messages

# Spring cloud Sleuth

---

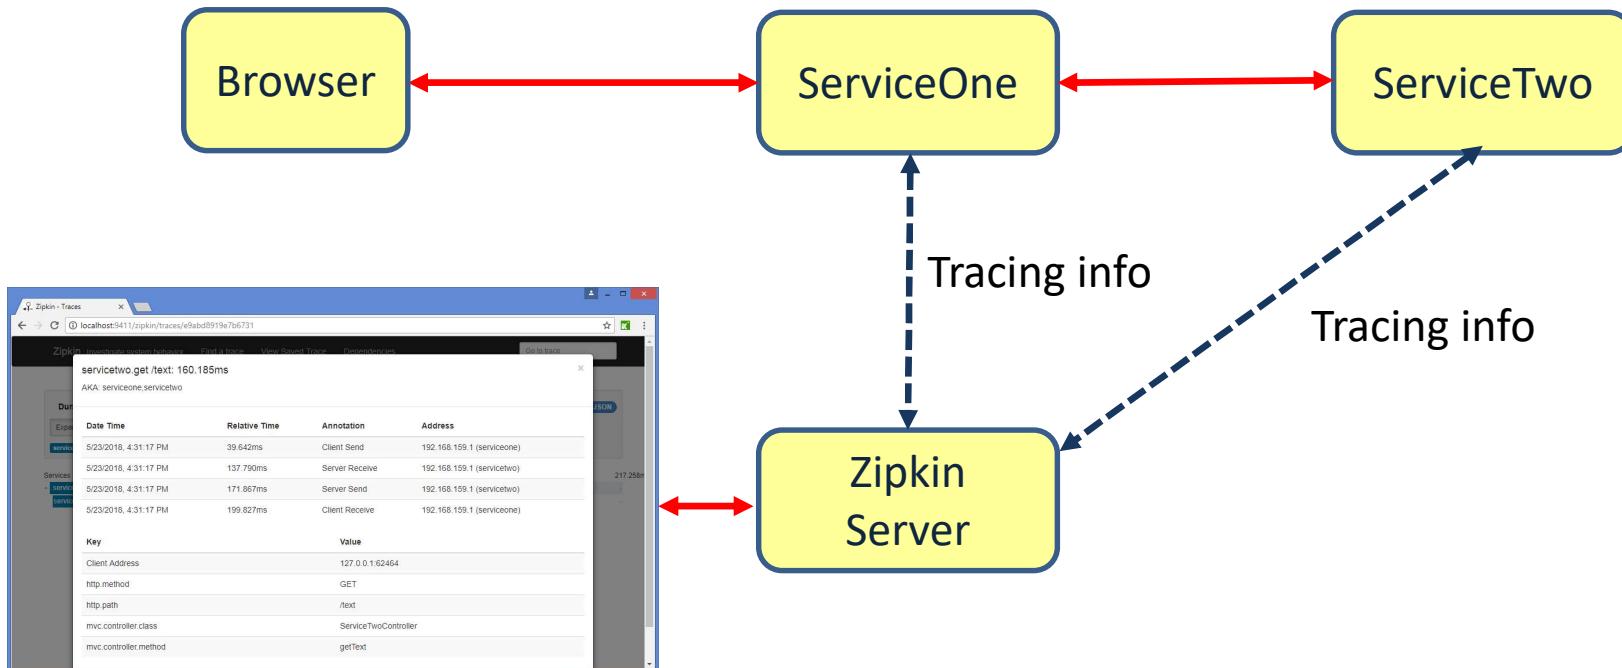
- Span: an individual operation
- Trace: a set of spans



# Zipkin

---

- Centralized tracing server
  - Collects tracing information
- Zipkin console shows the data



# Service1

```
@SpringBootApplication
public class Service1Application {
    public static void main(String[] args) {
        SpringApplication.run(Service1Application.class, args);
    }
}
```

```
@RestController
public class ServiceOneController {

    @Autowired
    RestTemplate restTemplate;

    @RequestMapping("/text")
    public String getText() {
        String service2Text = restTemplate.getForObject("http://localhost:9091/text",
                                                       String.class);
        return "Hello " + service2Text;
    }

    @Bean
    RestTemplate getRestTemplate() {
        return new RestTemplate();
    }
}
```

# Service1

---

## application.yml

```
server:
  port: 9090

spring:
  application:
    name: ServiceOne
  zipkin:
    base-url: http://localhost:9411/

  sleuth:
    sampler:
      probability: 1 #100% (default = 10%)
```

# Service1

---

## pom.xml

```
<dependency>
    <groupId>org.springframework.boot</groupId>
    <artifactId>spring-boot-starter-web</artifactId>
</dependency>
<dependency>
    <groupId>org.springframework.cloud</groupId>
    <artifactId>spring-cloud-starter-sleuth</artifactId>
</dependency>
<dependency>
    <groupId>org.springframework.cloud</groupId>
    <artifactId>spring-cloud-starter-zipkin</artifactId>
    <version>2.2.3.RELEASE</version>
</dependency>
```

# Service2

```
@SpringBootApplication
public class Service2Application {
    public static void main(String[] args) {
        SpringApplication.run(Service2Application.class, args);
    }
}
```

```
@RestController
public class ServiceTwoController {

    @RequestMapping("/text")
    public String getText() {
        return "World";
    }
}
```

# Service2

---

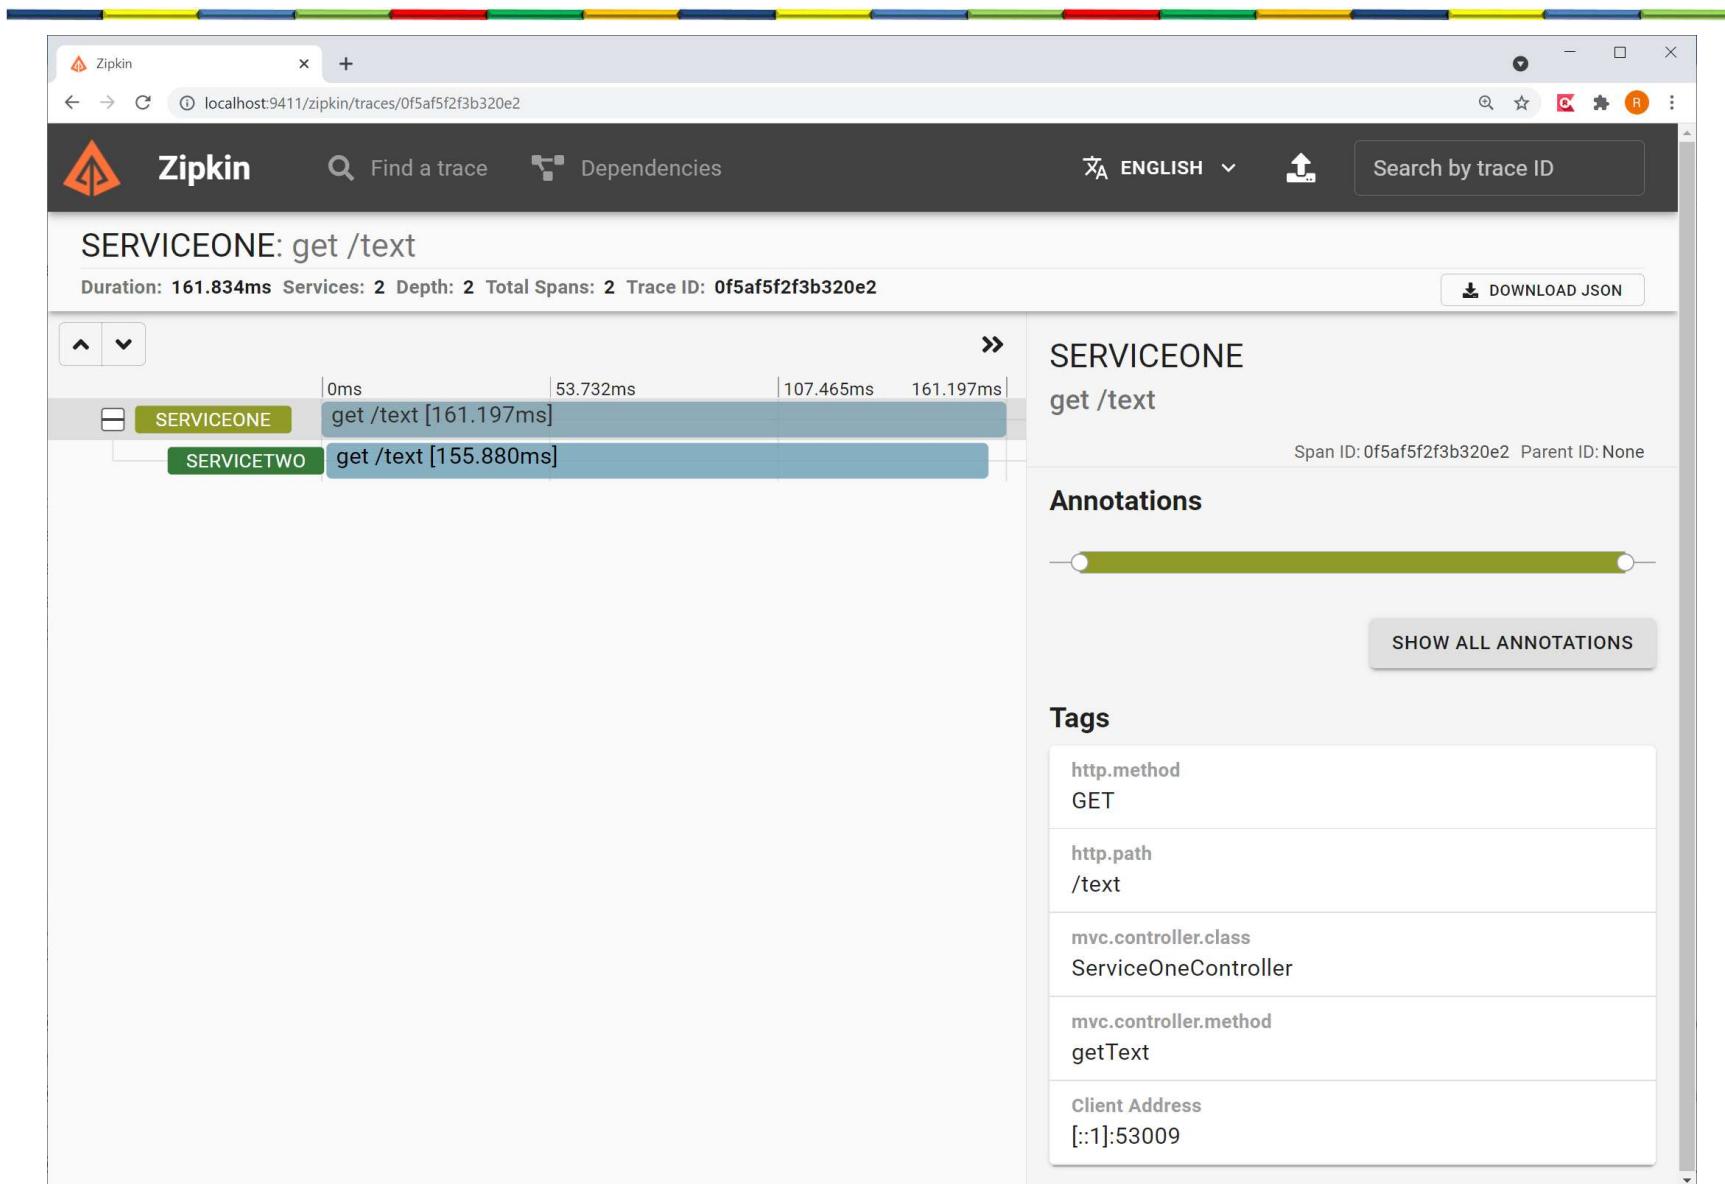
## application.yml

```
server:
  port: 9091

spring:
  application:
    name: ServiceTwo
  zipkin:
    base-url: http://localhost:9411/

sleuth:
  sampler:
    probability: 1 #100% (default = 10%)
```

# Zipkin console

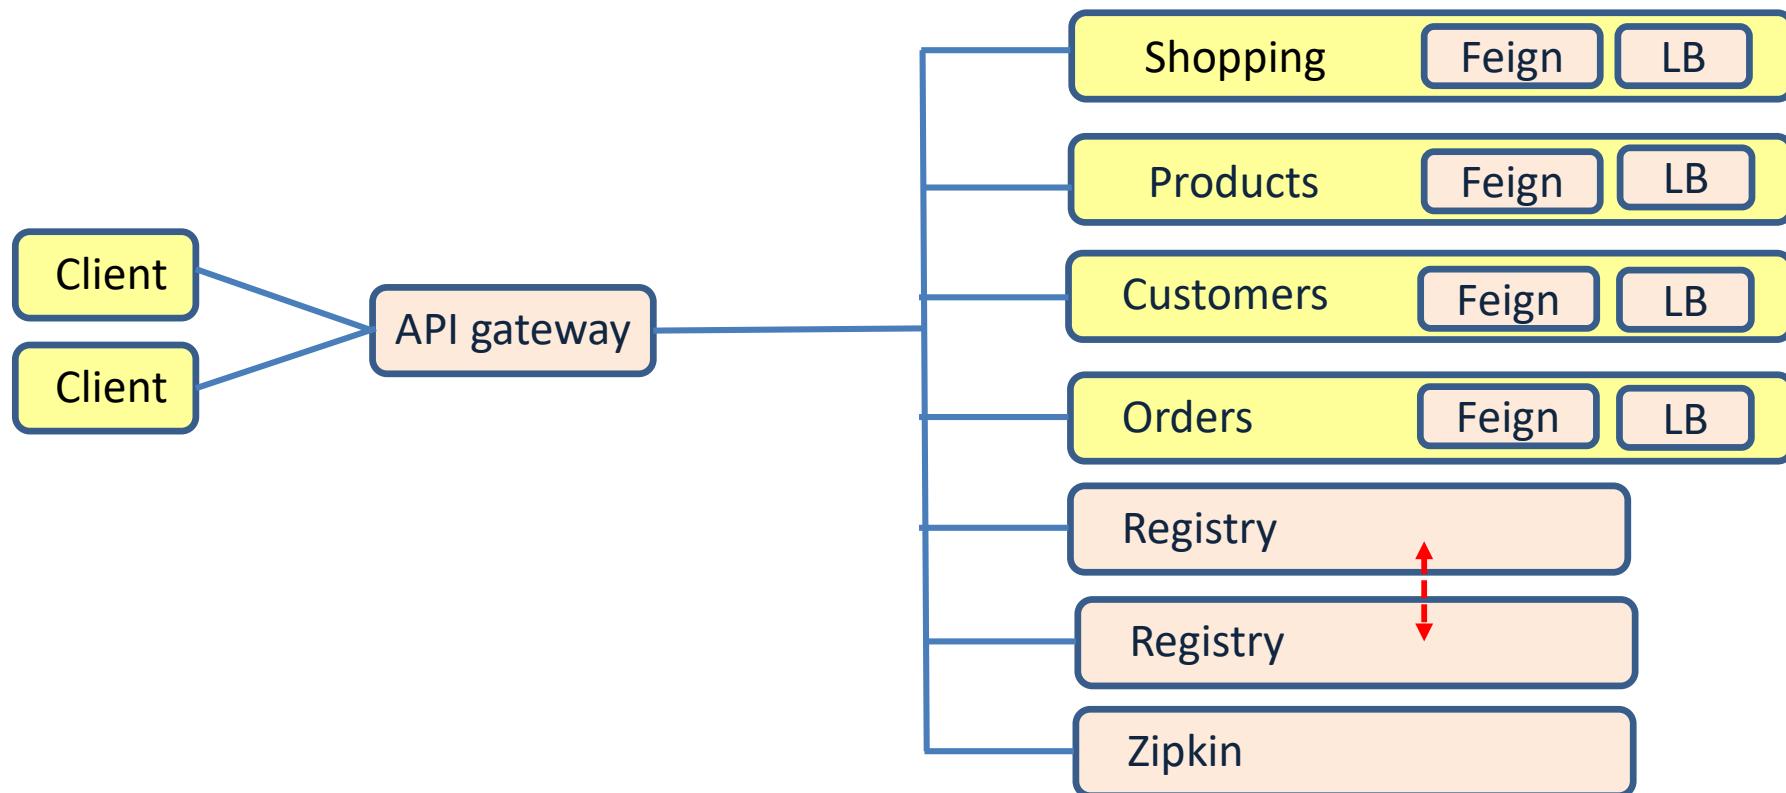


# Zipkin console

The screenshot shows a browser window for the Zipkin console at the URL `localhost:9411/zipkin/dependency?startTime=1627548020343&endTime=1627634420350`. The interface includes a header with the Zipkin logo, search bar, language selection (ENGLISH), and dependency navigation. Below the header, there are date range inputs for 'Start Time' (07/29/2021 03:40:20) and 'End Time' (07/30/2021 03:40:20), and a search button. A dropdown menu is open, showing 'Select...'. The main area displays a trace visualization with two nodes: 'serviceone' at the top and 'servicetwo' at the bottom, connected by a single vertical line.

# Implementing microservices

---



# Challenges of a microservice architecture

---

Challenge	Solution
Complex communication	Feign Registry API gateway
Performance	
Resilience	Registry replicas Load balancing between multiple service instances
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	Zipkin

# Main point

---

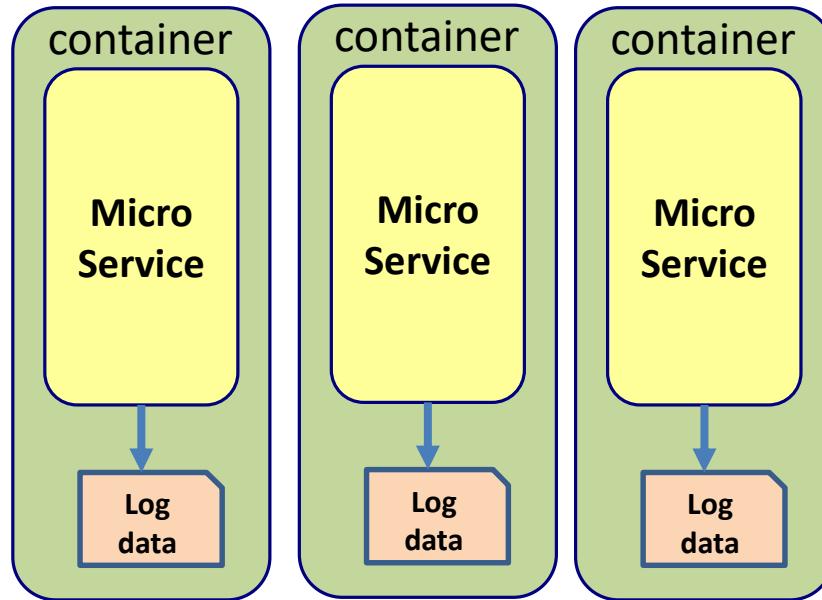
- We need zipkin in order to monitor and debug service-to-service communication
- The Unified Field is the field of perfection

# **DISTRIBUTED LOGGING**

# **ELK STACK**

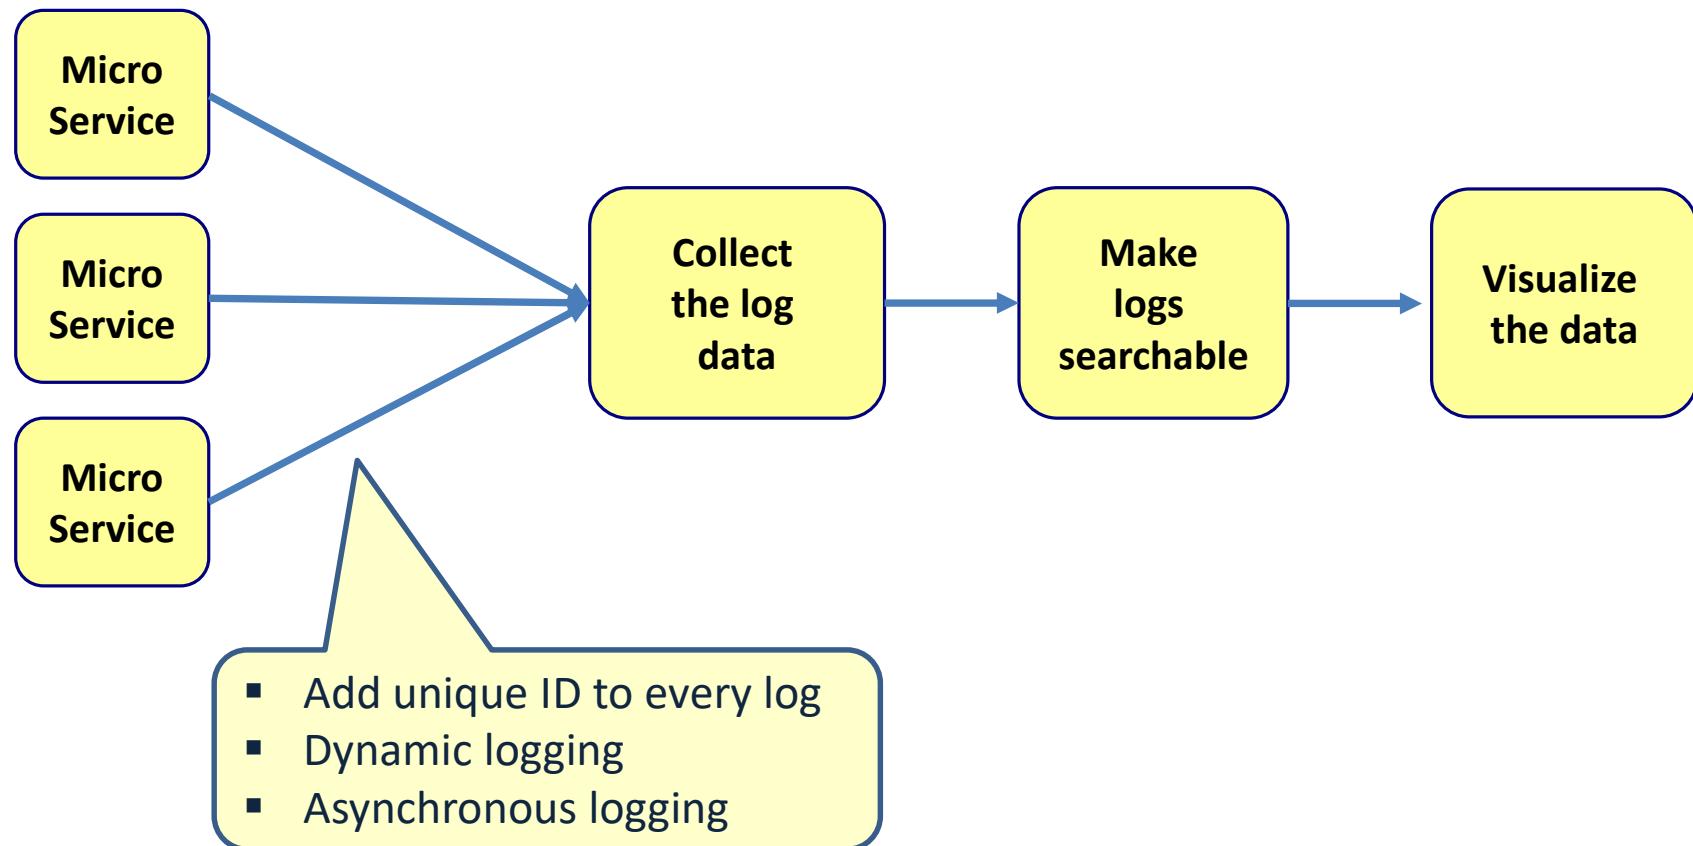
# The need for centralized logging

---

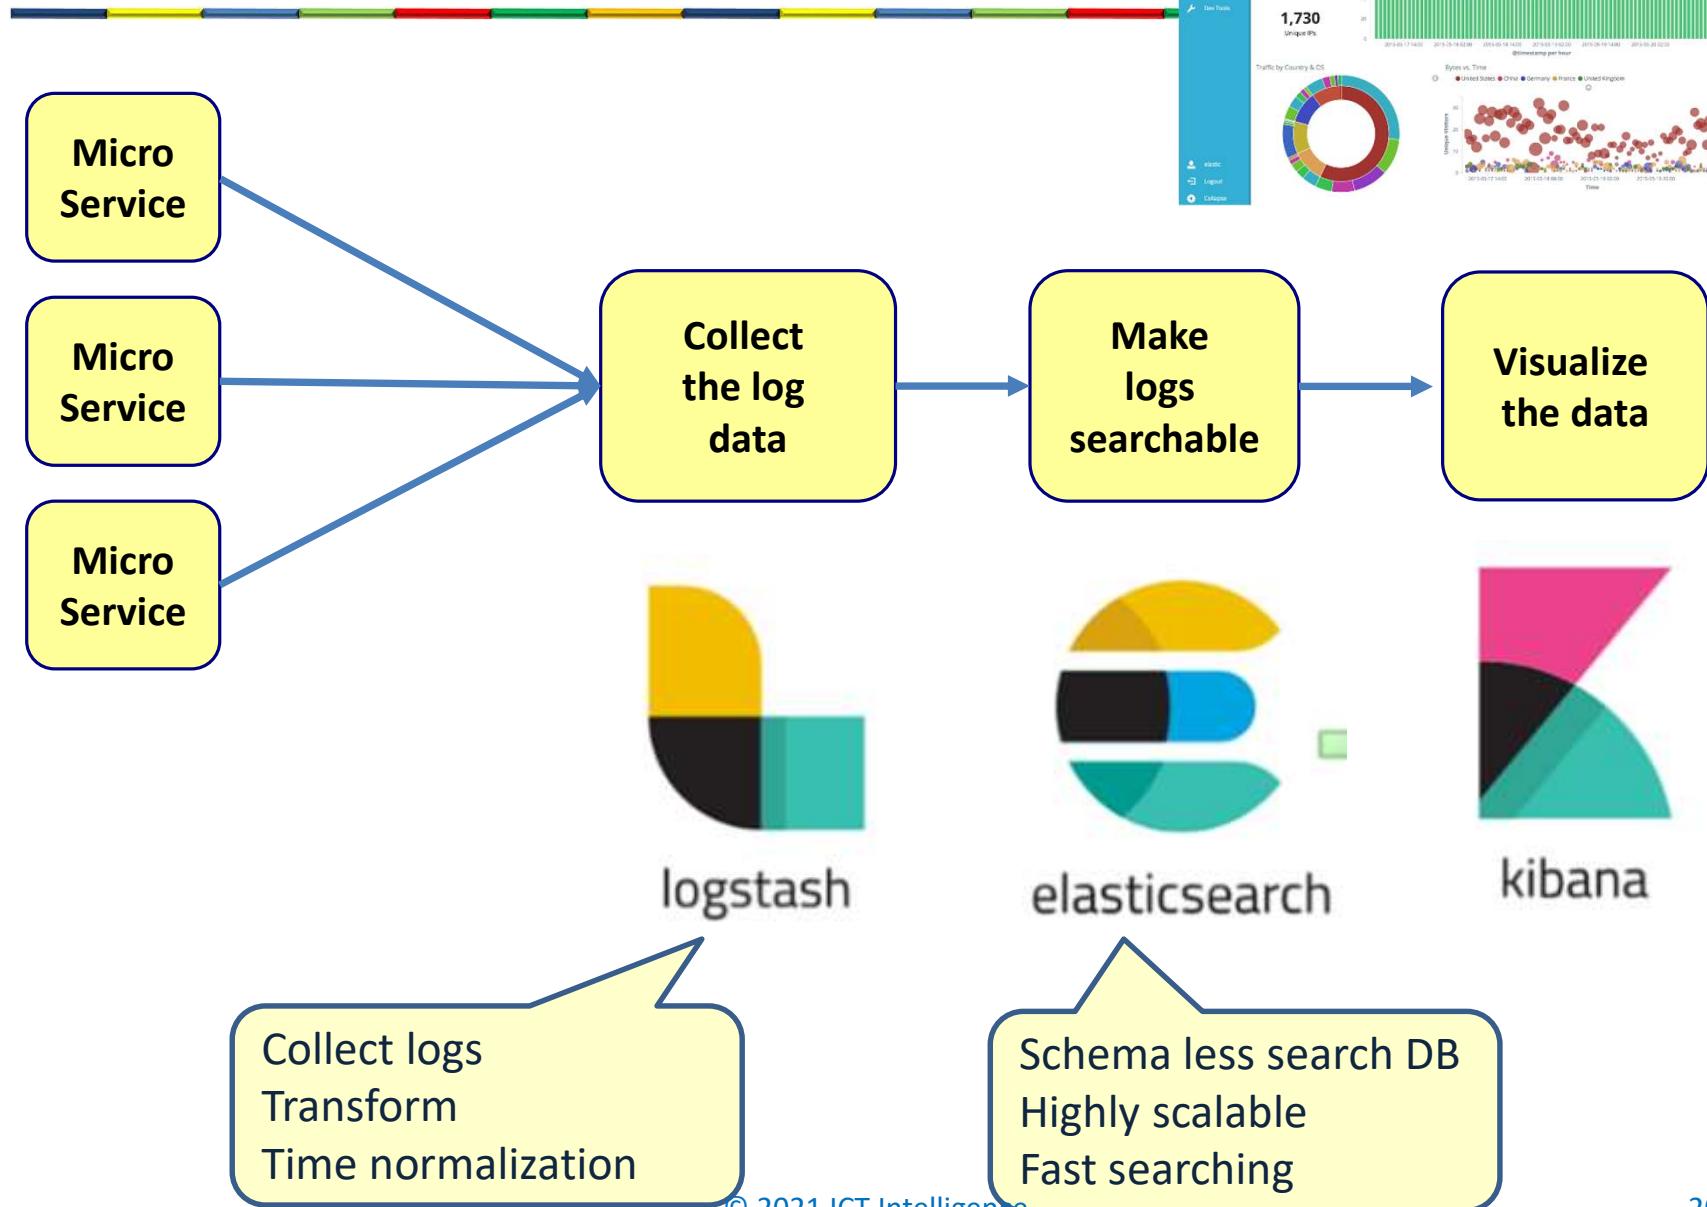


- Local logging does not work
  - Containers come and go
  - Containers have no fixed address

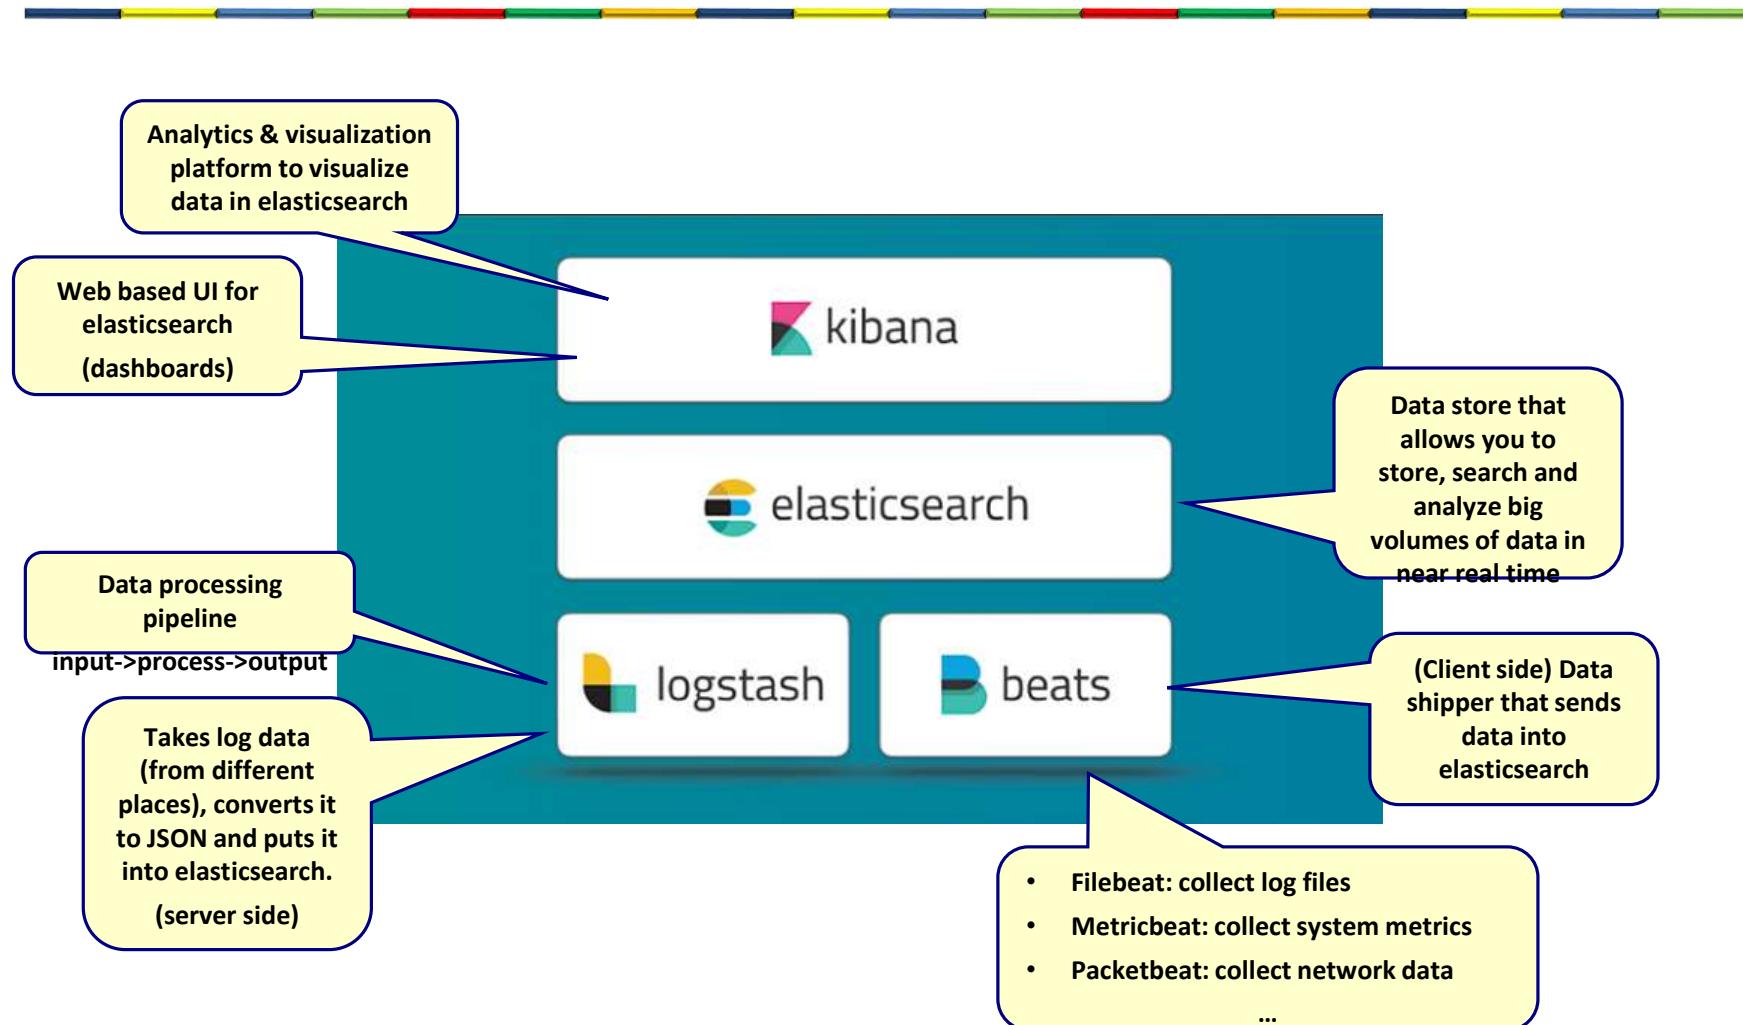
# Microservice logging architecture



# ELK stack



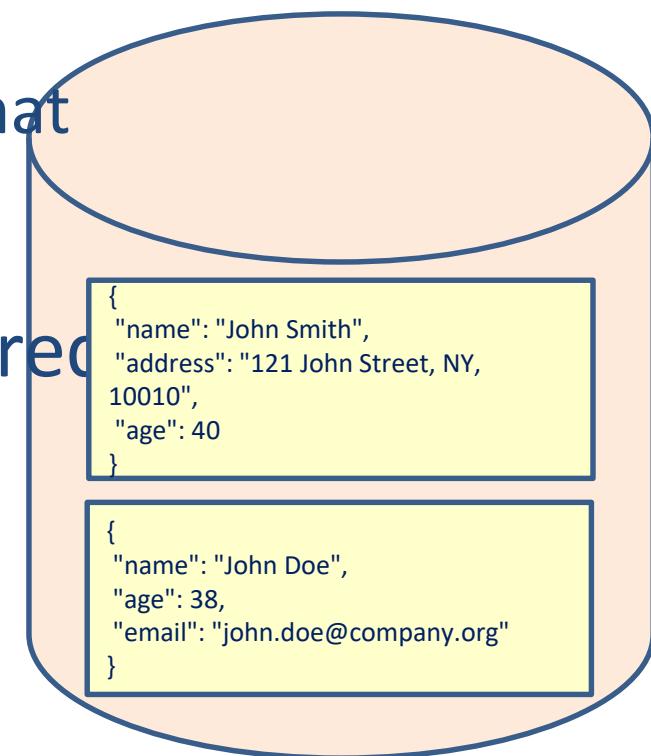
# Elastic stack components



# What is Elasticsearch?

---

- Database
  - Data is stored as documents
  - Data is structured in JSON format
- Full text search engine
- Analytics platform for structured data

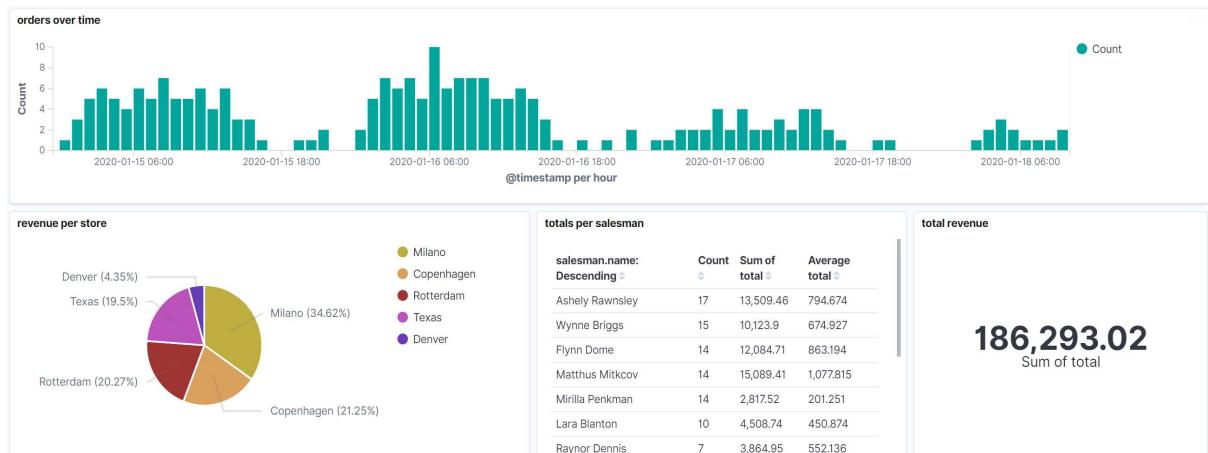


# KIBANA

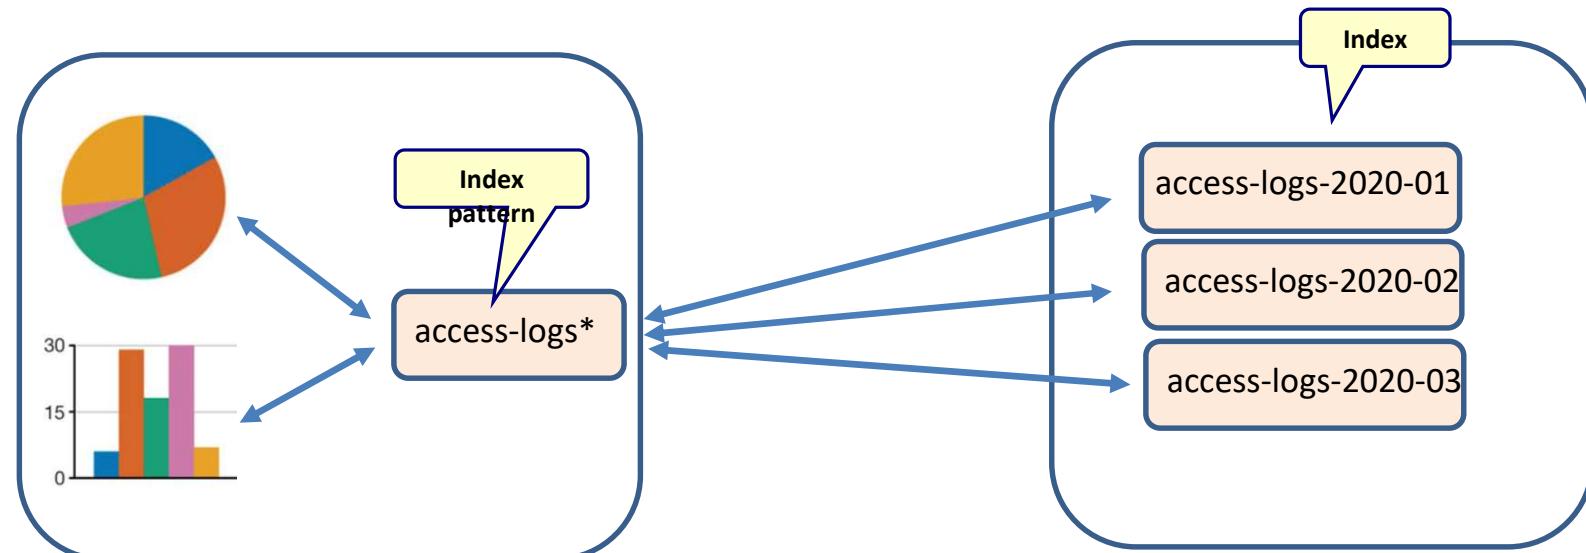
# Kibana



- Web UI on top of elasticsearch
- Has its own Kibana query language (KQL)
- Objects (Queries, visualizations, dashboards, etc.) are saved in elasticsearch

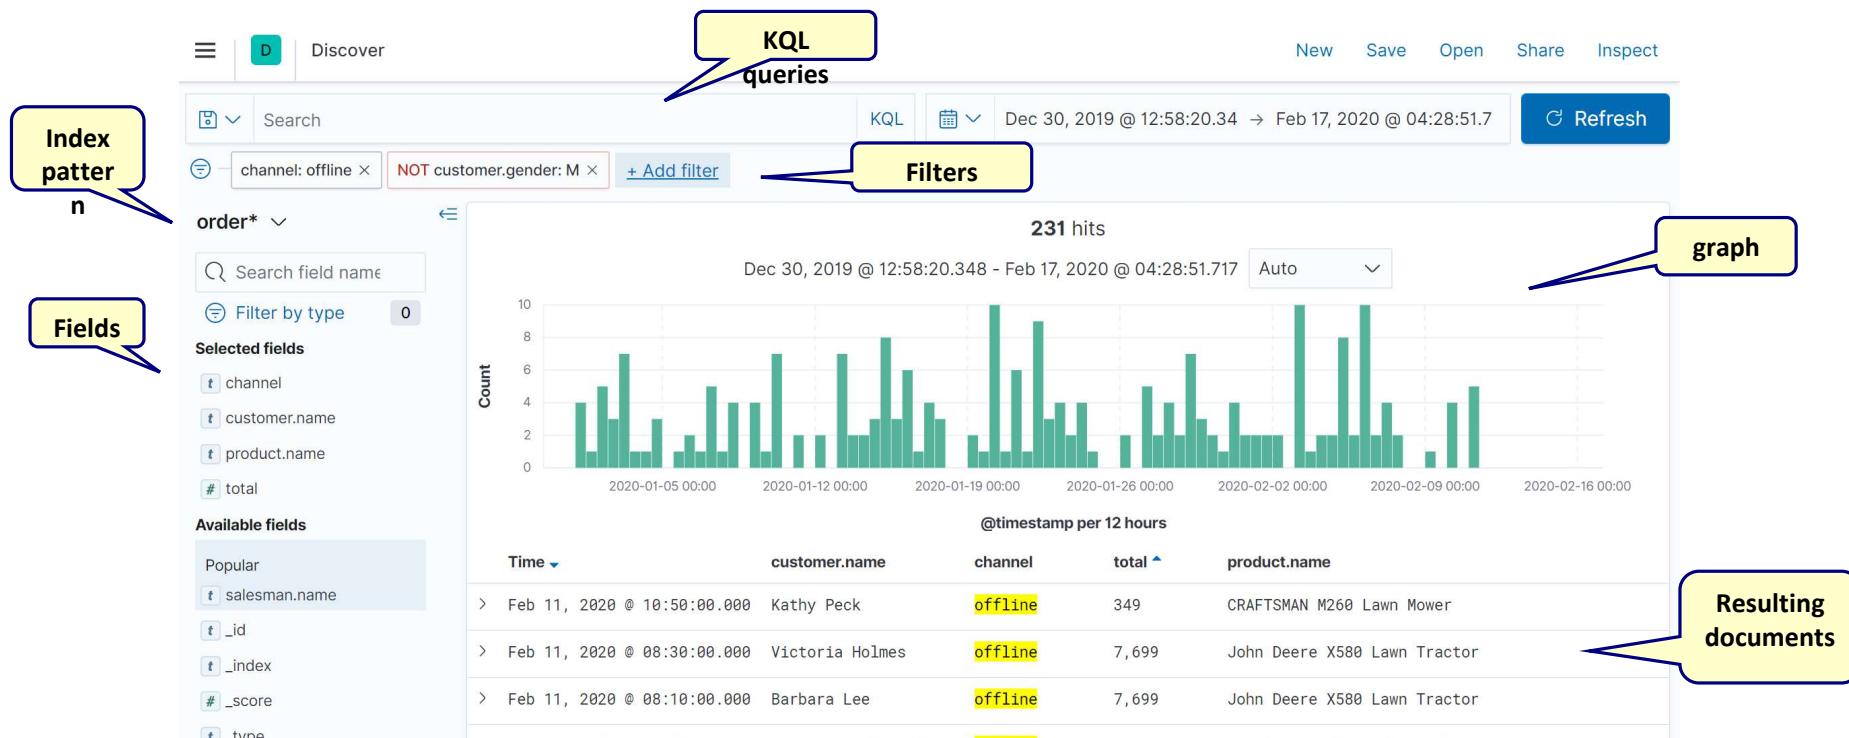


# Index patterns



# Discover app

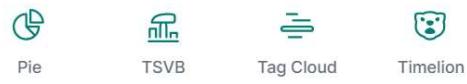
- Good for exploring and analyzing data



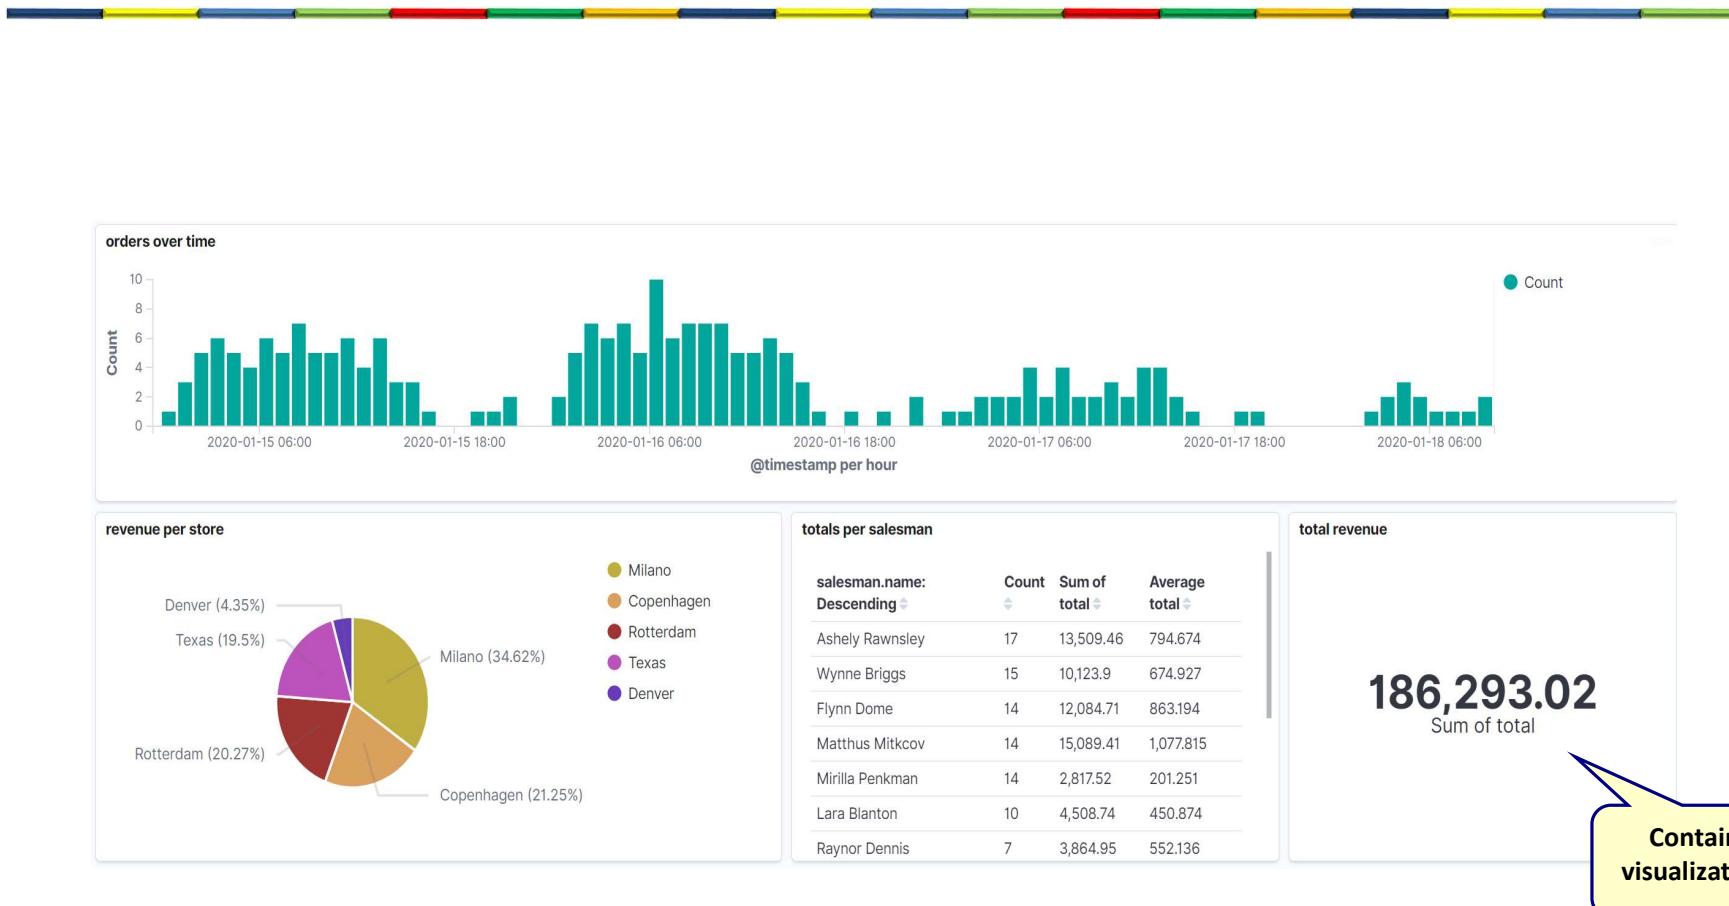
# Visualizations



New Visualization



# Dashboard

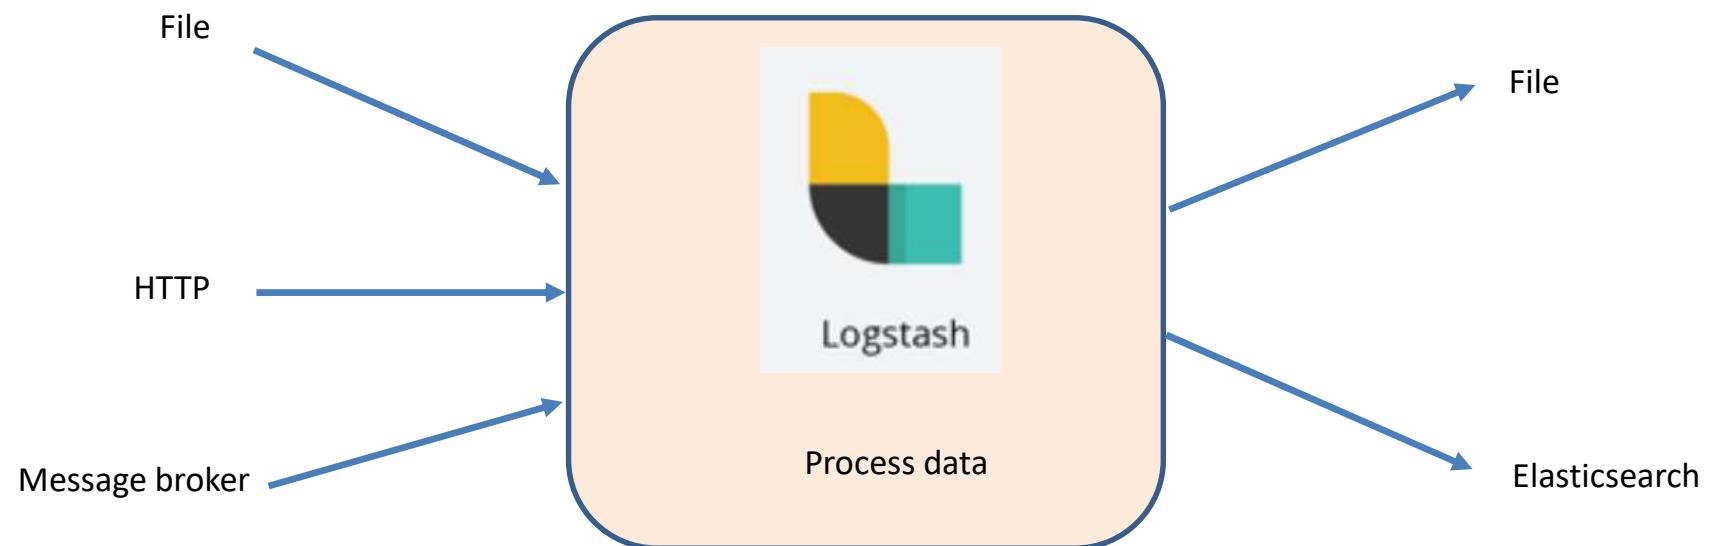


# **LOGSTASH**

# Logstash

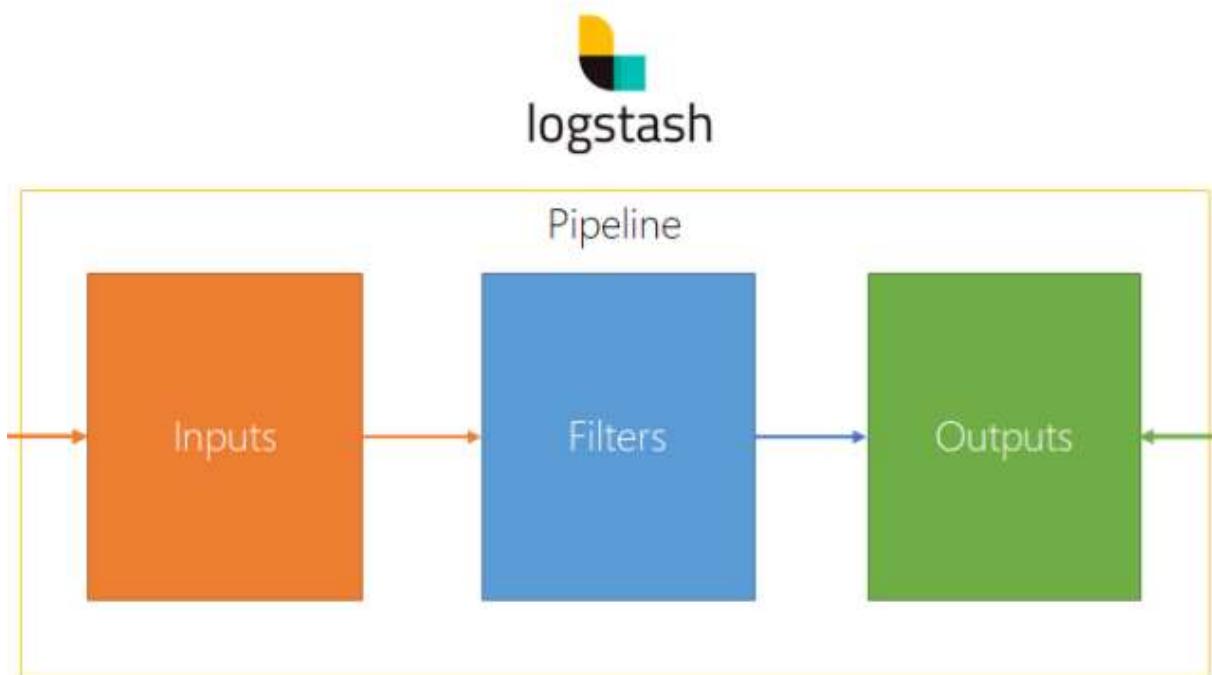
---

- Event processing engine



# How does Logstash work?

---



# Logstash configuration



# Logstash configuration

---

input.txt

Hello world

pipeline.conf

```
input {  
  file {  
    path => "C:/elasticsearchtraining/temp/input.txt"  
    start_position => "beginning"  
  }  
}  
  
output {  
  stdout {  
    codec => rubydebug  
  }  
  file {  
    path => "C:/elasticsearchtraining/temp/output.txt"  
  }  
}
```

output.txt

```
{  
  "host": "DESKTOP-BVHRK6K",  
  "@version": "1",  
  "path": "C:/elasticsearchtraining/temp/input.txt",  
  "message": "Hello world\r",  
  "@timestamp": "2021-01-16T13:52:32.726Z"  
}
```

Anytime this file changes, read from this file

Write the output to the console

Write the output to the specified file

# Logstash configuration

---

input.txt

Hi there

pipeline.conf

```
input {
  file {
    path => "C:/elasticsearchtraining/temp/input.txt"
      start_position => "beginning"
  }
}

filter {
  mutate {
    uppercase => ["message"]
  }
}

output {
  stdout {
    codec => rubydebug
  }
  file {
    path => "C:/elasticsearchtraining/temp/output.txt"
  }
}
```

output.txt

```
{
  "path": "C:/elasticsearchtraining/temp/input.txt",
  "message": "HI THERE\r",
  "host": "DESKTOP-BVHRK6K",
  "@version": "1",
  "@timestamp": "2021-01-16T14:17:10.537Z"
}
```

# Testing grok

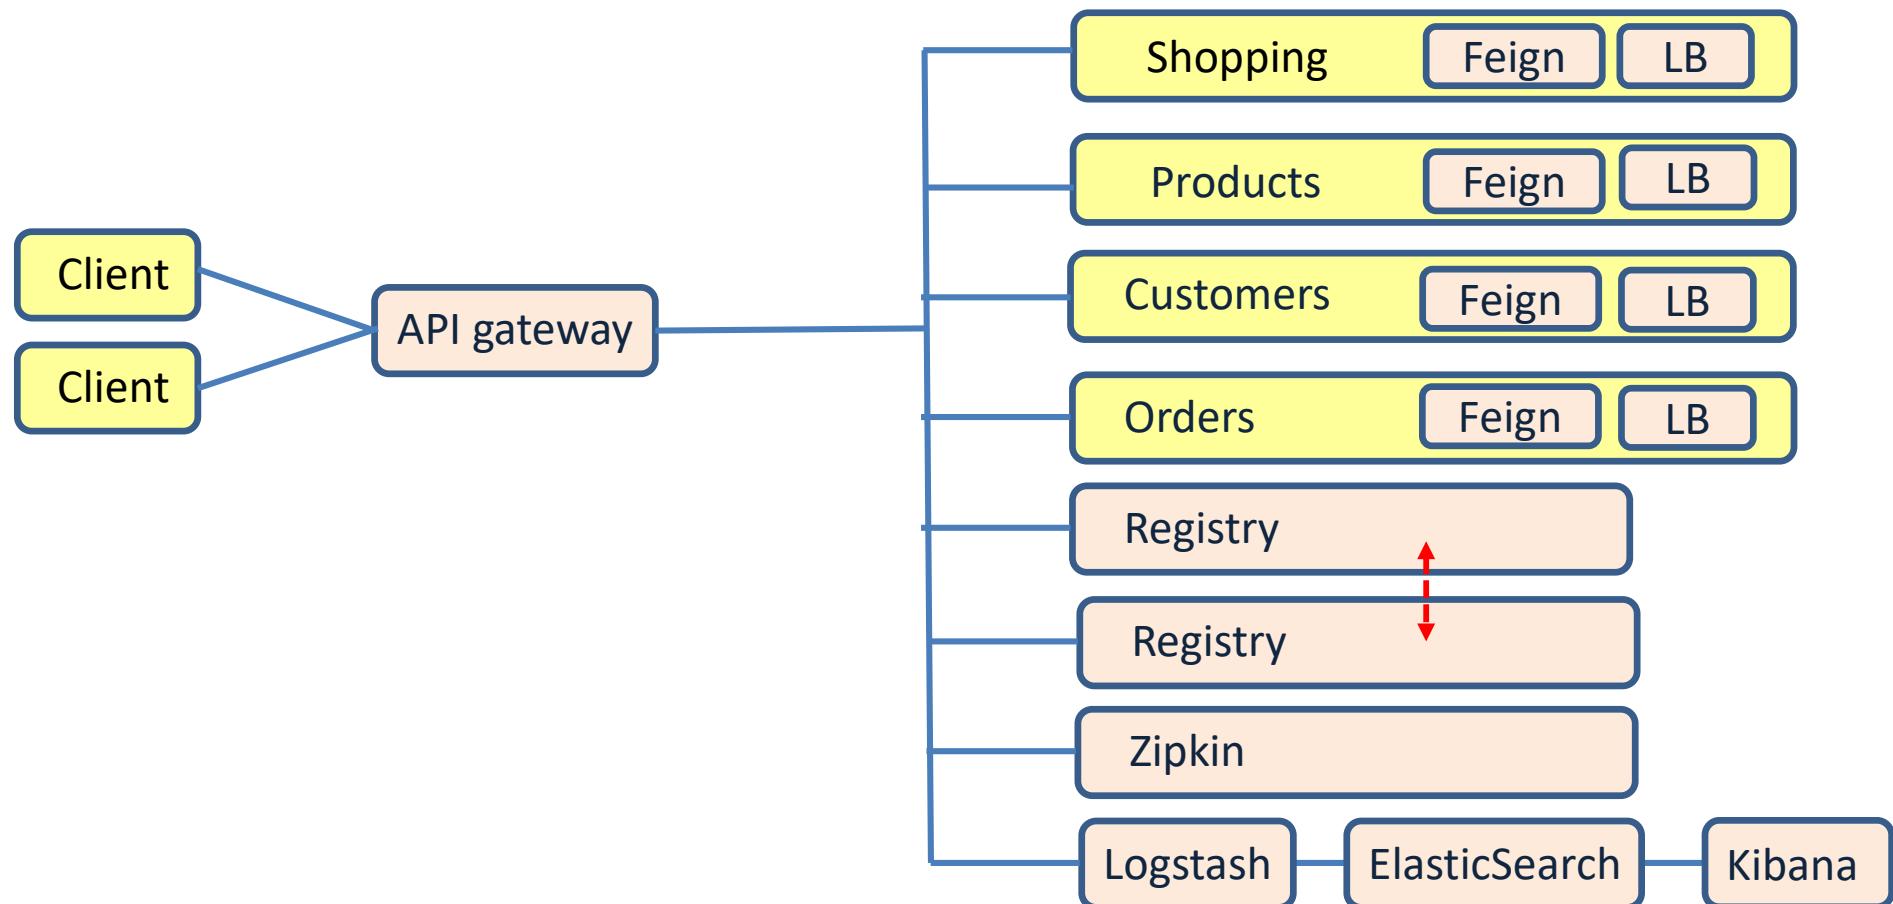
---

- <http://grokdebug.herokuapp.com/>

The screenshot shows the Grok Debugger interface. At the top, there's a header bar with a 'Not secure' warning and the URL 'grokdebug.herokuapp.com'. Below the header are tabs: 'Grok Debugger' (which is active), 'Debugger', 'Discover', and 'Patterns'. The main area displays a log entry: '129.10.12.121 GET /compute 2500 300'. Below this, a pattern is shown: '%{IP:client} %{WORD:method} %{URIPATHPARAM:request} %{NUMBER:bytes} %{NUMBER:duration}'. At the bottom of the interface, there are several checkboxes: 'Add custom patterns', 'Keep Empty Captures', 'Named Captures Only', and 'Singles'. A large text area on the right contains the generated Grok pattern:

```
{  
    "client": [  
        [  
            "129.10.12.121"  
        ]  
    ],  
    "IPV6": [  
        [  
            null  
        ]  
    ],  
    "IPV4": [  
        [  
            "129.10.12.121"  
        ]  
    ],  
    "method": [  
        [  
            "GET"  
        ]  
    ]  
}
```

# Implementing microservices



# Challenges of a microservice architecture

Challenge	Solution
Complex communication	Feign Registry API gateway
Performance	
Resilience	Registry replicas Load balancing between multiple service instances
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	Zipkin ELK

# **RESILIENCE**

The ability to recover from failures

# Fallacies of distributed computing

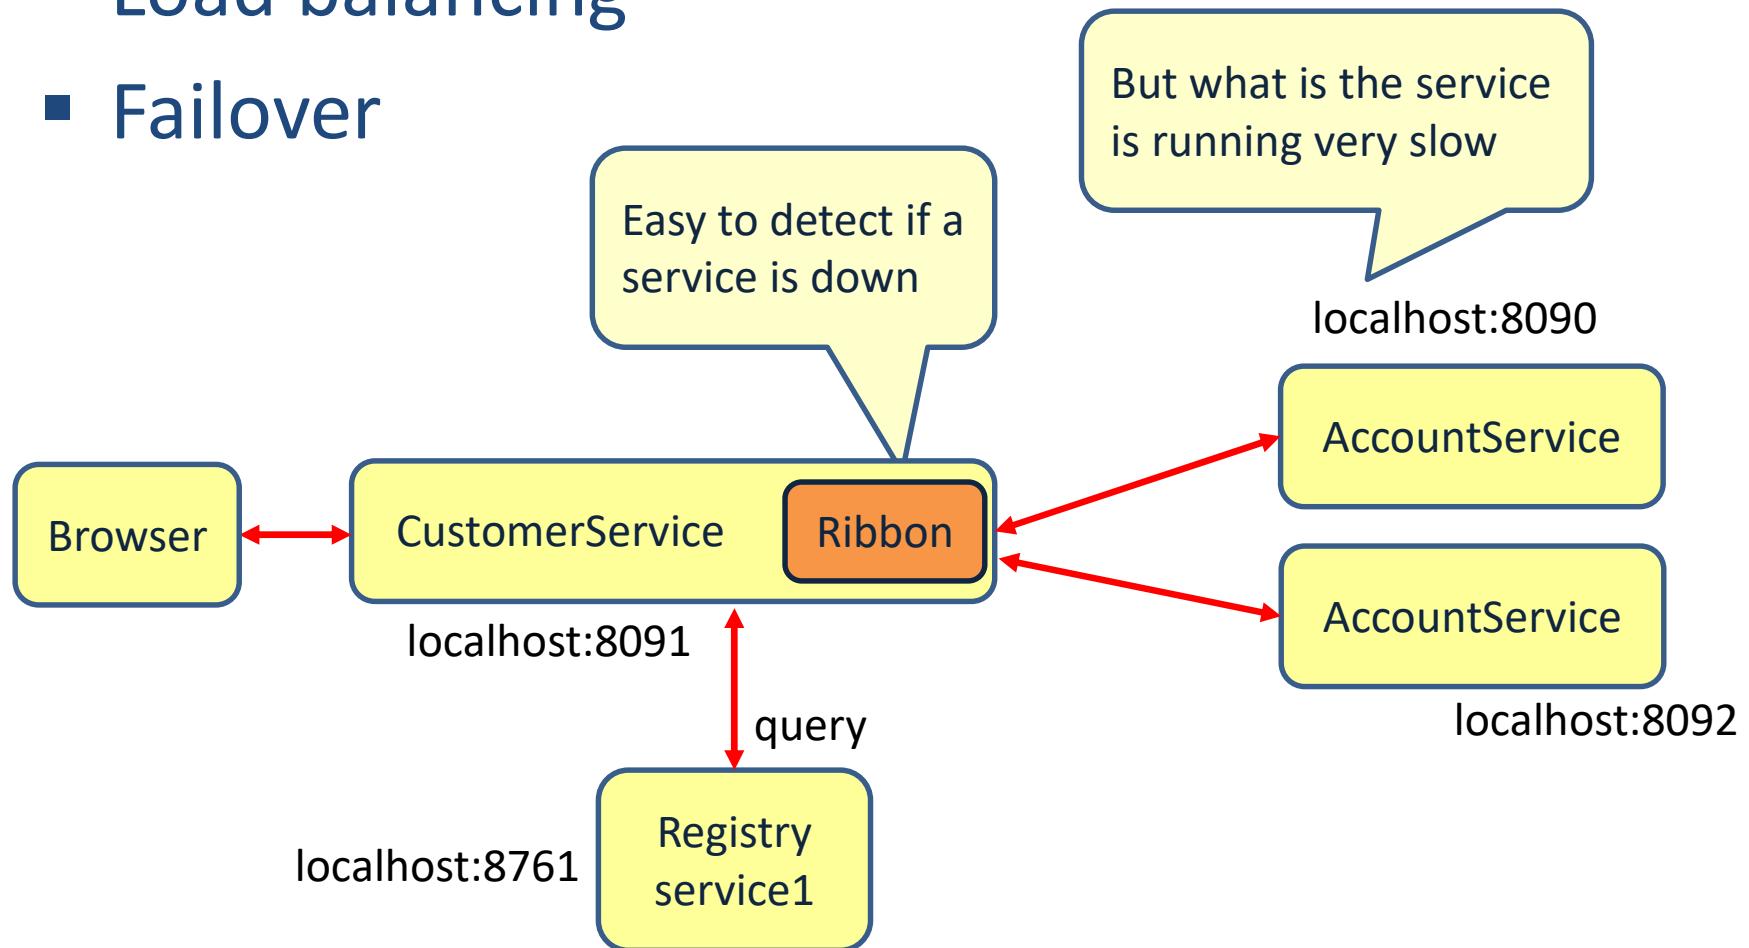
---

- The network is reliable
- Latency is zero
- Bandwidth is infinite
- The network is secure
- Topology doesn't change
- There is one administrator
- Transport cost is zero
- The network is homogeneous

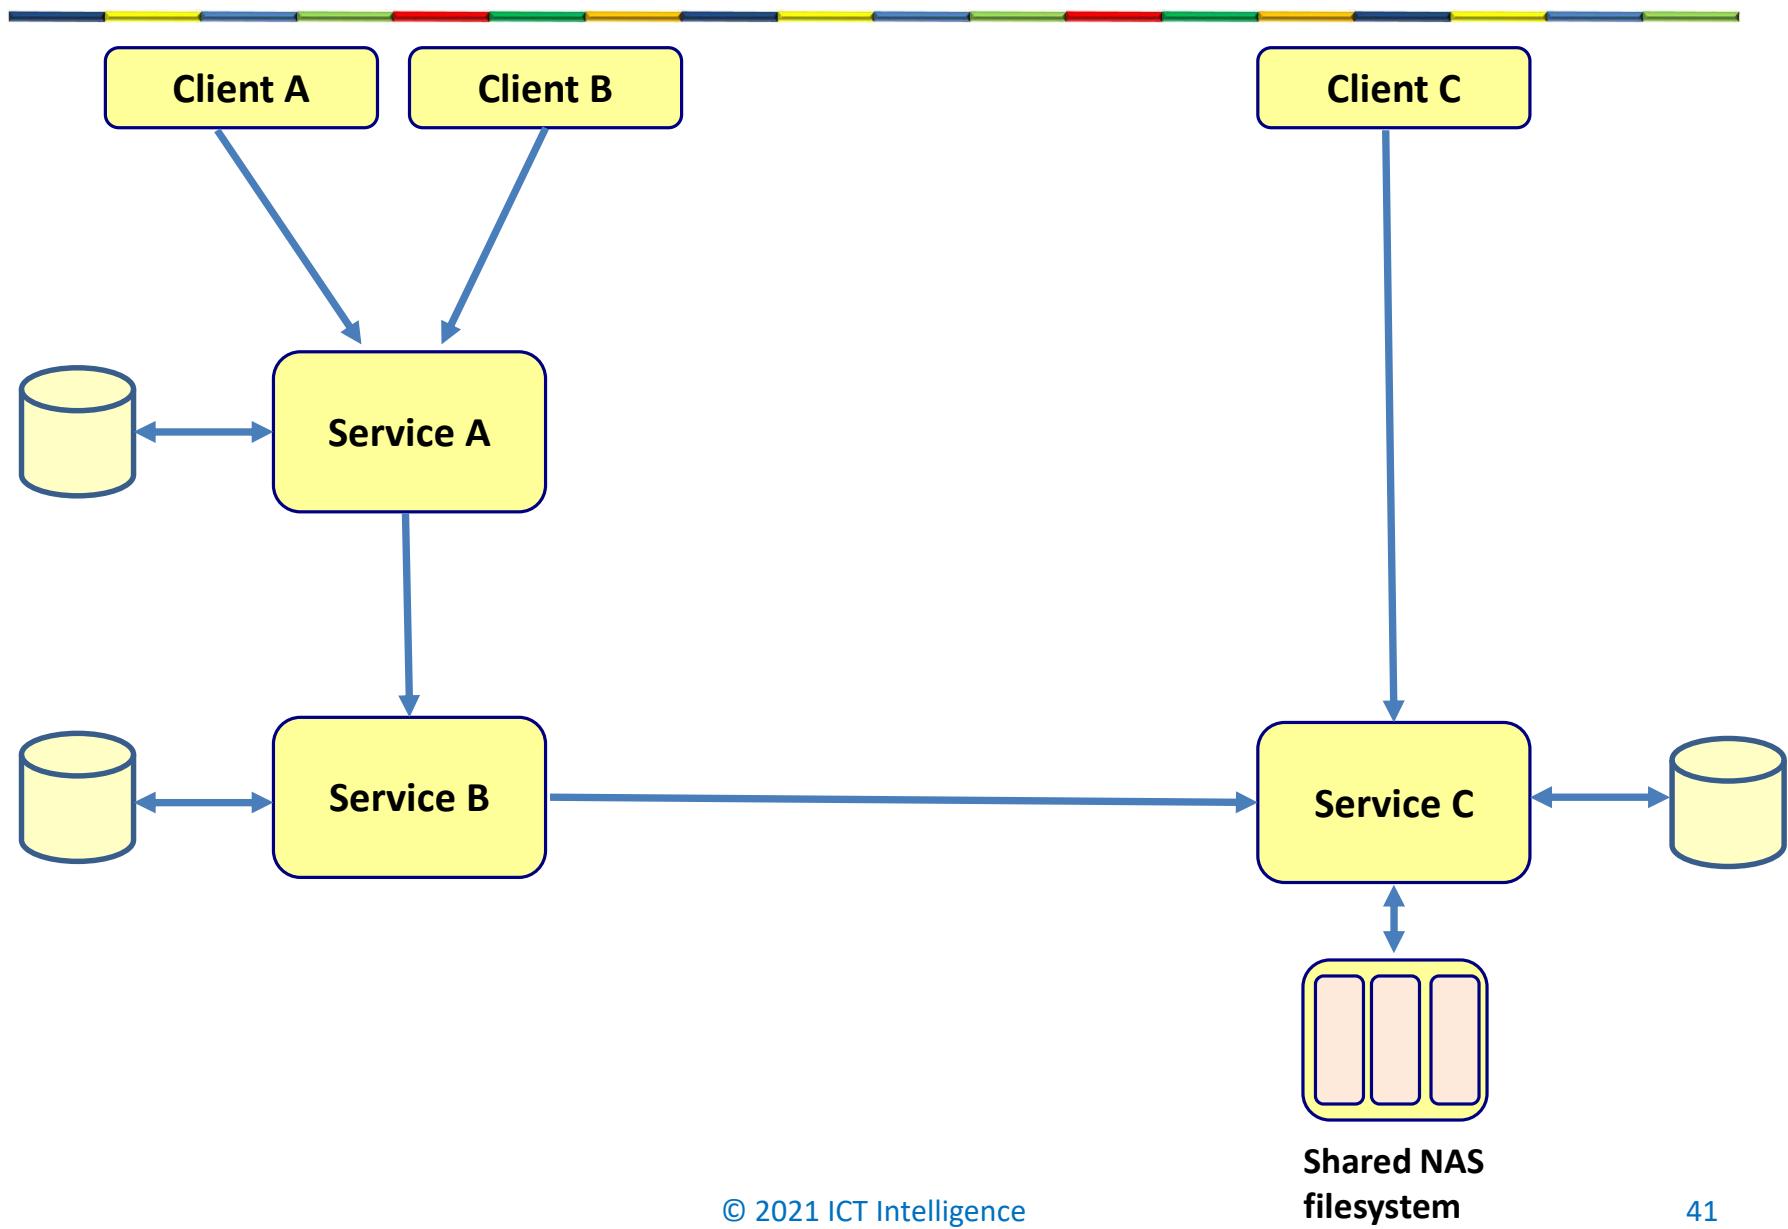
# Clustering

---

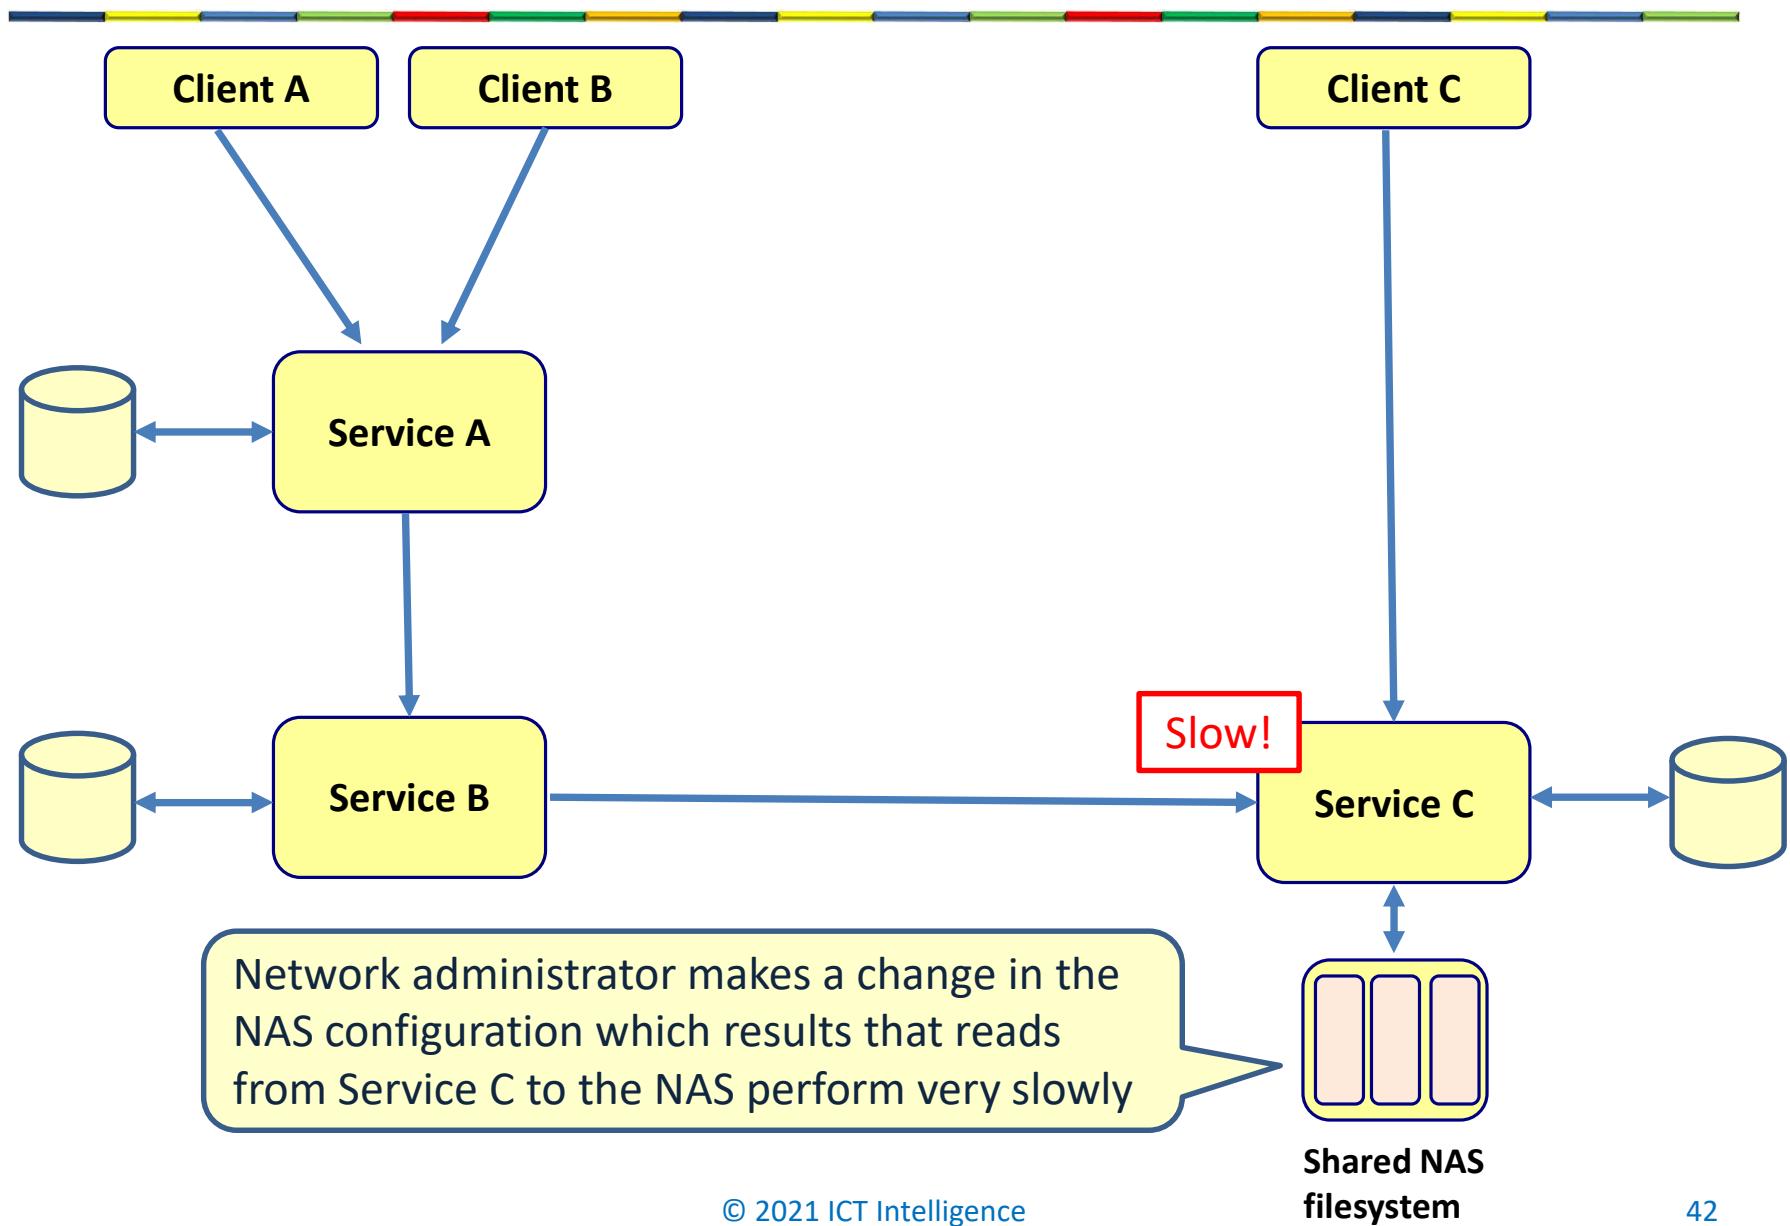
- Load balancing
- Failover



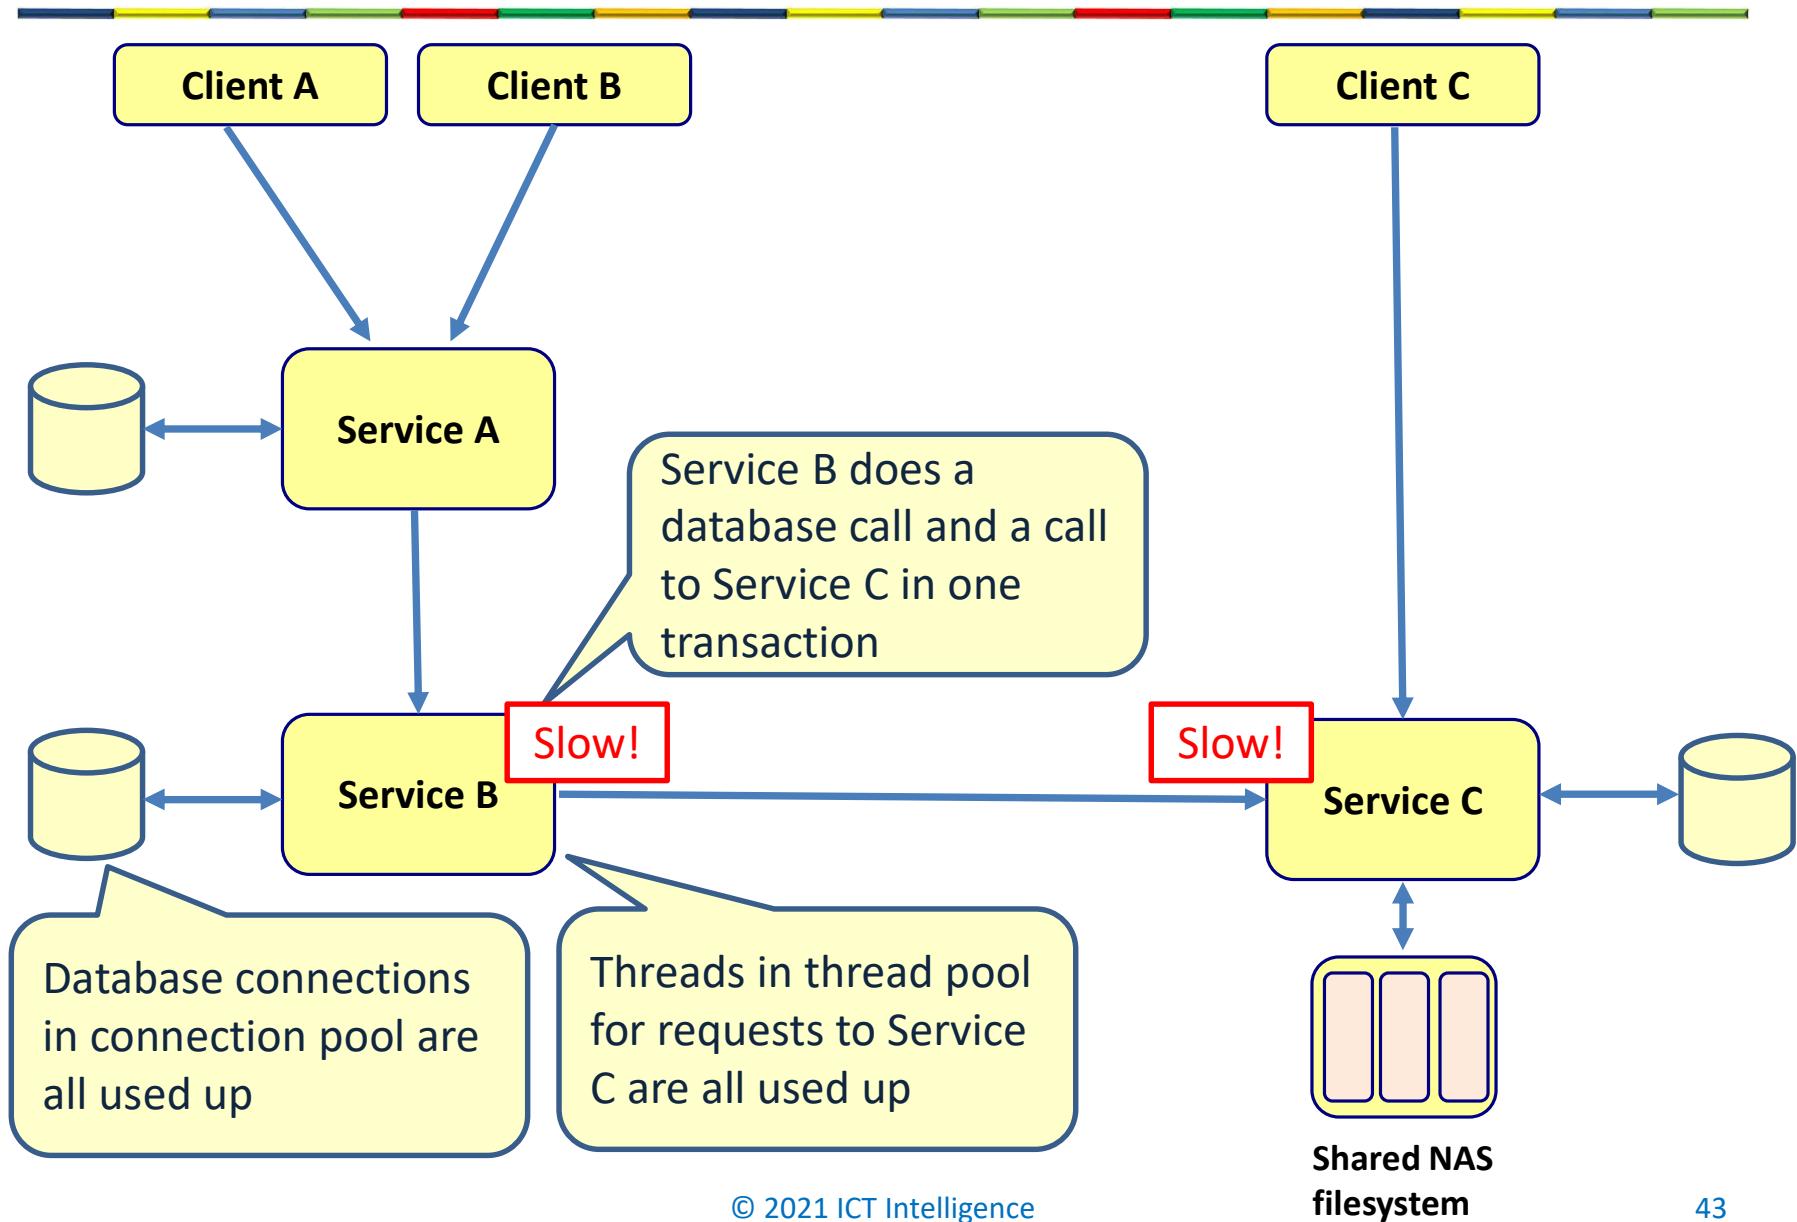
# Example



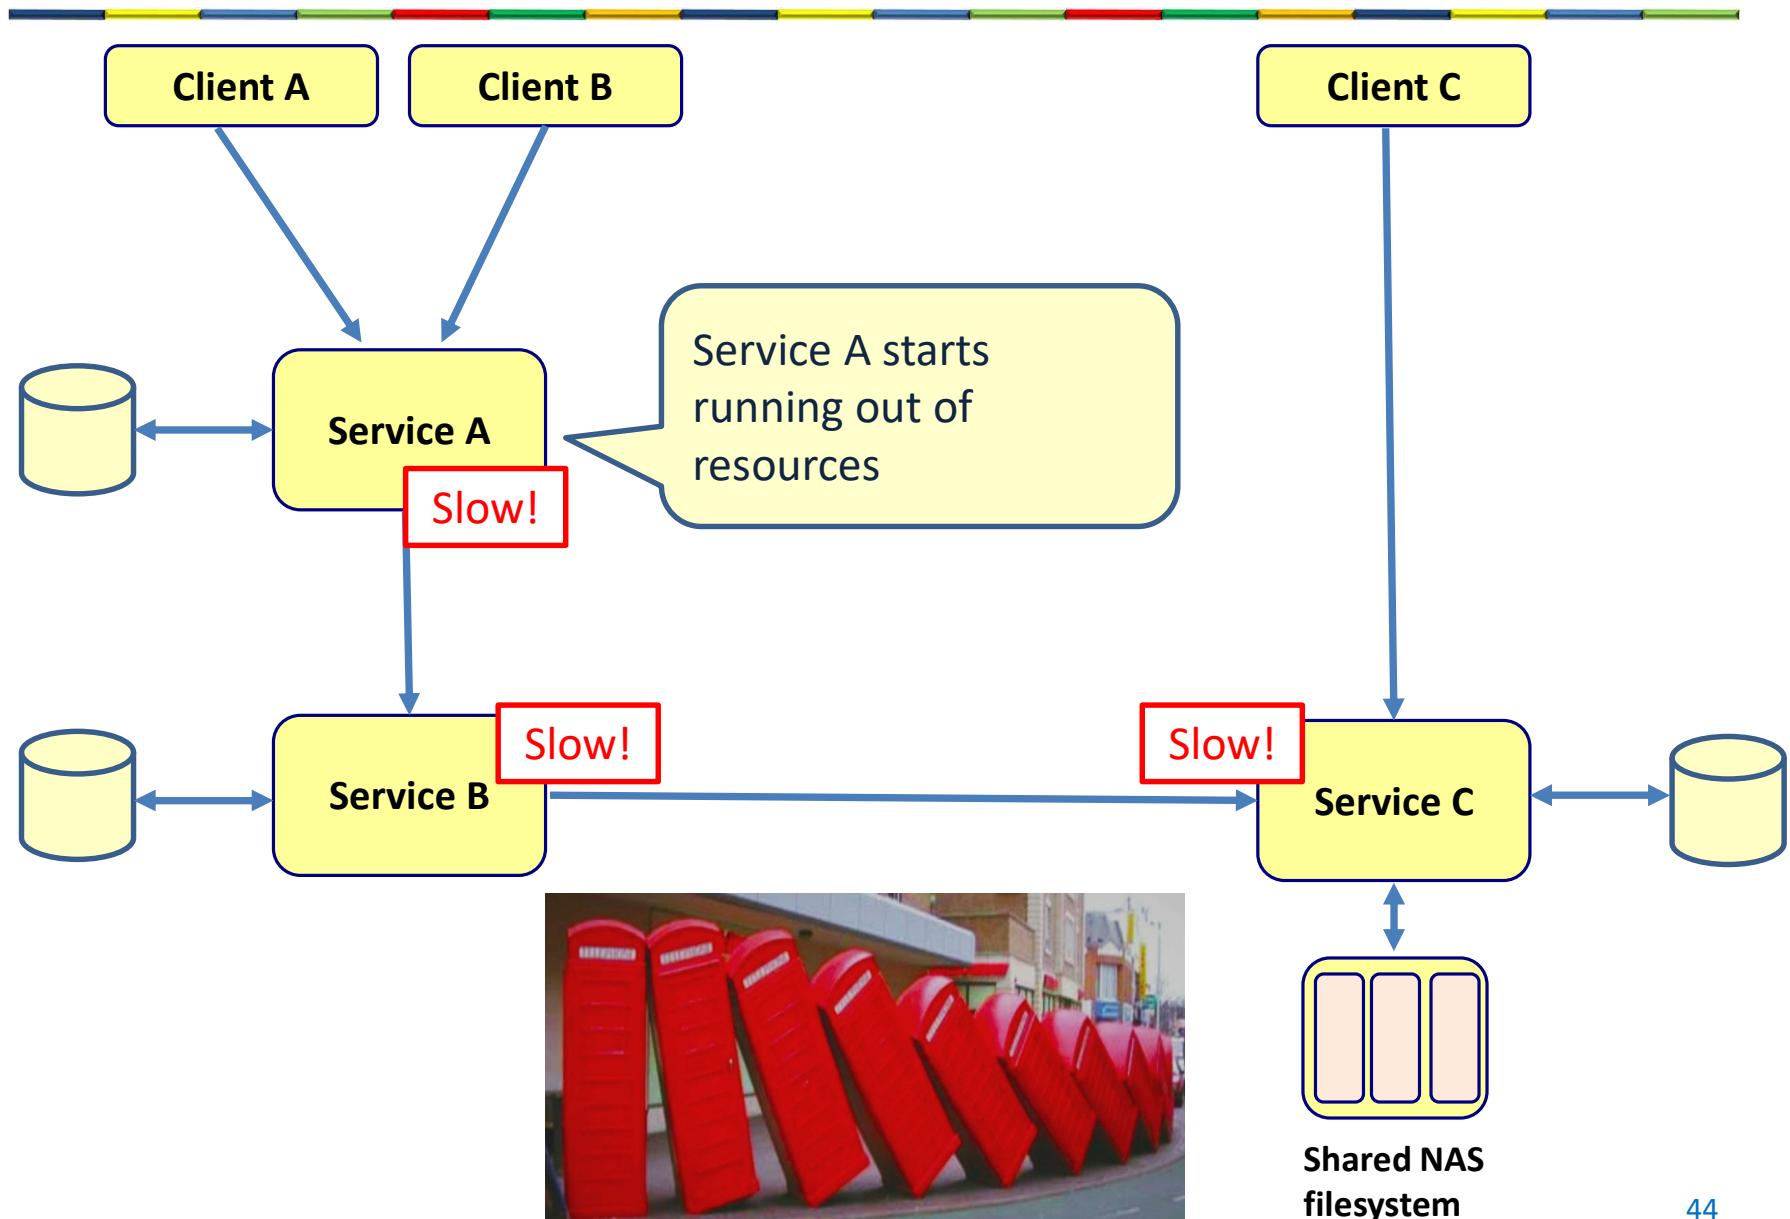
# Example



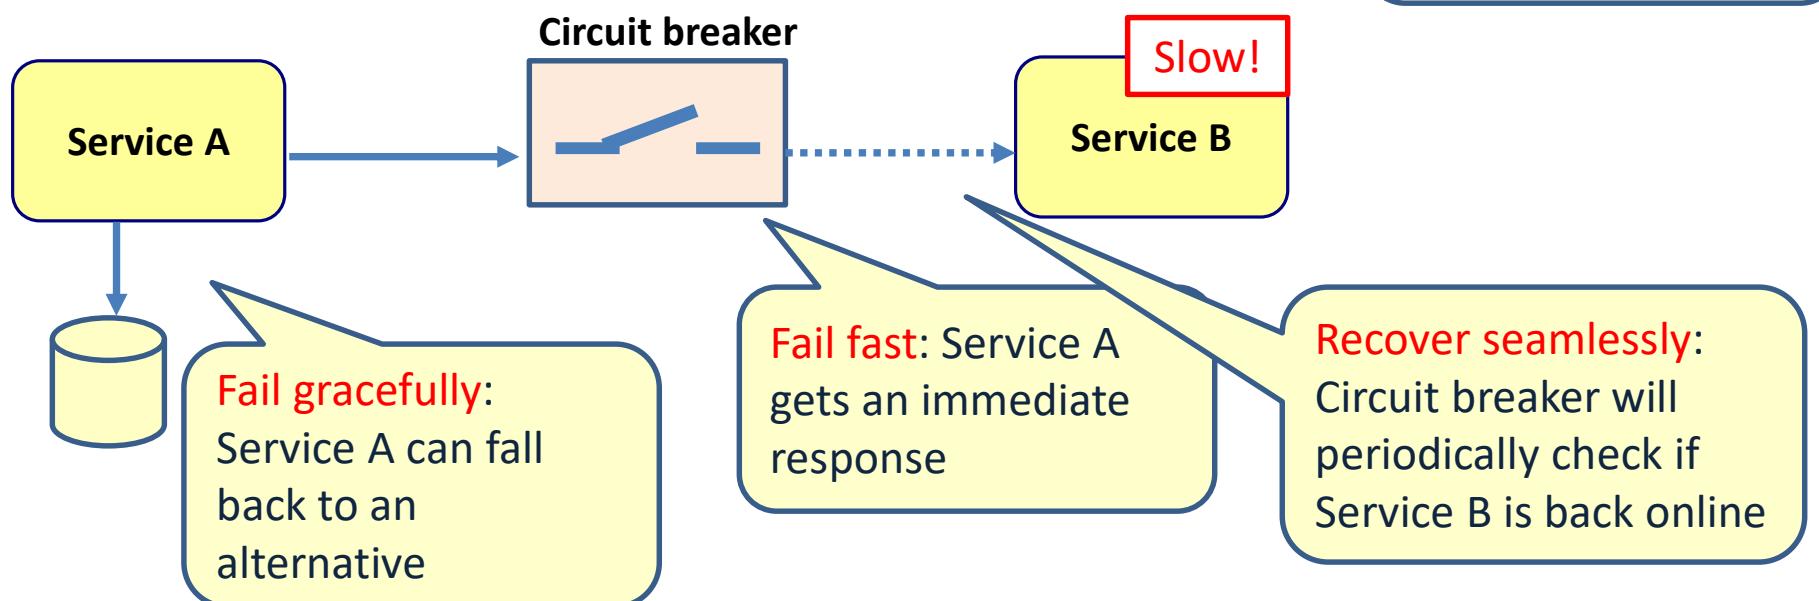
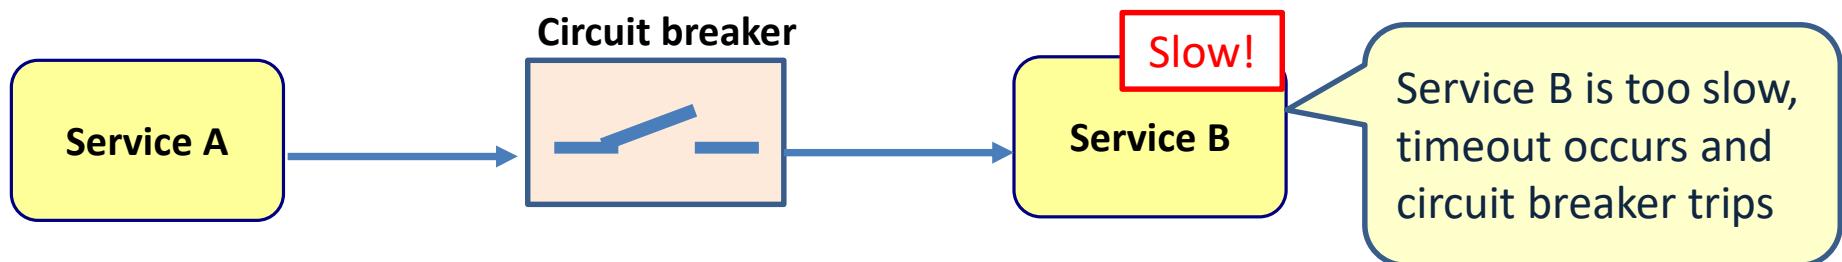
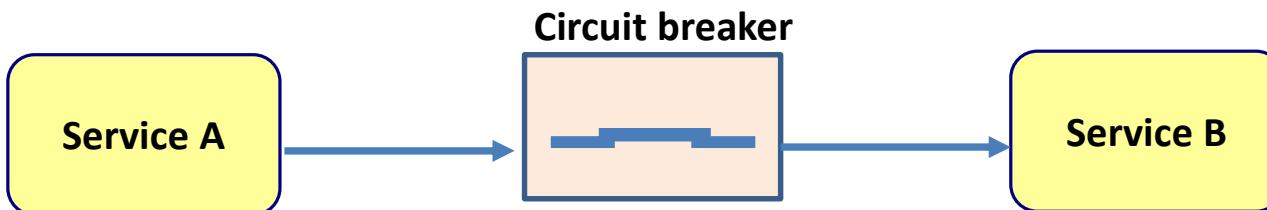
# Example



# Example



# Circuit breaker



# Main point

---

- A circuit breaker takes care that not the whole microservice architecture gets slow when one service becomes slow.
- Every relative part of creation is connected at the level of pure consciousness.

# **RESILIENCE: HYSTRIX**

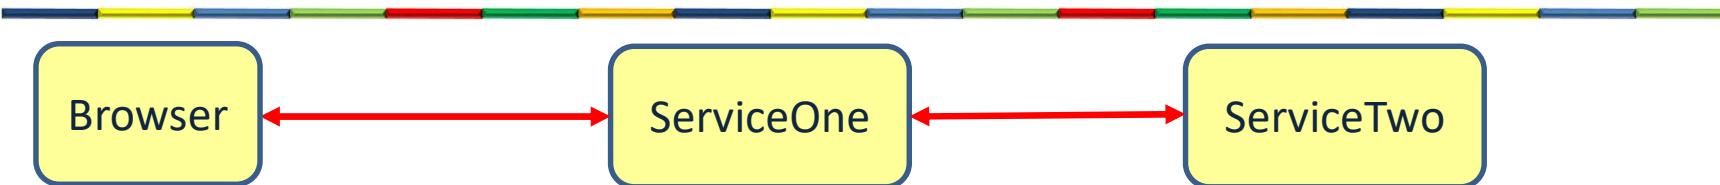
# Hystrix dependency



pom.xml

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-netflix-hystrix</artifactId>
</dependency>
```

# Enable Hystrix



```
@SpringBootApplication  
@EnableCircuitBreaker  
public class ServiceOneApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(Service2Application.class, args);  
    }  
}
```

Enable Hystrix

```
@SpringBootApplication  
@EnableCircuitBreaker  
public class ServiceTwoApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(Service2Application.class, args);  
    }  
}
```

Enable Hystrix

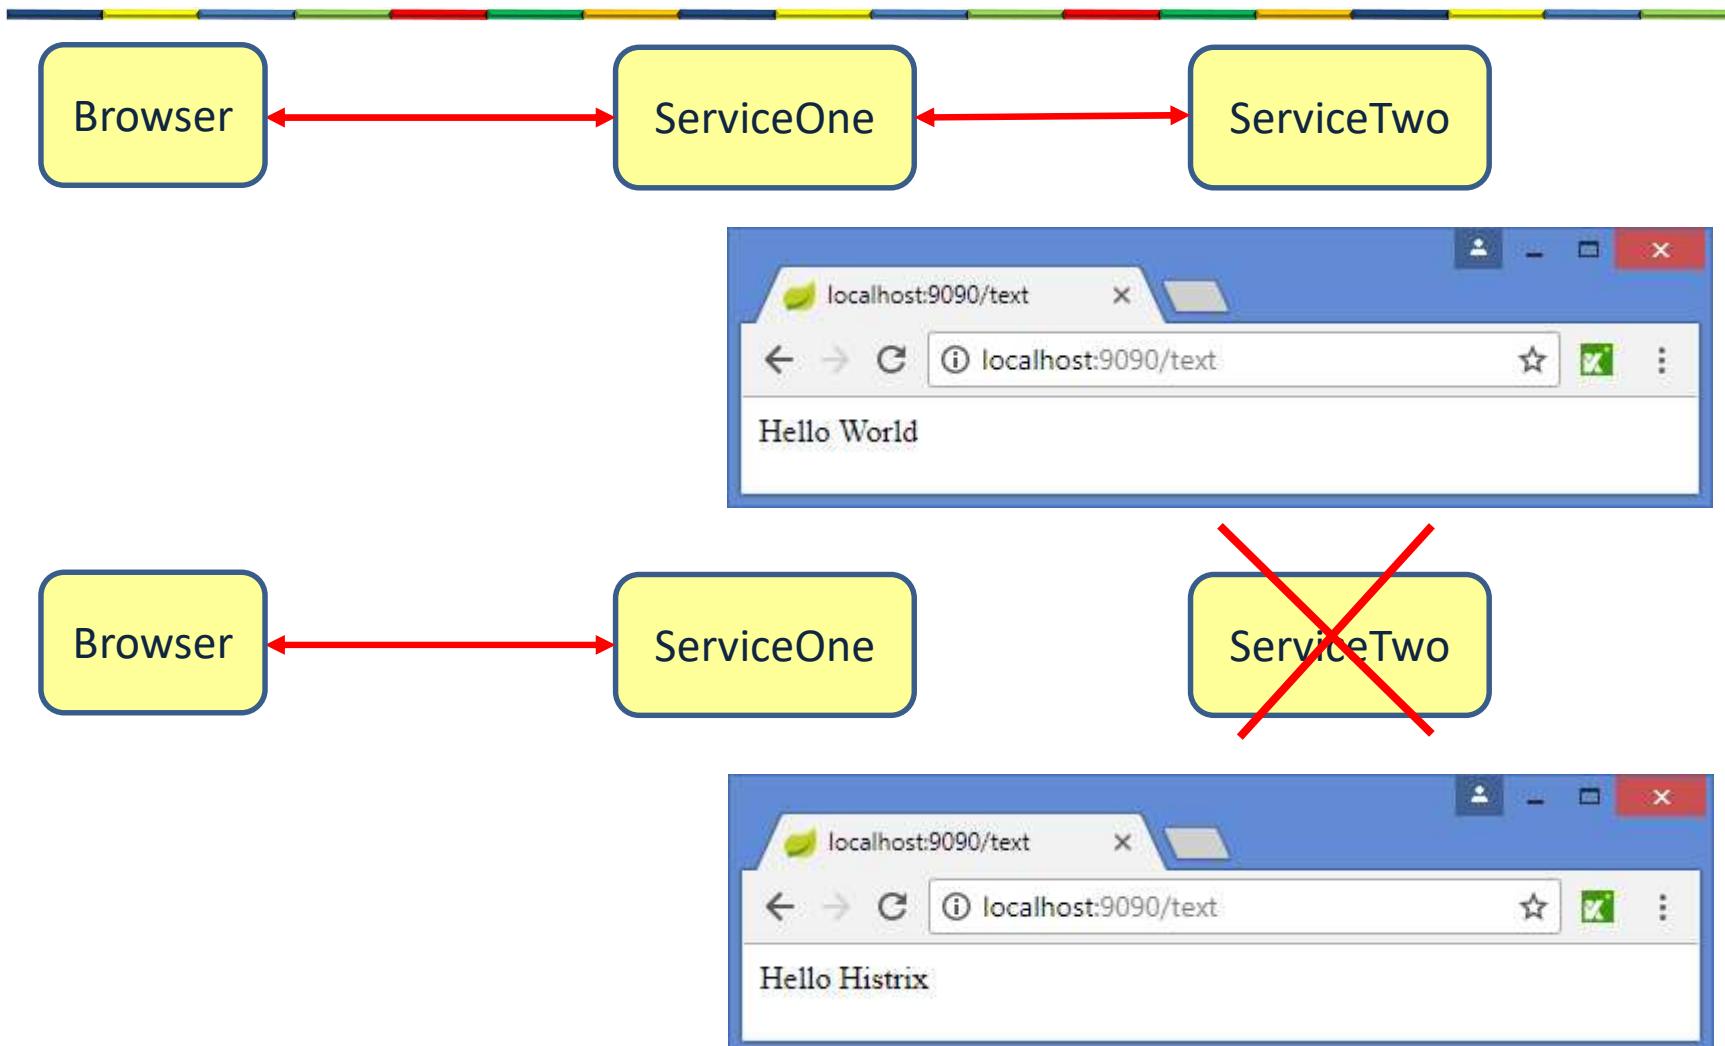
# Using the circuit breaker

```
public class ServiceOneController {  
  
    @Autowired  
    RestTemplate restTemplate;  
  
    @RequestMapping("/text")  
    @HystrixCommand(fallbackMethod = "getTextFallback")  
    public String getText() {  
        String service2Text = restTemplate.getForObject("http://localhost:9091/text",  
                                                       String.class);  
        return "Hello "+ service2Text;  
    }  
  
    public String getTextFallback() {  
        return "Hello Hystrix";  
    }  
  
    @Bean  
    RestTemplate getRestTemplate() {  
        return new RestTemplate();  
    }  
}
```

If this method throws an exception or takes longer than **2 seconds**, call the fallback method

Fallback method

# Using Histrix



# Setting the timeout

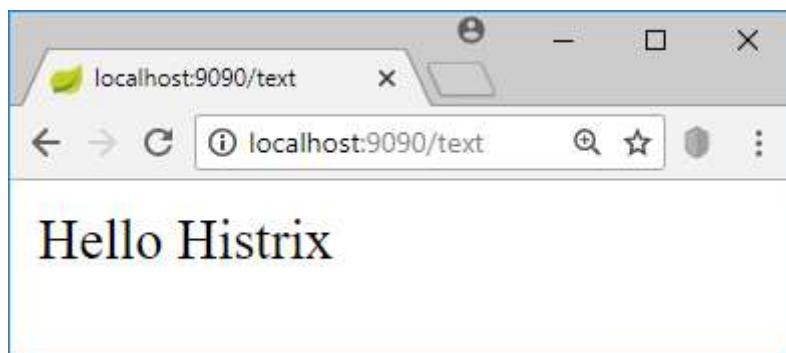
```
public class ServiceOneController {  
  
    @Autowired  
    RestTemplate restTemplate;  
  
    @RequestMapping("/text")  
    @HystrixCommand(fallbackMethod = "getTextFallback", commandProperties=  
    {@HystrixProperty(name="execution.isolation.thread.timeoutInMilliseconds",  
                     value="4000")})  
    public String getText() {  
        String service2Text = restTemplate.getForObject("http://localhost:9091/text",  
                                                       String.class);  
        return "Hello " + service2Text;  
    }  
  
    public String getTextFallback() {  
        return "Hello Hystrix";  
    }  
  
    @Bean  
    RestTemplate getRestTemplate() {  
        return new RestTemplate();  
    }  
}
```

Set timeout to 4 seconds

# Setting the timeout

```
@RestController  
public class ServiceTwoController {  
  
    @RequestMapping("/text")  
    public String getText() throws InterruptedException {  
        Thread.sleep(5000);  
        return "World";  
    }  
}
```

Sleep of 5 seconds



# **RESILIENCE: RESILIENCE4J**

# dependency



## pom.xml

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-circuitbreaker-resilience4j</artifactId>
  <version>1.0.2.RELEASE</version>
</dependency>
```

# Using the circuit breaker

```
@RestController
public class ServiceOneController {
    @Autowired
    RestTemplate restTemplate;

    @Autowired
    private CircuitBreakerFactory circuitBreakerFactory;

    @RequestMapping("/text")
    public String getText() {
        CircuitBreaker circuitBreaker = circuitBreakerFactory.create("circuitbreaker");

        String service2Text = circuitBreaker.run(() -> restTemplate.getForObject("http://localhost:9091/text",
            String.class), throwable -> getFallbackName());
        return "Hello " + service2Text;
    }

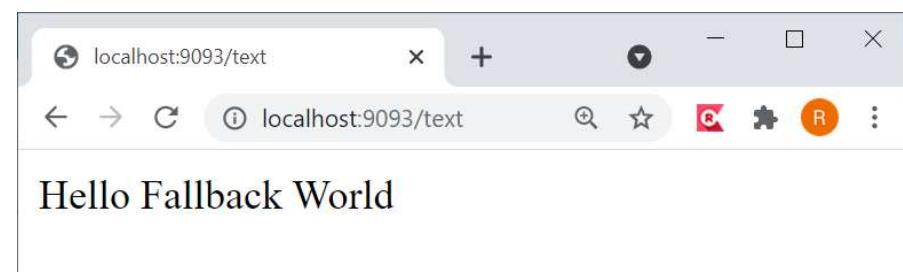
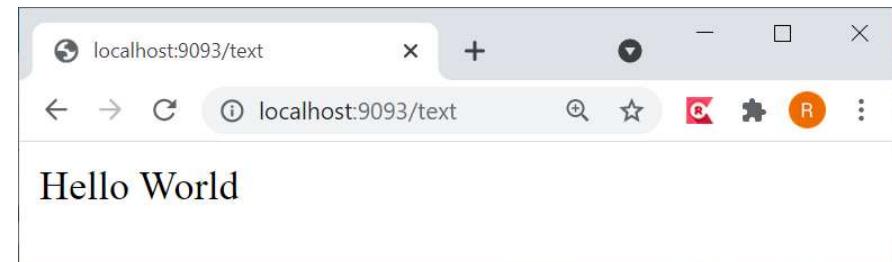
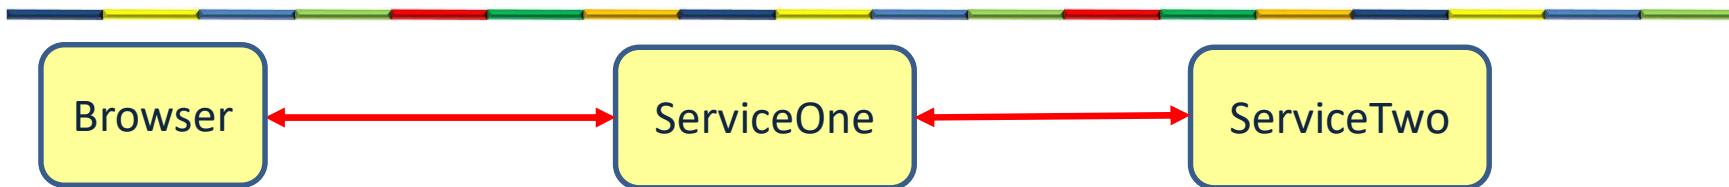
    private String getFallbackName() {
        return "Fallback World";
    }

    @Bean
    RestTemplate getRestTemplate() {
        return new RestTemplate();
    }
}
```

If this method throws an exception or takes longer than **2 seconds**, call the fallback method

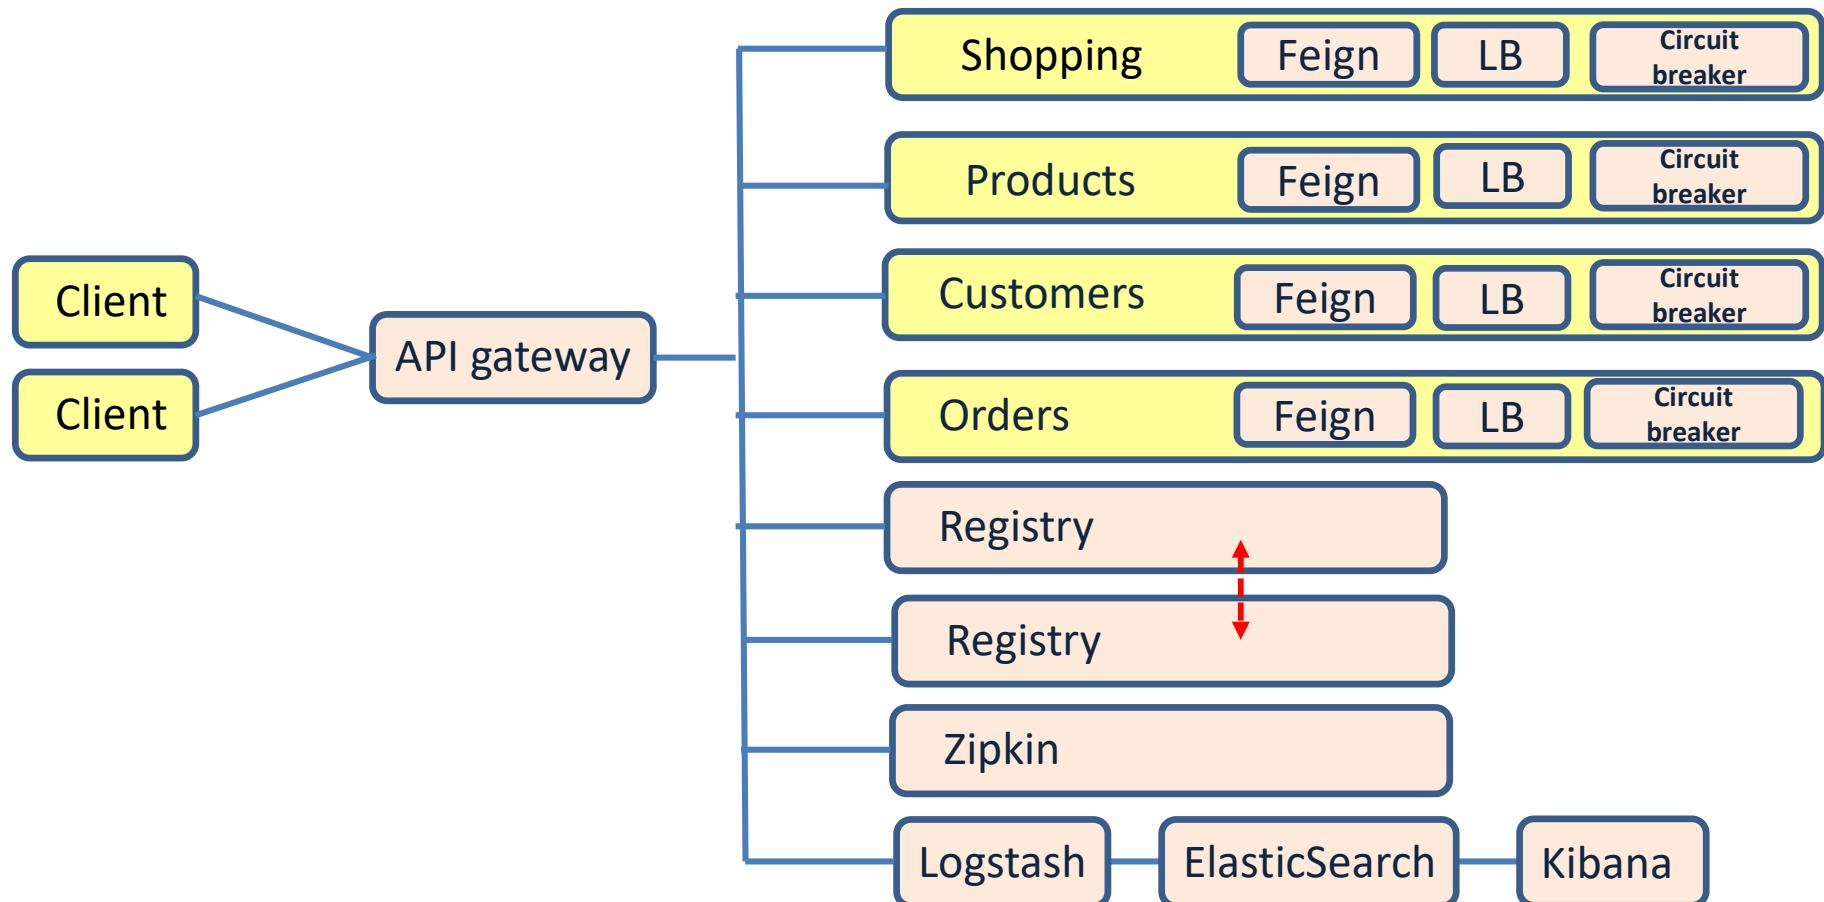
Fallback method

# Using Resilience4J



# Implementing microservices

---



# Challenges of a microservice architecture

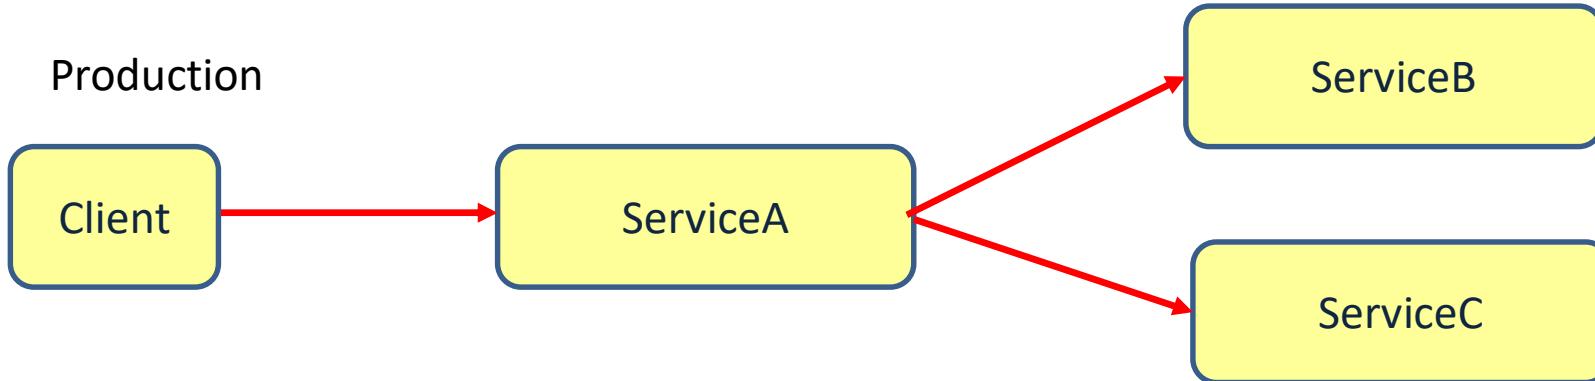
Challenge	Solution
Complex communication	Feign Registry API gateway
Performance	
Resilience	Registry replicas Load balancing between multiple service instances Circuit breaker
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	Zipkin ELK

# **SPRING CLOUD CONTRACT**

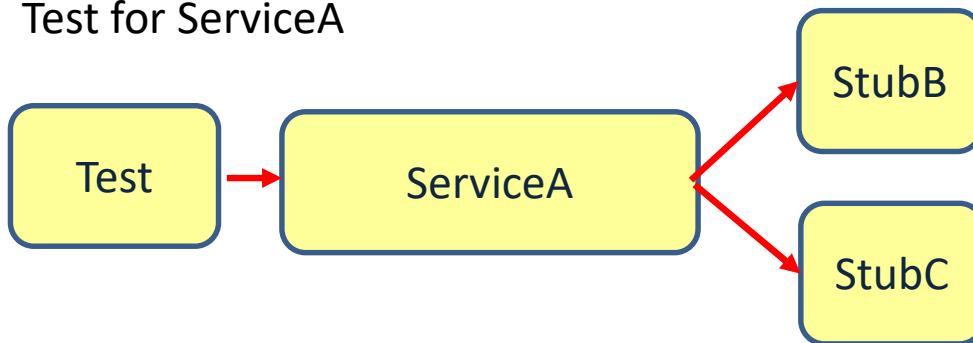
# How to test microservices

---

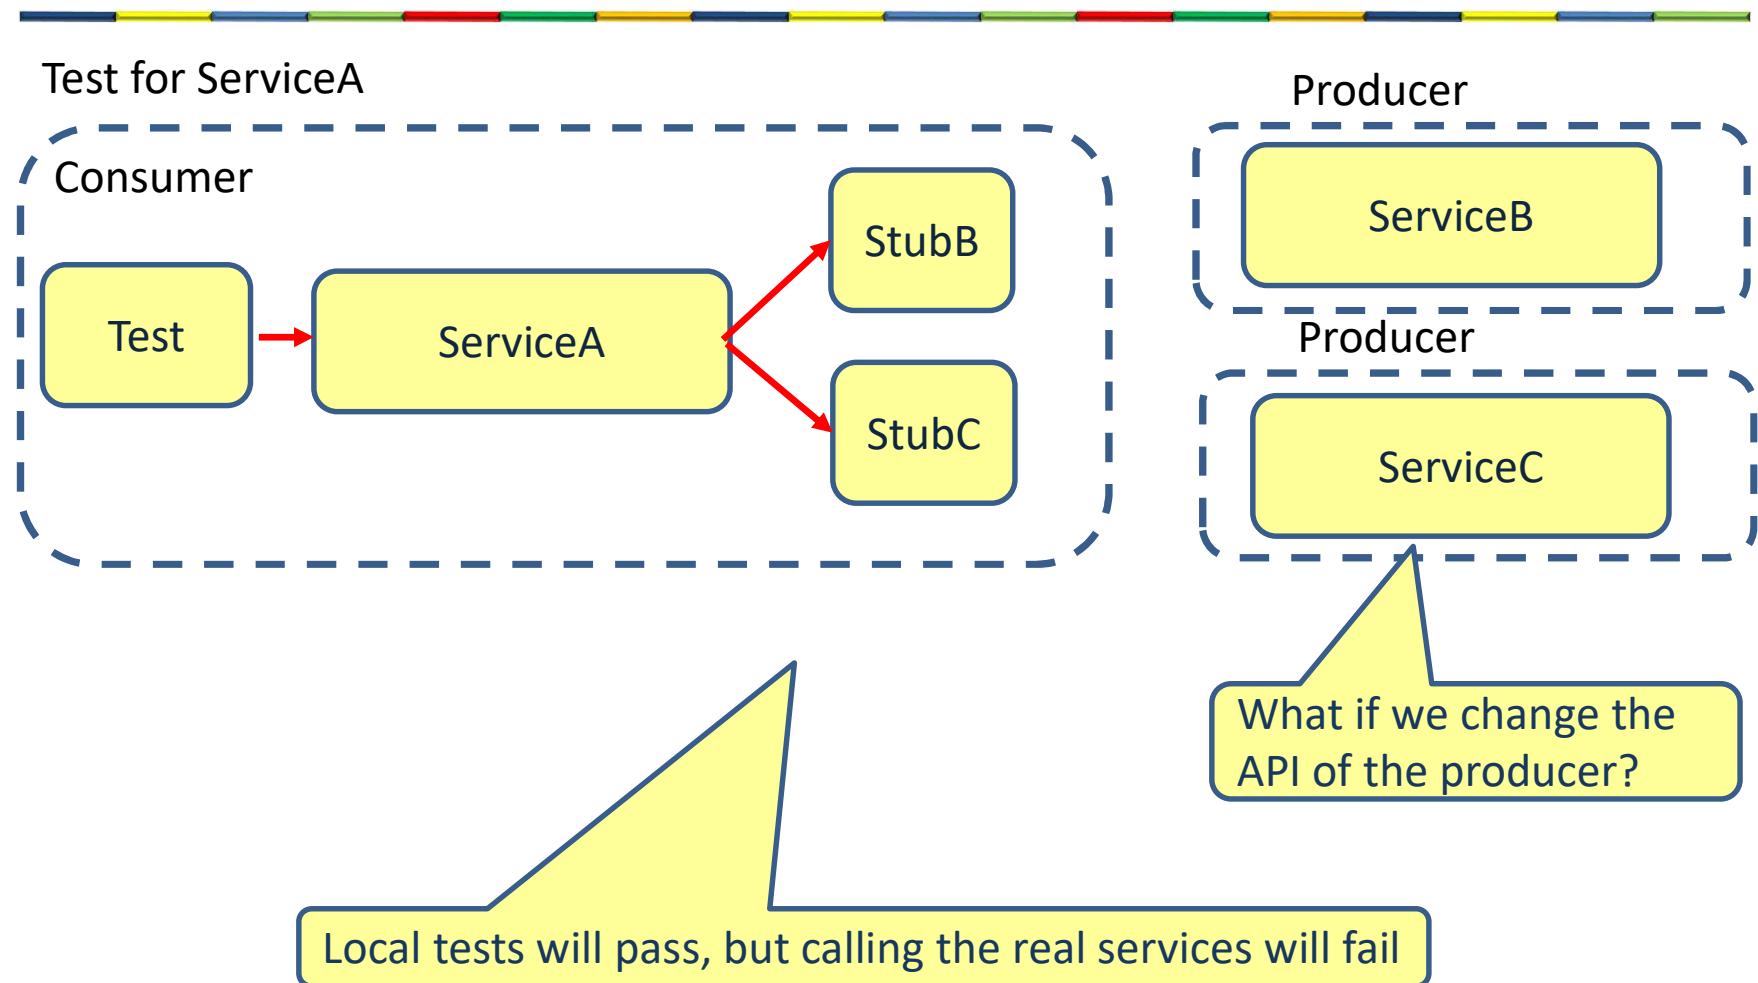
Production



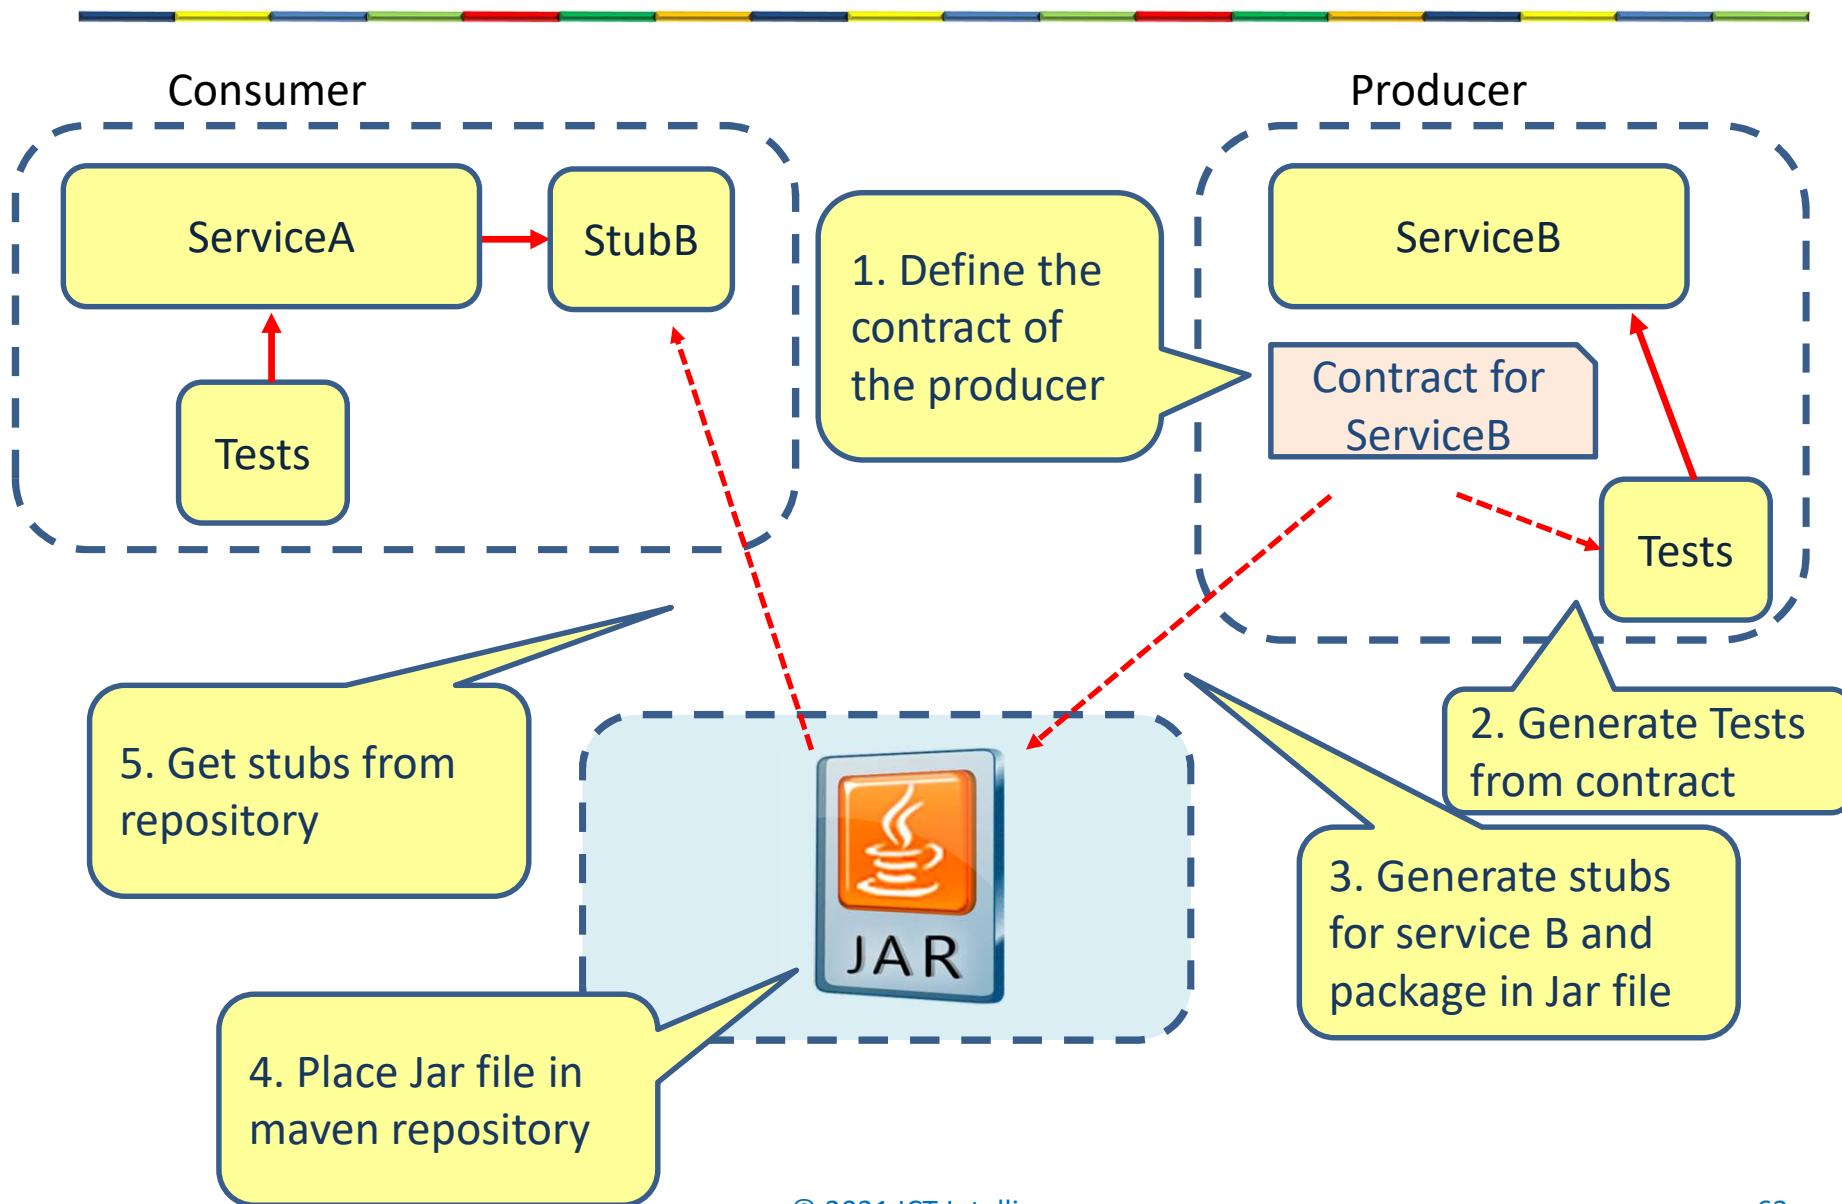
Test for ServiceA



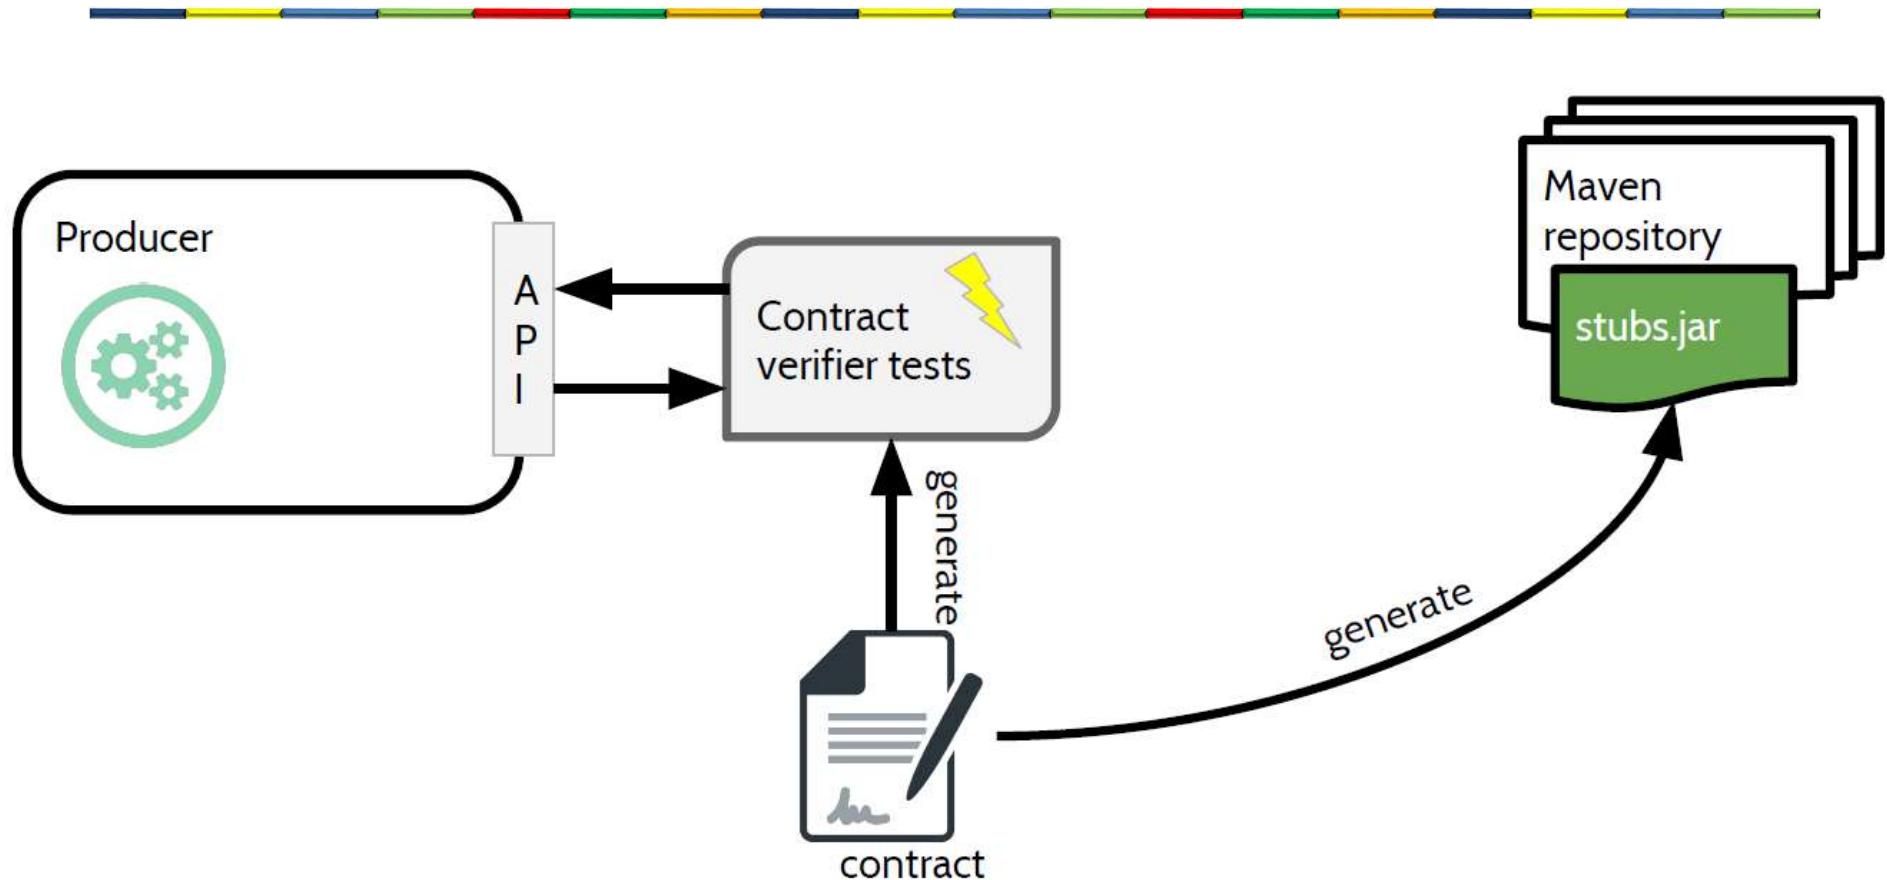
# Stubs live at the consumer



# Spring cloud contracts



# Producer



# Producer maven configuration

---

```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-contract-verifier</artifactId>
  <scope>test</scope>
</dependency>

<plugin>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-contract-maven-plugin</artifactId>
  <version>2.2.2.RELEASE</version>
  <extensions>true</extensions>
  <configuration>
    <baseClassForTests>service.BaseTestClass</baseClassForTests>
    <testFramework>JUNIT5</testFramework>
  </configuration>
</plugin>
```

# Producer

---

```
@RestController
public class EvenOddController {

    @GetMapping("/validate")
    public String evenOrOdd(@RequestParam("number") Integer number) {
        return number % 2 == 0 ? "Even" : "Odd";
    }
}
```

```
@SpringBootApplication
public class EvenoddServiceApplication {

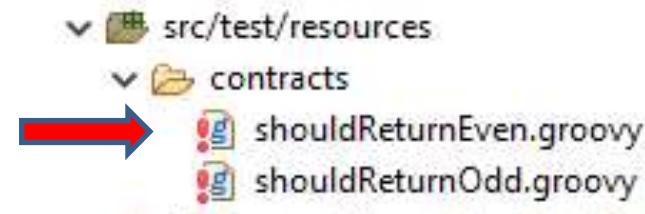
    public static void main(String[] args) {
        SpringApplication.run(EvenoddServiceApplication.class, args);
    }
}
```

# Producer contract 1

```
import org.springframework.cloud.contract.spec.Contract

Contract.make {
    description "should return even when number input is even"
    request{
        method GET()
        url("/validate") {
            queryParameters {
                parameter("number", "2")
            }
        }
    }
    response {
        body("Even")
        status 200
    }
}
```

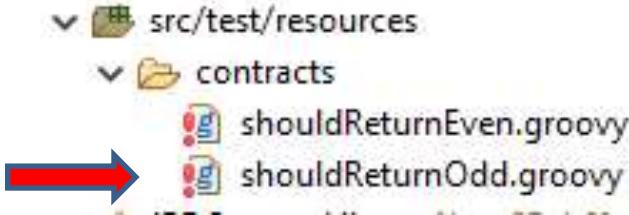
Contract in groovy



# Producer contract 2

```
import org.springframework.cloud.contract.spec.Contract

Contract.make {
    description "should return odd when number input is odd"
    request {
        method GET()
        url("/validate") {
            queryParameters {
                parameter("number", "1")
            }
        }
    }
    response {
        body("Odd")
        status 200
    }
}
```



# Producer: base test class

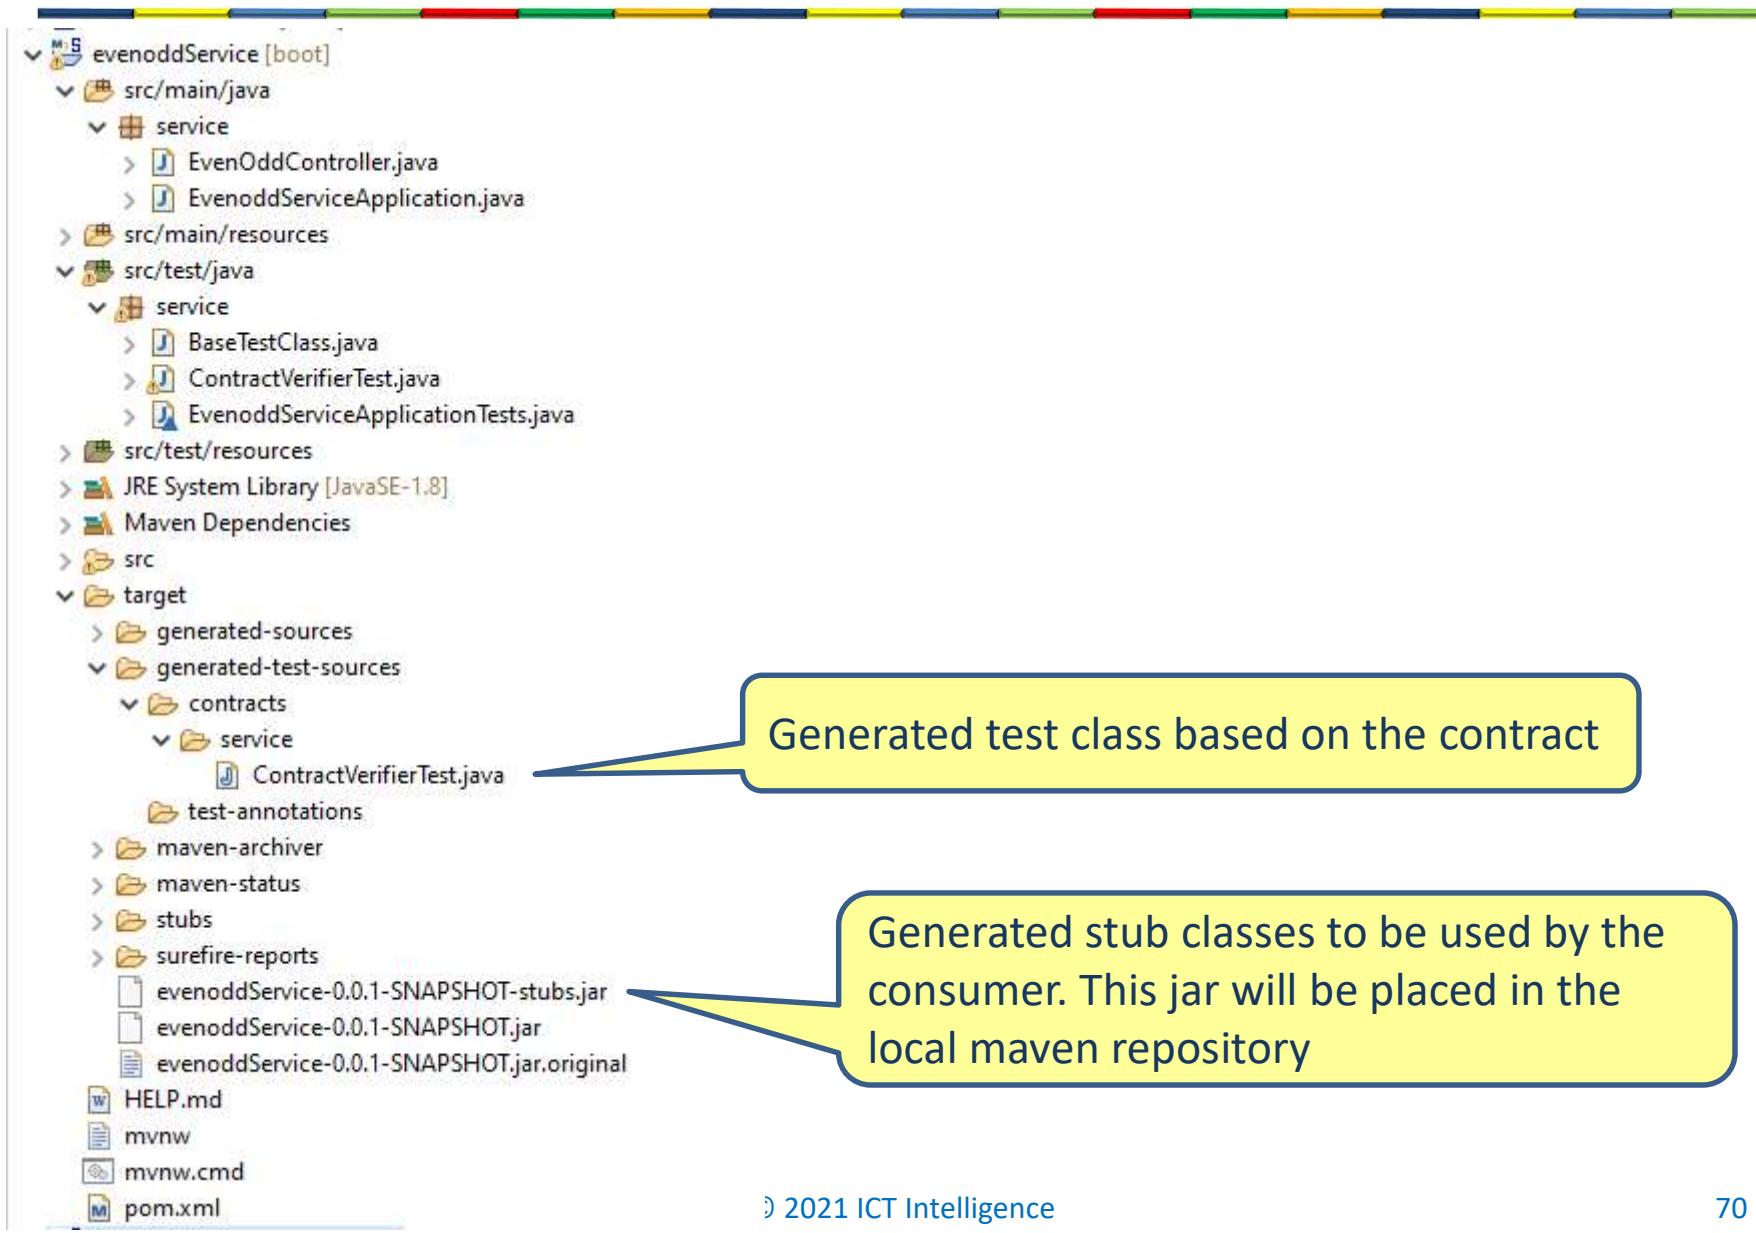
```
@SpringBootTest(webEnvironment = SpringBootTest.WebEnvironment.MOCK)
@DirtiesContext
@AutoConfigureMessageVerifier
public class BaseTestClass {

    @Autowired
    private EvenOddController evenOddController;

    @BeforeEach
    public void setup() {
        StandaloneMockMvcBuilder standaloneMockMvcBuilder
            = MockMvcBuilders.standaloneSetup(evenOddController);
        RestAssuredMockMvc.standaloneSetup(standaloneMockMvcBuilder);
    }
}
```

This is the base class for all to be generated test classes

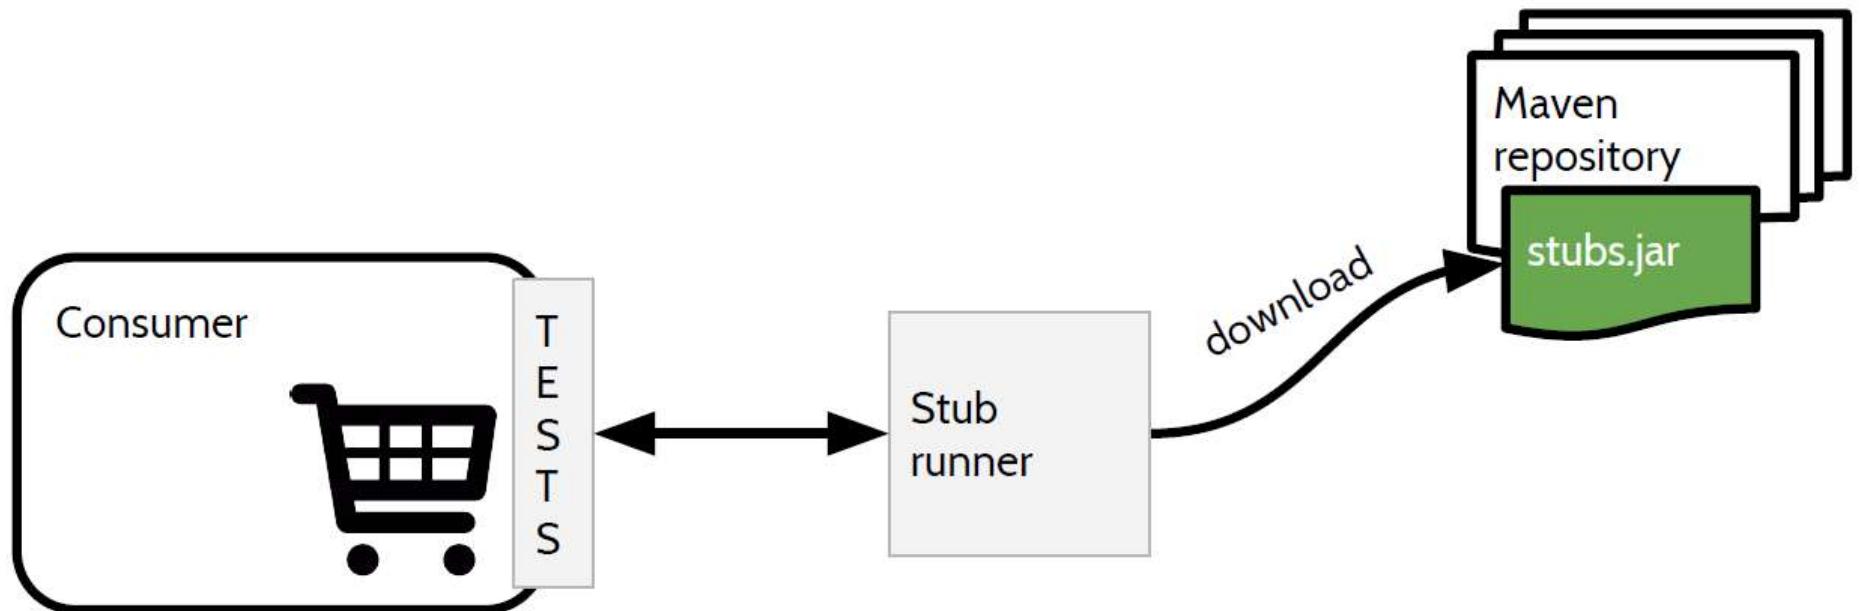
# After running maven install



# Spring cloud contract DSL

```
import org.springframework.cloud.contract.spec.Contract
Contract.make {
    description("GET employee with id=1")
    request {
        method 'GET'
        url '/employee/1'
    }
    response {
        status 200
        body("""
{
    "id": "1",
    "fname": "Jane",
    "lname": "Doe",
    "salary": "123000.00",
    "gender": "M"
}
""")
        headers {
            contentType(applicationJson())
        }
    }
}
```

# Consumer



```
<dependency>
  <groupId>org.springframework.cloud</groupId>
  <artifactId>spring-cloud-starter-contract-stub-runner</artifactId>
  <version>2.2.2.RELEASE</version>
  <scope>test</scope>
</dependency>
```

# Consumer

```
@RestController
public class MathController {

    @Autowired
    private RestTemplate restTemplate;

    @GetMapping("/calculate")
    public String checkOddAndEven(@RequestParam("number") Integer number) {
        HttpHeaders httpHeaders = new HttpHeaders();
        httpHeaders.add("Content-Type", "application/json");

        ResponseEntity<String> responseEntity = restTemplate.exchange(
            "http://localhost:8090/validate?number=" + number,
            HttpMethod.GET,
            new HttpEntity<>(httpHeaders),
            String.class);

        return responseEntity.getBody();
    }
}
```

# Consumer

---

```
@SpringBootApplication
public class MathServiceApplication {

    public static void main(String[] args) {
        SpringApplication.run(MathServiceApplication.class, args);
    }

    @Bean
    RestTemplate restTemplate() {
        return new RestTemplate();
    }
}
```

# Consumer test

Get the stubs from  
the local repository

```
@SpringBootTest(webEnvironment = SpringBootTest.WebEnvironment.MOCK)
@Autowired
@AutoConfigureMockMvc
@AutoConfigureJsonTesters
@AutoConfigureStubRunner(stubsMode = StubRunnerProperties.StubsMode.LOCAL,
    ids = "com.acme:evenoddService:+:stubs:8090")
public class MathControllerIntegrationTest {

    @Autowired
    private MockMvc mockMvc;

    @Test
    public void given_WhenPassEvenNumberInQueryParam_ThenReturnEven() throws Exception {
        mockMvc.perform(MockMvcRequestBuilders.get("/calculate?number=2")
            .contentType(MediaType.APPLICATION_JSON))
            .andExpect(status().isOk())
            .andExpect(content().string("Even"));
    }

    @Test
    public void given_WhenPassOddNumberInQueryParam_ThenReturnOdd() throws Exception {
        mockMvc.perform(MockMvcRequestBuilders.get("/calculate?number=1")
            .contentType(MediaType.APPLICATION_JSON))
            .andExpect(status().isOk())
            .andExpect(content().string("Odd"));
    }
}
```

# Consumer test

```
@SpringBootTest(webEnvironment = SpringBootTest.WebEnvironment.MOCK)
@AutowiredMockMvc
@AutoConfigureJsonTesters
@AutoConfigureStubRunner(stubsMode = StubRunnerProperties.StubsMode.LOCAL,
    ids = "com.acme:evenoddService:+:stubs:8090")
public class MathControllerIntegrationTest {
```

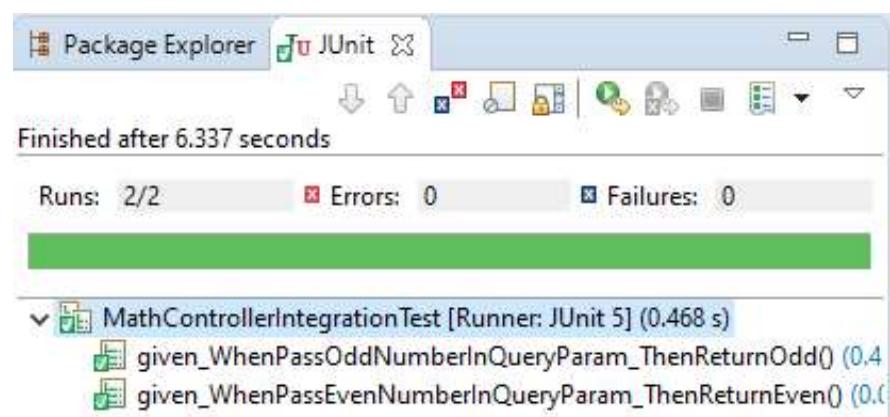
Group id

Artifact id

Version + means latest version

stubs

Port number to run the stubs on



# Challenges of a microservice architecture

Challenge	Solution
Complex communication	Feign Registry API gateway
Performance	
Resilience	Registry replicas Load balancing between multiple service instances Circuit breaker
Security	
Transactions	
Following the process	
Keep data in sync	
Keep interfaces in sync	Spring cloud contract
Keep configuration in sync	
Monitor health of microservices	
Follow/monitor business processes	Zipkin ELK

Lesson 9

# **MICROSERVICES**

# **SERVICE DEPLOYMENT**

# Service deployment

---

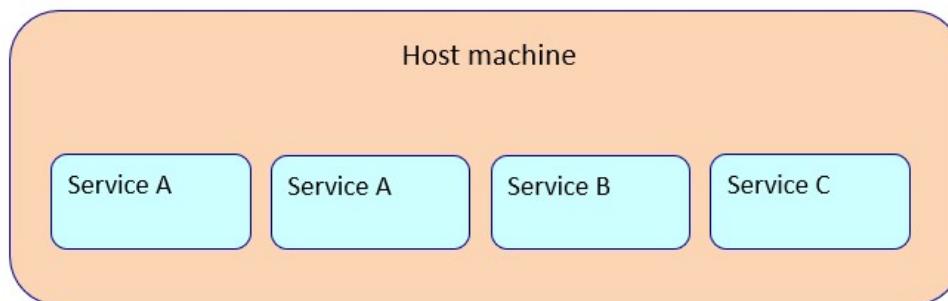
- Service are written using different languages, frameworks, framework versions
- Run multiple service instances of a service for throughput and availability
- Building and deploying should be fast
- Instances need to be isolated
- Constrain the resources a service may consume (CPU, memory, etc.)
- Deployment should be reliable

# Multiple service instances per host

---

- Benefits

- Efficient resource utilization
- Fast deployment



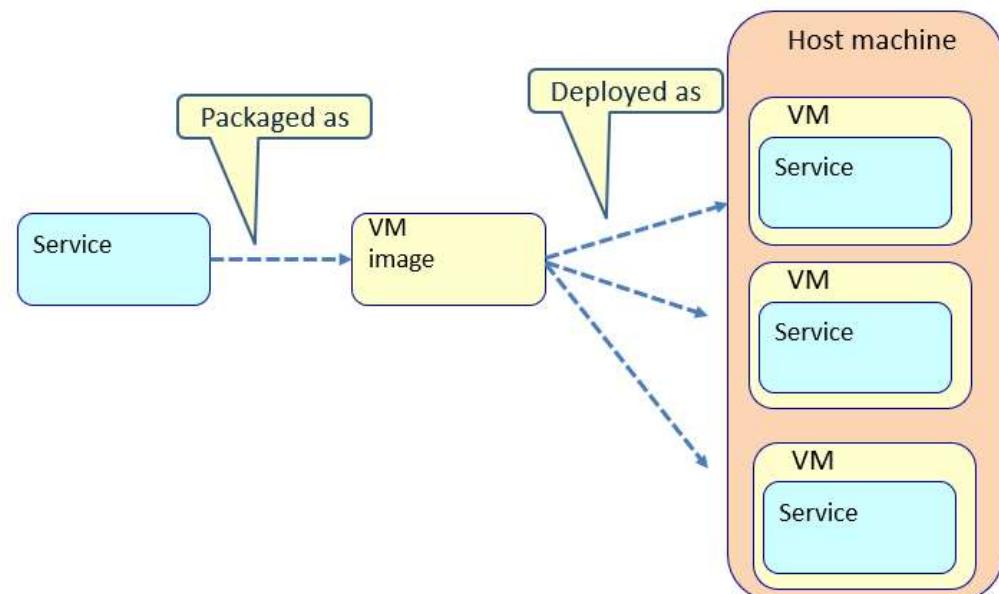
- Drawbacks

- Poor isolation
- Poor visibility of resource utilization
- Difficult to constrain resource utilization
- Risk of dependency version conflicts
- Poor encapsulation of implementation technology

# Service per VM

---

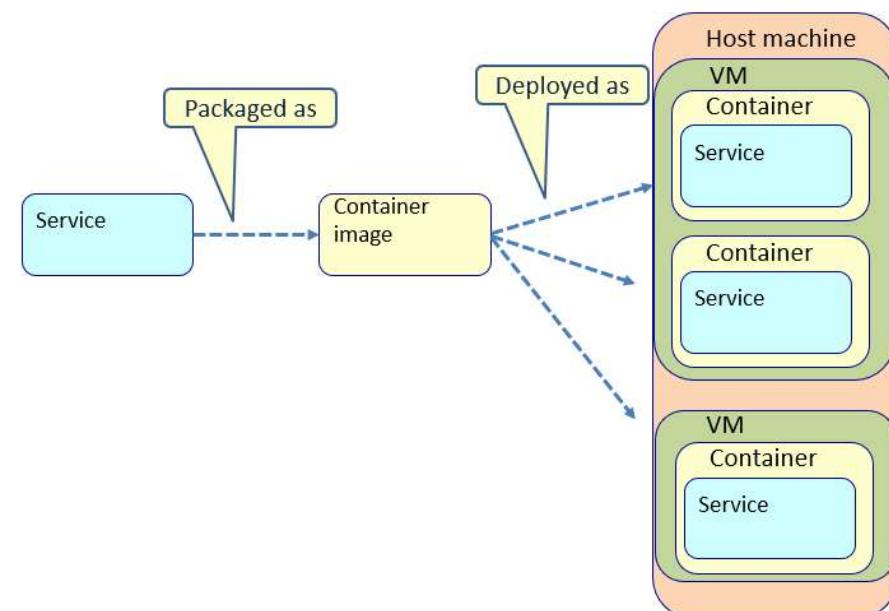
- Benefits
  - Great isolation
  - Great manageability
  - VM encapsulates implementation technology
  - Leverage cloud infrastructure for auto scaling/load balancing
- Drawbacks
  - Less efficient resource utilization
  - Slow deployment



# Service per container

---

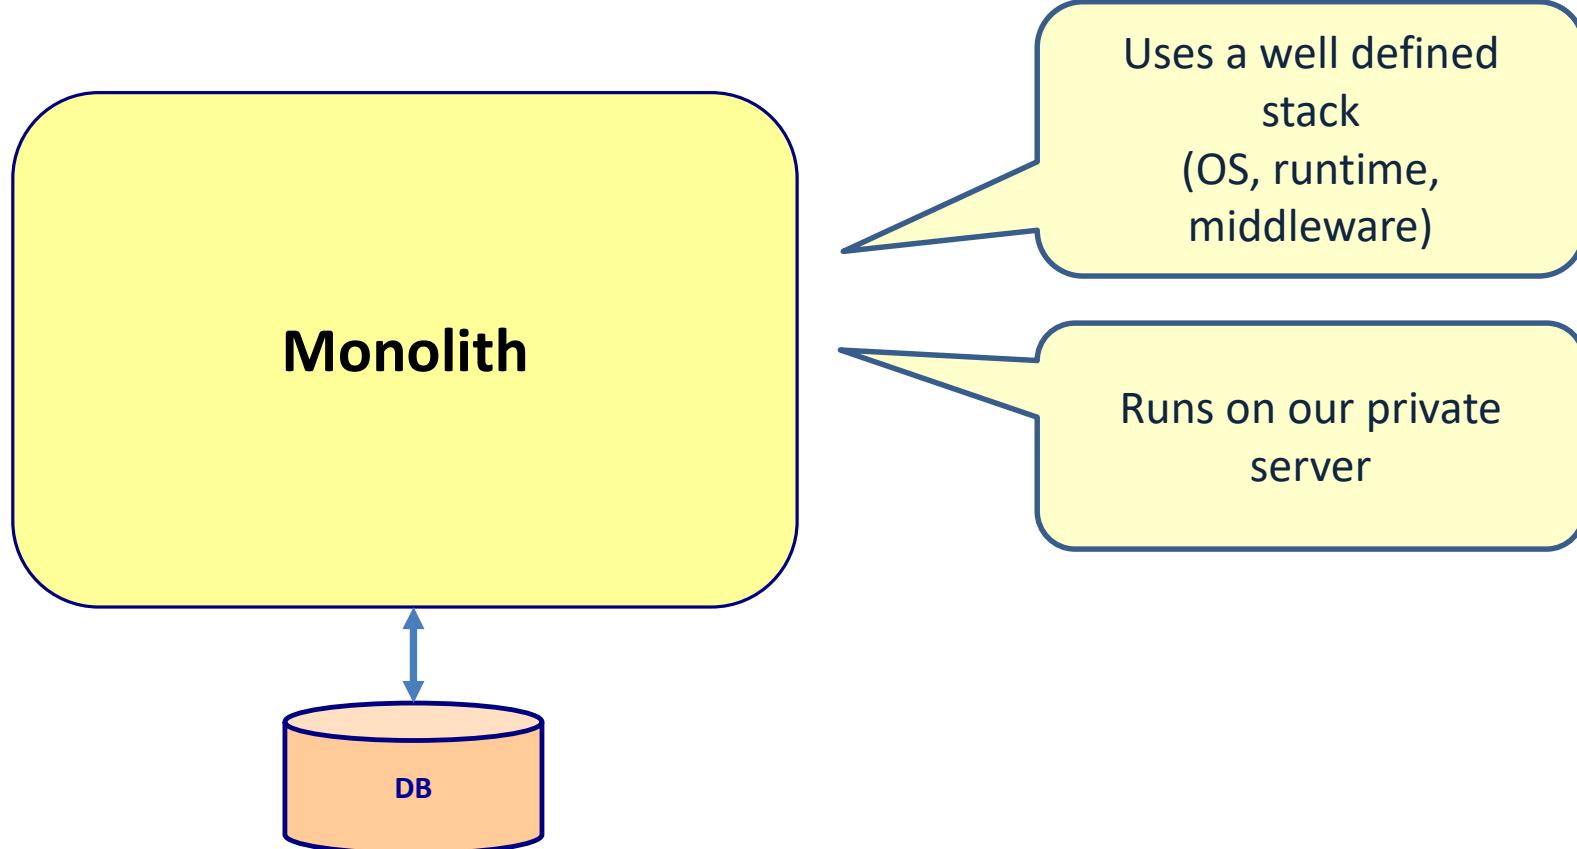
- Benefits
  - Great isolation
  - Great manageability
  - Container encapsulates implementation technology
  - Efficient resource utilization
  - Fast deployment
- Drawbacks
  - Technology is not as mature as VM's
  - Containers are not as secure as VM's



# **CONTAINERS**

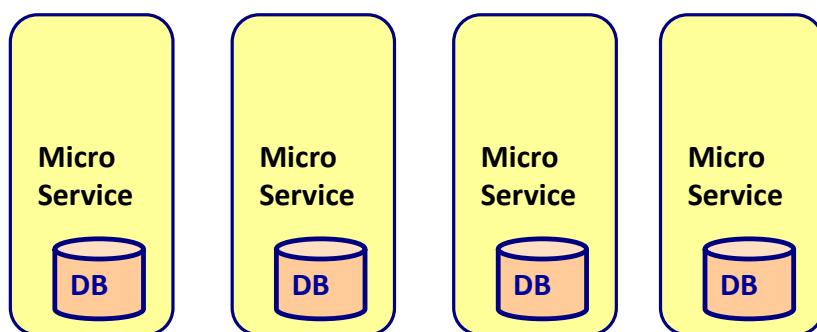
# Monolith

---



# Microservice

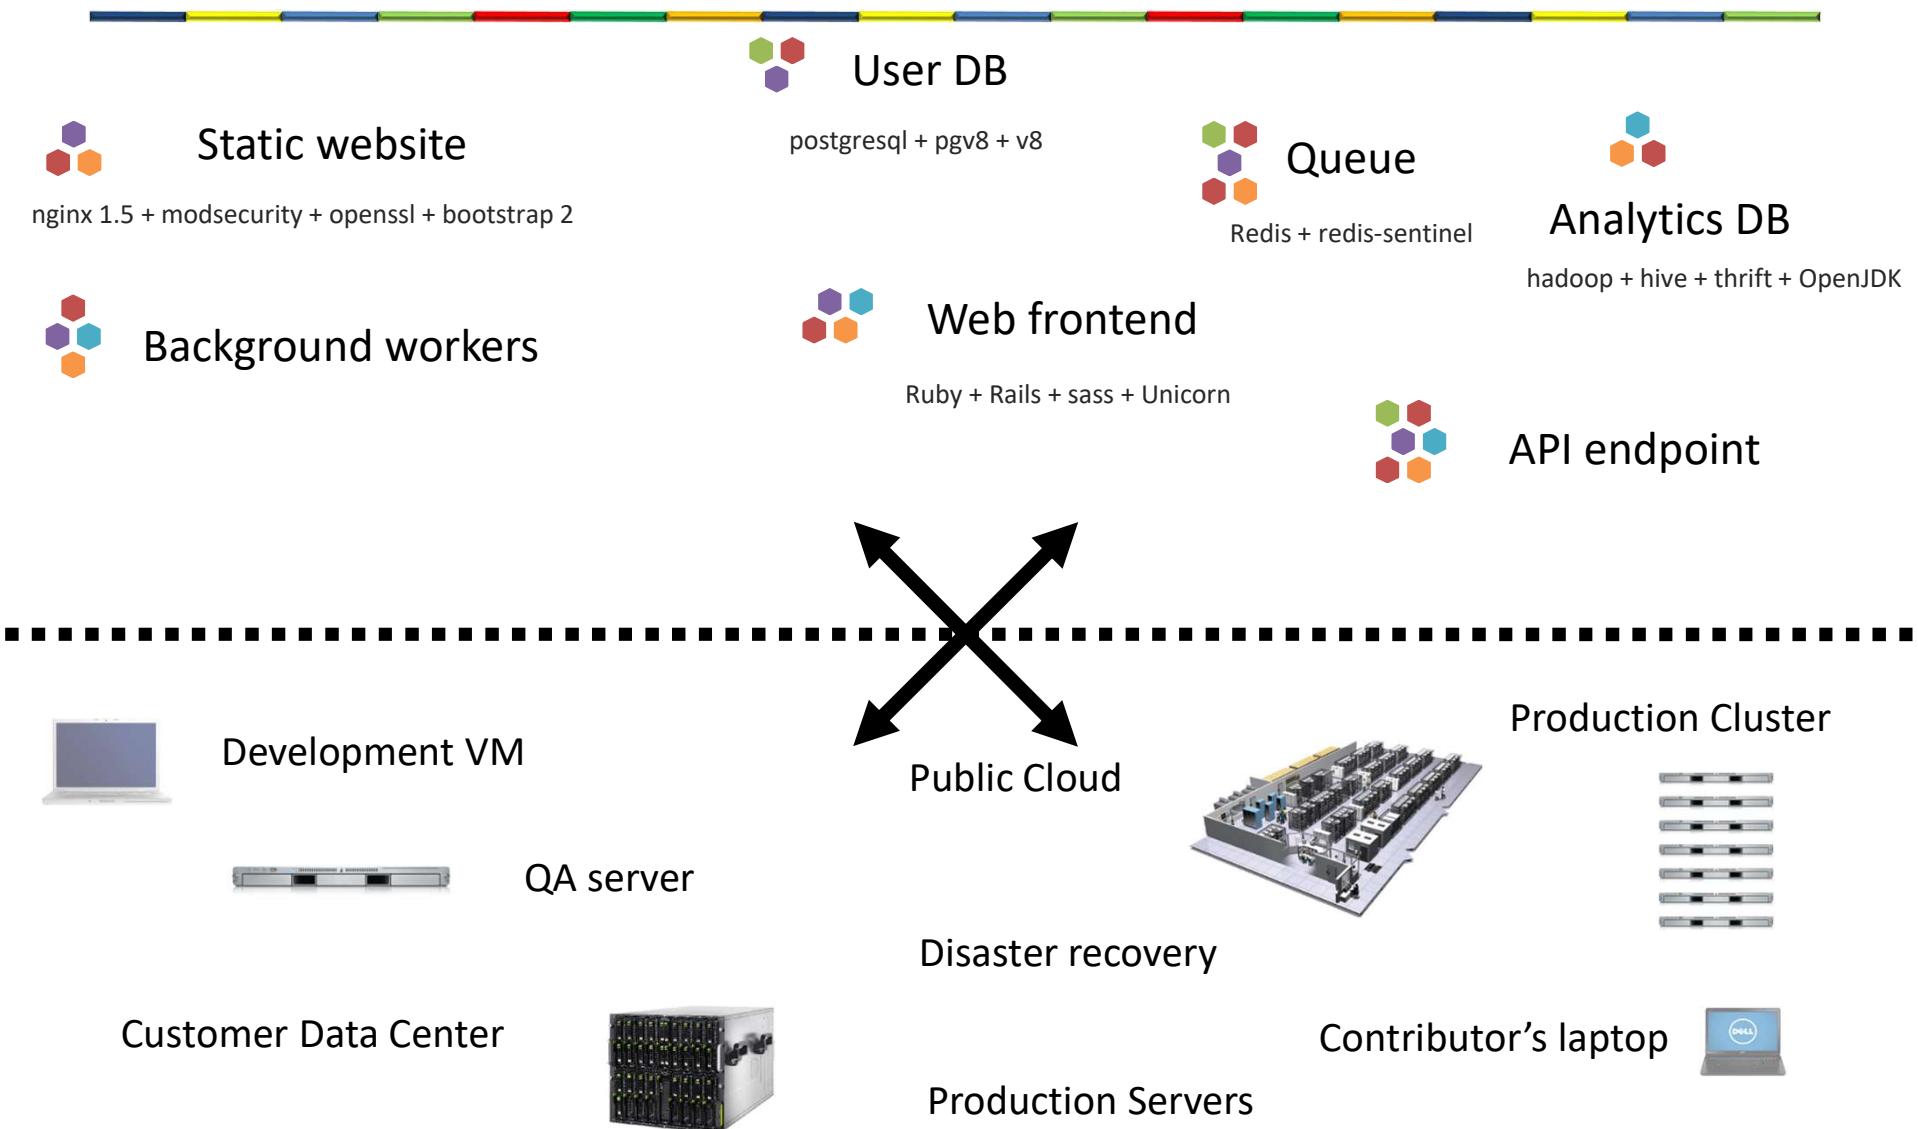
---



Microservices use different stacks. (OS, runtime, middleware)

Runs anywhere (private, cloud)

# The challenge

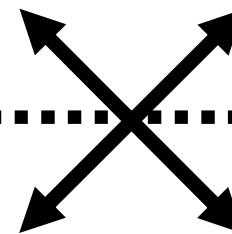


# Matrix from hell

	Development VM	QA Server	Single Prod Server	Onsite Cluster	Public Cloud	Contributor's laptop	Customer Servers
Static website	?	?	?	?	?	?	?
Web frontend	?	?	?	?	?	?	?
Background workers	?	?	?	?	?	?	?
User DB	?	?	?	?	?	?	?
Analytics DB	?	?	?	?	?	?	?
Queue	?	?	?	?	?	?	?



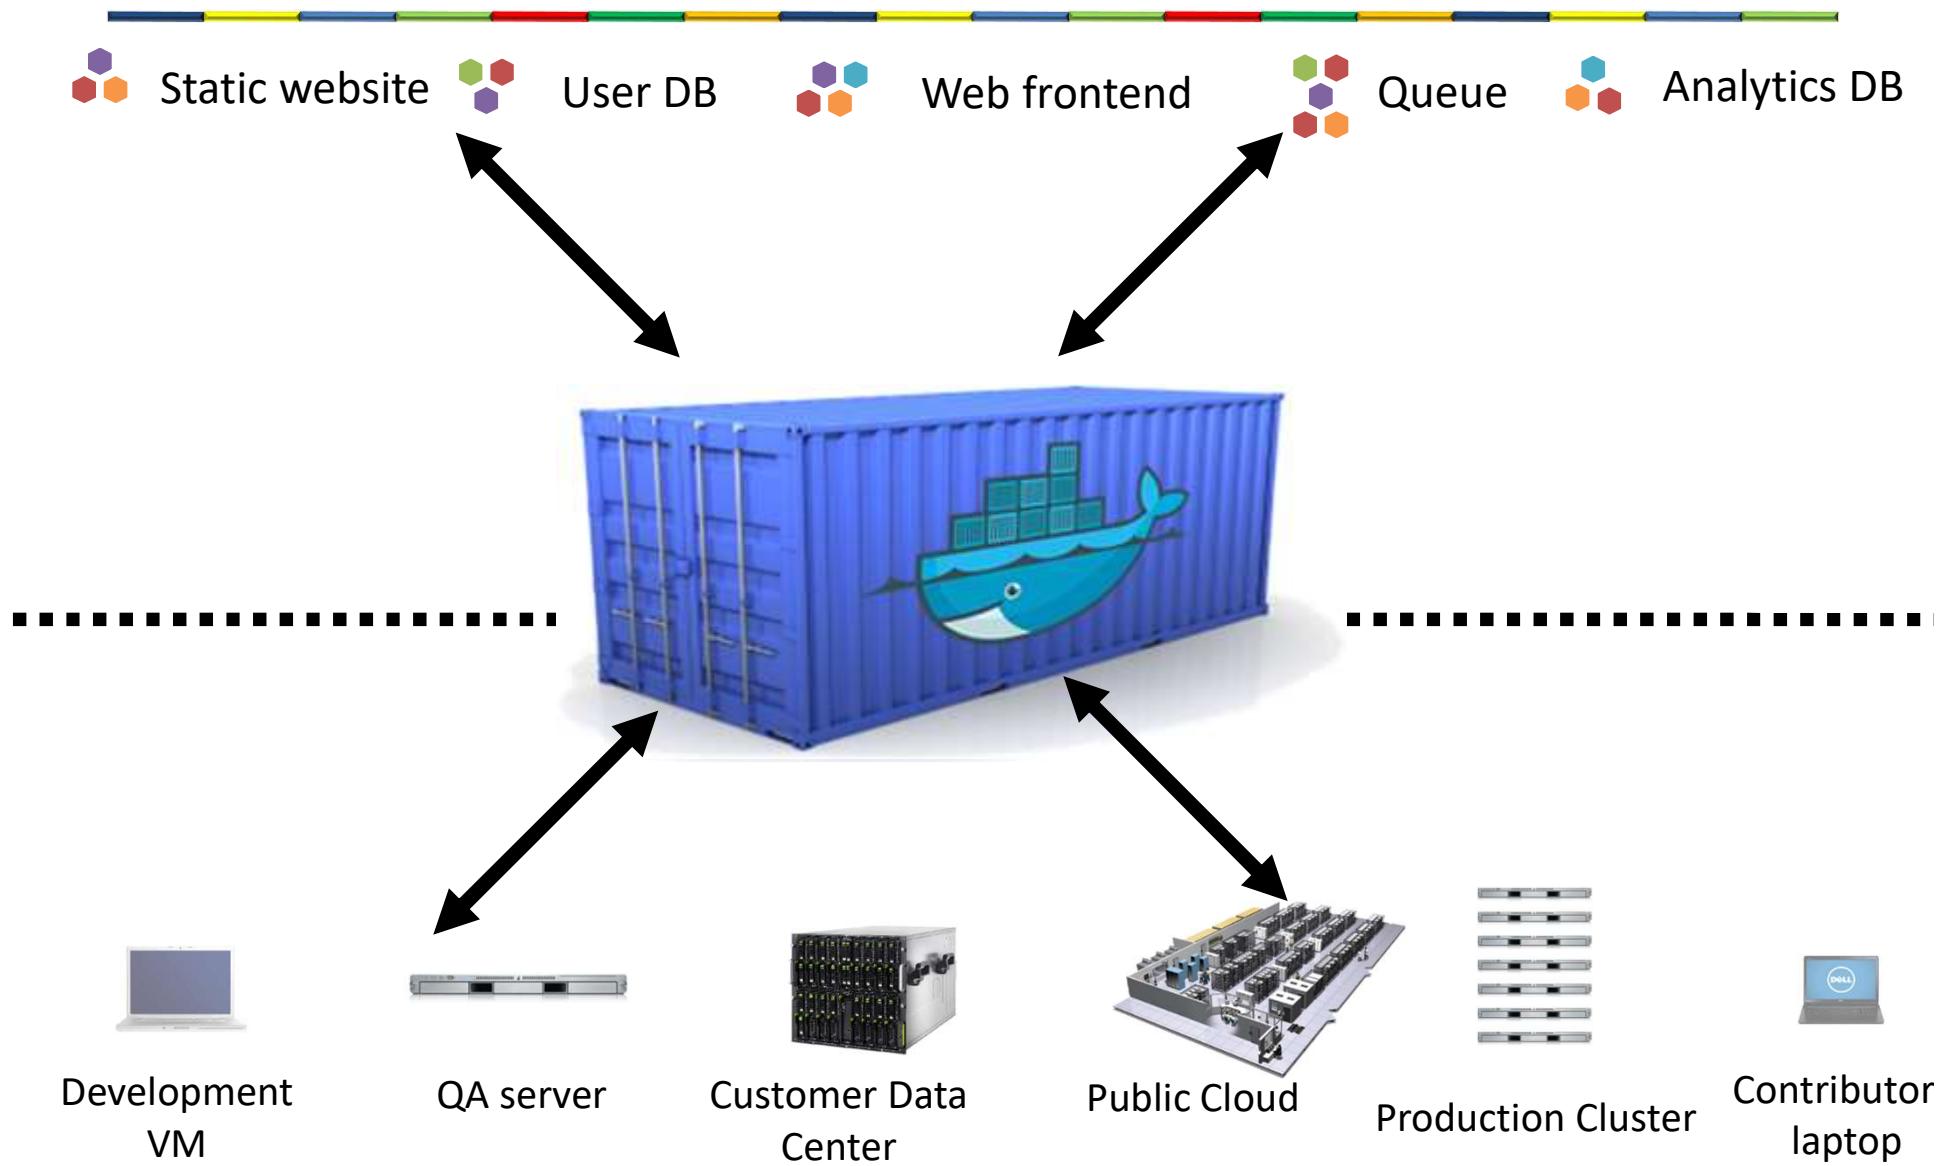
# Cargo Transport Pre-1960



# The solution



# Docker is a shipping container for code



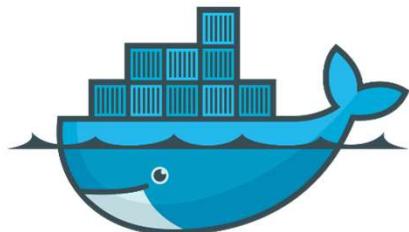
# Separation of concern

---

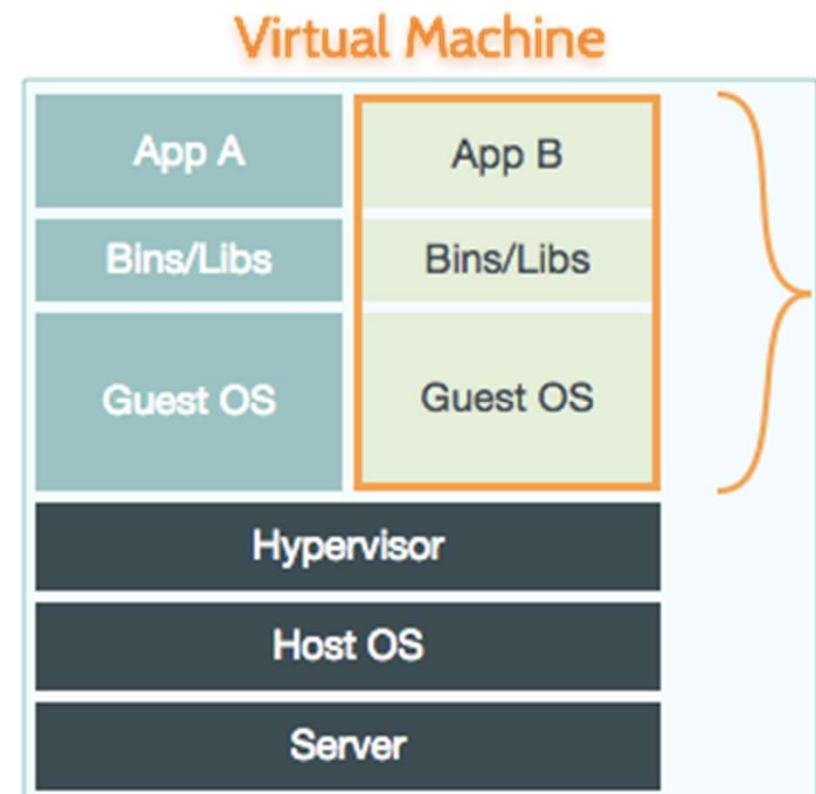
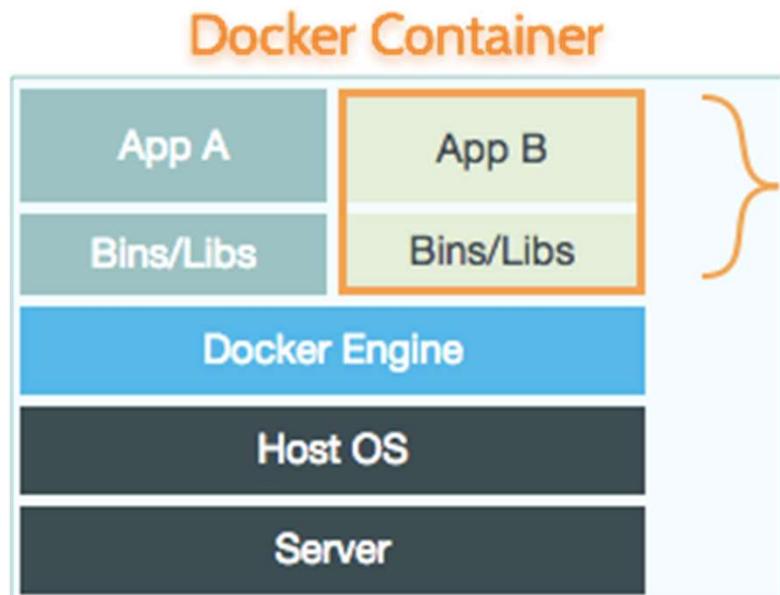
- Dan the Developer
  - Worries about what's "inside" the container
    - His code
    - His Libraries
    - His Package Manager
    - His Apps
    - His Data
- Oscar the Ops Guy
  - Worries about what's "outside" the container
    - Logging
    - Remote access
    - Monitoring
    - Network config
  - All containers start, stop, copy, attach, migrate, etc. the same way

# Containers

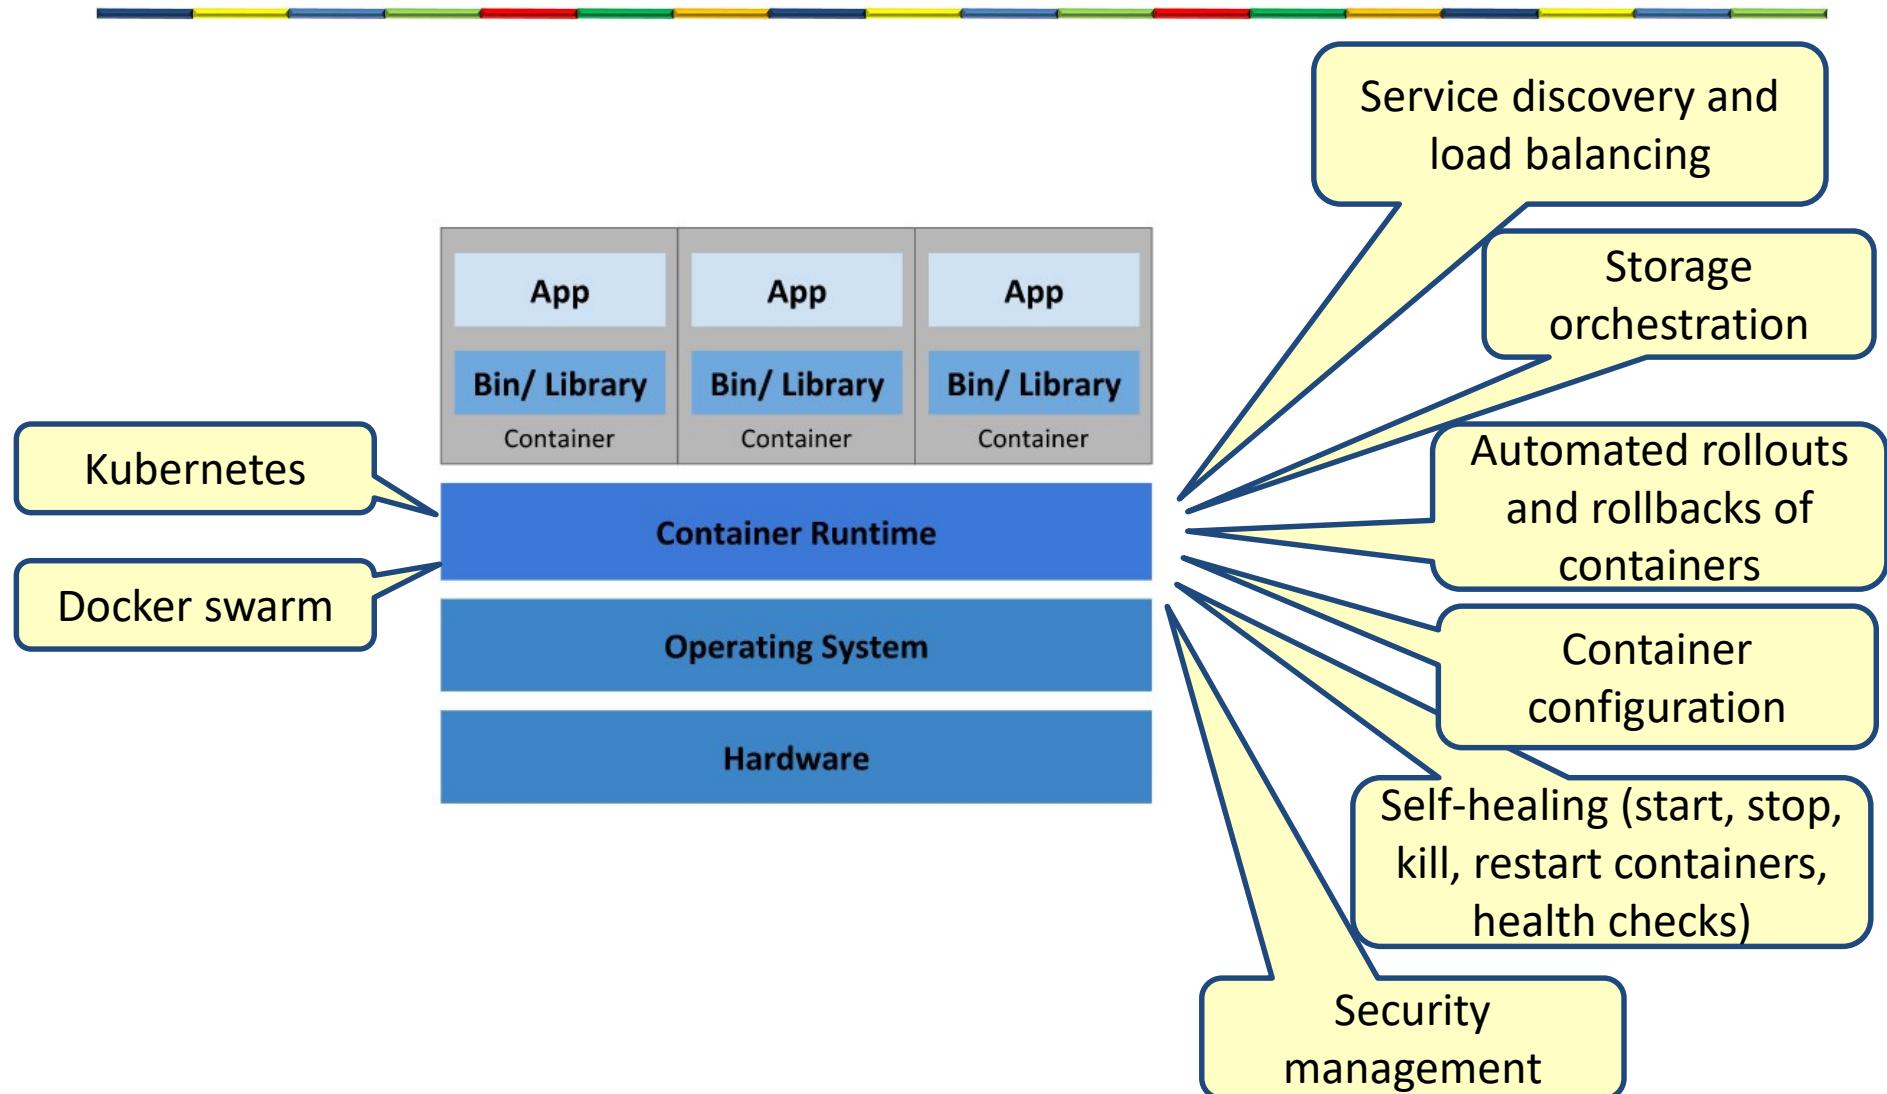
- Docker



docker



# Container management



# CQRS

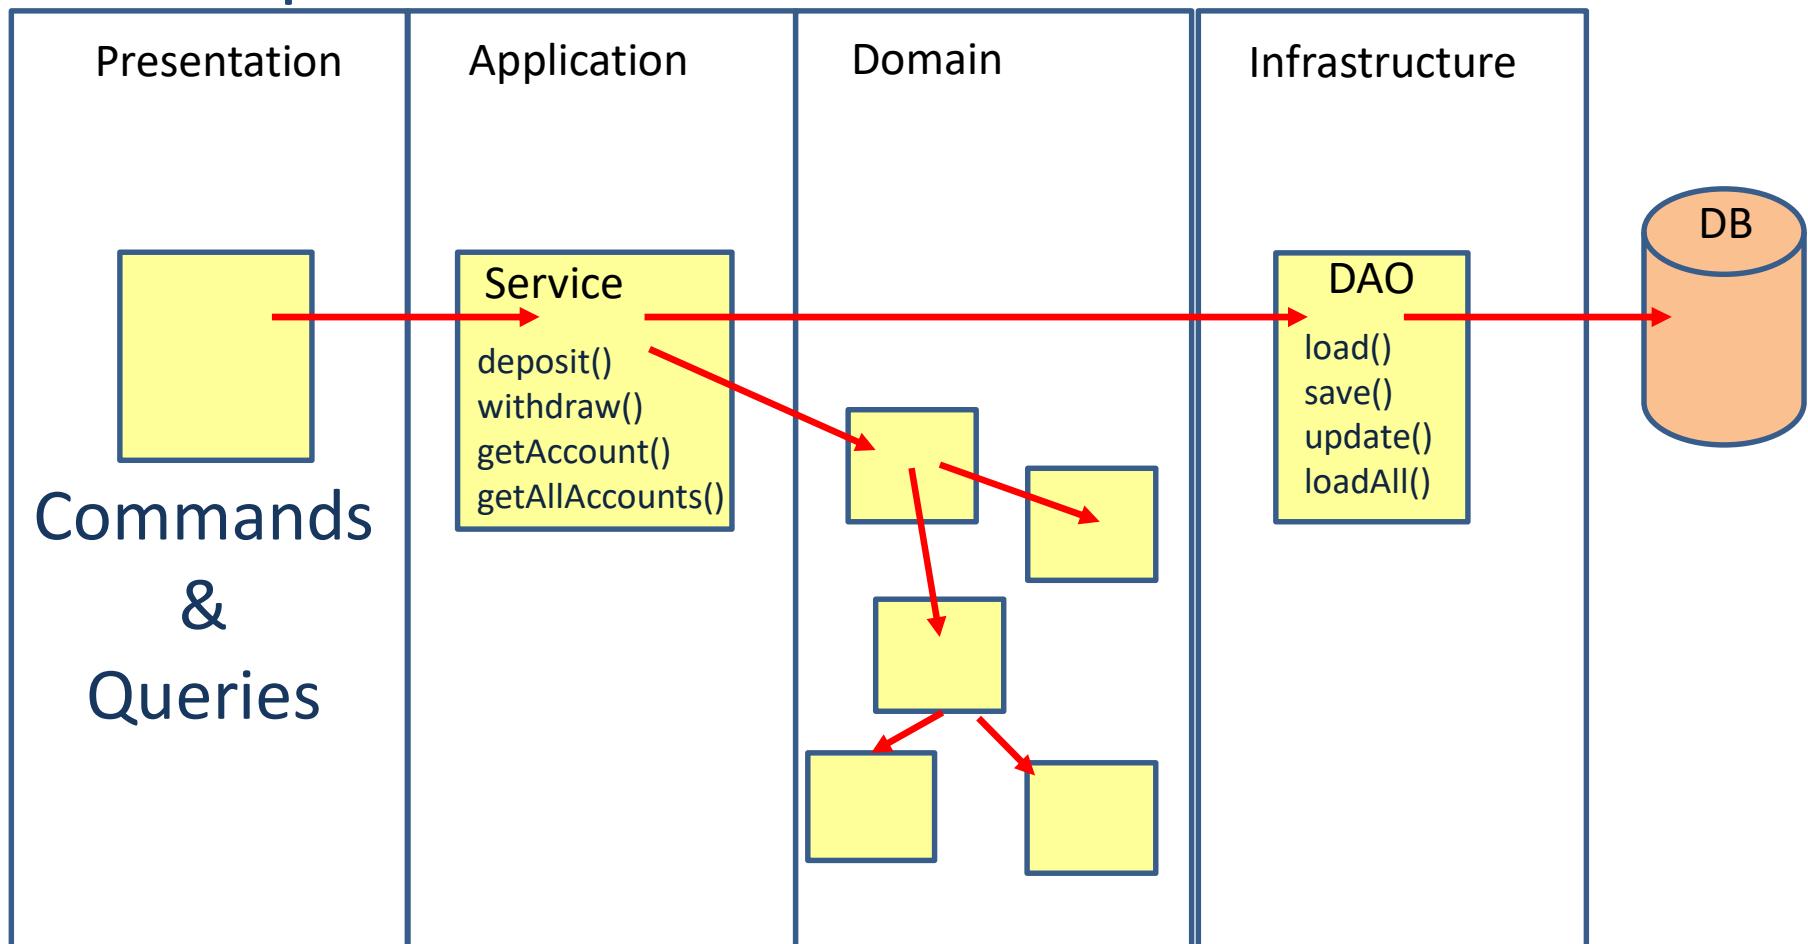
# Command Query Responsibility Segregation (CQRS)

---

- Separates the querying from command processing by providing two models instead of one.
  - One model is built to handle and process commands
  - One model is built for presentation needs (queries)

# Typical architecture

- One domain model that is used for commands and queries

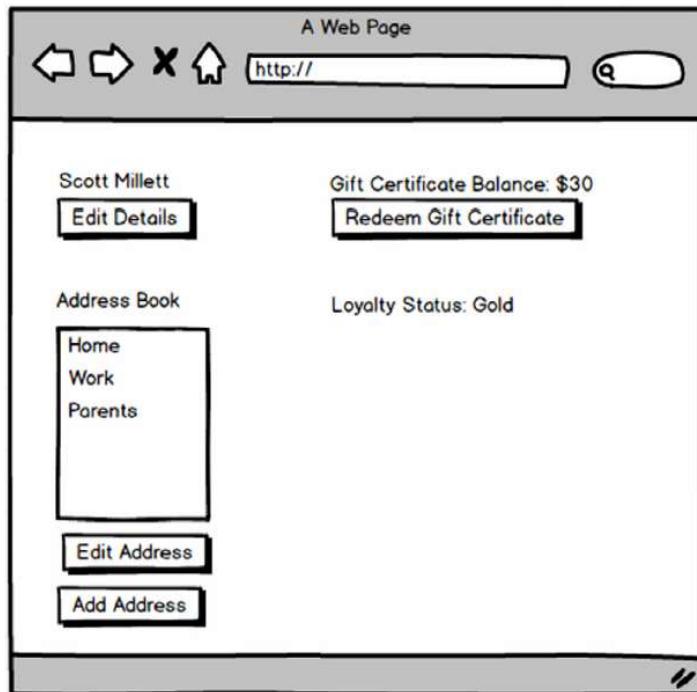


# One model for both commands and queries

---

- To support complex views and reporting
  - Required domain model becomes complex
  - Internal state needs to be exposed
  - Aggregates are merged for view requirements
  - Repositories often contain many extra methods to support presentation needs such as paging, querying, and free text searching
- Result: single model that is full of compromises

# Example of complex aggregates



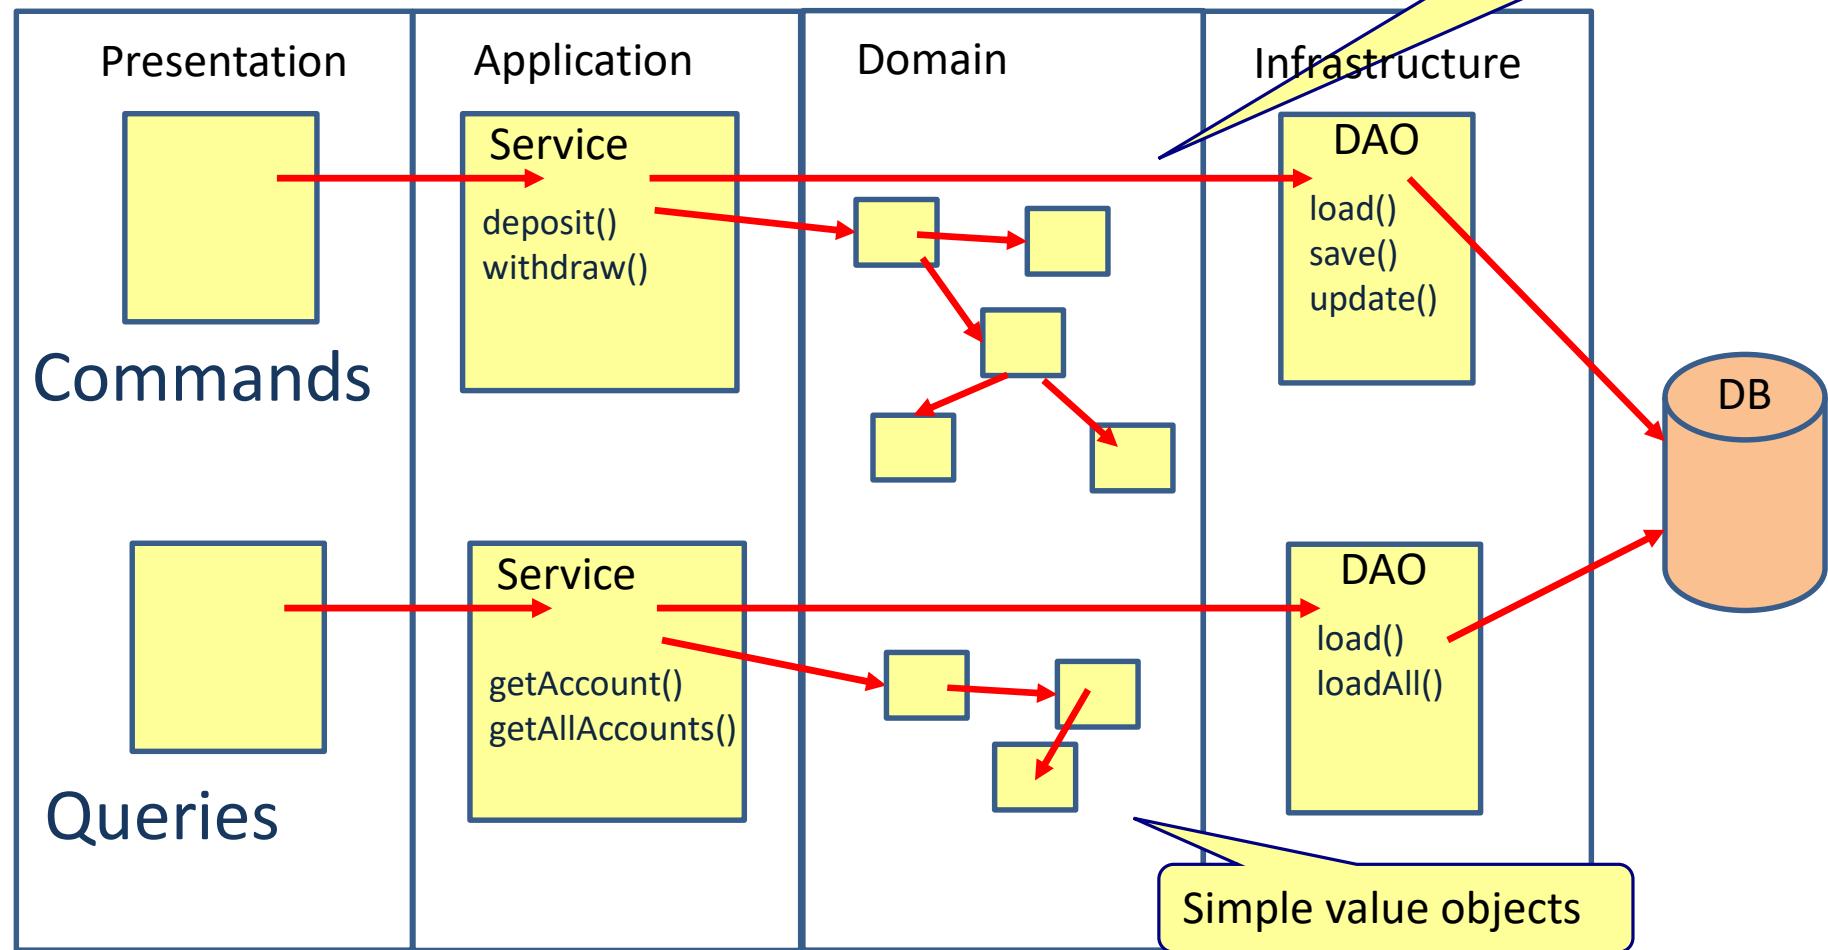
Complex aggregate  
because of UI needs

We don't need this  
complexity for commands

```
public class Customer
{
    // ...
    public ContactDetails ContactDetails { get; private set; }
    public LoyaltyStatus LoyaltyStatus { get; private set; }
    public Money GiftCertBalance { get; private set; }
    public IEnumerable<Address> AddressBook { get; private set; }
}
```

# CQRS

- Domain model is used for commands
- View model is used for queries



# 2 services instead of one

---

## Traditional service

```
CustomerService
void MakeCustomerPreferred(CustomerId)
Customer GetCustomer(CustomerId)
CustomerSet GetCustomersWithName(Name)
CustomerSet GetPreferredCustomers()
void ChangeCustomerLocale(CustomerId, NewLocale)
void CreateCustomer(Customer)
void EditCustomerDetails(CustomerDetails)
```



## Service with CQRS

```
CustomerWriteService
void MakeCustomerPreferred(CustomerId)
void ChangeCustomerLocale(CustomerId, NewLocale)
void CreateCustomer(Customer)
void EditCustomerDetails(CustomerDetails)

CustomerReadService
Customer GetCustomer(CustomerId)
CustomerSet GetCustomersWithName(Name)
CustomerSet GetPreferredCustomers()
```

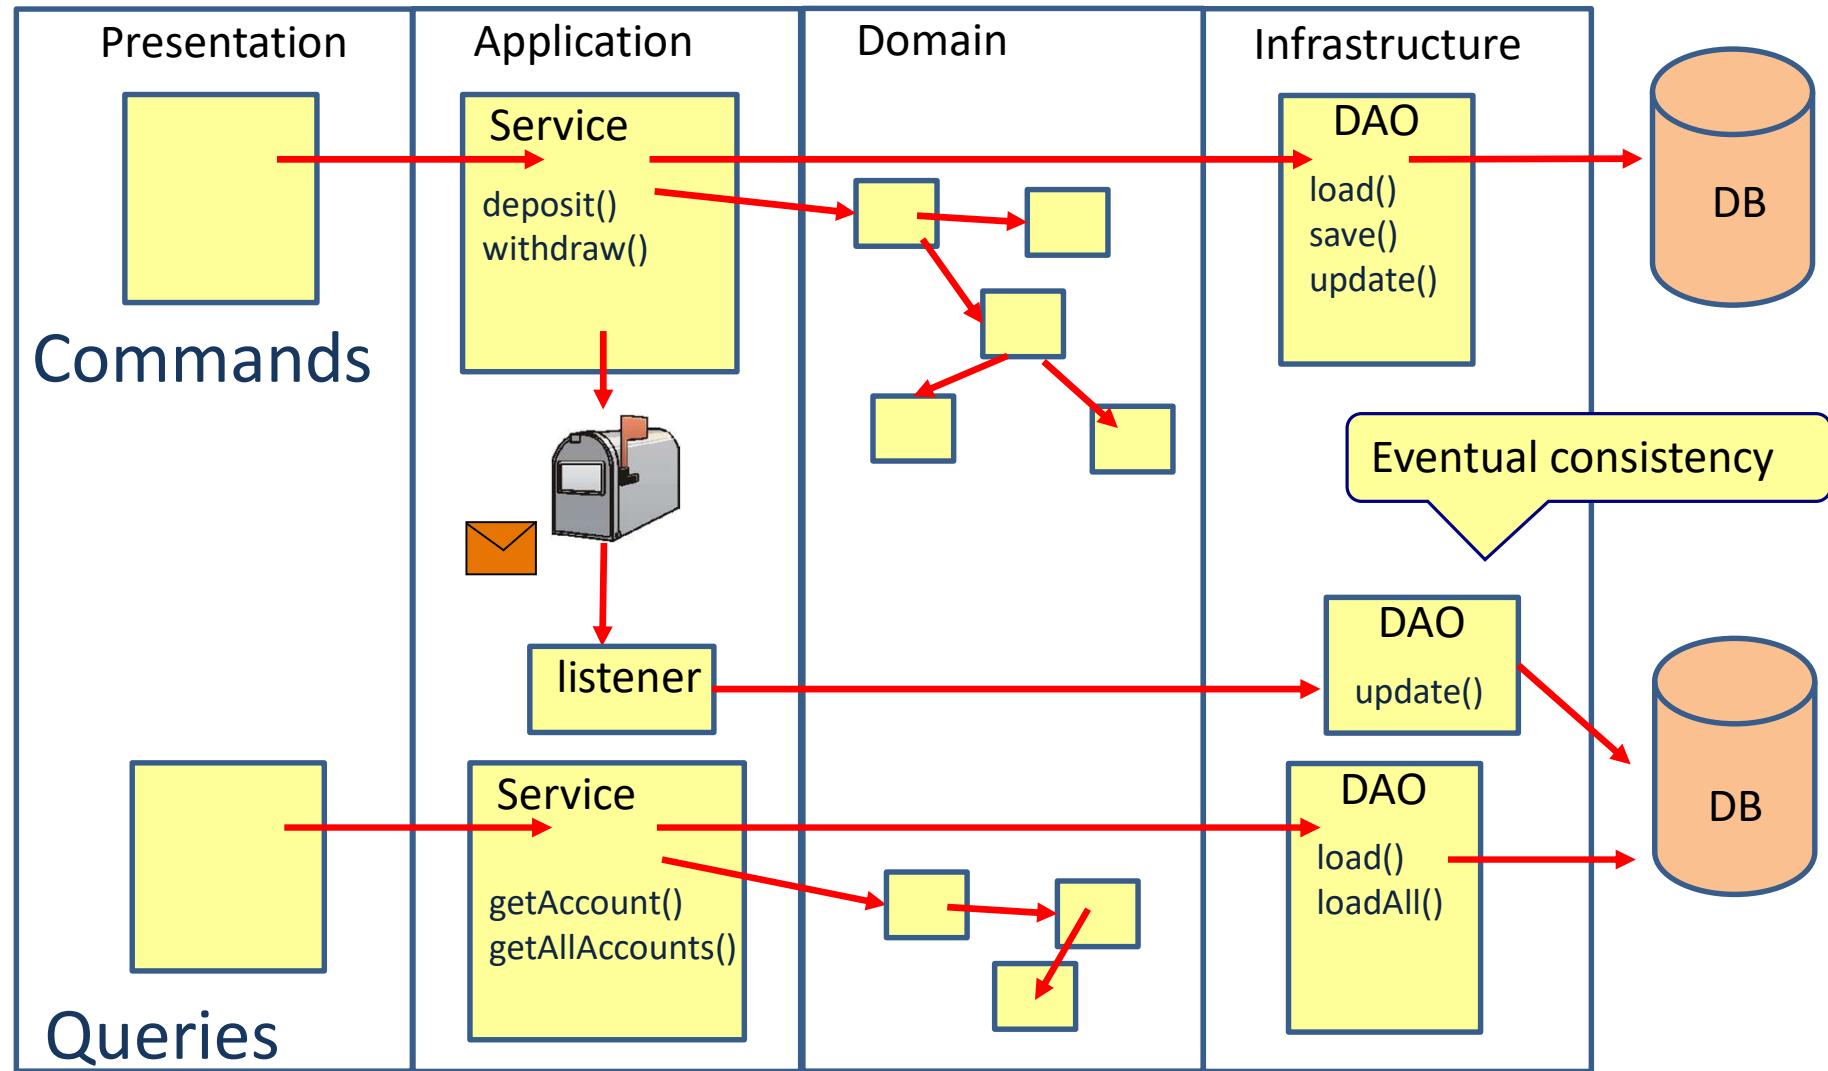
# Architectural properties

---

- Command and query side have different architectural properties
- Consistency
  - Command: needs consistency
  - Query: eventual consistency is mostly OK
- Data storage
  - Command: you want a normalized schema (3<sup>rd</sup> NF)
  - Query: denormalized (1<sup>st</sup> NF) is good for performance (no joins)
- Scalability
  - Command: commands happens not that often. Scalability is often not important.
  - Query: queries happen very often, scalability is important

# Eventual consistency

- Views will become eventual consistent



# CQRS advantages

---

- The query side is very simple
  - No transactions
- The query side can be optimized for performance
  - Fast noSQL database
  - Caching is easy
- Queries and commands can be scaled independently

# When to use CQRS?

---

- When queries and commands have different scaling requirements
- When read performance is critical
- When your screens start to look very different than your tables
- When you apply event sourcing

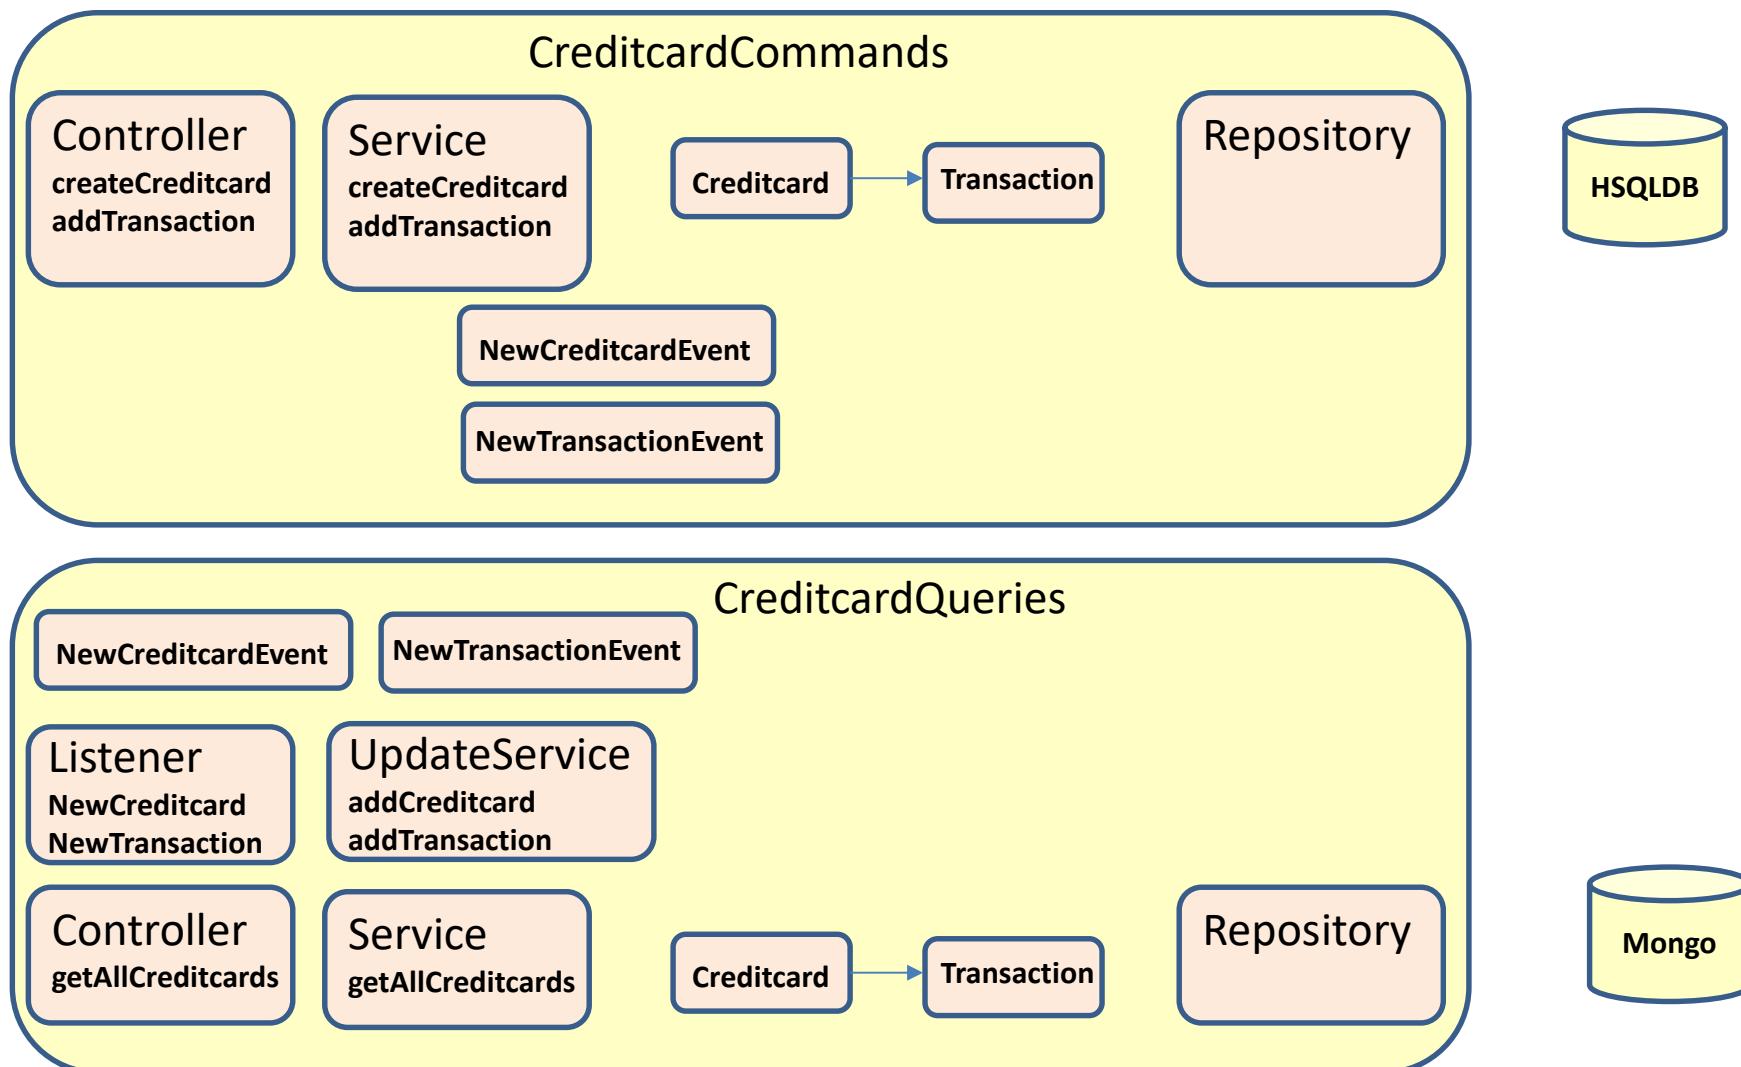
# When not to use CQRS

---

- Systems with simple CRUD functionality
- When you need strict consistency (instead of eventual consistency)

# CQRS example

---



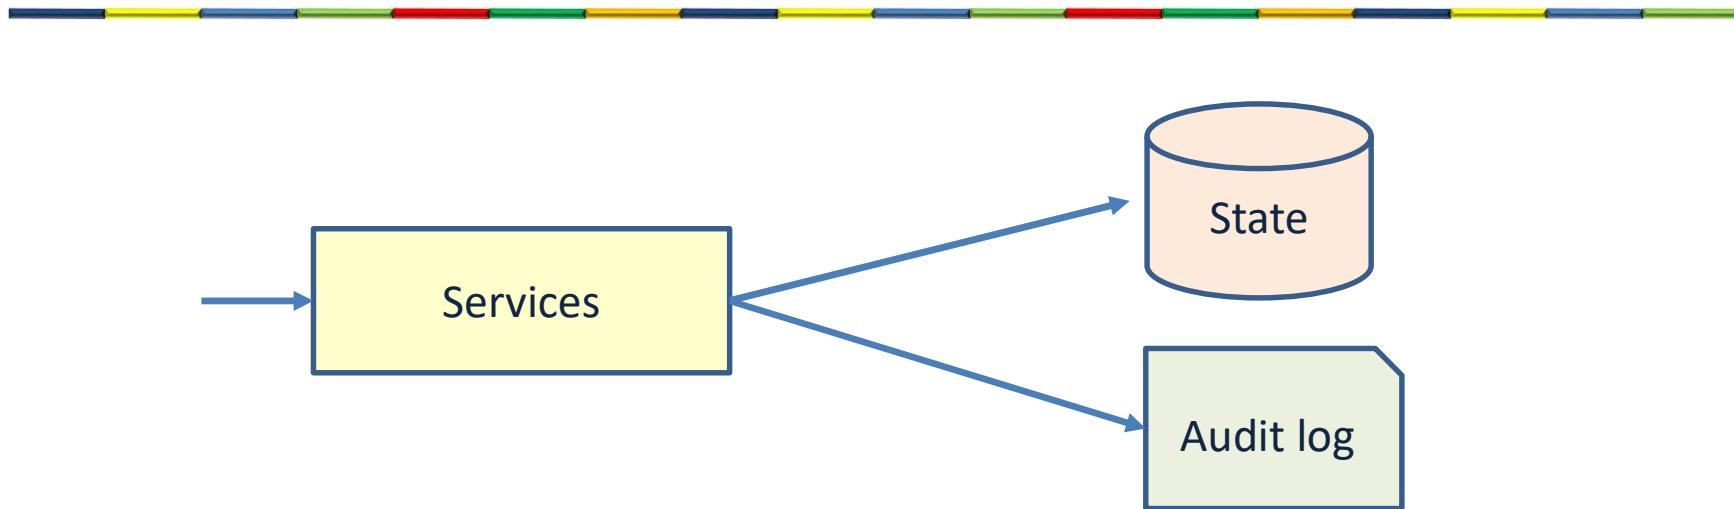
# **EVENT SOURCING**

# State based persistance

---

- You capture where you are, but not how you got there
- You don't know what happened in the past
- You cannot fix bad state because of bugs in the code

# Audit log

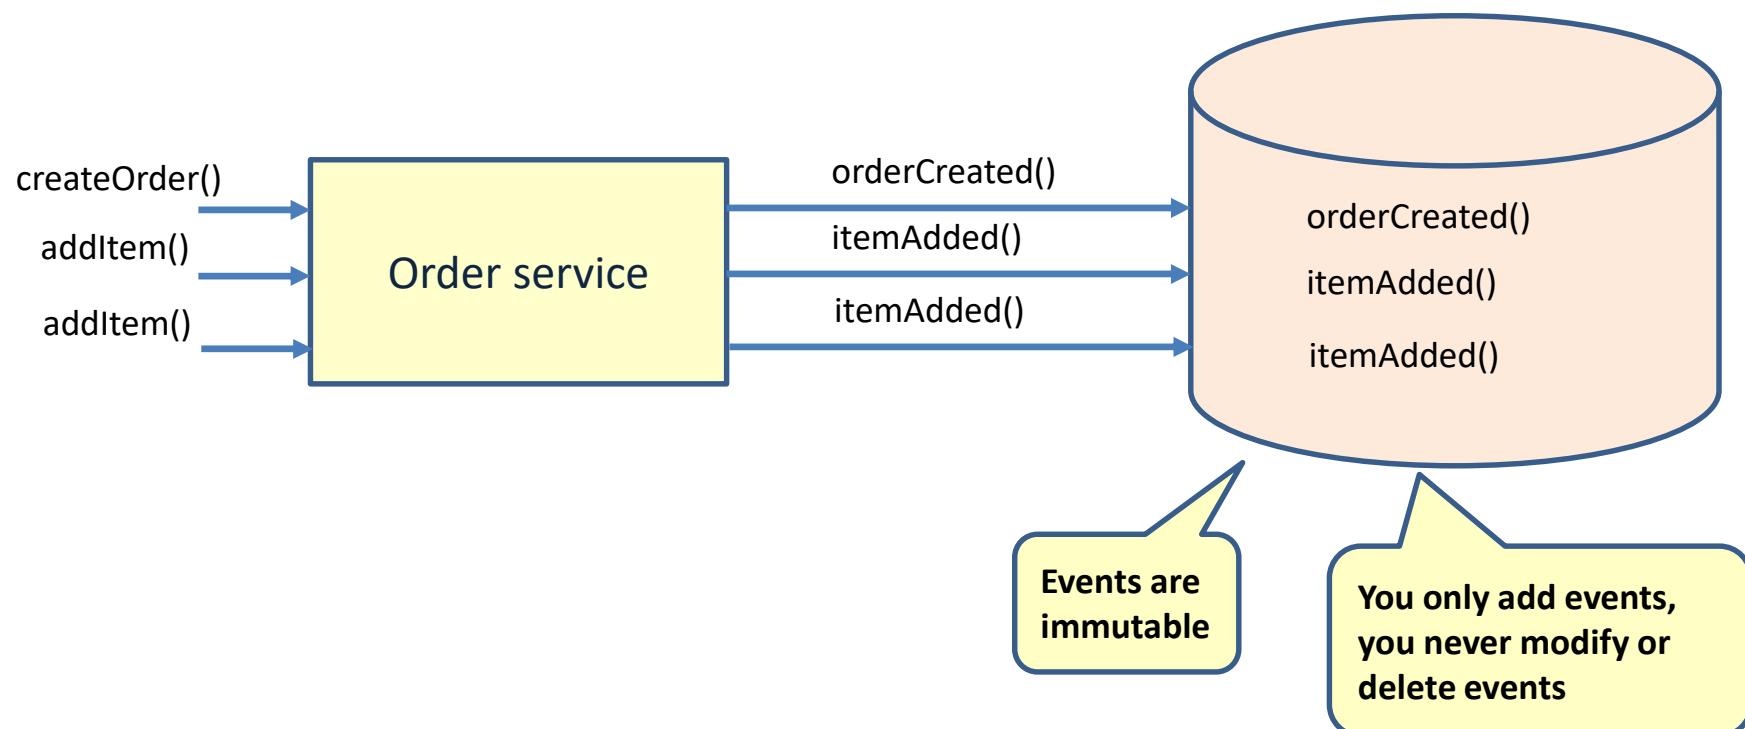


- Now we know how we got there
- What if the audit log gets out of sync with the state due to a bug?
- Which is the source of truth? (Audit log)

# Event sourcing

---

- Only capture how we got there



# Event sourcing

---

- Bank account
  - Store all deposits and withdrawals
- Phone account
  - Store all phone calls
- Version control systems
  - Each commit or change is an event
- DBMS
  - Databases keep a transaction log of all inserts, deletes and updates

# Storing state and storing events

---

- Store state

ID	status	data...
101	accepted	...

Store entity data

- Event sourcing

Entity ID	Entity type	Event ID	Event type	Event data...
101	Order	901	OrderCreated	...
101	Order	902	OrderApproved	...
101	Order	903	OrderShipped	...

Store state changing events

# Store the state of a system

---



## Structural representation

List of ordered goods

Payment information

Shipping information

# Store the events of a system

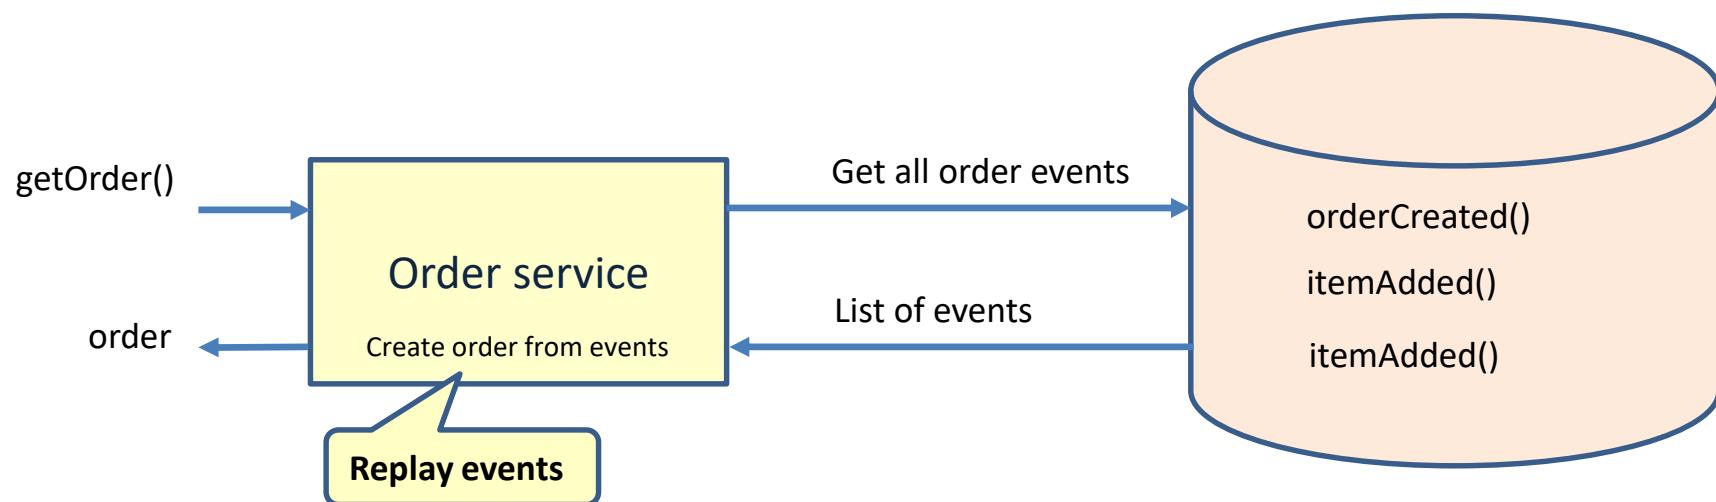
---



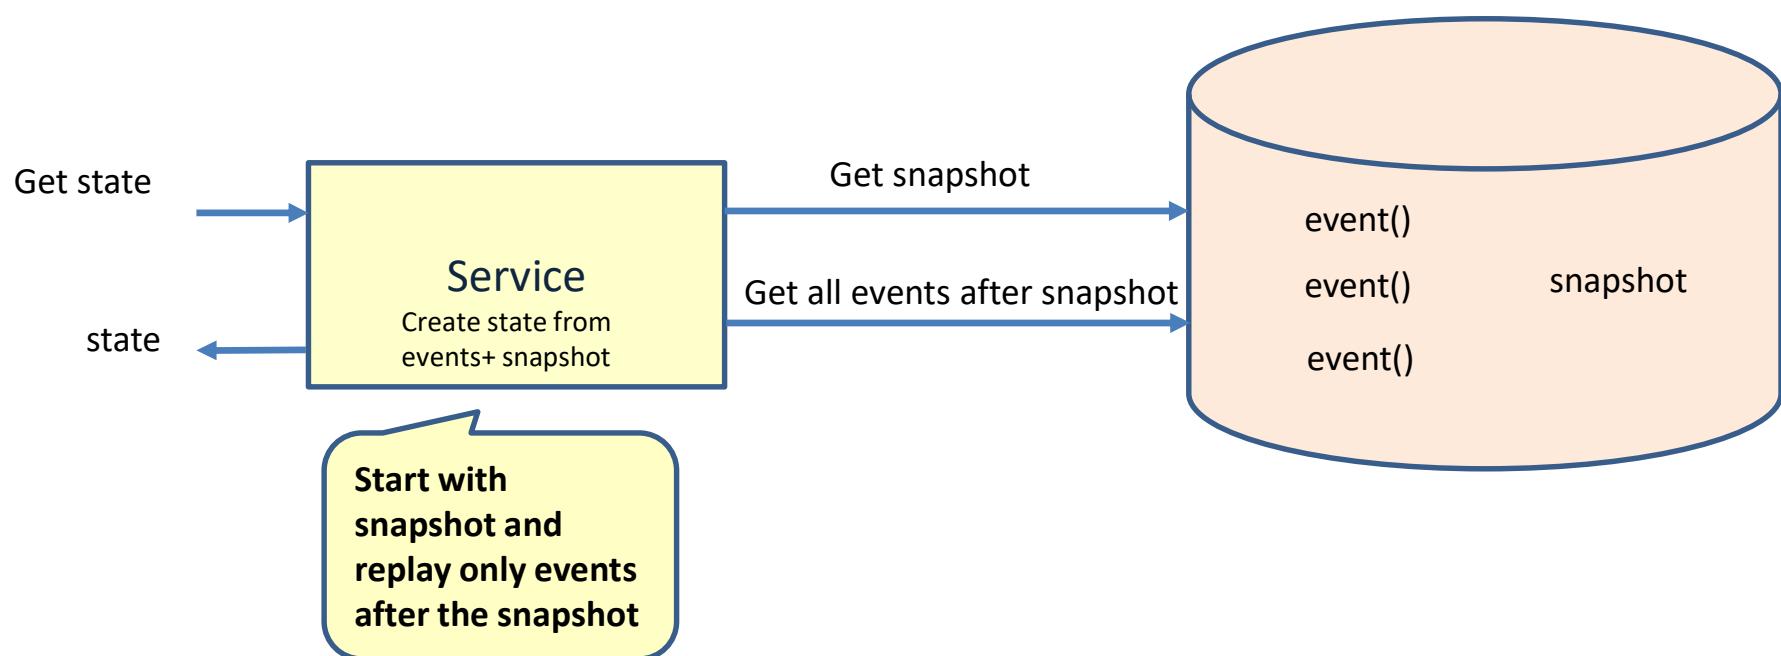
## Event representation

-  Add item #1
-  Add item #2
-  Add payment info
-  Update item #2
-  Remove item #1
-  Add shipping info

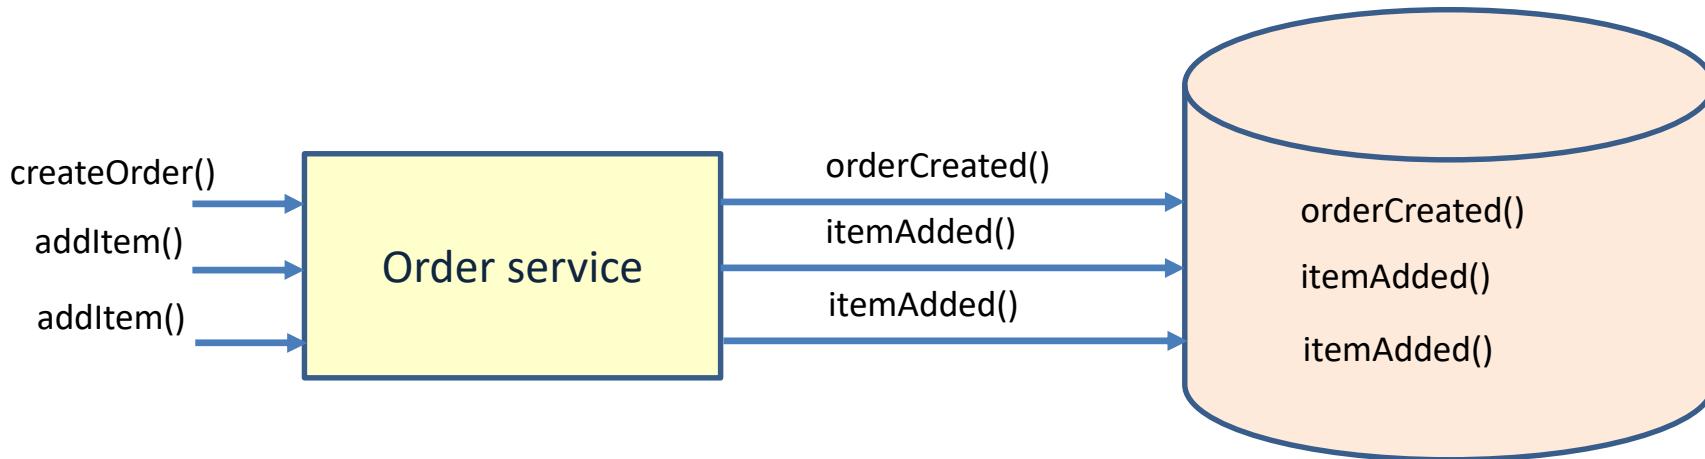
# How do we get the state?



# Snapshots

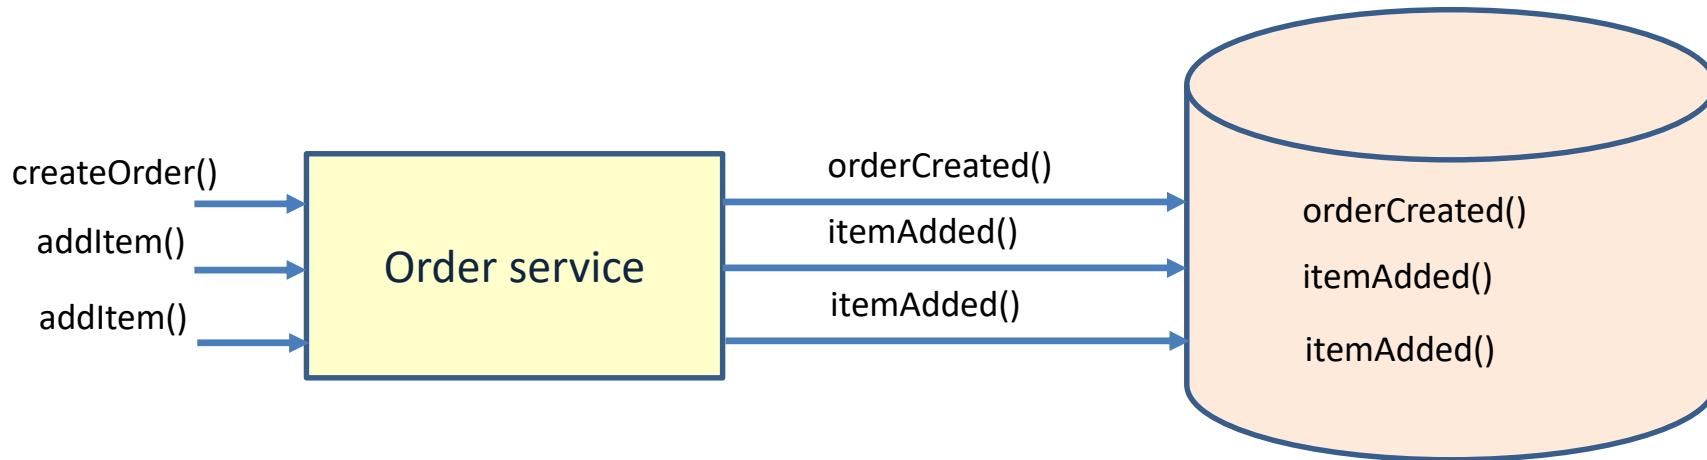


# Event sourcing advantages



- You don't miss a thing
  - Business can analyze history of events
  - Bugs can be solved easier
- Creates a built-in audit log
- Allows for rewinding or undo changes
- Append is usually more efficient in databases
- No transactions needed

# Event sourcing disadvantages



- You cannot easily see what your state is
  - You always have to replay the events first (lower read performance)
  - Always use CQRS when you apply event sourcing
- You need more storage
- What if events have to change?
  - We need event versioning
  - The logic needs to support all versions

# When to consider event sourcing?

---

- When you need an audit log
- When you want to derive additional business value from the event history
- When you need fast performance on the command side

# **CENTRALIZED CONFIGURATION SERVICE**

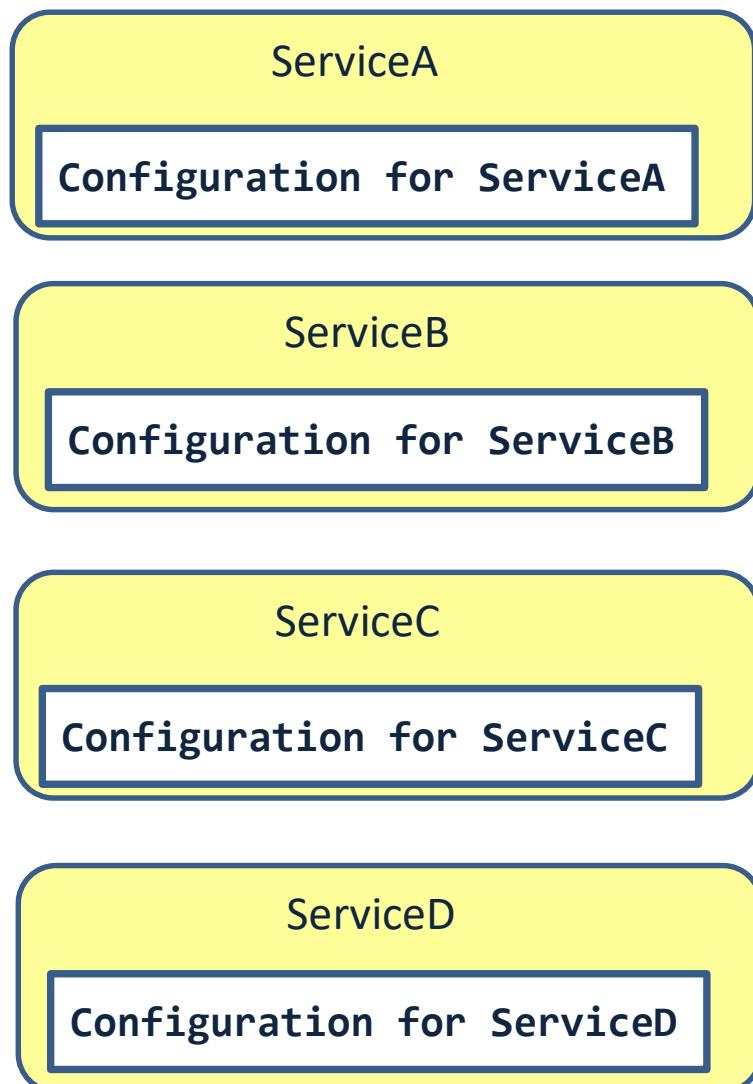
# Configuration in microservices

---

- Remove settings from code
- Change runtime behavior
- Enforce consistency across elastic services

# Local configuration

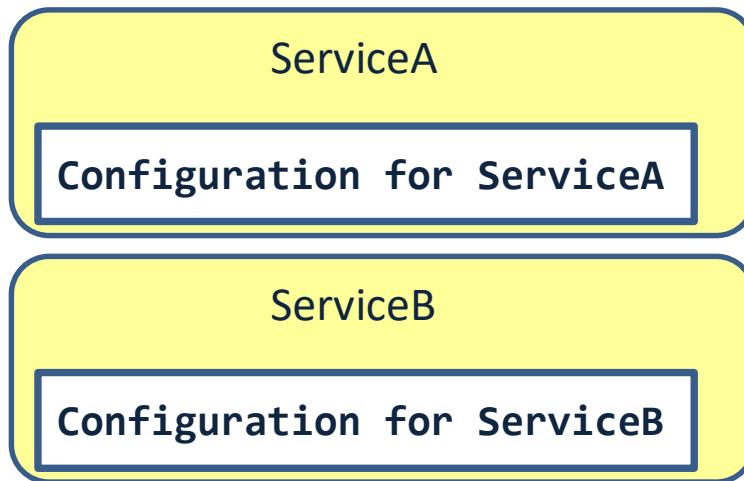
---



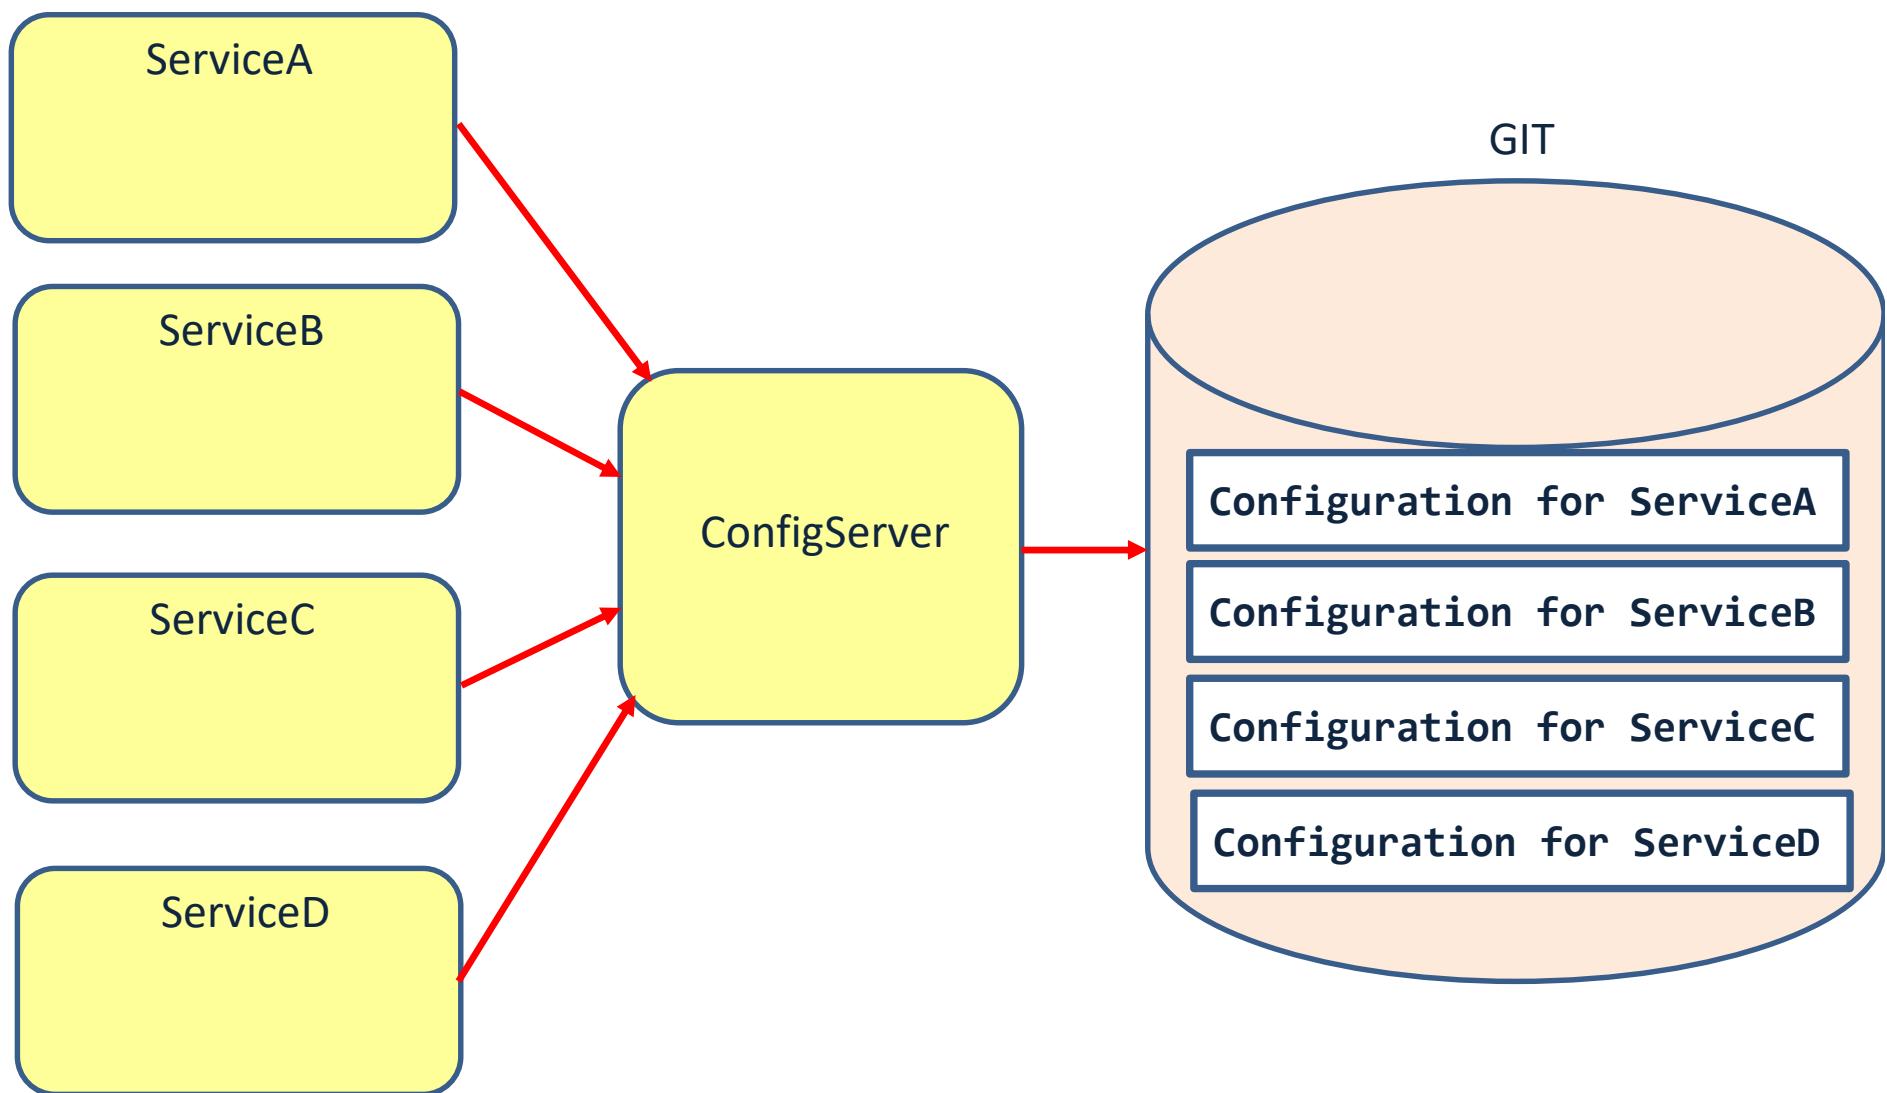
# Local configuration challenges

---

- When we change the configuration we need to rebuild and redeploy the application
- Configuration may contain sensitive information
- Some of the properties are the same among services: lots of duplication



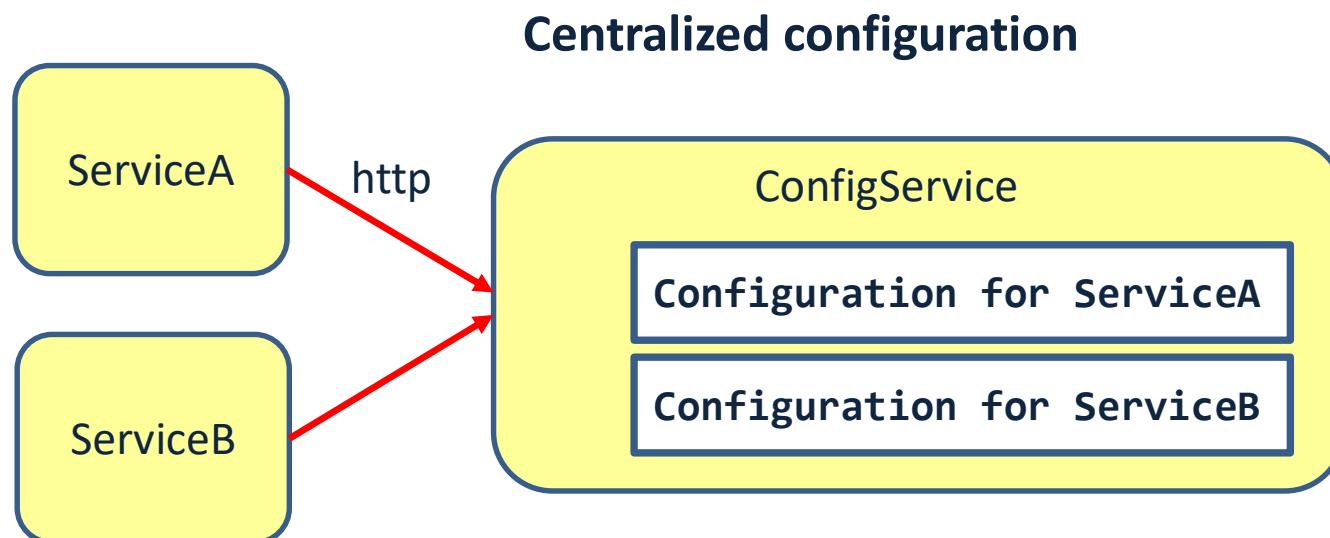
# Spring cloud config server



# Spring cloud config

---

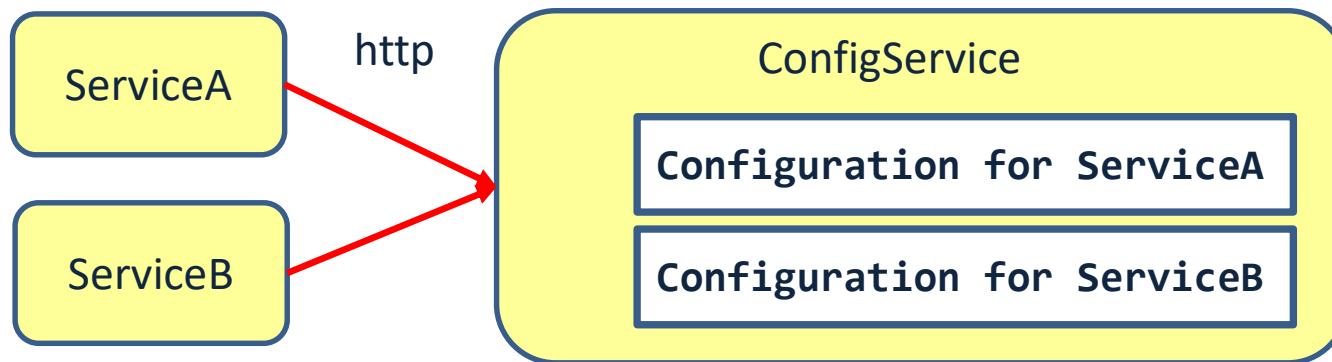
- HTTP access to centralized configuration



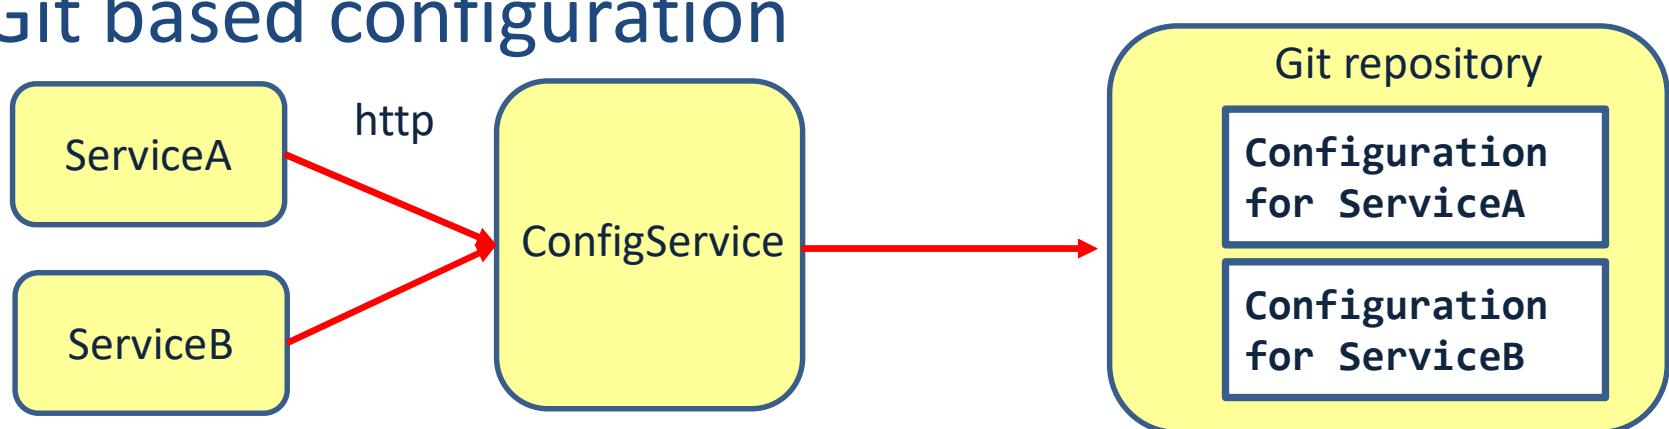
# Spring cloud config

---

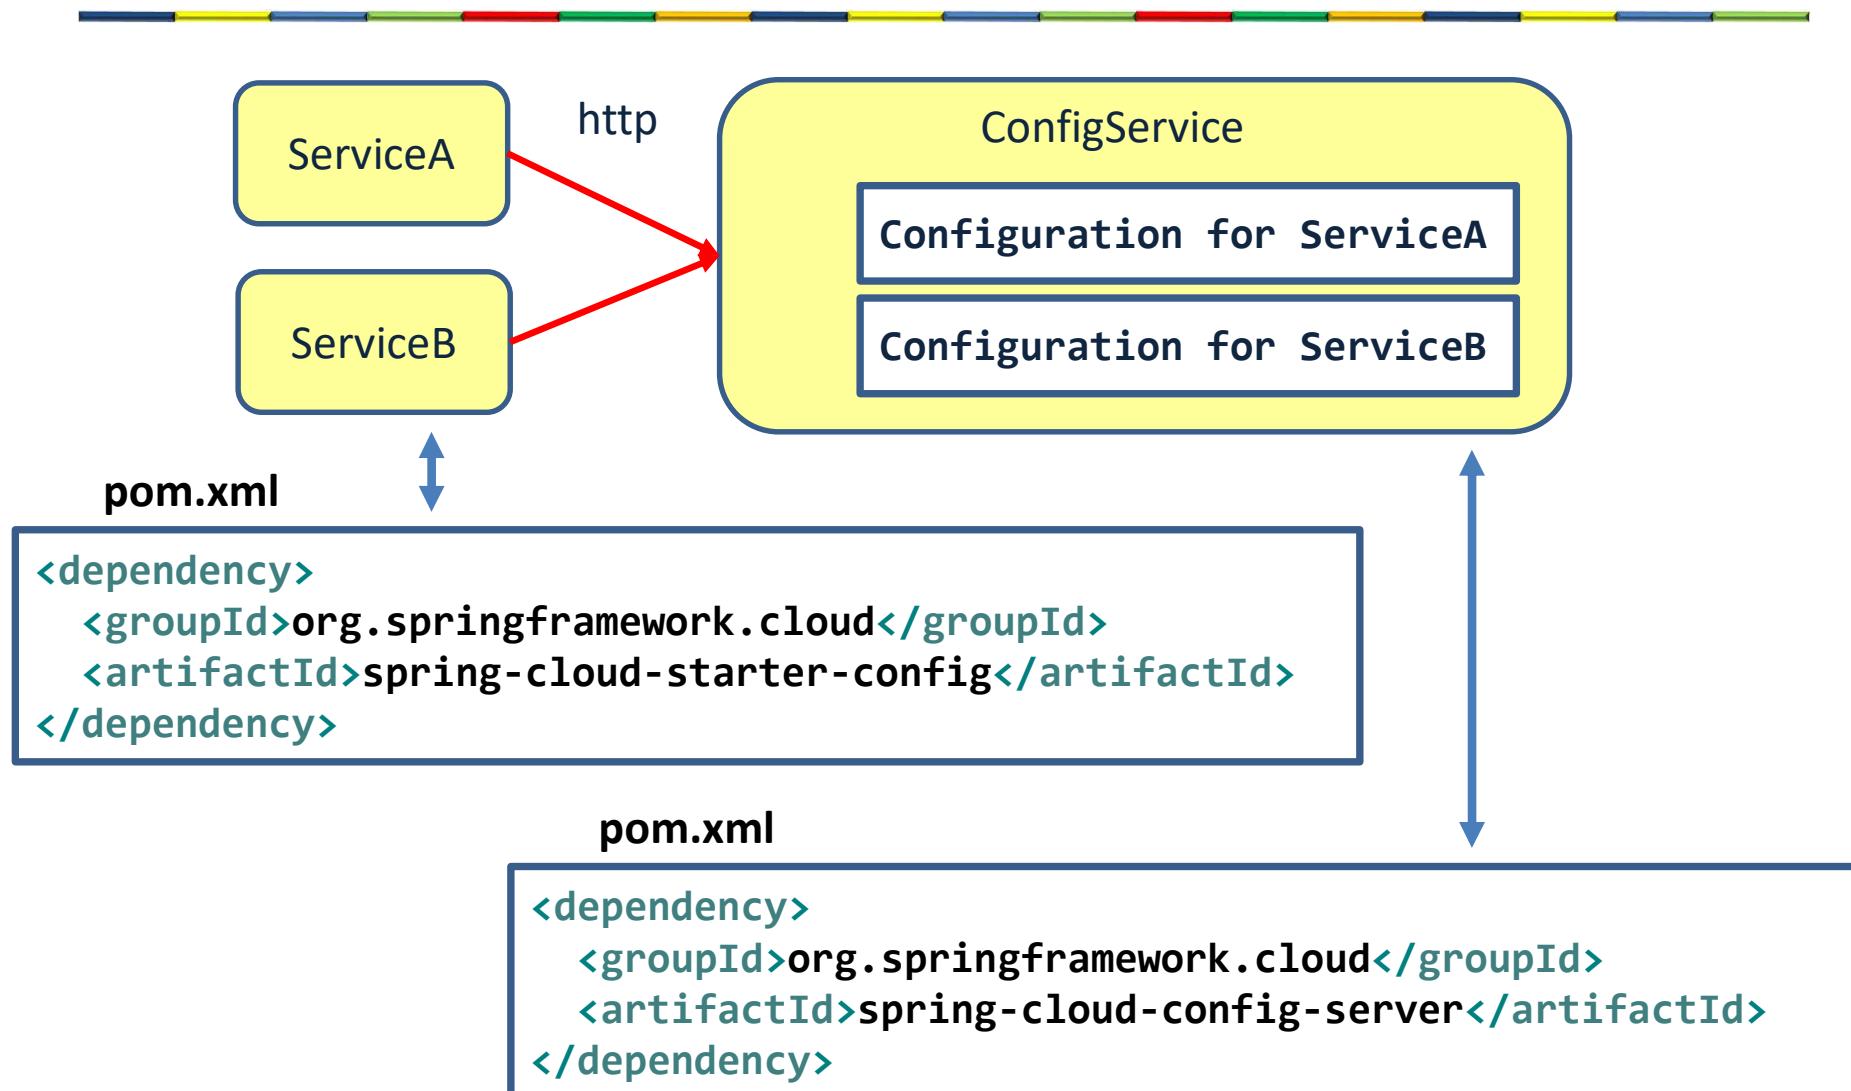
- File based configuration



- Git based configuration



# Spring cloud config example



# Configuration server

```
import org.springframework.boot.SpringApplication;
import org.springframework.boot.autoconfigure.SpringBootApplication;
import org.springframework.cloud.config.server.EnableConfigServer;

@SpringBootApplication
@EnableConfigServer
public class ConfigServiceApplication {

    public static void main(String[] args) {
        SpringApplication.run(ConfigServiceApplication.class, args);
    }
}
```

application.properties

```
spring.profiles.active=native
server.port=8888
```

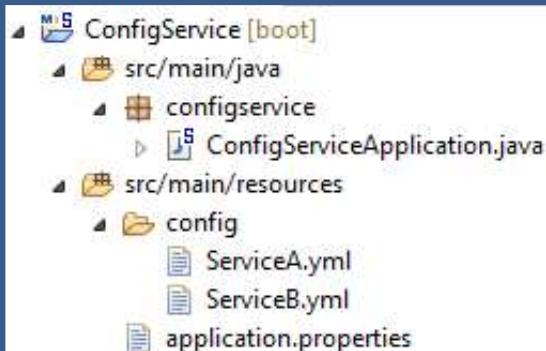
Do not use GIT,  
but local files

config/ServiceA.yml

```
greeting: Hello from Service A
```

config/ServiceB.yml

```
greeting: Hello from Service B
```

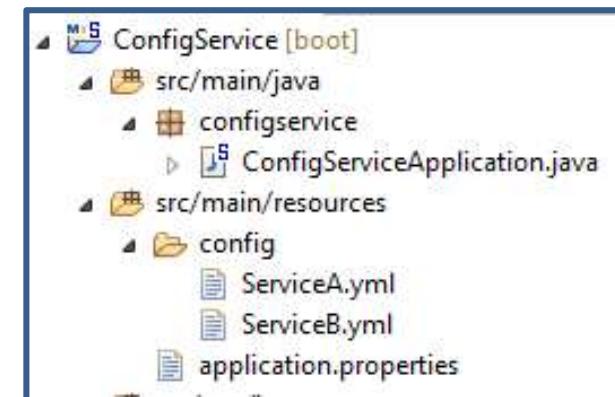


# Configuration server



A screenshot of a web browser window. The address bar shows "localhost:8888/ServiceA/default". The page content displays a JSON object:

```
{"name": "ServiceA", "profiles": [{"default"], "label": null, "version": null, "state": null, "propertySources": [{"name": "classpath:/config/ServiceA.yml", "source": {"greeting": "Hello from Service A"}}]}]
```



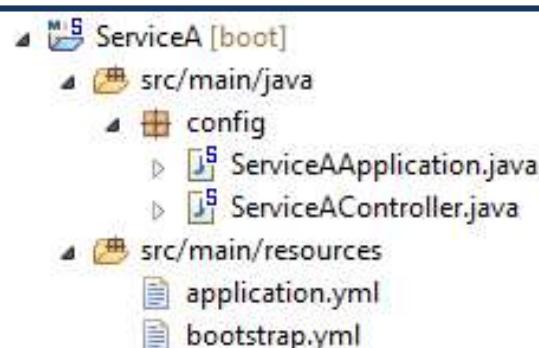
A screenshot of a web browser window. The address bar shows "localhost:8888/ServiceB/default". The page content displays a JSON object:

```
{"name": "ServiceB", "profiles": [{"default"], "label": null, "version": null, "state": null, "propertySources": [{"name": "classpath:/config/ServiceB.yml", "source": {"greeting": "Hello from Service B"}}]}]
```

# Configuration client: ServiceA

```
@SpringBootApplication  
public class ServiceAApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(ServiceAApplication.class, args);  
    }  
}
```

```
@RestController  
public class ServiceAController {  
    @Value("${greeting}")  
    private String message;  
  
    @RequestMapping("/")  
    public String getName() {  
        return message;  
    }  
}
```



**application.yml**

```
server:  
  port: 8090
```

**bootstrap.yml**

```
spring:  
  application:  
    name: ServiceA  
  cloud:  
    config:  
      url: http://localhost:8888
```

# Spring cloud applications

---

- 2 configuration files

- bootstrap.yml

- Is loaded before applications.yml
    - Is needed when configuration is stored on a remote config server
    - Contains
      - The name of the application
      - Location of the configuration server

- applications.yml

- Contains standard application configuration

```
spring:          bootstrap.yml
  application:
    name: ServiceA
  cloud:
    config:
      url: http://localhost:8888
```

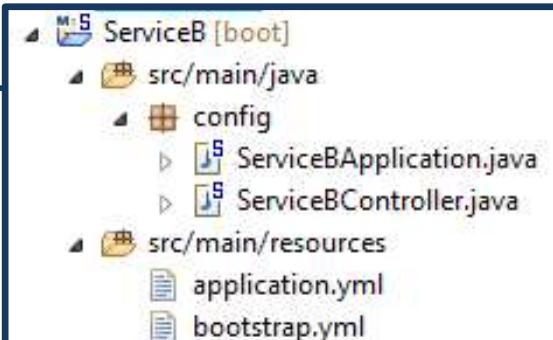
**application.yml**

```
server:
  port: 8090
```

# Configuration client: ServiceB

```
@SpringBootApplication  
public class ServiceBApplication {  
  
    public static void main(String[] args) {  
        SpringApplication.run(ServiceBApplication.class, args);  
    }  
}
```

```
@RestController  
public class ServiceBController {  
    @Value("${greeting}")  
    private String message;  
  
    @RequestMapping("/")  
    public String getName() {  
        return message;  
    }  
}
```



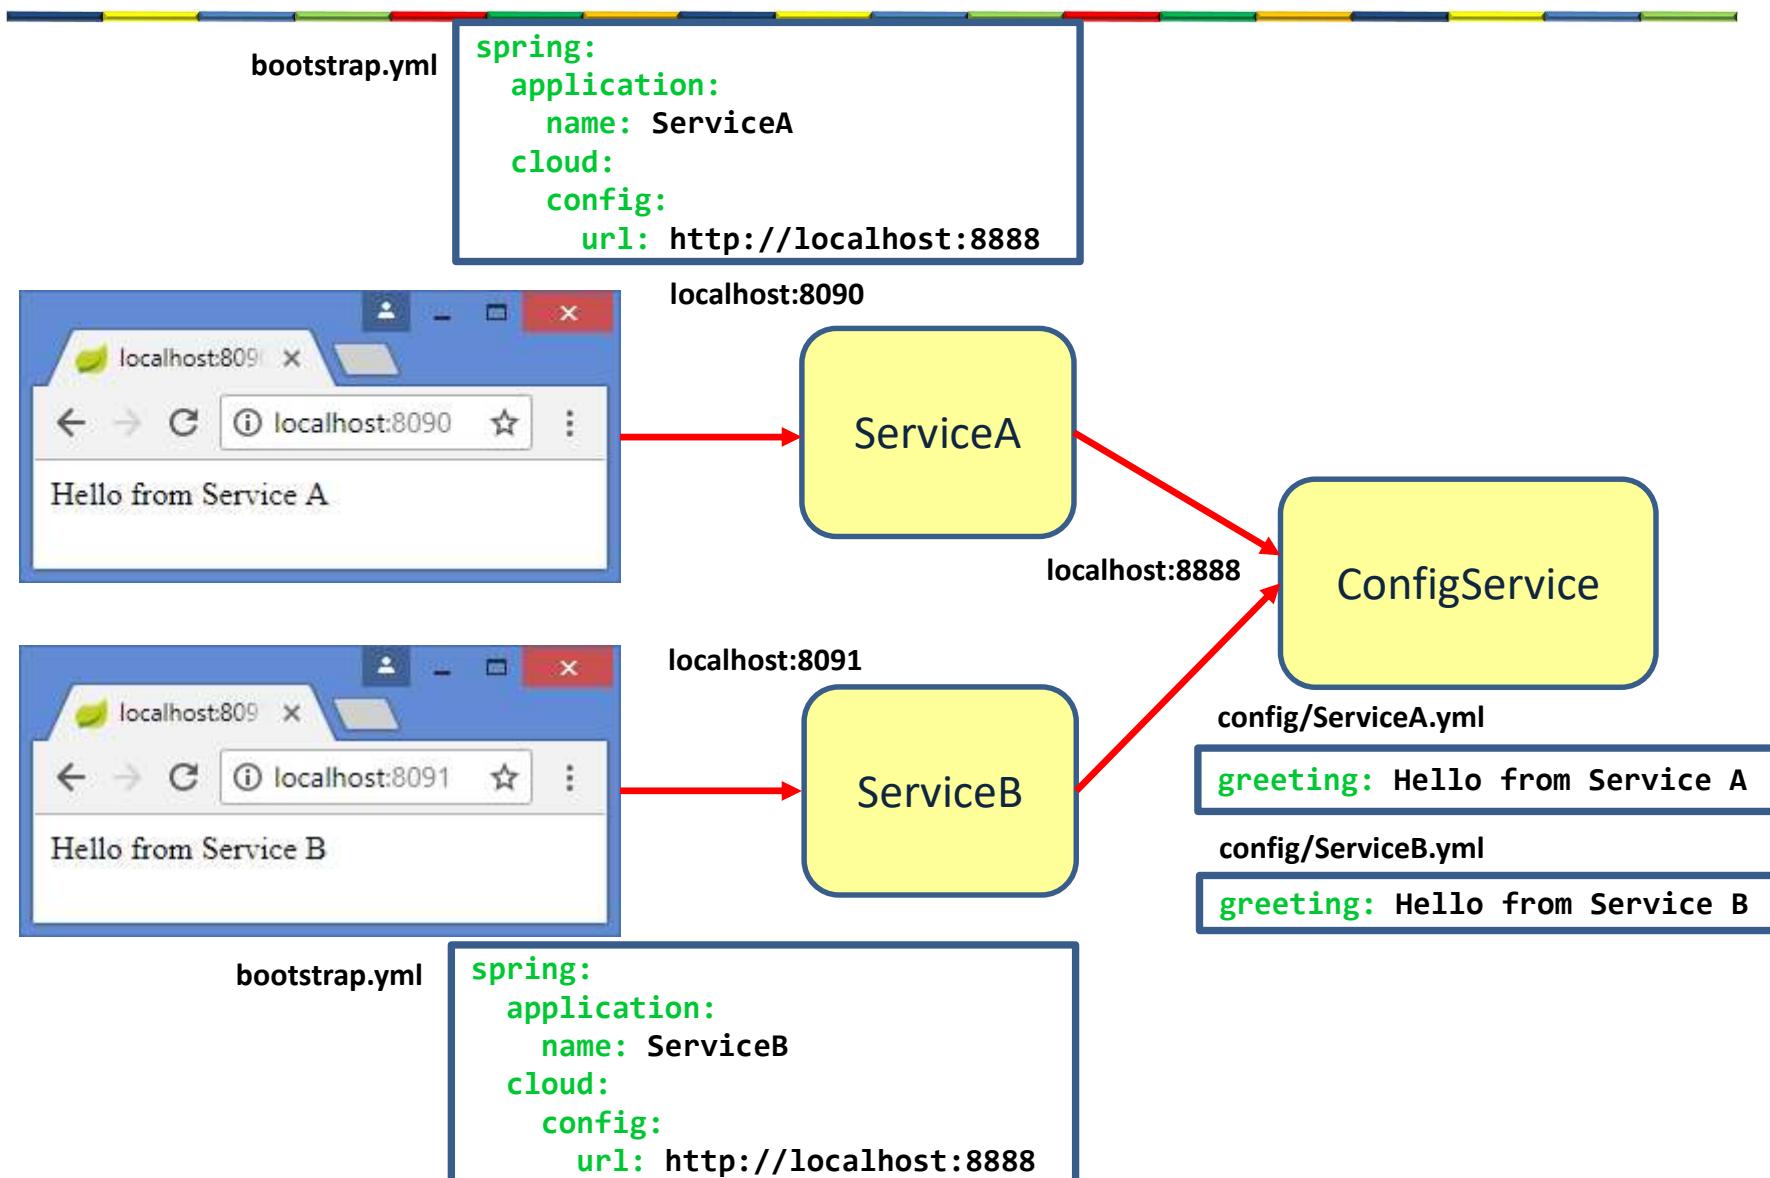
**application.yml**

```
server:  
  port: 8091
```

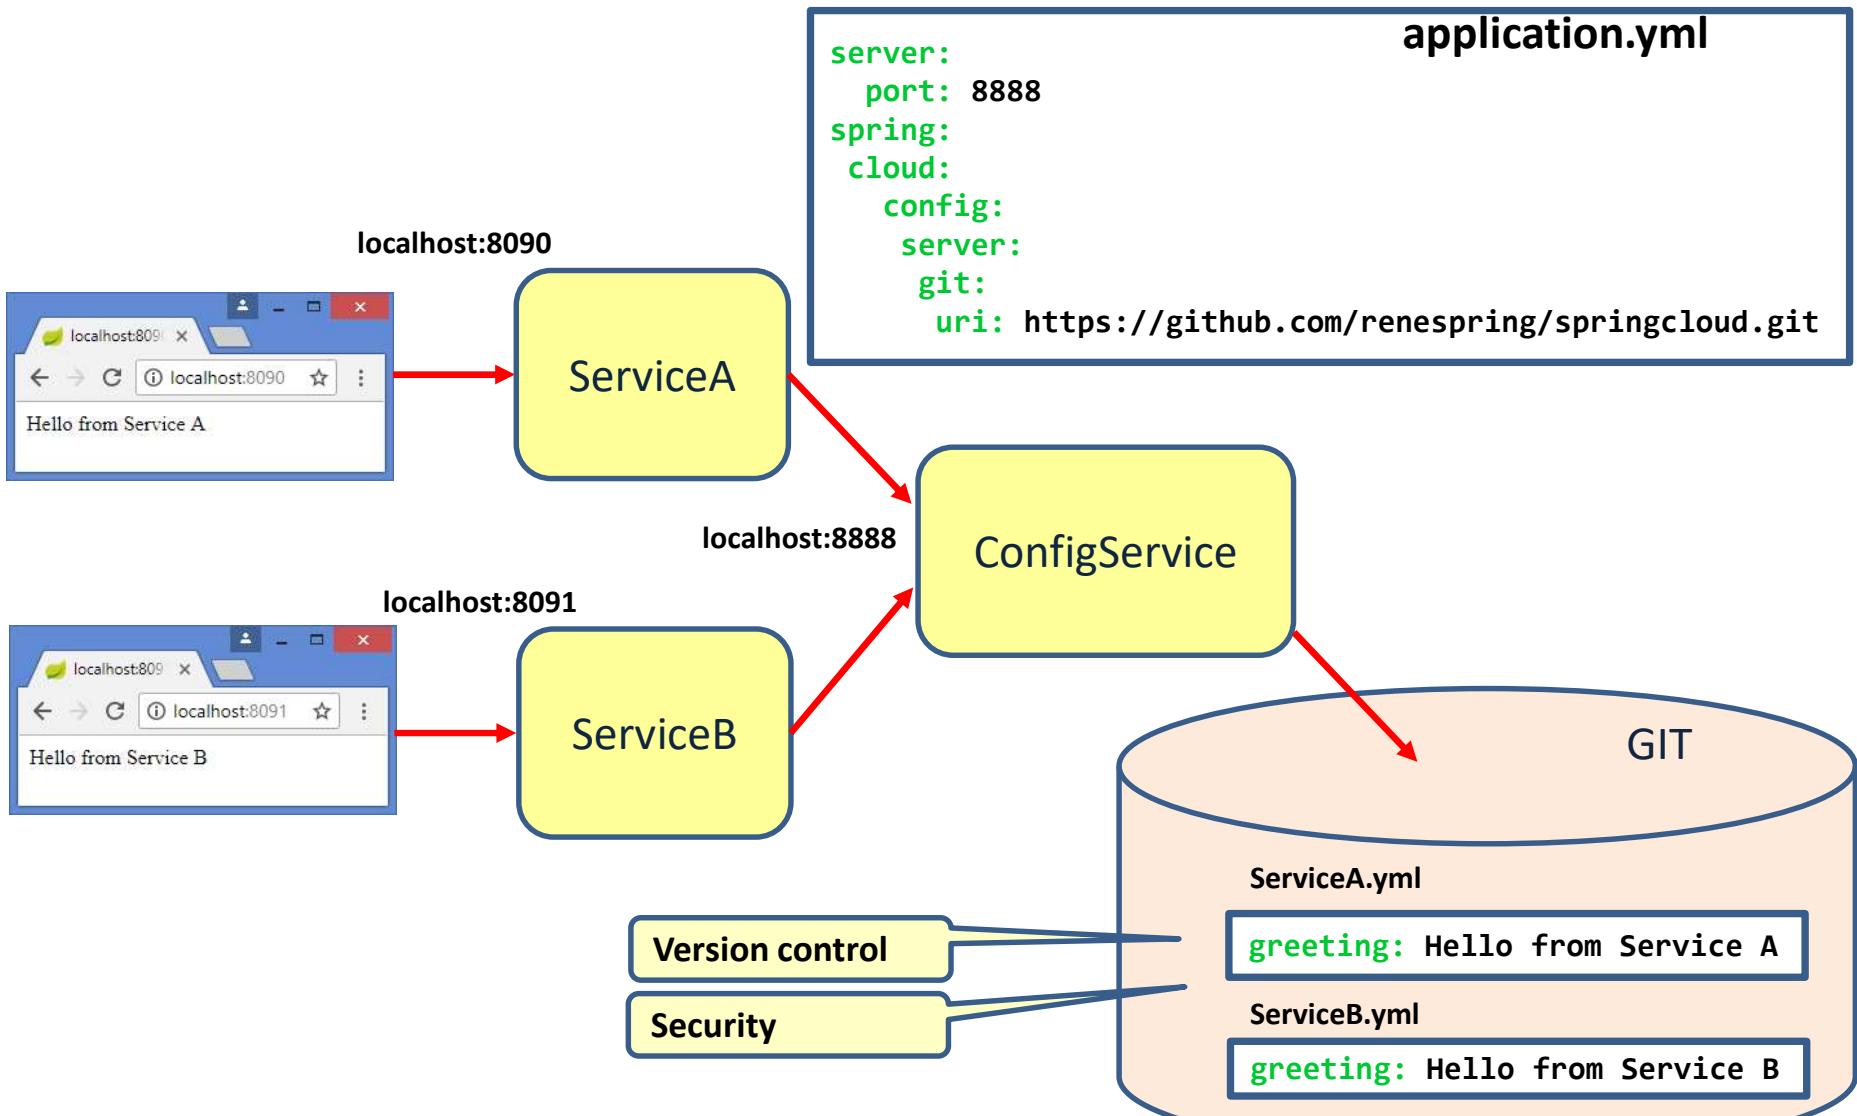
**bootstrap.yml**

```
spring:  
  application:  
    name: ServiceB  
  cloud:  
    config:  
      url: http://localhost:8888
```

# Use of the Config Server



# Using Git (or GitHub)



# Refreshing configuration

---

- Option 1: use `@RefreshScope` and “`/actuator/refresh`” event

```
    @RestController  
    @RefreshScope  
    public class ServiceAController {  
        @Value("${greeting}")  
        private String message;  
  
        @RequestMapping("/")  
        public String getName() {  
            return message;  
        }  
    }
```

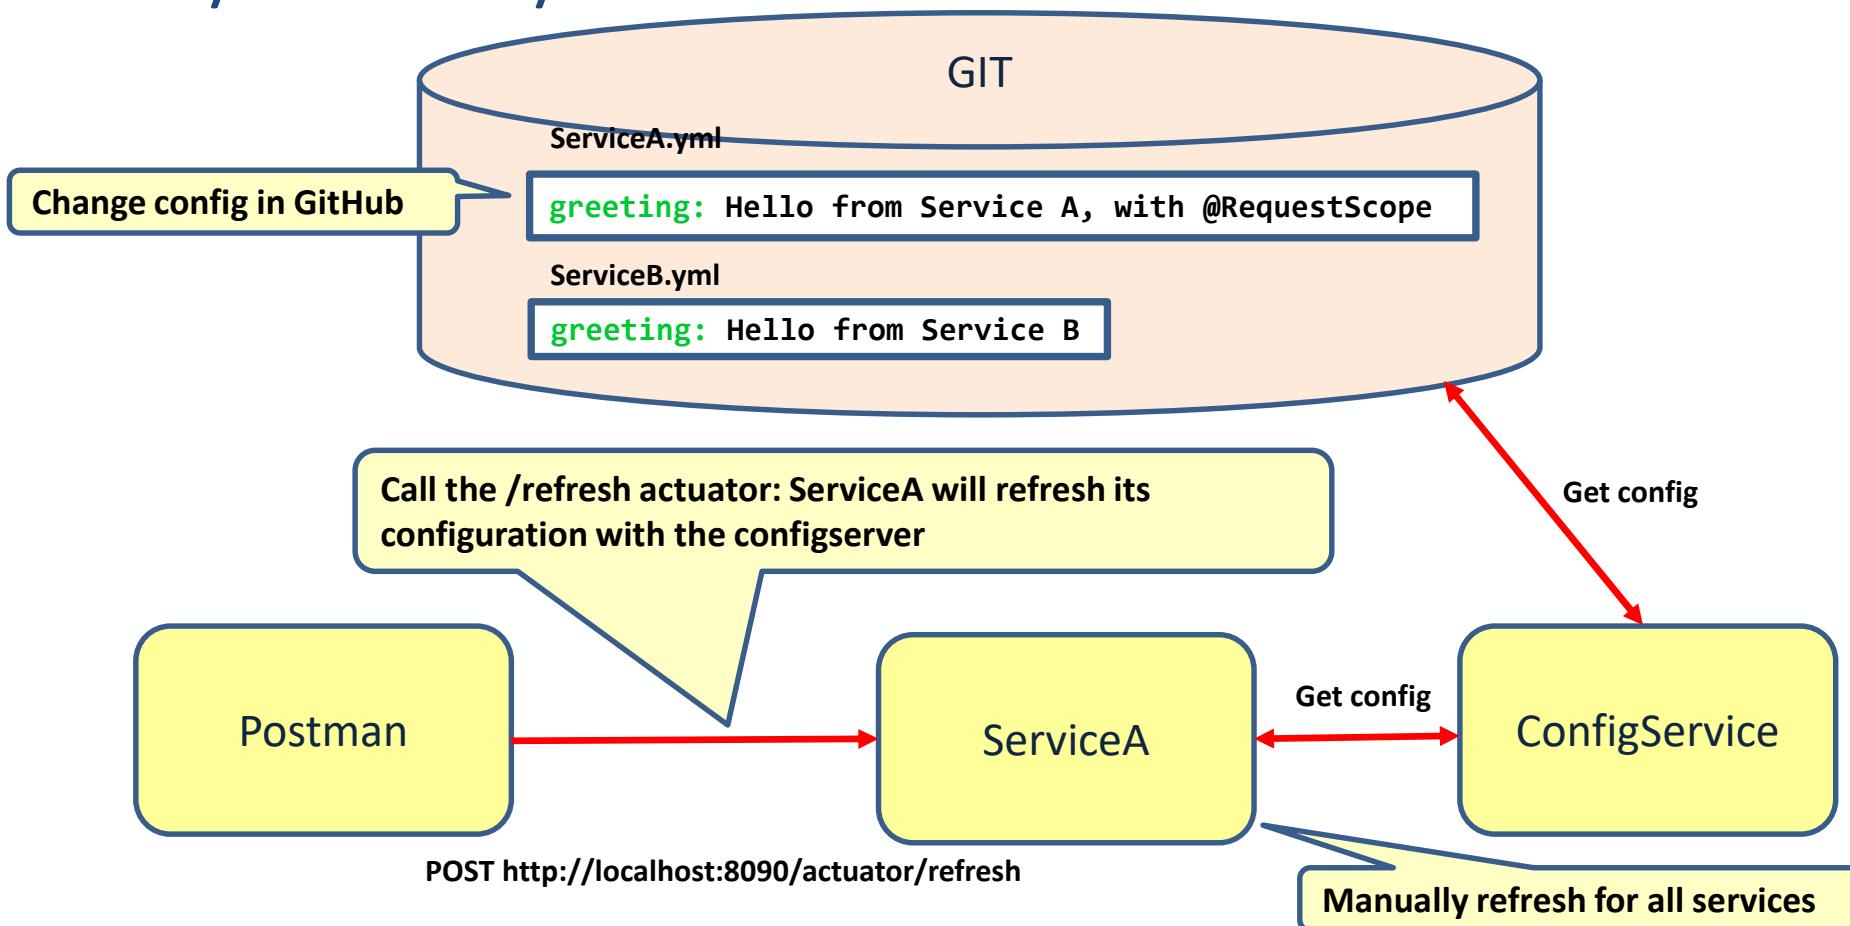
`@RefreshScope`

```
server:  
port: 8888  
  
spring:  
cloud:  
config:  
server:  
git:  
uri:  
https://github.com/renespring/springcloud.git  
  
management:  
endpoints:  
web:  
exposure:  
include: refresh
```

Expose the `/refresh` actuator

# Refreshing configuration

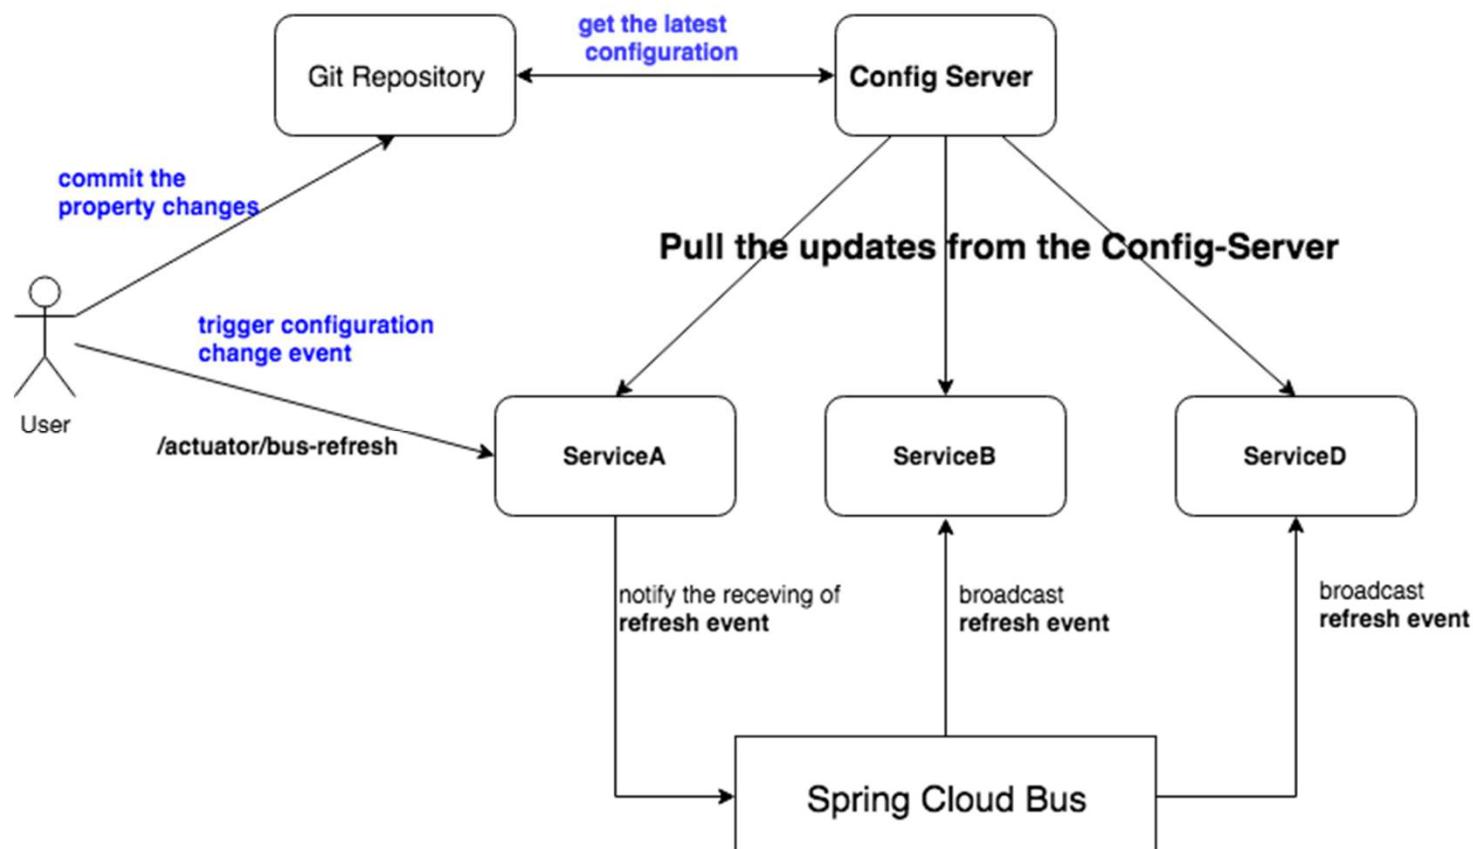
- Option 1: use `@RefreshScope` and “`/actuator/refresh`” event



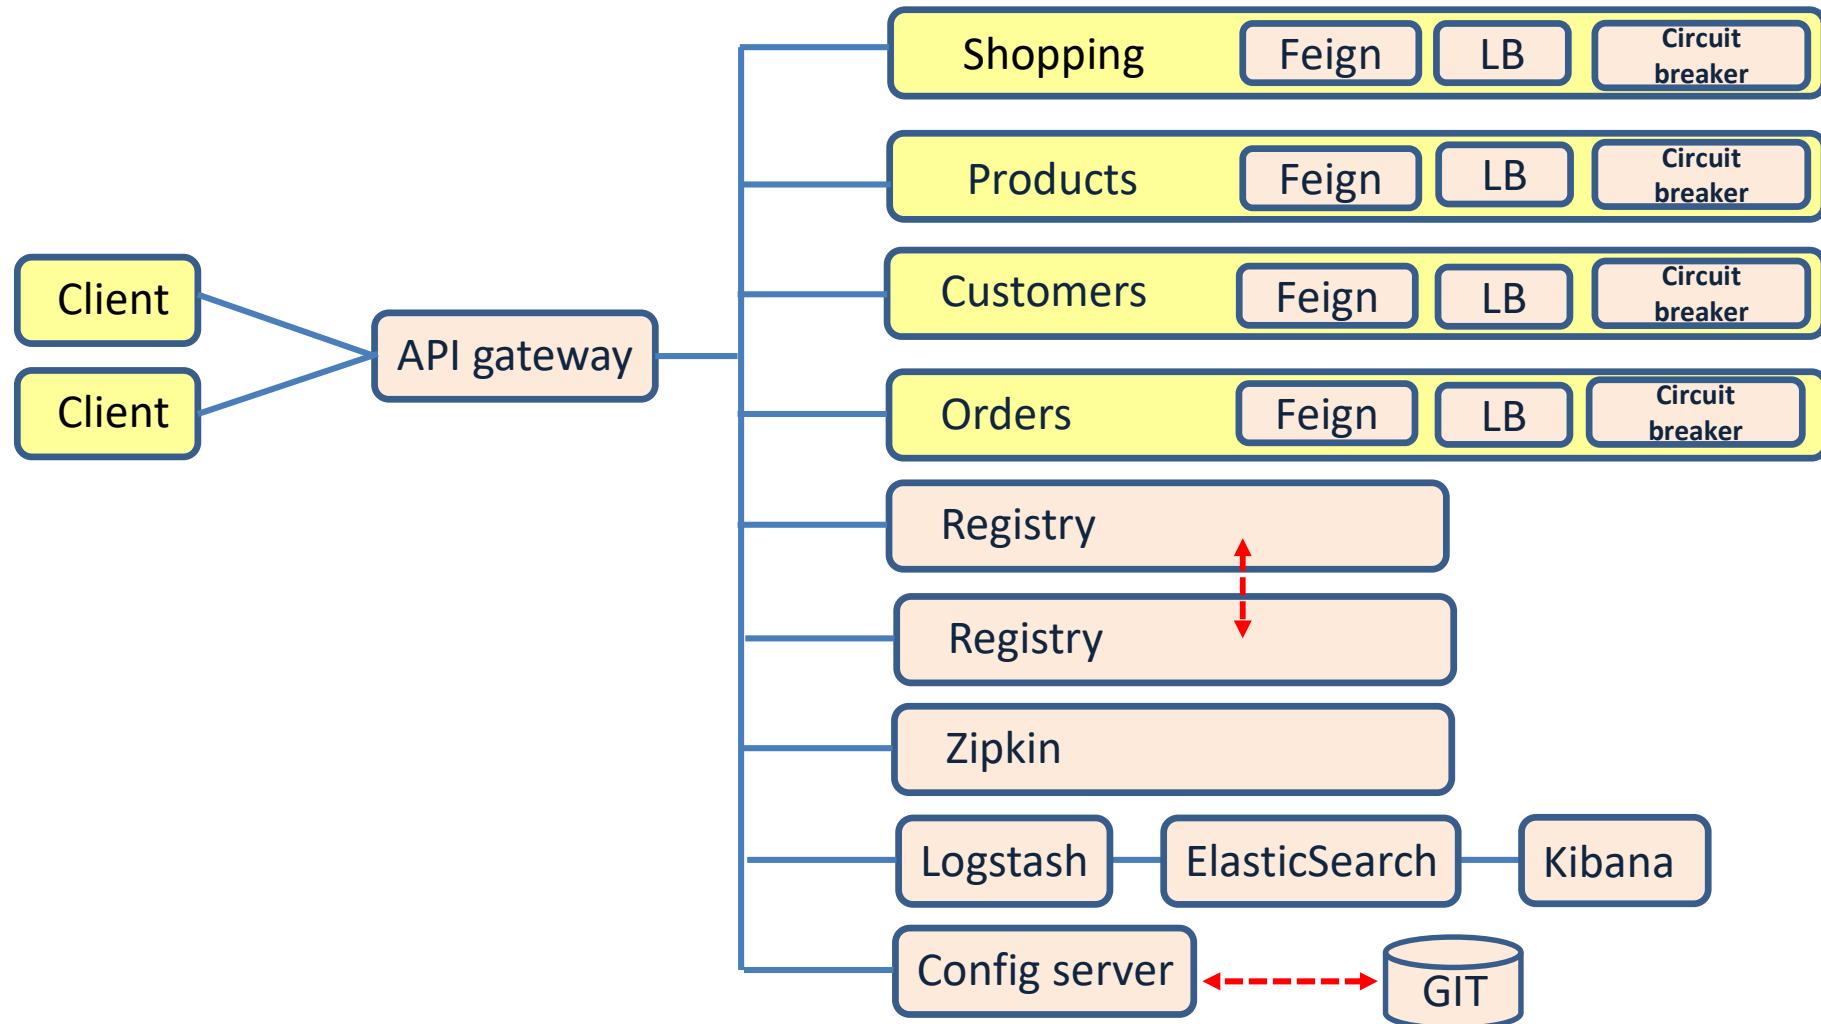
# Refreshing configuration

---

- Option 2: use Spring cloud bus for broadcasting refresh events



# Implementing microservices



# Challenges of a microservice architecture

Challenge	Solution
Complex communication	Feign Registry API gateway
Performance	
Resilience	Registry replicas Load balancing between multiple service instances Circuit breaker
Security	
Transactions	
Keep data in sync	
Keep interfaces in sync	Spring cloud contract
Keep configuration in sync	Config server
Monitor health of microservices	ELK + beats
Follow/monitor business processes	Zipkin ELK